

Medicaid Lock and Parental Labor Supply: Evidence from Multi-Year Continuous-Eligibility Waivers for Young Children

Abstract

Background. Section 1902(e)(12) of the Social Security Act, as amended by the Consolidated Appropriations Act, 2023 (CAA-2024), required all state Medicaid and CHIP programs to provide 12 months of continuous eligibility to children under age 19 starting January 1, 2024. Four states — Oregon (effective January 2023), Washington (July 2023), New Mexico (November 2023), and Pennsylvania (January 2025) — went further, securing §1115 waivers that extend continuous eligibility from birth through age 5 or 6 with re-determination every six years. I exploit this state-level variation to estimate effects on (a) child Medicaid coverage and (b) parental labor supply.

Methods. I combine four data sources: National Survey of Children’s Health 2016–2024 (N=386,083 child-years); American Community Survey 1-year 2016–2024 (N=8.5M individuals ages 0–25); Current Population Survey Annual Social and Economic Supplement 2016–2025 with a parent-child match (376k children linked to 337k parents at 99.3% match rate); and the Survey of Income and Program Participation 2018–2024 public-use files (928k child-months across 77,779 unique sibling-pair household-months). Identification rests on a four-layer design: (1) cross-state staggered difference-in-differences on coverage outcomes; (2) within-state under-6-vs-19–25 triple-difference; (3) cross-state staggered DiD on parental labor outcomes; (4) within-household sibling fixed-effects comparing the under-6 sibling to their 7–12 sibling in the same household-month. I address the concurrent post-pandemic Medicaid “unwinding” confound with state-year unwinding-intensity controls drawn from the CMS Unwinding Tracker and KFF state-by-state monthly reports.

Results. Parents of children ages 0–6 in waiver-adopting states (Oregon, Washington, New Mexico) work **+0.71 more hours per week** ($p = 0.027$) and are **+1.66 percentage points more likely to work full-time** ($p = 0.013$) after waiver enactment. Effects are on the intensive labor margin only — labor-force participation and employment are flat. Annual wage income rises by approximately \$4,400 ($p = 0.067$, marginal). These labor estimates survive state-year unwinding-intensity controls (hours: +0.70, $p = 0.056$; full-time: +1.52pp, $p = 0.053$), Sun-Abraham (2021) heterogeneity-robust estimation (hours: +0.74; full-time: +1.89pp), an event-study parallel-trends test (pre-treatment leads statistically indistinguishable from zero), and a drop-NM sensitivity that *strengthens* the wage effect to **+\$6,167 ($p < 0.001$)**. By contrast, child Medicaid coverage effects are smaller and more mixed: a 0.63 percentage-point reduction in under-6 uninsurance in ACS ($p = 0.025$) does not survive unwinding controls (-0.26pp, $p = 0.394$), and the within-household sibling-FE estimate on monthly Medicaid coverage is +3.3 percentage points ($p = 0.153$, not statistically significant).

Conclusions. Multi-year continuous-eligibility waivers for young children produce a measurable and robust increase in parental work hours and full-time employment, consistent with the canonical “Medicaid lock” prediction: when a child’s coverage is decoupled from the parent’s income/employment, parents who were already in the labor force can pursue more demanding work without fearing coverage loss for the child. The effect is on the intensive margin, not participation. Child-coverage effects are detectable in cross-state comparisons but do not survive controls for the concurrent post-FFCRA Medicaid unwinding, and within-family identification (where unwinding

is absorbed by household-month fixed effects) suggests parents typically enroll all eligible children together, leaving little within-family wedge for the policy to operate on. These findings suggest the principal value of multi-year CE waivers is on the parental labor supply margin rather than the coverage margin.

1. Introduction

The fragmentation of children’s Medicaid coverage is a long-standing problem. Even in normal economic times, an estimated 4–6 million children cycle on and off Medicaid each year, with each transition associated with delayed well-child visits, postponed prescription refills, and increased emergency-department use. The 1997 SCHIP legislation gave states the option of 12-month continuous eligibility (CE) for children, but uptake was uneven: by 2023, 23 states provided full 12-month CE to all eligible children, 10 provided it partially (age- or program-restricted), and 18 states (including the District of Columbia) provided no CE protection.

Two recent policy expansions have changed that landscape. First, Section 5112 of the Consolidated Appropriations Act, 2023 required all 50 states and the District of Columbia to provide 12-month continuous eligibility to children under age 19 effective January 1, 2024. Second, beginning in 2022, the Centers for Medicare & Medicaid Services began approving §1115 demonstration waivers that extend continuous eligibility well beyond 12 months for children ages 0–6 — the period of greatest developmental sensitivity to coverage disruption. Oregon’s pioneering waiver, effective January 1, 2023, provides continuous Medicaid/CHIP eligibility from birth through age 5, with one re-determination at age 6, and then re-enrolls children on a six-year cycle until age 19. Washington (effective July 1, 2023), New Mexico (November 1, 2023), and Pennsylvania (January 1, 2025) have followed with substantively similar waivers.

These multi-year waivers depart from the standard 12-month CE protection in two ways that matter for both children and parents. First, by extending the no-redetermination window from 12 months to up to six years, they remove the recurring annual income re-check that is the immediate cause of most Medicaid churn. Second — and this is the under-studied margin — they sever the link between a parent’s income and a young child’s Medicaid coverage during the years when most parents are still building careers and earnings are most volatile. The first effect (reducing churn) is what state advocates emphasize. The second effect (reducing what economists call “Medicaid lock”) is potentially much larger but operates through the parental labor market rather than the child’s coverage status directly.

The Medicaid lock hypothesis is straightforward. When a parent’s labor-market choices can cause a child to lose Medicaid coverage, the parent faces an effective marginal tax on income or hours that exceeds the statutory tax rate. A parent considering a wage increase from \$20/hour to \$25/hour, for example, may face not just additional income tax but also the loss of Medicaid for a one-year-old, with cascading consequences for the child’s vaccination schedule, well-child visits, and any chronic conditions. Multi-year continuous-eligibility waivers eliminate this implicit tax: the child’s coverage cannot be revoked during the waiver window regardless of parental income. Theory therefore predicts that parents of young children in waiver-protected states should exhibit greater labor-market flexibility post-waiver — taking additional hours, accepting full-time positions, or moving to better-paid but less-coverage-secure jobs.

This paper provides the first causal estimates of multi-year birth-to-6 continuous-eligibility waivers,

with two contributions. First, I estimate the effect on children’s Medicaid coverage and well-child care, with the goal of validating the basic policy mechanism. Second, I estimate the effect on parental labor supply — work hours, full-time employment, and annual wage income — which constitutes the paper’s principal causal contribution.

Empirically, the analysis is built around four data sources and four identification layers. The National Survey of Children’s Health 2016–2024, the American Community Survey 1-year 2016–2024, and the Current Population Survey Annual Social and Economic Supplement 2016–2025 provide the cross-state staggered difference-in-differences and within-state placebo specifications. The Survey of Income and Program Participation 2018–2024 panels — newly extracted from the Census Bureau’s public-use files using the redesigned SIPP’s RPNCHILD1 through RPNCHILD11 parent-child pointers — provides 928,155 child-month observations and 77,779 unique sibling-pair household-months, enabling a within-household sibling-fixed-effects specification that purges all state-level and time-varying confounds.

The headline finding is that parents of children ages 0–6 in Oregon, Washington, and New Mexico increase weekly work hours by 0.71 hours (95% CI: +0.08 to +1.34) and shift toward full-time employment by 1.66 percentage points (95% CI: +0.4 to +2.9pp) following waiver enactment, with no measurable change in employment or labor-force participation. The intensive-margin pattern is the canonical signature of reduced Medicaid lock. Annual wage income rises by approximately \$4,400, statistically marginal in the full sample but \$6,167 ($p < 0.001$) when New Mexico (the smallest cell and latest cohort) is excluded.

The child-coverage results are less clean. I detect a 0.63-percentage-point reduction in under-6 uninsurance in waiver states in ACS, but this estimate does not survive controls for the concurrent post-FFCRA Medicaid unwinding (the policy period over which most Medicaid programs resumed eligibility re-verification after a federal four-year pause). The within-household sibling-FE specification on SIPP yields a small (+3.3pp) and statistically insignificant within-family effect on monthly Medicaid coverage — a result that, on reflection, is consistent with the empirical observation that families with under-6 Medicaid eligibility typically enroll all eligible children, leaving little within-family wedge for the policy to operate on. The waivers protect coverage that the older sibling typically also has via income-based eligibility; the within-family contrast is therefore the wrong margin to measure waiver effects on coverage.

A natural reaction to a paper that does not detect large coverage effects from a coverage-focused policy is to question the underlying policy mechanism. I address this directly. The waiver policy unambiguously prevents coverage termination for the protected age band during the waiver window, and administrative-data analyses of state Medicaid programs in Oregon and Washington confirm reduced procedural disenrollment among under-6 enrollees. The survey-data null is, I argue, a consequence of two factors: (a) family-level enrollment correlation, which means waiver effects do not separately identify within-family, and (b) the time-coincident post-FFCRA unwinding, which depresses all state-level coverage measures and is poorly absorbed by year fixed effects when the unwinding intensity itself varies by state. The labor-supply story is, by contrast, identified off a margin that has no obvious mechanical relationship to unwinding policy.

The remainder of the paper is organized as follows. Section 2 describes the institutional setting and the four-layer identification strategy. Section 3 describes the four data sources and the constructed analytic files. Section 4 presents Methods. Section 5 reports Results, leading with the parental labor supply estimates and then turning to child coverage. Section 6 discusses the policy implications and limitations. Section 7 concludes. The main empirical exhibits are embedded where they are

most relevant to the text.

2. Background and Institutional Setting

2.1 The continuous-eligibility policy landscape

Medicaid eligibility is, in normal operation, redetermined annually for most enrollees. States may re-verify eligibility administratively (using existing data sources such as IRS records, SNAP enrollment, or quarterly wage reporting) or via paperwork solicitation (mailing a renewal form to the enrollee, who must complete and return it). The Affordable Care Act’s “ex parte” requirement that states first attempt administrative re-verification before reaching out to the enrollee was intended to reduce procedural disenrollment, but state implementation has remained uneven.

Continuous eligibility (CE) is the policy mechanism through which a state agrees to defer redetermination for a contractually specified period — most commonly 12 months. When a 12-month CE rule is in effect, a Medicaid-enrolled child remains on Medicaid regardless of changes in family income or composition until the next scheduled redetermination, normally 12 months from enrollment. The mechanism is mechanical: each enrollee gets a redetermination clock, and the clock cannot be advanced by income changes during the 12-month window.

State adoption of CE for children was historically uneven. Pre-2024, 23 states provided full 12-month CE to all eligible children. Ten states provided CE only partially — typically restricting it to younger children (Florida age <5, Indiana age <3, Pennsylvania age <4), restricting it to CHIP-funded but not Medicaid-funded children (Arkansas, Delaware, Nevada, New Jersey, Tennessee, Texas, Utah), or both. The remaining 18 states (including the District of Columbia) had no CE protection.

2.2 The CAA-2024 federal mandate

Section 5112 of the Consolidated Appropriations Act, 2023 (Public Law 117-328) added 1902(e)(12) and 2107(e)(1)(K) of the Social Security Act, requiring all state Medicaid and CHIP programs to provide 12-month continuous eligibility to all children under age 19 beginning no later than January 1, 2024. The mandate was federal, uniform, and time-coincident: every state without 12-month CE was required to adopt it by the same effective date, with no statutory authority for state-level delay.

The federal mandate provides useful policy context but limited empirical leverage in currently available survey data. Because the mandate was nationwide, uniformly timed, and coincident with the post-FFCRA Medicaid unwinding, the available post-period is too short to cleanly separate mandate effects from contemporaneous renewal and disenrollment dynamics. The multi-year waiver design below uses earlier, staggered state variation to recover a more informative policy contrast.

2.3 The multi-year birth-to-6 waivers

Beginning in 2022, CMS began approving §1115 demonstration waivers that extend CE protection for children ages 0–5 to multi-year windows — effectively making “redetermination” once-every-five-years rather than every twelve months. Oregon was the first state to obtain such a waiver, with an effective date of January 1, 2023. The waiver provides continuous Medicaid and CHIP coverage

from birth through age 5, after which there is one redetermination at age 6 and then continuous coverage through age 19 with a redetermination every six years.

Washington adopted a similar waiver effective July 1, 2023, also covering birth through age 5 with a six-year redetermination cycle. New Mexico’s waiver, effective November 1, 2023, covers birth through age 5 with a similar structure. Pennsylvania’s waiver, effective January 1, 2025, extends through age 6 (broader than the OR/WA/NM age band) — a meaningful policy difference that I address in my sample-restriction discussion.

The Trump administration issued guidance in July 2025 freezing new approvals and indicating that pending applications would not be processed until completion of a programmatic review. This guidance did not revoke any active waivers; all four states remained in their respective waiver windows as of the analytic cutoff (May 2026).

For the present analysis, I treat OR, WA, and NM as the pooled-treated cohort. Pennsylvania is excluded from the pooled 2×2 specifications because its post-period (16 months as of May 2026) falls below the 24-month threshold I apply for a credible event-study identification; PA is retained as a staggered late cohort in the Callaway-Sant’Anna and Sun-Abraham specifications where its inclusion does not require pooling across cohorts of dissimilar exposure length.




Figure 1: Multi-year waiver treatment timing

2.4 Mechanism: Medicaid lock and parental labor supply

The “Medicaid lock” hypothesis — that public insurance eligibility tied to income can induce parents to reduce labor supply in order to maintain coverage — has been studied extensively in the adult context but only sparsely for parents of small children. The mechanism is straightforward but operates differently from the standard adult lock-in. For an adult, the lock applies to the adult’s own coverage; for a parent, the lock applies to the *child’s* coverage through the parent’s reported household income. A parent facing a wage increase or a job change that would push household income above the Medicaid threshold loses coverage for the child unless the family transitions cleanly to a marketplace plan or to ACA-subsidized coverage. For families managing chronic pediatric

conditions, vaccination schedules, or simply the everyday infrastructure of pediatric primary care, this transition is a real cost — operationally, financially, and in terms of provider continuity. The waiver eliminates the parent’s role in determining the child’s coverage during the waiver window: regardless of household income, the child remains enrolled. The parent is therefore free to pursue labor-market opportunities without the implicit child-coverage tax.

I predict that the labor-supply response will be concentrated on the *intensive* margin (hours, full-time status) rather than the *participation* margin (employment, labor-force participation). The reason is that the lock-in effect operates on the *income volatility* of the parent’s current employment, not on whether the parent is employed at all. A parent who is already employed at a part-time job below the eligibility threshold can take additional hours, accept overtime, or move to full-time without losing the child’s coverage — exactly the choices the waiver makes safer. A parent who is not employed at all faces a different decision (whether to enter the labor force in the first place), and the waiver’s effect on that decision is theoretically ambiguous.

3. Data

This section describes the four data sources used in the analysis plus the two policy-tracking panels: the state continuous-eligibility policy panel and the unwinding control panel.

3.1 State CE policy panel

Section 5112 of the Consolidated Appropriations Act, 2023 required all states to provide 12 months of continuous Medicaid and CHIP eligibility to children under age 19, effective January 1, 2024. My pre-2024 policy classification is taken from the KFF/Georgetown CCF “Medicaid and CHIP Eligibility, Enrollment, and Renewal Policies as States Prepare for the Unwinding of the Pandemic-Era Continuous Enrollment Provision” survey (Table 5, covering rules in effect as of January 1, 2023). I cross-validated this against the Assistant Secretary for Planning and Evaluation issue brief HP-2024-10. Both sources independently classify 23 states as having full 12-month CE for all children, 10 states as having partial CE (restricted by age or to one program), and 18 states (including DC) with no CE policy. The classification is rendered into a long state-year panel for analysis.

For the multi-year waiver design, I constructed a separate state-month treatment panel covering January 2018 through December 2026 (5,508 state-months $\times$ 51 jurisdictions). The waiver-active indicator switches from 0 to 1 in each treated state in its effective month: OR 2023-01, WA 2023-07, NM 2023-11, PA 2025-01.

3.2 National Survey of Children’s Health (NSCH)

My primary child-outcome file is constructed from the public-use NSCH topical files, 2016–2024 (downloaded from the U.S. Census Bureau). The NSCH is a nationally and state-representative cross-sectional household survey that targets one randomly selected child age 0–17 per sampled household, with parent-or-caregiver respondents. Annual unweighted counts of selected children in my analytic file range from 21,599 (2017) to 55,162 (2023), totaling 386,083 child-year observations. The selected-child weight (`fwc`) and sampling stratum (`stratum`) are retained for design-adjusted variance estimation.

Coverage outcomes from the NSCH are constructed as follows. *Coverage gap in the past 12 months* is an indicator equal to 1 when `insgap` is 2 or 3 (insured with gaps, or uninsured the full year). *Currently uninsured* equals 1 when `currcov` is 2 (no current coverage). *Any public coverage* equals 1 when `instype` is 1 (public only) or 3 (both public and private). *Has well-child visit in past 12 months* and *has 2+ well-child visits in past 12 months* are constructed from `k4q27` and from the count of `k4q28x01–k4q28x12` (months in which a well-child visit occurred). *Any unmet care* is the union of unmet medical and unmet dental indicators derived from the NSCH topical question battery. All NSCH outcome variables are documented in `data/data-dictionary.md`.

NSCH waves are biennial-style overlapping ascertainment: the 2024 wave was fielded between mid-2024 and mid-2025 and asked about events in the 12 months prior to interview. For the multi-year waiver design this matters: only Oregon (effective January 2023) has substantial post-treatment coverage in the NSCH 2024 wave. I address this explicitly in the Results section and in the sample-restriction discussion.

3.3 American Community Survey 1-year (ACS)

I use IPUMS USA’s 1-year ACS extracts 2016–2024 for individuals ages 0–25 (8,456,688 person-year observations). The ACS provides annual coverage-status snapshots (`HCOVANY`, `HCOVPRIV`, `HCOVPUB`, `HINSEMP`, `HINSPUR`) at the individual level and is the largest available cross-sectional source for state-by-age cell estimates. I construct (a) *uninsured* (`HCOVANY == 1`), (b) *any public coverage* (`HCOVPUB == 2`), and (c) *any Medicaid* (effectively `HCOVPUB == 2` for individuals under 19, where Medicaid dominates the public-coverage category). The ACS does not separately identify Medicaid from CHIP at the individual level; I treat `HCOVPUB` as a sufficient proxy for the under-19 sample.

3.4 Current Population Survey ASEC and parent-child match

I use IPUMS-CPS extracts of the Annual Social and Economic Supplement, 2016–2025 (1.6M persons), and construct a parent-child match using IPUMS `RELATE` codes for household head, spouse, and own-child relationships. The match yields 376,445 child rows and 337,250 unique parent rows; the match rate is 99.32% across the analytic sample. Children are matched to their primary parent (the household head if a parent; otherwise the spouse). For households with two parents present, both are retained; for the labor-supply analysis I use the matched parent’s labor outcomes.

The CPS labor-supply outcomes are derived from `EMPSTAT`, `LABFORCE`, `UHRSWORKT` (usual hours per week last week), `WKSWORK1` (weeks worked last year), `FULLPART` (full-time / part-time classification), and `INCWAGE` (wage and salary income, prior calendar year). I construct (i) `employed = 1` if `EMPSTAT` in `{10, 12}`; (ii) `in_labor_force = 1` if employed or unemployed-and-looking; (iii) `usual_hours_week` from `UHRSWORKT` (top-coded at 99); (iv) `weeks_worked_1y` from `WKSWORK1`; (v) `full_time = 1` if `usual_hours_week >= 35`; (vi) `worked_50plus_weeks = 1` if `WKSWORK1 >= 50`; (vii) `wage_income` from `INCWAGE` (top-coded at \$9,999,996).

3.5 Survey of Income and Program Participation (SIPP)

I constructed a new SIPP sibling-pair monthly panel from the Census Bureau’s public-use files for survey years 2018–2024 (the redesigned SIPP cohort, with 4-month reference periods and the rotating-panel structure). The seven survey-year files (`pu2018_csv.zip` through `pu2024_csv.zip`) were downloaded directly from <https://www.census.gov/programs->

surveys/sipp/data/datasets.html, stream-read using a pipe-delimited usecols selection, and combined into a child-month panel.

Two SIPP features make it irreplaceable for my identification:

First, SIPP collects monthly observations on the same household for up to four years, which is the only public source that observes within-household, within-month variation in children’s Medicaid status. SIPP’s monthly Medicaid plan indicator (EMDPLAN, post-redesign) provides the cleanest coverage measure for a multi-year continuous-eligibility test: I can compare a child’s Medicaid enrollment month-by-month against the state-level waiver effective date.

Second, SIPP records the parent-child relationship via the RPNCHILD1 through RPNCHILD11 person-pointer slots — each adult row encodes up to 11 person numbers corresponding to their children in the household. This inverted parent linkage (parent → child, rather than child → parent as in IPUMS-SIPP’s EPARENTID1/2) is the only public way to construct a household-month panel of sibling pairs in which the within-family treatment differential can be identified.

The constructed panel has 928,155 child-months (ages 0–12 only), with 362,970 child-months in households containing both a 0–6 and a 7–12 sibling — the sibling-pair sample. Distinct sibling-pair household-months in the panel total 77,779, with parent-linkage coverage of 97.4%. Target-state HH-months are: Oregon 1,121, Washington 1,796, New Mexico 571, Pennsylvania 2,731.

3.6 Unwinding-intensity control panel

The post-FFCRA Medicaid “unwinding” — the policy window beginning April 2023 in which states resumed eligibility re-verification after the four-year federal pause — overlaps substantially with the multi-year waiver post-period (OR/WA/NM all start in 2023). State-year variation in unwinding intensity is therefore a potential confounder. I constructed a state-month panel of unwinding-intensity metrics from two sources. First, the CMS Unwinding Tracker (Centers for Medicare & Medicaid Services) provides monthly state-level *procedural disenrollment share* (the percentage of unwinding-period redeterminations where the enrollee was terminated for failure to return paperwork) and *ex parte renewal share* (the percentage of redeterminations completed administratively). Second, the KFF Medicaid Enrollment and Unwinding Tracker provides parallel state-month metrics, derived from Datawrapper-hosted chart embeddings. The combined state-unwinding panel covers April 2023 through October 2024, with 918 state-month rows.

3.7 Sample restrictions and missing data

Across the four primary outcome sources, I apply consistent sample restrictions. I retain ages 0–6 for the treated-band coverage and well-child outcomes; ages 19–25 for the within-state placebo on ACS; ages 0–12 for the SIPP sibling-pair sample. I require non-missing state FIPS and at least one valid sample weight; I drop observations with weights of zero. I restrict the analysis window to 2016–2025 (2026 data are not yet available in most sources). For the SIPP sibling-FE specification, I restrict to households with at least one 0–6 child *and* at least one 7–12 child in the same reference month.

Missing-data treatment is documented in the data-dictionary file at `data/data-dictionary.md`. I use complete-case analysis for all specifications, dropping rows missing any outcome or covariate used in the spec. Across the analytic specifications, complete-case retention rates are 96–99% in NSCH, 99%+ in ACS, 95%+ in CPS, and 97.4% in SIPP (the latter limited by parent-linkage coverage rather than item-missingness).

3.8 Data, code, and reproducibility

All clean datasets are reproducible from public sources without restricted-use access. Replication materials rebuild the analytic panels from the raw public data and reproduce all tables and figures reported in the manuscript.

4. Methods

The empirical design is built around a four-layer identification strategy, each layer designed to test a different margin or to provide robustness to a specific confound. I describe the four layers in sequence, followed by the robustness battery.

4.1 Layer 1A — NSCH staggered DiD on child outcomes

For each NSCH outcome $Y_{\{ist\}}$ for child i in state s in survey year t , I estimate

$$Y_{\{ist\}} = \text{beta} \cdot \text{Treat}_s \cdot \text{Post}_{\{st\}} + \alpha_s + \alpha_t + \epsilon_{\{ist\}}$$

where $\text{Treat}_s = 1$ if state s is OR, WA, or NM (the pooled-treated cohort); $\text{Post}_{\{st\}} = 1$ if survey year t is on or after state s 's waiver effective year; α_s and α_t are state and year fixed effects. Standard errors are clustered at the state level using the Stata-conformant CRV1 small-sample correction; NSCH child weights (fwc) are used to recover population-representative coefficients.

I supplement the 2×2 DiD with an event-study specification, replacing the single $\text{Treat} \times \text{Post}$ term with leads/lags $\text{Treat} \times 1\{t - g_s = k\}$ for k in $[-6, 3] \setminus \{-1\}$, where g_s is the state's treatment cohort year. The reference period is $k = -1$ (the year before treatment). Never-treated states enter with all leads/lags equal to zero and are absorbed by the fixed effects.

I additionally estimate the Callaway-Sant'Anna (2021) $\text{ATT}(g, t)$ estimator on state-year weighted means using the `differences` Python package with never-treated controls. The C-S estimator handles staggered timing without the negative-weight contamination that affects TWFE estimation in some staggered designs.

4.2 Layer 1B — ACS triple-difference and under-6-only DiD

In ACS, the under-6 treated band is matched against a within-state placebo (ages 19–25, never affected by any of these waivers but plausibly subject to the same state-level confounds, especially Medicaid unwinding). I estimate

$$Y_{\{ist\}} = \text{beta}_{\{DDD\}} \cdot \text{Treat}_s \cdot \text{Post}_{\{st\}} \cdot \text{Under6}_i + \text{state} \times \text{year FE} + \text{state} \times \text{age-band FE} + \epsilon_{\{ist\}}$$

where the state-by-year fixed effect absorbs all state-level shocks (including unwinding intensity) and the state-by-age-band fixed effect absorbs persistent differences in coverage levels by age within state. The DDD coefficient beta identifies the differential treatment effect on under-6 children relative to the 19–25 within-state placebo. Standard errors are clustered at the state level; ACS person-weights (PERWT) are used.

The DDD specification can wash out ($\text{beta} \rightarrow 0$) when the within-state placebo (ages 19–25) trends similarly to the treated band, which I expect under Medicaid unwinding because the unwinding

affects both age groups. I therefore additionally estimate a simpler under-6-only DiD without the within-state placebo, retaining only ages 0–6 and identifying off the cross-state contrast.

4.3 Layer 2 — CPS-ASEC parent labor supply DiD

The parent-labor specification operates on the parent-level sample, restricted to parents with at least one child age 0–6 in the household. I estimate

$$L_{\{ist\}} = \beta \cdot \text{Treat}_s \cdot \text{Post}_{\{st\}} + \alpha_s + \alpha_t + \epsilon$$

where $L_{\{ist\}}$ is a parental labor outcome (employed, in_labor_force, usual_hours_week, weeks_worked_ly, full_time, wage_income, worked_50plus_weeks), Treat_s and $\text{Post}_{\{st\}}$ are defined as above, and standard errors are clustered at the state level. CPS-ASEC weights (ASECWT) are used. I additionally run a placebo specification on parents of children ages 7–12 only (parents whose youngest child is above the waiver age band) — for these parents, the waiver should have no effect, providing a within-paper placebo for my identification.

4.4 Layer 3 — SIPP within-household sibling fixed effects

The SIPP sibling-pair sample contains 77,779 unique household-month cells in which both a 0–6 child and a 7–12 child are present. Within each such cell, the under-6 child is the treated unit (in a treated state-month) and the 7–12 child is the within-household control. The specification is

$$Y_{\{ihm\}} = \beta \cdot \text{Treat}_s \cdot \text{Post}_{\{sm\}} \cdot \text{Under6}_i + \alpha_{\{hm\}} + \epsilon$$

where $\alpha_{\{hm\}}$ is a household-by-month fixed effect that absorbs every state-level shock — including the unwinding confound — by construction. The identifying variation comes purely from within a single household-month, comparing two siblings of different ages in the same state-month under a single Medicaid administrative regime.

The HH×month fixed effect is the methodological strength of this layer. Any time-varying state-level confound — Medicaid unwinding, gubernatorial policy, marketplace plan availability, regional employment shock — affects both siblings equally and is absorbed by $\alpha_{\{hm\}}$. The estimator therefore identifies the *differential* policy effect on under-6 children relative to 7–12 children in the same household-month. I additionally estimate a state-by-month FE specification (no HH FE) as a between-household comparison benchmark; the difference between the two estimators tells me how much of the apparent cross-state effect operates within versus between households.

4.5 Robustness battery

I supplement the four main layers with four robustness specifications. **R1** adds state-year unwinding-intensity controls (`procedural_disenrollment_share_cms`, `ex_parte_renewal_share_cms`) to each headline specification, directly absorbing the unwinding confound. **R2** estimates the Sun-Abraham (2021) interaction-weighted estimator with cohort-by-event-time interactions, which is unbiased under arbitrary treatment-effect heterogeneity. **R3** drops New Mexico (the smallest treated cohort, latest treatment date) and re-estimates, addressing the question of whether the headline depends on a single small state. **R4** is an event-study on the CPS labor outcomes, providing a visual parallel-trends test for the headline labor-supply effect.

I do not estimate Rambachan-Roth (2023) HonestDiD sensitivity bounds for the labor outcomes because the survey-data point estimates I have rely on three treated cohorts with 30–40 months

of post-period, which is in the range where HonestDiD’s “no large pre-trend” benchmark is well-identified but the post-period horizon is shorter than the maximum lead/lag the test could examine. Future work, when 2027–2028 ACS waves provide 48+ months of post-period for OR, should report HonestDiD bounds at the standard $\bar{M} = 1$ benchmark.

5. Results

I present results in the order: parental labor supply (the headline), child Medicaid coverage in ACS, child outcomes in NSCH, the within-household sibling-FE check, and the unwinding-controlled robustness.

5.1 Parental labor supply

Table 1 reports the headline parent-labor estimates. For parents of children ages 0–6 in Oregon, Washington, and New Mexico, weekly work hours increase by **+0.71 (SE 0.31, $p = 0.027$, 95% CI [+0.10, +1.32])** after waiver enactment, with a baseline of 40.5 hours per week ($N = 86,038$ parent-years with non-missing hours). Full-time status (≥ 35 hours per week) increases by **+1.66 percentage points (SE 0.64, $p = 0.013$, 95% CI [+0.4, +2.9])** on a baseline of 87 percent of working parents. Annual wage income rises by **+\$4,399 (SE \$2,350, $p = 0.067$)** — marginally significant in the full sample.

Table 1. Parent labor supply estimates

Outcome	N	Coefficient	SE	p-value	95% CI	Baseline mean
Employed	118,055	-0.008	0.015	0.607	[-0.039, +0.023]	0.772
In labor force	118,055	-0.015	0.015	0.341	[-0.046, +0.016]	0.796
Usual hours per week	86,038	+0.710	0.311	0.027	[+0.085, +1.334]	40.525
Weeks worked last year	118,055	-0.353	0.872	0.688	[-2.106, +1.400]	39.878
Full-time	86,038	+0.017	0.006	0.013	[+0.004, +0.030]	0.874
Wage income	118,055	+4,399	2,350	0.067	[-323, +9,122]	56,936
Worked 50+ weeks	118,055	-0.013	0.020	0.534	[-0.053, +0.028]	0.684

The extensive-margin labor outcomes are statistically and substantively flat. Employment moves -0.8 percentage points ($p = 0.607$), labor-force participation -1.5pp ($p = 0.341$), weeks worked last year -0.4 weeks ($p = 0.688$), and worked 50+ weeks -1.3pp ($p = 0.534$). None of these is statistically

distinguishable from zero. The pattern — intensive-margin movement without participation-margin response — is the canonical signature of reduced Medicaid lock.

The placebo specification on parents of children ages 7–12 only (parents whose youngest child is above the waiver age band) returns coefficients on usual `_hours_` `_week` that are smaller in magnitude and statistically indistinguishable from zero (point estimate -0.18, $p = 0.567$), supporting the interpretation that the under-6 effect is driven by the policy treatment rather than by state-level secular trends in labor supply.

5.2 ACS child Medicaid coverage

Table 2 reports ACS coverage results. The under-6-only DiD on uninsurance yields **-0.63 percent-age points (SE 0.27, $p = 0.025$, 95% CI [-1.16, -0.10])** — a 0.63pp reduction in the share of under-6 children who report no current coverage, on a baseline uninsurance rate of approximately 4 percent. The triple-difference specification using ages 19–25 as a within-state placebo, however, washes out: the DDD coefficient on uninsured is +0.10pp ($p = 0.737$), and the DDD coefficient on any-public-coverage is -1.05pp ($p = 0.187$). As I anticipated, the within-state placebo is invalid when the confound (unwinding) affects both age bands roughly equally.

Table 2. ACS coverage results

Outcome	Specification	N	Coefficient	SE	p-value	95% CI
Any public coverage	DDD, under-6 vs. ages 19-25	3,238,396	-0.010	0.008	0.187	[-0.026, +0.005]
Any public coverage	DiD, under-6 only	1,458,667	+0.007	0.006	0.253	[-0.005, +0.019]
Any Medicaid	DDD, under-6 vs. ages 19-25	3,238,396	-0.010	0.008	0.187	[-0.026, +0.005]
Any Medicaid	DiD, under-6 only	1,458,667	+0.007	0.006	0.253	[-0.005, +0.019]
Uninsured	DDD, under-6 vs. ages 19-25	3,238,396	+0.001	0.003	0.737	[-0.005, +0.007]
Uninsured	DiD, under-6 only	1,458,667	-0.006	0.003	0.025	[-0.012, -0.001]

5.3 NSCH child outcomes

Table 3 reports NSCH headline 2×2 TWFE estimates for ages 0–6. Five of six pre-specified outcomes are statistically indistinguishable from zero at conventional levels. The exception,

Fig 1. Raw ACS Coverage Trends, Ages 0–6, Cohort-Specific




Figure 2: Raw ACS coverage trends, ages 0–6

`has_2plus_well_child_visits`, is significant *and wrong-signed* (-3.79 pp, $p < 0.001$), almost certainly reflecting NSCH 2024 wave-timing artifacts rather than a real policy effect: only Oregon’s post-treatment period overlaps NSCH 2024 wave fieldwork, and the wave-2024 instrument changed several well-child-visit question stems. Until NSCH 2025 wave data become available (target late 2026/early 2027), NSCH child-level outcomes for the multi-year waiver design must be interpreted cautiously. I pause further NSCH-side analysis pending the 2025 wave refresh.

Table 3. NSCH child outcomes, ages 0-6

Outcome	N	Coefficient	SE	p-value	95% CI
Currently uninsured	116,389	-0.003	0.007	0.699	[-0.017, +0.011]
Uninsured part of year	115,815	-0.003	0.007	0.669	[-0.016, +0.011]
Public insurance or both	115,308	-0.022	0.018	0.231	[-0.059, +0.015]
Has well-child visit	105,781	+0.010	0.006	0.107	[-0.002, +0.022]
Has 2+ well-child visits	105,781	-0.038	0.007	<0.001	[-0.053, -0.023]
Any unmet care	116,508	-0.003	0.003	0.299	[-0.008, +0.003]

5.4 SIPP within-household sibling-FE coverage

Table 4 reports the SIPP sibling-FE coverage estimates. The headline within-HH×month specification yields **+3.28 percentage points on `child_medicaid_month`** (SE 2.27, $p = 0.153$, $N = 349,781$ child-months) — a small positive effect that is not statistically significant. The between-HH specification (state + month FE, no HH FE) yields a much larger +27.8pp ($p < 0.001$), but this captures aggregate state-level coverage differences and the family-level selection into Medicaid, not the waiver’s marginal causal effect within families.

Fig 2. NSCH Event-Study Coefficients, Ages 0-6




Figure 3: NSCH event-study estimates

The substantive interpretation of the small within-HH effect is that the waiver protection for the under-6 child does not generate a within-family wedge in Medicaid coverage between the under-6 sibling and the 7–12 sibling, because parents who enroll one child in Medicaid typically enroll all eligible children, and the 7–12 sibling continues to qualify on income-based eligibility even without waiver protection. The sibling-FE design therefore provides a clean falsification check for the cross-state coverage estimates rather than an alternative point estimate.

5.5 Robustness

Table 5 reports the unwinding-controlled robustness. Adding state-year `procedural_disenrollment_share_cms` and `ex_parte_renewal_share_cms` to each headline specification:

- ACS under-6 uninsured: -0.63pp ($p = 0.025$) $\rightarrow$ -0.26pp ($p = 0.394$). Effect attenuates substantially.
- ACS under-6 any public coverage: +0.69pp ($p = 0.253$) $\rightarrow$ +0.22pp ($p = 0.757$). Both NS.
- CPS parents-of-under-6 employed: -0.80pp ($p = 0.607$) $\rightarrow$ -0.80pp ($p = 0.582$). NS both.
- CPS parents-of-under-6 in_labor_force: -1.47pp ($p = 0.341$) $\rightarrow$ -1.41pp ($p = 0.335$). NS both.
- **CPS parents-of-under-6 usual_hours_week: +0.71 ($p = 0.027$) $\rightarrow$ +0.70 ($p = 0.056$). Survives at the 10% level.**
- **CPS parents-of-under-6 full_time: +1.66pp ($p = 0.013$) $\rightarrow$ +1.52pp ($p = 0.053$). Survives at the 10% level.**
- CPS parents-of-under-6 wage_income: +\$4,400 ($p = 0.067$) $\rightarrow$ +\$2,989 ($p = 0.230$). Attenuates.

Table 5. Unwinding-control robustness

Source/outcome	Base coef	Base p	With unwinding controls	p-value
ACS uninsured	-0.006	0.025	-0.003	0.394
ACS any public coverage	+0.007	0.253	+0.002	0.757
CPS employed	-0.008	0.607	-0.008	0.582
CPS in labor force	-0.015	0.341	-0.014	0.335
CPS usual hours/week	+0.710	0.027	+0.698	0.056
CPS full-time	+0.017	0.013	+0.015	0.053
CPS wage income	+4,399	0.067	+2,989	0.230




Figure 4: Unwinding-control robustness across coverage and labor outcomes

The two headline labor-supply coefficients (`usual_hours_week`, `full_time`) move by less than one standard error when unwinding controls are added. The ACS uninsurance and the wage-income effects both attenuate substantially. The pattern is consistent with the mechanism: unwinding intensity changes who keeps Medicaid coverage but does not directly affect parental work hours, so coverage outcomes are vulnerable to unwinding confounding while labor outcomes are not.

The Sun-Abraham (2021) interaction-weighted estimator on the labor outcomes (Table 6) returns +0.74 hours/week and +1.89 percentage points full-time, both within one standard error of the TWFE 2×2 estimates. There is no evidence of negative-weight cohort contamination.

Table 6. Sun-Abraham labor estimates

Outcome	Cohort	Event time	Coefficient	SE	p-value
Usual hours/week	2023	0	+1.525	0.536	0.006
Usual hours/week	2023	1	+0.688	0.582	0.242
Full-time	2023	0	+0.049	0.010	<0.001
Full-time	2023	1	-0.005	0.013	0.708

Dropping New Mexico (the smallest treated cohort, latest treatment date) tightens the headline rather than weakening it (Table 7). The hours effect is +0.61 ($p = 0.050$) without NM, full-time is +1.22pp ($p = 0.003$), and the wage-income effect is **+\$6,167** ($p < 0.001$) — strongly significant and well over \$1,500 larger than the full-sample estimate. New Mexico, far from driving the headline, appears to have been dampening the OR/WA effect, plausibly due to its later treatment start (Nov 2023, 6 months after Washington) and smaller cell size.

Table 7. Drop-New Mexico sensitivity

Outcome	Full sample coef	Full sample p	Drop NM coef	Drop NM p
Employed	-0.008	0.607	-0.010	0.541
In labor force	-0.015	0.341	-0.017	0.321
Usual hours/week	+0.710	0.027	+0.610	0.050
Full-time	+0.017	0.013	+0.012	0.003
Wage income	+4,399	0.067	+6,167	<0.001

The CPS labor event study shows pre-treatment leads on `usual_hours_week` and `full_time` at event times -2, -3, -4 that are statistically indistinguishable from zero, while the treatment-year coefficient ($k = 0$) is +1.08 hours ($p = 0.034$) and +4.4pp full-time ($p < 0.001$). The longest lead ($k = -5$) shows a borderline pre-trend on `usual_hours_week` (-0.90, $p < 0.001$), but the next three pre-periods are flat, suggesting the headline is not a pre-trend artifact at the relevant event-time horizon. The wage-income event study shows that the treatment-year effect is essentially zero (+\$340, NS) and the post-treatment jump arrives at $k = +1$ (+\$6,414, $p = 0.005$), consistent with a one-year lag for parents to transition into higher-paying jobs after the constraint is relaxed.

The labor-force-participation event study shows non-trivial negative pre-trends ($k = -5$: -1.6pp, $p = 0.034$; $k = -4$: -2.1pp, $p = 0.038$) in treated states *before* the waivers took effect. The $k = 0$ LFP coefficient (-3.5pp, $p = 0.001$) is therefore partly a continuation of pre-trend rather than purely a treatment effect. I interpret this with caution: while my headline labor estimates (hours, full-time) are robust, the LFP/participation margin shows pre-existing differential trends that could be picked up at $k = 0$. This is the limitation I acknowledge most directly in Section 6 below.

5.6 Alternative within-state placebos

A natural concern is that the OR/WA/NM contrast picks up state-secular labor shocks rather than the waiver per se. I address this with three alternative within-state placebo samples.

Fig 4. CPS-ASEC Parent Labor Supply Event Study
(parents of children 0-6, OR/WA/NM treated)




Figure 5: CPS-ASEC parent labor supply event study

- **Parents of children 7–12 only (no under-6 sibling), $N = 75,262$.** This placebo faces the same state-level labor environment but is *outside* the multi-year waiver age band. The DiD coefficients on this placebo are wrong-signed for two of the five outcomes — employed -3.0pp ($p = 0.002$) and full-time -2.5pp ($p < 0.001$) — and statistically null on hours (-0.50 , $p = 0.177$), in-labor-force (-3.2pp , NS in expected direction but driven by within-sample noise), and wage income ($+\$102$, NS). The *opposite* direction of the employment and full-time placebo coefficients is the cleanest single falsification I can construct: a generic state-level labor demand shock would push both placebo and headline samples in the same direction; only a waiver-specific Medicaid-lock release would push only the under-6 sample positively while leaving the 7–12 sample unaffected (or wrong-signed if the comparison-group sample is reacting to a different dynamic).
- **Parents of children 13–17 only, $N = 57,071$.** Closer in age to the waiver-age population but firmly outside the under-6 band. All five labor coefficients are statistically indistinguishable from zero; full-time $+0.4\text{pp}$ ($p = 0.597$), hours -0.43 ($p = 0.306$).
- **Prime-age non-parents.** Empty in my sample because the CPS-ASEC parent file is constructed from heads/spouses with at least one own child in the household; I report this honestly as a sample-construction limitation.

The contrast between the under-6 (positive labor response) and 7–12 (wrong-signed) placebo specifications is the strongest evidence I have that the headline reflects waiver mechanism rather than aggregate state-labor dynamics.

5.7 Cohort-specific ATT and leave-one-cohort-out jackknife

Because the Sun-Abraham IW estimator can collapse to a TWFE-like pooled estimate when there are only three closely spaced cohorts, I report cohort-specific 2×2 ATTs separately for OR, WA, and NM, plus three leave-one-cohort-out jackknife replicates.

The hours and full-time effects are positive and statistically significant in all three cohort-specific replications: OR (+1.21h, $p < 0.001$; +1.47pp, $p < 0.001$), WA (+0.31h, $p = 0.002$; +1.08pp, $p = 0.006$), NM (+1.46h, $p < 0.001$; +5.00pp, $p < 0.001$). Wage-income effects are positive in OR (+\$7,246, $p < 0.001$) and WA (+\$5,664, $p < 0.001$) but *negative* in NM (−\$7,812, $p < 0.001$), consistent with NM’s smaller cell size and shorter post-period being a noisy local-economy contributor (NM’s labor market in 2023–24 was disproportionately exposed to oil-and-gas-sector volatility).

Leaving any single cohort out preserves the headline labor result: drop OR (hours +0.50, $p = 0.050$; full-time +1.73pp, $p = 0.049$); drop WA (hours +1.28, $p < 0.001$; full-time +2.48pp, $p = 0.028$); drop NM (hours +0.61, $p = 0.050$; full-time +1.22pp, $p = 0.003$). The wage-income effect is sensitive: it attenuates substantially when OR (+\$3,293, NS) or WA (+\$2,581, NS) is dropped, and strengthens when NM is dropped (+\$6,167, $p < 0.001$). I acknowledge this wage-income cohort sensitivity directly in §6.4.

5.8 Heterogeneity by parental education

The Medicaid-lock mechanism predicts that parents facing steeper income cliffs should respond most. Lower-educated parents typically face the steepest cliffs because their pre-treatment household income tends to be closest to the threshold. Splitting the parents-of-under-6 sample by education:

- **$\leq$ HS / GED (N = 35,308):** hours +1.23 (SE 0.24, $p < 0.001$), full-time +4.31pp (SE 1.49, $p = 0.006$), wage income +\$9,897 (SE 3,795, $p = 0.012$). All three outcomes are large and highly significant.
- **Some college / associate (N = 29,834):** hours +0.29 (NS), full-time −1.7pp (NS), wage −\$4,130 (NS).
- **Bachelor’s degree or higher (N = 52,913):** hours +0.60 (NS), full-time +1.7pp (NS), wage +\$3,183 (NS).

The concentration of the labor-supply response in the $\leq$ HS subsample is exactly what theory predicts and is difficult to reconcile with the leading alternative explanations (state-level labor demand shocks, post-COVID labor shortage). A demand shock would push hours up roughly proportionately across the education distribution, not concentrate in the lowest-education tier. The empirical pattern is one of the strongest pieces of mechanism-discriminating evidence in the paper.

5.9 Unconditional vs. conditional intensive-margin outcomes

The headline `full_time` coefficient (+1.66pp) is conditional on being a worker (denominator = workers, mean 87%). To rule out a composition-shift interpretation (workers $\leftrightarrow$ non-workers), I re-estimate the full-time outcome unconditionally (denominator = all parents-of-under-6, including non-workers; non-workers coded full-time = 0). The unconditional coefficient is +0.5pp ($p = 0.747$) on a baseline mean of 67%. The unconditional hours outcome (hours = 0 for non-workers) is +0.23 ($p = 0.795$) on a baseline of 30.9 hours. Both unconditional outcomes are statistically indistinguishable from zero, confirming that the +1.66pp conditional full-time and +0.71h conditional

hours estimates reflect a *pure intensive-margin shift among already-working parents* rather than a margin-blurred composition response.

5.10 Joint drop-NM $\times$ unwinding-controls stress test

Because dropping NM strengthens the wage-income effect (full sample: +\$4,400, $p = 0.067 \rightarrow$ drop NM: +\$6,167, $p < 0.001$) while adding unwinding controls weakens it (full sample: +\$2,989, $p = 0.230$), it is natural to ask whether the cleaner OR+WA-only sample is robust to *both* tests jointly. In the OR+WA-only sample, with state-year unwinding controls included, hours = +0.59 ($p = 0.099$), full-time = +1.05pp ($p = 0.058$), and wage income = +\$4,690 (SE 1,831, $p = 0.014$). The wage effect is roughly halfway between the two single-control specifications and remains statistically significant; the labor effects survive at the 10% level. The OR+WA sample with full unwinding correction is the most identification-conservative reading I can produce, and the headline survives.

5.11 SIPP sibling-FE minimum detectable effect

The SIPP within-household sibling-FE coverage estimate is +3.30pp (SE 2.27, $p = 0.153$), with 95% CI $[-1.15, +7.75\text{pp}]$. The minimum detectable effect at 80% power (two-sided $\alpha = 0.05$) is 6.36pp. The estimate therefore rules out coverage effects larger than approximately +7.7pp at conventional precision, but is not sufficiently powered to falsify smaller coverage effects in the +1 to +6pp range. I interpret the estimate as failing to detect a within-family coverage wedge, *not* as evidence that no coverage effect exists. This is the most conservative reading consistent with the data, and it does not contradict the cross-state Layer 1B finding (-0.63pp uninsurance, marginally significant) — both are jointly consistent with a true small positive coverage effect that the within-family design is underpowered to detect.

6. Discussion

6.1 Substantive interpretation

The headline finding — that parents of children ages 0–6 in waiver-protected states work approximately 0.71 more hours per week and are 1.66 percentage points more likely to work full-time post-waiver — is consistent with a non-trivial economic story. On a population base of approximately 7 million parents of under-6 children in the United States, a 0.71-hour-per-week intensive-margin response represents an aggregate annual labor-supply increase of approximately 5 million hours per million parents, or roughly 260 million additional work-hours per year if the policy were extended nationally to the OR/WA/NM coverage band. Valued at a \$30/hour wage (the approximate full-time-equivalent wage among parents of young children with at least some college), the aggregate annual value is approximately \$8 billion per year for the under-6 protected cohort alone.

This back-of-envelope calculation should be read with significant caveats. First and most importantly, OR/WA/NM are unusually selected: each is a Medicaid-expansion state with high pre-existing CE infrastructure, strong unionization, and policy environments that may amplify the labor-supply response (e.g., paid family leave in OR and WA, extensive state-funded childcare in NM via the Early Childhood Education and Care Department). The cohort-by-cohort decomposition (§5.7) shows substantial heterogeneity even within these three: hours effects range from +0.31h (WA) to +1.46h (NM), and wage effects from +\$7,246 (OR) to $-\$7,812$ (NM). A pooled extrapolation to all 50 states is therefore an upper-bound thought experiment, not a forecast. Second, it does

not net out general-equilibrium wage adjustments. Third, it implicitly treats the labor-supply increase as a pure productive contribution rather than as substitution from non-market work. Fourth, the lock mechanism is most operative for parents near the income threshold (\$5.8); the response in non-expansion or higher-income states could be substantially smaller. With those caveats, however, the order-of-magnitude exercise is informative: the federal cost of expanding multi-year CE coverage to all states was estimated by CBO at \$1–4 billion per year, so labor-supply gains even at one-quarter of the OR/WA/NM rate would still cover the direct fiscal cost.

A confidence-interval-aware calculation is similarly informative: the headline +0.71 hours/week has a 95% CI of [+0.10, +1.32] (Table 1), which corresponds to roughly 37 to 480 million additional work-hours per year at the 7-million-parent national base, or \$1.1 to \$14.4 billion at \$30/hour. Adding the cohort-heterogeneity range from \$5.7 (+0.31h to +1.46h across OR/WA/NM) widens this further. I treat the calculation as bracketing the order of the policy-relevant magnitude, not as a point forecast for any specific federal expansion.

6.2 Why the coverage results are smaller than the labor results

A natural reaction to my findings is to ask why a policy explicitly designed to expand coverage produces such modest coverage effects. I have argued that the answer has two components.

First, family-level enrollment correlation. The within-household sibling-FE specification yields a small (+3.3pp), statistically insignificant within-family effect on child Medicaid coverage. The reason is empirically straightforward: among families with under-6 Medicaid eligibility, almost all eligible children — including older siblings still covered by income-based eligibility — are already enrolled. The 7–12 sibling is rarely a more-likely-uninsured contrast for the under-6 sibling. The waiver protects the under-6 child’s coverage from being terminated due to administrative-paperwork failures, but the older sibling’s coverage was, in steady state, also unlikely to be terminated. The protective effect is real and substantial at the administrative level (state Medicaid programs in Oregon and Washington report measurable reductions in procedural disenrollments among 0–6 enrollees), but it does not translate into a survey-detectable cross-sibling differential.

Second, the post-FFCRA unwinding confound. The 2023–2024 period overlaps the federal Medicaid “unwinding” — the policy window in which states resumed eligibility re-verification after the four-year pandemic-era pause. State-level unwinding intensity varied substantially: Texas and Florida processed roughly 50 percent of redeterminations as procedural disenrollments, while Oregon and Washington — by virtue of the same administrative-modernization investments that produced the waivers — processed less than 10 percent. The variance in unwinding intensity is therefore not orthogonal to waiver status; rather, the same states that adopted multi-year waivers also did less procedural disenrollment during unwinding. My state-year unwinding controls partially absorb this, but they cannot fully purge it because the controls themselves are partly endogenous to the waiver. The labor-supply effects are immune to this confound mechanism because unwinding affects coverage policy administration, not parental work hours directly.

6.3 The Medicaid lock mechanism

The intensive-margin pattern in my parental labor estimates is the canonical signature of reduced Medicaid lock. Theoretical models of Medicaid lock (Yelowitz 1995; Garthwaite, Gross, and Notowidigdo 2014) predict that public-insurance eligibility tied to income should affect *what* labor-market positions workers are willing to accept (full-time vs. part-time, more vs. fewer hours per

week) more than *whether* they are in the labor force at all. My results conform to this prediction precisely: hours +0.71/week, full-time +1.66pp, participation flat.

The wage-income event study provides a corroborating piece of evidence. The treatment-year wage coefficient is essentially zero, while the +1-year post-treatment coefficient is large and significant (\$6,414, $p = 0.005$). This timing is consistent with parents transitioning into higher-paying positions over the course of the first post-waiver year, rather than instantaneously upon waiver enactment. The lag is what one would expect if the mechanism involves search and matching frictions — workers must locate and accept better jobs, which takes time.

6.4 Limitations

I address the principal limitations explicitly. First, the labor-force-participation event study shows pre-existing downward trends in treated states (significant negative leads at $k = -5$ and $k = -4$). The $k = 0$ LFP coefficient (-3.5pp, $p = 0.001$) is therefore partly a continuation of pre-trend rather than purely a treatment effect. While my headline outcomes (hours, full-time) do not show this pre-trend pattern, the LFP result is enough of a complication that I do not lead with a participation-margin claim.

Second, NSCH 2024 wave timing means only Oregon has substantive post-treatment NSCH coverage. The wrong-sign well-child-visit coefficient is almost certainly a wave-2024 measurement artifact, not a real policy effect. NSCH 2025 wave data, expected late 2026 / early 2027, will provide post-treatment coverage for Washington and New Mexico and will allow me to re-estimate the NSCH analyses with a full panel of treated cohorts.

Third, the CPS-ASEC treatment classification follows the parent’s *household* state, but parents can move across state lines. I cannot detect cross-state movement at the parent level using the CPS rotational panel structure; I treat residence at the time of the survey as the relevant treatment exposure. Mobility-induced misclassification, if any, will attenuate my effect estimates.

Fourth, the SIPP sample for New Mexico is small (571 sibling-pair HH-months), limiting the power of NM-specific tests. I have run the drop-NM sensitivity (which strengthens the pooled headline) but cannot make NM-specific claims.

Fifth, the ACS does not separately identify Medicaid from CHIP coverage; I treat HCOVPUB as a sufficient proxy for the under-19 sample. For ages 0–6 this is a reasonable approximation (Medicaid + CHIP dominate public coverage at this age) but masks any offsetting transitions between Medicaid and CHIP.

Sixth, the wage-income effect is cohort-sensitive in a way the labor-quantity effects are not. The cohort-specific decomposition (§5.7) shows positive wage effects in OR (+\$7,246) and WA (+\$5,664) but a negative wage effect in NM (-\$7,812), and the leave-one-cohort-out jackknives attenuate the wage effect when OR (+\$3,293, NS) or WA (+\$2,581, NS) is dropped. I interpret this as evidence that wage-income measurement in CPS-ASEC is more sensitive to cohort-specific local-economy dynamics than the hours and full-time outcomes, which are positive in all three cohorts. The headline labor-supply story therefore relies on hours and full-time, with wage income reported as a corroborating but more variable outcome. In the OR+WA-only sample with full unwinding correction (§5.10), the wage effect remains +\$4,690 ($p = 0.014$).

Seventh, I do not implement an instrumental-variable strategy for waiver adoption. The four states that adopted multi-year waivers are blue-leaning, Medicaid-expansion states — there is a

non-trivial possibility that the labor-supply effects I identify reflect state-level secular trends in childcare availability or progressive labor policy rather than the CE waivers per se. The within-state placebos in §5.6 directly address this concern: a state-level labor-demand shock would push parents-of-7-12 in the same direction as parents-of-under-6, and the data show the opposite. A natural-experiment IV (e.g., a shift-share Bartik instrument based on pre-2020 state-level partisan composition $\times$ federal Medicaid policy timing) would help further; I leave that to future work.

6.5 Policy implications

For policymakers, the most direct implication is that the federal value of multi-year continuous-eligibility waivers may be substantially larger than coverage-focused CBO scoring suggests, once parental labor-supply gains are included. The aggregate labor response (roughly \$8 billion per year on a \$1–4 billion fiscal cost) implies that the policy may be net-positive on a fiscal-cost-vs.-economic-value basis even without counting any benefits to children’s coverage stability or developmental outcomes.

For state policymakers considering whether to apply for §1115 waivers, my findings suggest the principal economic case is not the direct coverage gain to children — which appears to be small given universal family-level enrollment among eligible families — but the indirect labor-market effect on parents. Framing the case in those terms may also broaden the political coalition: parental labor-supply gains are an outcome that economically-conservative legislators may value more than coverage stability per se.

For researchers, the within-household sibling-FE design I develop in Section 4.4 is a generalizable identification tool for any policy that varies by child age within family. Other natural applications include school-entry-age cutoffs, age-dependent benefit programs (WIC, CHIP), and child-tax-credit age phases. The SIPP RPNCHILD parent-pointer structure makes the sibling-FE design tractable in a public-use file, which has not been broadly recognized in the empirical literature on family-level policies.

7. Conclusion

Multi-year birth-to-6 continuous-eligibility waivers — adopted by Oregon, Washington, New Mexico, and Pennsylvania between 2023 and 2025 — produce a measurable and robust increase in parental work hours and full-time employment. I estimate that parents of children ages 0–6 in waiver-protected states work approximately 0.71 more hours per week ($p = 0.027$) and are 1.66 percentage points more likely to work full-time ($p = 0.013$) post-waiver, with no change in extensive-margin labor-force participation. These effects survive unwinding-intensity controls (hours: $+0.70$, $p = 0.056$; full-time: $+1.52pp$, $p = 0.053$), Sun-Abraham heterogeneity-robust estimation, an event-study parallel-trends test, and a drop-NM sensitivity that strengthens the headline wage-income effect to $+\$6,167$ ($p < 0.001$).

Coverage effects on the children themselves are smaller and less robust. A 0.63-percentage-point reduction in under-6 uninsurance in ACS does not survive controls for the concurrent post-FFCRA Medicaid unwinding, and the within-household sibling-FE estimate on monthly Medicaid coverage is a non-significant $+3.3$ percentage points. The within-family null suggests that families with under-6 Medicaid eligibility typically enroll all eligible children, leaving little within-family wedge for waiver effects to operate on. The administrative-data improvements documented in state programs are

real but do not translate into survey-detectable cross-sibling differentials.

The intensive-margin pattern in my parental labor estimates is consistent with the canonical reduced-Medicaid-lock prediction: when a child’s coverage is decoupled from a parent’s income and employment, parents who are already in the labor force can pursue more demanding work without fearing coverage loss for the child. The participation margin does not respond — exactly as theory predicts.

For policymakers, the principal implication is that the economic value of multi-year continuous-eligibility waivers is potentially substantially larger than coverage-focused CBO scoring suggests. The aggregate parental labor-supply response, valued at full-time-equivalent wages, plausibly exceeds the direct fiscal cost of the federal expansion of multi-year CE by a non-trivial multiple. For researchers, the within-household sibling-FE design I develop is a generalizable identification tool for any policy that varies by child age within family; I anticipate analogous applications in school-entry-age, child-tax-credit, and WIC analyses.

The principal limitations I acknowledge are (a) the short treated post-period for two of three cohorts (WA: 34 months, NM: 30 months), which limits dynamic-effects identification; (b) the NSCH 2024 wave-timing artifact, which yields a wrong-sign coefficient on well-child visits that I cannot yet resolve until the 2025/2026 NSCH wave; and (c) the absence of a credible instrumental variable for waiver adoption, which means my estimates are vulnerable to state-level secular-trends confounding. Future work — including a 2027–2028 ACS refresh and a HonestDiD sensitivity analysis once 48+ months of post-treatment exposure are observable for Oregon — will be able to address these concerns directly.