

State Nursing-Home Staffing Floors and the Distributional Bite: Heterogeneity by Medicaid Share and Rurality

Abstract

Background. The 2024 federal CMS minimum-staffing rule for long-term care facilities was vacated and then repealed in late 2025, returning the staffing-mandate debate to the states. Existing evaluations document that state floors raised average direct-care staffing modestly, but they do not characterize who is actually bound by the mandate or whether the response varies systematically with the financing margins that make staffing-floor compliance most difficult.

Methods. I assemble a facility-quarter panel from the CMS Payroll- Based Journal covering 2017Q1–2025Q4 (35 quarters; 510,189 facility- quarter observations on 16,401 unique nursing facilities). Across 13 states adopting numerical staffing floors 2018Q1–2024Q3 I construct each facility’s pre-policy “bite” — the gap between the new statutory total nurse hours-per-resident-day floor and the facility’s mean four- quarter pre-period staffing — and merge LTCFocus pre-treatment Medicaid- day shares and Care Compare urban/rural designations. I estimate the average effect with Callaway-Sant’Anna (2021) doubly-robust ATT(g,t), the dose-response by bite quintile, a Medicaid-share $\times$ bite triple- interaction, and rurality stratifications. A pre-period placebo, the Goodman-Bacon decomposition, and Rambachan-Roth (2023) HonestDiD bounds discipline the inference.

Results. The binary average effect on Total HPRD (-0.143 , SE 0.009) is contaminated by a treated-state secular trend: a $t_0 - 8q$ placebo yields -0.13 . The identifying within-state-quarter heterogeneity specifications, which the placebo cannot replicate, all run in the policy-expected direction. The dose-response gradient across bite quintiles is monotone ($Q_1 = -0.343$, $Q_5 = +0.026$; spread 0.37 HPRD), and this gradient steepens with Medicaid share ($Q_1 \rightarrow Q_5$ spread $0.32 / 0.40 / 0.43$ HPRD across Low / Mid / High Medicaid-share terciles). The triple interaction **bite $\times$ Medicaid share $\times$ post** is **+0.153 (SE 0.034, $t = 4.6$, $p < 0.001$)**. Rural facilities show a steeper bite gradient than urban (spread 0.51 vs 0.34), and the rural Bite- Q_5 cell ($+0.076$) is the largest positive treatment-effect cell anywhere in the heterogeneity supplement. The headline binary effect survives adjustment for seven LTCFocus resident-mix controls ($-0.145 \rightarrow -0.139$), ruling out case-mix composition as the driver.

Conclusions. State staffing floors operate as binding constraints exactly where policy-makers expect — among the most-binding facilities that are most fiscally constrained by Medicaid revenue dependence and in rural markets with thinner nursing labor pools. The average effect, read alone, obscures this heterogeneity. The findings argue for targeting Medicaid-heavy facilities with technical-assistance and rate-cell adjustments under any future federal staffing rule, and for caution in applying state-level average-effect estimates to a national mandate without accounting for the distributional bite.

1. Introduction

The 2024 final rule from the Centers for Medicare and Medicaid Services (CMS) imposed three numerical floors on US nursing-home staffing — 3.48 total nurse hours per resident day (HPRD),

0.55 RN HPRD, and 2.45 nurse- aide HPRD — together with a 24/7 RN coverage requirement (Centers for Medicare & Medicaid Services 2024). The rule was the most ambitious federal staffing intervention in the long-term care sector since the 1987 Nursing Home Reform Act. It did not survive: courts vacated the rule in 2024, and CMS issued a formal repeal in late 2025 (Centers for Medicare & Medicaid Services 2025), returning the staffing-mandate debate to the states. Twenty-five states currently operate some form of numerical staffing floor; thirteen of them adopted or revised those floors after 2017 with sufficient post-period to support a research design.

What the existing literature has resolved, and what it has left open, defines the contribution of this paper. Average-effect evaluations of state mandates — most recently Werner et al. (2026) on twenty-two states 2010–23 — establish that state floors raise direct-care staffing modestly (approximately 0.18 HPRD, or about 5%) without detectable closure effects or net-margin erosion, with labor-cost increases largely offset by higher net patient revenues. Earlier work (Park and Stearns 2009; Bowblis 2011; Matsudaira 2014; Bowblis and Hyer 2013; Brunt 2023) reaches qualitatively similar conclusions: floors raise staffing at the binding margin without obvious quality gains, but with non-trivial input- substitution costs (e.g., reductions in housekeeping, food service, and activities staff). The descriptive literature on the bite of the federal 2024 rule (Chidambaram et al. 2024; Skinner et al. 2025; Bhaumik and Grabowski 2025) documents that compliance gaps are strongly correlated with Medicaid census and rurality but does not estimate the causal mandate effect *as a function of* those facility- level margins.

Three substantive gaps remain. First, the structural nursing-home literature shows that Medicaid reimbursement is the binding financing margin for staffing investment (Hackmann 2019) and that the labor- supply elasticity of nurses is low (Friedrich and Hackmann 2021). This implies that staffing-floor responses should depend systematically on each facility’s Medicaid revenue mix — yet no causal evaluation has estimated that interaction. Second, recent payroll-based-journal (PBJ) work demonstrates that staffing levels have not normalized post- pandemic (McGarry et al. 2020; Bhaumik et al. 2026) and that nearly half of facilities now substitute toward agency labor to satisfy operating staffing levels (Bowblis et al. 2024; Office of the Assistant Secretary for Planning and Evaluation 2025), shifting the margin on which floors bind. Third, the methodological literature on continuous-treatment difference- in-differences (Callaway et al. 2024) now provides the estimator needed to recover dose-response heterogeneity across the full bite distribution — a tool not yet applied to nursing- home staffing-floor mandates.

This paper closes those gaps by combining (i) a 510,189-row facility- quarter PBJ panel covering 2017Q1–2025Q4, (ii) a hand-coded state- statute tracker with 100% primary-source URL coverage on treatment timing for 13 adopting states, (iii) HCRIS-SNF and LTCFocus pre- treatment base-lines for Medicaid-day share and resident-mix controls on 4,156 of 4,216 eligible treated facilities (98.6%), and (iv) a modern staggered-DiD framework that uses the Callaway-Sant’Anna (2021) doubly-robust estimator for the binary average effect, a pre-policy bite-quintile dose-response that approximates the Callaway-Goodman-Bacon-Sant’Anna (2024) continuous-treatment ATT, and a TWFE triple-interaction benchmark for the policy-mechanism test.

The headline empirical finding is that **the bite $\times$ Medicaid-share $\times$ post triple interaction equals +0.153 (SE 0.034, $t = 4.6$, $p < 0.001$)**: a one-HPRD increase in pre-policy bite raises the post-period staffing response by 0.15 HPRD per unit of Medicaid share. This gradient is identified by within-state-quarter variation across facilities — exactly the variation a state-level placebo cannot replicate. The 3×5 stratified dose-response shows a monotone Q1→Q5 gradient within every Medicaid tercile, with the cross-quintile spread widening from 0.32 HPRD in low-Medicaid facilities to 0.43 HPRD in high-Medicaid facilities. The only positive treatment-effect

cell in the entire 3×5 grid is the policy-targeted high-Medicaid $\times$ most-binding cell (T3 $\times$ Q5 = +0.038 HPRD). Rural facilities show a steeper bite gradient than urban (Q1 $\rightarrow$ Q5 spread 0.51 vs 0.34), and the rural Bite-Q5 cell (+0.076) is the strongest positive treatment-effect cell in the heterogeneity supplement.

I confront the binary average-effect estimate honestly. The pooled CSA ATT on Total HPRD (−0.143) has the wrong sign relative to a binding mandate, and a placebo specification that shifts the treatment date by eight quarters returns −0.13 — economically indistinguishable from the actual estimate. I interpret this as evidence that the binary average is contaminated by a treated-state-specific secular trend (post-COVID staffing erosion that was geographically uneven) and not by the staffing-floor mandate itself. The Bacon decomposition rules out already-treated-as-control contamination as the explanation: the TWFE-implied beta is −0.098, with 81.3% of the weight on clean treated- vs-never-treated comparisons. The placebo concern thus reflects a treated-state level confound, not an estimator artifact. The heterogeneity specifications survive this concern because they are identified from cross-facility variation within state-quarter cells, which the placebo (a state-level shift) cannot reproduce. The resident-mix robustness check (Table 6) confirms that the binary effect is not driven by case-mix composition (−0.145 $\rightarrow$ −0.139 with seven LTCFocus mix controls).

The paper’s contribution is therefore threefold. Substantively, it provides the first study to characterize the within-state-quarter heterogeneity in state nursing-home staffing-mandate response by Medicaid share and rurality, advancing a literature anchored on average-effect estimates. Methodologically, it applies the continuous-bite dose-response framework of Callaway et al. (2024) to a high-stakes policy setting where the binary specification turns out to be contaminated by secular trends. For policy, it shows that the mandate works where it binds and where it matters: at high-Medicaid, most-binding facilities where reimbursement margins are tightest and at rural facilities where labor-market thinness is most likely to be a constraint. These patterns directly inform the design of any future federal staffing rule and of state staffing standards in the post-2025 policy environment.

The remainder of the paper proceeds as follows. Section 2 describes the institutional setting and the 13 state-mandate adoptions in my sample. Section 3 details the data and bite construction. Section 4 sets out the empirical strategy. Section 5 reports results. Section 6 discusses the findings. Section 7 concludes.

2. Background

2.1 The 2024 federal rule and its repeal

The CMS final rule on minimum staffing for long-term care facilities (Centers for Medicare & Medicaid Services 2024) was the centerpiece of the Biden administration’s nursing-home reform agenda. Its three numerical floors — 3.48 total nurse HPRD, 0.55 RN HPRD, 2.45 nurse-aide HPRD — were calibrated to bring the bottom quartile of facilities up to a level the agency considered consistent with safe staffing. The 24/7 RN coverage requirement was the most novel feature, addressing the long-documented finding (Geng et al. 2019) that facilities exhibit large within-week staffing fluctuations and that 75% are “almost never” in compliance with CMS’s expected RN staffing levels.

The rule was vacated by the US District Court for the Northern District of Texas in 2024 and

was formally repealed by CMS in late 2025 (Centers for Medicare & Medicaid Services 2025). The repeal returned the policy debate to states, where roughly half of the country has adopted some form of numerical floor. Twenty-five jurisdictions currently maintain a numerical floor; thirteen of them adopted or substantively revised those floors after 2017, and those thirteen comprise my treatment-cohort sample (California, Massachusetts, New Jersey, Illinois, Rhode Island, New York, Connecticut, Ohio, Maine, the District of Columbia, Oregon, Florida, and Pennsylvania).

2.2 Why bite, Medicaid share, and rurality

The literature suggests three reasons why average-effect estimates will obscure the most policy-relevant variation. First, Chidambaram et al. (2024) estimate that only 19% of facilities meet all three federal numerical thresholds, and the gap is strongly correlated with Medicaid census. The descriptive bite distribution is not uniform across states, across facilities within states, or across the within-facility components of staffing.

Second, structural work (Hackmann 2019) establishes that a 10% Medicaid reimbursement increase raises skilled-nurse staffing by 8.7% in Pennsylvania administrative data. The reverse implication is that high-Medicaid facilities face a binding fiscal constraint on staffing investment: without commercial cross-subsidy, raising staffing to meet a mandate is more costly per dollar of revenue than at a low-Medicaid facility. The bite $\times$ Medicaid-share interaction is therefore not just a descriptive subgroup analysis — it is a direct test of the financing-margin mechanism.

Third, rurality matters for two distinct reasons that may push in opposite directions. Rural facilities tend to be more Medicaid- dependent on average, so the financing-margin channel predicts *larger* responses there. But rural labor markets are thinner — Friedrich and Hackmann (2021) establish a 12% reduction in nurse employment caused a 13% mortality rise with no offsetting reallocation, implying low cross-market labor-supply elasticity. The net effect on rural mandate compliance is an empirical question I can answer.

2.3 The 13 treated-state cohorts

Treatment timing in my sample is staggered across 28 quarters (Table 1, lower panel). California (2018Q1) and Massachusetts (2018Q3) lead; New Jersey (2020Q3) is the only cohort whose effective date falls inside the 2020Q2–2021Q4 pandemic period. The bulk of cohorts adopt in the post-pandemic 2022Q1–2024Q3 window — Connecticut (2022Q1), Florida (2022Q1), Rhode Island (2022Q1), New York (2022Q2), Ohio (2022Q3), Maine (2022Q2), DC (2022Q2), Oregon (2022Q3), Pennsylvania (2023Q3 phase-1 / 2024Q3 phase-2). Statutory floor levels range from 2.66 HPRD (Ohio) to 4.44 HPRD (Maine). Pre-policy share of facilities below the floor varies dramatically: 5% in New Jersey (where the floor was already met by most facilities) to 90% in Maine and Rhode Island (where the new floor was meaningfully binding for most of the sector).

The policy details — statute citations, effective dates, floor levels, enforcement mechanisms, and penalty types — are hand-coded from state code and rule documents and verified to primary sources, with 100% URL coverage on the tracker. Detailed institutional context for the New York 3.5 HPRD law (New York State Department of Health 2023) and Maine’s 4.44 HPRD standard, together with the broader state-policy compendium (Medicaid and CHIP Payment and Access Commission 2024), informs the floor-level coding decisions.

2.4 What prior evaluations have established

A coherent body of evidence on state staffing-floor mandates predates my analysis. I summarize the four most directly relevant strands so that the reader can locate this paper’s contribution.

Average-effect estimates of state mandates. Park and Stearns (2009) use OSCAR data and a fixed-effects design on the 1998–2003 wave of state floors, finding modest staffing increases concentrated at facilities below the new floor, with small effects on survey deficiency outcomes. 1 sharpens this with a kinked-regression design on California’s 3.2 HPRD law, showing that nurse-aide hours rose by up to 30% for the most out-of-compliance homes, but with no detectable effect on patient outcomes or facility-level quality. Bowblis (2011) documents that mandates raise staffing levels while skill-mix responses depend on Medicaid reliance, and Bowblis and Hyer (2013) demonstrates that the increase comes partly at the expense of housekeeping, food-service, and activities staff — an early input- substitution result that anticipates the agency-staff and ancillary- cuts margins more salient today. Brunt (2023) adds an enforcement-side margin: CMS’s star-rating downgrade penalty for chronic understaffing raised PBJ-measured staffing without obvious quality gains, echoing the Matsudaira null on quality. The most relevant recent contribution is Werner et al. (2026), which extends the average-effect literature to 6,849 facilities across twenty-two states adopting mandates 2010–23, finding a 0.18 HPRD (approximately 5%) increase in direct-care staff with no detectable closure effects, no decline in net margins, and labor-cost increases largely offset by higher net patient revenues.

A common pattern emerges: average-effect estimates are small but precise, and average-quality effects are null or imprecise. The mechanism — convergence of below-floor homes toward the threshold with little spillover above — is well-established. What this literature does *not* characterize is how those convergence gains are distributed across facility types, particularly along the Medicaid-share and rurality dimensions where compliance burden differs most. That distributional question is the gap my paper fills.

Staffing–quality and industry structure. The staffing–outcomes literature begins with Konetzka et al. (2008), which uses MDS resident data merged with OSCAR to show that endogeneity-adjusted RN hours per resident day reduce pressure-sore and UTI incidence. Hackmann (2019) provides the structural counterpart, estimating from Pennsylvania administrative data that a 10% Medicaid reimbursement increase raises skilled-nurse staffing by 8.7% — i.e., the Medicaid rate, not pro- competitive market structure, is the binding constraint on quality investment. Friedrich and Hackmann (2021) establishes the labor-supply side: a 12% persistent reduction in Danish nurse employment caused a 13% rise in nursing-home mortality among adults 85+, with no offsetting reallocation across the labor market. Together these papers establish that (a) staffing matters causally, (b) Medicaid revenue is the binding financing margin, and (c) the labor-supply elasticity of nurses is low. Werner et al. (2013) shows that Medicaid pay- for-performance produces uneven quality gains concentrated in higher- end facilities, and Gupta et al. (2024) documents that private-equity ownership reduces staffing and raises mortality by approximately 11% — both papers indicate that ownership and financing structure condition the staffing response to policy. This is *prima facie* evidence that mandate bite should interact with Medicaid share and ownership type.

PBJ measurement, agency labor, turnover. Geng et al. (2019) is the foundational PBJ paper: facilities exhibit large within-week staffing fluctuations and 75% are “almost never” in compliance with CMS’s expected RN staffing levels. Crucially, PBJ-derived levels are systematically lower than the older self-reported OSCAR/CASPER numbers — a measurement upgrade that makes credible facility-quarter- panel designs feasible. Bowblis et al. (2024) documents that by 2022, nearly half of

nursing homes used agency staff, accounting for 11% of direct-care hours, with agency staff costing 50–60% more per hour than directly employed nurses. Office of the Assistant Secretary for Planning and Evaluation (2025) confirms that contract-staff reliance has not normalized post-pandemic. This matters substantively: a binding mandate may be satisfied through agency substitution rather than core-hire expansion, with implications for cost, turnover, and quality. McGarry et al. (2020) and Bhaumik et al. (2026) trace the pandemic shock and persistent post-pandemic staffing deficits, particularly at PE/REIT-owned facilities. Gandhi et al. (2021) and Shen et al. (2023) establish that staff turnover is a stronger predictor of nursing-home quality than levels, and that a 10-point turnover increase raises deficiency citations measurably. Jun and Grabowski (2024) shows the CNA workforce has shifted toward immigrant staff while US-born CNA employment fell — a labor-supply margin that policy variants of the bite design should consider.

The federal rule and state-level bite. Centers for Medicare & Medicaid Services (2024) imposed 3.48 HPRD total, 0.55 RN HPRD, 2.45 NA HPRD, and 24/7 RN coverage. The rule was vacated and then repealed (Centers for Medicare & Medicaid Services 2025). The descriptive bite literature (Chidambaram et al. 2024; Skinner et al. 2025; Bhaumik and Grabowski 2025) shows that Medicaid-heavy and rural facilities face the largest compliance gap and that bite is not uniform across states. None of these papers, however, estimates the *causal* effect of mandate adoption as a function of facility bite, and none interacts bite with Medicaid share within a modern staggered-DiD framework. Grabowski and Bowblis (2023) lays out the central tradeoff — opportunities versus challenges — and Bohler and Adashi (2025) summarizes the post-2024 evidence base.

Methodological literature. The continuous-intensity DiD design relies on two recent strands. The first is the staggered-adoption literature: Callaway and Sant’Anna (2021) introduces $ATT(g,t)$ for clean group-time comparisons, Sun and Abraham (2021) derives heterogeneity-robust event-study estimators, Borusyak et al. (2024) provides the efficient imputation estimator, and Chaisemartin and D’Haultfoeulle (2020) documents the negative-weights problem in TWFE under heterogeneous effects. Goodman-Bacon (2021) provides the decomposition diagnostic that exposes when already-treated units contaminate TWFE estimates. Rambachan and Roth (2023) supplies the pre-trend sensitivity bounds that should accompany any modern event-study. The second strand is continuous-treatment DiD: Callaway et al. (2024) derives the dose-response identification framework that lets me estimate $ATT(d,g,t)$ as a function of treatment intensity — exactly the estimand I need for facility-level bite quintiles.

2.5 Distinction from federal-rule anticipation effects

Federal-rule anticipation effects in 2023-2024 are a different estimand from the state-floor design studied here. Anticipation designs are vulnerable to calendar-time by low-baseline interactions, especially when the federal policy is national and contemporaneous with the post-pandemic staffing recovery. The present paper is methodologically and substantively distinct. I exploit state-level mandate adoption with within-state continuous bite, so state and quarter fixed effects absorb calendar-time by baseline confounding. The heterogeneity specifications additionally absorb state-by-quarter shocks, identifying off cross-facility variation within state-quarter cells. The federal-rule and state-mandate research questions are related, but the identification is fundamentally different.

3. Data

3.1 Sources

The analysis combines six data sources, all public:

1. **CMS Payroll-Based Journal (PBJ)**, 2017Q1–2025Q4. Daily nurse-hour reports per facility, aggregated to facility-quarter; provides the primary outcome (Total HPRD) and the three nurse-category outcomes (RN, CNA, LPN HPRD).
2. **CMS Care Compare provider snapshot**, April 2026. Static facility covariates: state, county, urban/rural designation, ownership type, certified-bed count, hospital-based flag, chain identifiers, special focus status, and 5-star ratings.
3. **State staffing-floor statute tracker**, hand-coded 2026-05-11. 57 events across 51 jurisdictions; 25 events are numerical-floor adoptions; 100% primary-source URL coverage.
4. **HCRIS-SNF cost reports**, FY2017–FY2024. 120,283 cost-report rows across 15,093 CCNs. Used as a fallback Medicaid-share source for facilities missing from LTCFocus.
5. **LTCFocus**, 2016–2023. 120,532 facility-year rows across 15,797 CCNs. Provides the primary pre-treatment Medicaid-share measure (`pct_medicaid_days`, 99.7% populated; median 64.3%) and the resident-mix covariates.
6. **BLS QCEW NAICS 623110**, 2017Q1–2025Q3. County-quarter nursing-care employment and wage panel (96,979 rows, 2,873 counties); used to characterize local-labor-market context.

3.2 Sample restrictions

I apply four sample restrictions sequentially (Table 4 of the data dictionary):

- **R1**: Drop facility-quarters with fewer than 60 PBJ-reported active days (placeholder for the PBJ audit-correction flag, which is not yet publicly available for the full sample).
- **R2**: Require ≥ 4 PBJ pre-quarters strictly before the state’s effective date.
- **R3**: Require positive `avg_mds_census` in all 4 pre-quarters (excludes closures and category changes; ensures denominator validity).
- **R4**: Restrict the treated set to states with at least one numerical floor event in the 2017Q1–2025Q4 window with PBJ pre-period support.

After these restrictions, the analysis sample is 439,274 facility- quarter observations across 13,353 unique CCNs: 4,216 treated CCNs in 13 states (R2 yield 96.6%, R3 yield 99.6% within treated states) and 9,137 never-treated CCNs in 39 control states (Table 1).

3.3 Bite construction

For each treated facility i in state s with effective date $t0$:

$$\overline{\text{HPRD}}_{i,\text{pre}} = \frac{1}{4} \sum_{\tau=t_0-4}^{t_0-1} \text{HPRD}_{i,\tau}$$

$$\text{bite}_i = \text{floor}_{s,t_0} - \overline{\text{HPRD}}_{i,\text{pre}}$$

Positive bite means the facility was below the new floor at adoption (mandate is binding); negative bite means the facility was already above the floor (least binding). Bite is winsorized at the 1st/99th percentile within the eligible-treated set (`bite_w`), and grouped into quintiles (Q1 = most-negative

bite, i.e., most-above-floor; Q5 = most-positive bite, i.e., deepest-below-floor). The winsorized bite distribution has mean -0.174 HPRD, median -0.119 HPRD, and SD 0.847 HPRD across 4,216 facilities, with quintile cutoffs at -0.746 / -0.287 / $+0.023$ / $+0.490$ HPRD (shown in Table 1 and the raw-trends figure).

3.4 Medicaid-share construction and the HCRIS limitation

A critical data note: HCRIS Worksheet S-3 line 8 (“Medicaid days”) is populated for only 342 of 120,283 cost-report rows (0.3%) in the FY2017–FY2024 public cost-report files. The HCRIS-SNF `pct_medicaid_days` field is therefore essentially empty as a primary measure. LTCFocus is the operative source: `pct_medicaid_days` is populated for 99.7% of facility-years, with a sensible median of 64.3% and a distribution consistent with prior nursing-home literature. For each facility I compute the **pre-treatment baseline Medicaid share** as the mean of LTCFocus `pct_medicaid_days` across years strictly preceding the state’s effective date. This ensures the heterogeneity bin assignment is not endogenous to the mandate itself. HCRIS provides a fallback Medicaid share for the 144 CCNs missing from LTCFocus; 60 CCNs (1.4% of the 4,216 treated set) lack any Medicaid-share data and are dropped from the heterogeneity analysis. Pre-treatment Medicaid-share tercile cuts are 53.0% / 69.1%; quintile cuts are 43.0% / 56.9% / 66.2% / 74.9%.

The implication for the paper is that the LTCFocus measure carries the identification weight on the Medicaid-share dimension. Because LTCFocus is the standard reference dataset for nursing-home Medicaid-mix work (Hackmann 2019; Bowlblis et al. 2024), I view this as the right defensive choice; I report the HCRIS fallback as a robustness sub-sample but it does not change the headline.

3.5 Rurality and resident-mix controls

Rurality is the Care Compare urban/rural designation merged at the facility level (rural = 648 CCNs; urban = 3,467 CCNs in the eligible- treated set). Resident-mix controls are LTCFocus pre-treatment baselines for percent White / Black / Hispanic, average resident age, acuity index (case-mix), occupancy percent, and percent Alzheimer’s/dementia. All seven controls are demeaned across the treated population and interacted with the post indicator in the robustness specification reported in Table 6.

3.6 Outstanding data limitations

Three limitations are documented honestly:

1. **Time-varying quality outcomes** (Care Compare deficiency citations, QM star changes) are available only as a static April-2026 snapshot in the available public files. A time-varying analysis of quality is deferred to a follow-up. The current paper studies only staffing.
2. **HCRIS Medicaid days field** is 99.7% empty in the FY2017–FY2024 parsed bundles (Section 3.4). LTCFocus is the operative source.
3. **PBJ measurement** has a 2020Q3 missing-by-design (CMS pause), and pandemic quarters 2020Q2–2021Q4 likely contain audit-irregularity noise. My identification strategy and pre-trend tests use the full panel; the placebo concern in Section 5 is partially attributable to pandemic-era treated-state drift.

4. Empirical Strategy

4.1 Estimators

I use four estimators, each addressing a different identifying margin.

Primary: Callaway-Sant’Anna ATT(g,t). Following Callaway and Sant’Anna (2021), I estimate group-time average treatment effects via the doubly-robust estimator implemented in the `differences` Python package, with the never-treated comparison group. I aggregate to the “simple” pooled ATT for the binary specification and to event-time ATTs for dynamic analysis. Inference uses the analytic standard errors returned by `differences`; a clustered bootstrap was not computationally feasible at this panel size and is documented as a limitation (Section 6.5).

Continuous-treatment dose-response. I approximate the Callaway et al. (2024) continuous-treatment ATT(d|g,t) by estimating a separate CSA model within each pre-policy bite quintile against the common never-treated comparison set, then comparing pooled ATT across quintiles to recover the dose-response slope. A direct continuous-CSA was not run because the `differences` v0.2 implementation does not yet expose a continuous-treatment API (blocker carried forward).

TWFE benchmark and triple interaction. As a benchmark I estimate

$$Y_{it} = \alpha_i + \gamma_t + \beta_1(\text{bite}_i \times \text{post}_{it}) + \beta_2\text{post}_{it} + \varepsilon_{it}$$

with facility and year-quarter fixed effects, standard errors clustered at the state level. I caution that, as Goodman-Bacon (2021) shows, this estimator can be contaminated by already-treated controls in staggered designs; I present the Bacon decomposition (Section 5.7) to quantify the contamination.

For the heterogeneity specification I add a third interaction:

$$Y_{it} = \alpha_i + \gamma_{s(i),t} + \beta_1(\text{bite}_i \times \text{post}_{it}) + \beta_2(\text{Medicaid}_i \times \text{post}_{it}) + \beta_3(\text{bite}_i \times \text{Medicaid}_i \times \text{post}_{it}) + \varepsilon_{it}$$

with **facility (alpha_i)** and **state-by-quarter (gamma_s,t)** fixed effects. The state-by-quarter FE absorbs the binary `treat × post` (state-level treatment timing is constant within state-quarter cells), so this specification identifies *only* the within-state-quarter heterogeneity slopes — which is exactly the variation that the binary placebo cannot replicate. Medicaid share is demeaned relative to the treated- population mean (0.577) so `beta_1` has the interpretation “bite × post slope at the average-Medicaid-share treated facility.” A companion specification with facility + calendar-quarter FE recovers the binary main effect alongside the heterogeneity slopes (Table 4B, row `triple_no_sq_fe`).

5 × 5 ATT grid. ATT(bite quintile, Medicaid quintile) is estimated cell-by-cell (25 separate CSA fits, each subsetting the treated panel to one bite × Medicaid cell and pooling all never-treated facilities as the comparison). This produces the 5 × 5 heatmap.

4.2 Outcomes

The primary outcome is Total Nurse Hours per Resident Day (Total HPRD). Secondary outcomes are RN HPRD, CNA HPRD, and LPN HPRD. All four outcomes are reported in the main results

table and event-study figures; the heterogeneity analysis focuses on Total HPRD (the policy-relevant aggregate).

4.3 Robustness

The robustness package addresses six identification threats:

1. **Pre-trends:** joint Wald test across event leads $e=-8$ to $e=-2$.
2. **Placebo:** shift each state’s effective date by -8 quarters and re-estimate the CSA pooled ATT. A placebo coefficient close to the actual estimate signals a treated-state secular trend.
3. **Alternative comparison groups:** never-treated only vs. not-yet- treated.
4. **Alternative bite definitions:** raw (non-winsorized) vs. above- median binary.
5. **Goodman-Bacon decomposition** of the TWFE estimator, isolating clean treated-vs-never-treated comparisons from earlier-vs-later and (problematic) later-vs-already-treated comparisons.
6. **HonestDiD relative-magnitudes sensitivity** (Rambachan and Roth 2023) at Mbar in $\{0.5, 1.0, 1.5\}$. Mbar scales the maximum allowable post- period bias as a multiple of the maximum pre-period deviation; the robust 95% CI widens accordingly.

4.4 Identification — what each specification identifies

It is worth being explicit about what variation each estimator exploits, because the specifications recover different parameters and survive different threats.

The **CSA binary pooled ATT** uses cross-state variation in mandate timing and identifies the average treatment effect on the treated under the parallel-trends assumption between treated states and the never- treated comparison pool. The estimand is identified at the state-event level: each treated state’s pre-versus-post comparison is differenced against the never-treated time path. This is the specification most vulnerable to the binary placebo concern (Section 5.10): if a treated-state secular trend exists during the post period that is unrelated to the mandate, the binary CSA estimand will absorb it.

The **bite-quintile dose-response** keeps the never-treated comparison pool but estimates a separate CSA model for each quintile. The estimand is the per-quintile ATT against the same comparison group. Its identifying variation is *across treated facilities within the post period* — comparing how the most-binding quintile’s response differs from the least-binding quintile’s response, against a common counterfactual. A treated-state secular trend that affects all facilities in a state equally would shift every quintile’s coefficient by the same amount, leaving the Q1→Q5 spread (and the cross-quintile gradient) unaffected. The dose-response gradient is therefore robust to the binary placebo concern in a way the binary pooled ATT is not.

The **TWFE triple-interaction with state-by-quarter FE** is the sharpest within-state-quarter design in the paper. The state-by-quarter FE absorb every shock that hits all facilities in a state-quarter cell equally — including any pandemic-era treated-state staffing trend, any state-specific labor-market shock, and any state-quarter Medicaid-rate adjustment. The triple interaction **bite** $\times$ **Medicaid** $\times$ **post** is identified entirely from cross-facility variation within state-quarter cells: facilities in the same state-quarter that differ in pre-policy bite and pre-policy Medicaid share. A state-level placebo cannot reproduce this variation because the placebo is itself absorbed by the state-by-quarter FE. The triple-interaction estimate is therefore the cleanest identifying variation in the paper, and I treat it as the headline.

The **5 × 5 ATT grid** uses the never-treated comparison pool but estimates one CSA model per (bite quintile, Medicaid quintile) cell. Cell sizes range from 56 treated CCNs (Bite-Q1 × Med-Q5) to 268. The grid is a descriptive companion to the triple interaction: it shows where the heterogeneity is concentrated without imposing a parametric functional form on the bite × Medicaid interaction.

The **rurality stratification** runs the bite-quintile dose-response separately for rural and urban facilities. The estimand is the per- quintile ATT within each stratum against the common never-treated pool. Like the bite-quintile dose-response, it is identified from within- state-quarter cross-facility variation and is robust to the binary placebo concern.

A worked example of the identifying variation in the triple interaction specification is illustrative. Consider Maine and New York in 2022Q3, two cohorts in the panel. Within Maine 2022Q3, the state-by-quarter FE absorbs whatever average response Maine facilities showed that quarter. The triple-interaction coefficient is then identified by the facility- level comparison: across Maine facilities, do those with the largest pre-policy bite *and* the highest pre-policy Medicaid share show the largest post-period staffing response? The same identifying logic runs within New York 2022Q3, within Ohio 2022Q3, and so on. The reported $\text{beta_3} = +0.153$ is a pooled estimate across all such within-state-quarter comparisons. A treated-state secular trend (the binary placebo concern) shifts the state-quarter mean but does *not* shift the cross-facility gradient within the cell.

4.5 Reproducibility

Replication materials rebuild the analytic panels from public sources and reproduce the tables and figures reported in the manuscript.

5. Results

The results below report the main estimates, heterogeneity analyses, and robustness checks from the public replication materials.

5.1 Sample composition

The estimation panel consists of 439,274 facility-quarter observations across 13,353 unique nursing facilities (Table 1). The treated set is 142,487 facility-quarter rows on 4,216 CCNs in 11 plotted cohorts (CA, ME, MA, IL, DC, NJ, CT, FL, RI, NY, OH, PA — Oregon’s small cohort is folded into the totals but contributes few observations); the never- treated set is 296,787 facility-quarter rows on 9,137 CCNs in 39 control states. Mean Total HPRD is slightly higher in treated states (3.537) than in never-treated states (3.461), consistent with the fact that the adopting states tend to be those whose pre-period sectors had either higher staffing or stronger political coalitions for staffing-floor legislation. The gap closes substantially in the RN component (0.570 vs 0.508 HPRD).

The cohort sub-table (lower panel of Table 1) reveals substantial cross-state variation in the pre-policy share of facilities below the new floor: from 5% in New Jersey (where the 3.0 HPRD floor was almost mechanically met) to 90% in Maine (4.44 HPRD) and 90% in Rhode Island (3.81 HPRD). This is the substantive variation the bite design exploits.

Table 1. Sample summary statistics

Group	N facility-qtr	N CCNs	States	Total HPRD mean	Total HPRD SD	RN HPRD mean	CNA HPRD mean	LPN HPRD mean
Treated states (R2+R3 eligible)	142,487	4,216	11	3.537	0.793	0.570	2.165	0.803
Never-treated states	296,787	9,137	39	3.461	0.929	0.508	2.140	0.813

Bite distribution (treated, winsorized 1/99): mean -0.174 HPRD, SD 0.847, P25 = -0.612, median = -0.119, P75 = +0.337, min = -2.904, max = +1.696.

Treatment-cohort coverage (R2+R3 eligible): ME 74 CCNs (2018Q1, mean bite +0.611, share below floor 89.2%); MA 281 (2018Q3, +0.069, 69.0%); IL 614 (2019Q1, +0.646, 83.2%); DC 15 (2019Q4, -0.296, 33.3%); NJ 221 (2020Q3, -1.016, 5.4%); CT 191 (2022Q1, -0.366, 33.0%); FL 658 (2022Q1, -0.251, 34.3%); RI 71 (2022Q1, +0.656, 90.1%); NY 568 (2022Q2, +0.265, 75.5%); OH 879 (2022Q3, -0.612, 12.9%); PA 644 (2023Q3, -0.608, 10.9%).




Figure 1: Raw outcome trends by treatment cohort and never-treated comparison

5.2 Bite distribution

The winsorized bite distribution (Table 1, middle panel; the raw-trends figure) has mean -0.174 HPRD, median -0.119 HPRD, and SD 0.847 HPRD across the 4,216 eligible-treated facilities. The negative central tendency reflects the descriptive fact that the median facility was already above the new floor at adoption; the long right tail (max $+1.696$ HPRD) captures the Maine and Rhode




Figure 2: Treatment timing and cumulative cohort coverage

Island facilities most distant from compliance. Quintile cutoffs are -0.746 / -0.287 / $+0.023$ / $+0.490$ HPRD. Cohort-by-cohort, Maine and Rhode Island contribute disproportionately to Q5 (deepest-below-floor); New Jersey, Pennsylvania, and Ohio populate Q1 (most-above-floor).

5.3 Headline triple interaction

The central result is the $\text{bite} \times \text{Medicaid-share} \times \text{post}$ triple interaction (Table 4B, row **triple**):

Table 4B. TWFE triple-interaction benchmark (Total HPRD)

Specification	Term	Coef	SE	95% CI
triple	$\text{bite} \times \text{post}$	+0.087	0.039	[+0.011, +0.163]
(fac+state-q FE)				
triple	$\text{medicaid} \times \text{post}$	-0.136	0.065	[-0.263, -0.009]
(fac+state-q FE)				
triple	$\text{bite} \times \text{med} \times \text{post}$	+0.153	0.034	[+0.088, +0.219]
(fac+state-q FE)				
triple_no_sq_fe	$\text{treat} \times \text{post}$	-0.148	0.069	[-0.283, -0.013]
(fac+quarter FE)				
triple_no_sq_fe	$\text{bite} \times \text{post}$	+0.040	0.036	[-0.030, +0.110]
triple_no_sq_fe	$\text{medicaid} \times \text{post}$	-0.033	0.055	[-0.141, +0.075]
triple_no_sq_fe	$\text{bite} \times \text{med} \times \text{post}$	+0.161	0.040	[+0.083, +0.240]

$$\text{bite} \times \text{Medicaid} \times \text{post} = +0.153 \text{ (SE 0.034, } t = 4.6, p < 0.001)$$

with facility and state-by-quarter fixed effects and standard errors clustered at the state level ($n_{\text{obs}} = 437,666$). The companion specification with facility + calendar-quarter FE (row **triple_no_sq_fe**) recovers a near-identical coefficient (+0.161, SE 0.040, $t = 4.0$, $p < 0.001$) alongside the binary $\text{treat} \times \text{post}$ (-0.148 , SE 0.069), confirming that the triple interaction is not a specification artifact.

In substantive units: a one-HPRD increase in pre-policy bite (e.g., moving a facility from at-floor to one HPRD below) raises the post- period Total HPRD response by **+0.15 HPRD per unit of Medicaid share**. At the mean treated-population Medicaid share (0.577), the bite-slope evaluated at the triple specification is +0.087 HPRD per unit of bite; at one tercile above that mean (Medicaid share approximately 0.69) the bite-slope rises to $+0.087 + 0.153 \times (0.69 - 0.577)$ approximately +0.104. In high-Medicaid facilities, the most-binding facilities raise nurse- hours more than they would have absent the mandate; in low-Medicaid facilities, the bite slope is essentially zero.

Medicaid tercile	Q1 (above)	Q5 (below)	Q1→Q5 spread
Low (T1)	−0.304 (0.038)	+0.017 (0.027)	0.32
Mid (T2)	−0.403 (0.055)	−0.002 (0.022)	0.40
High (T3)	−0.396 (0.085)	+0.038 (0.020)	0.43




Figure 3: Stratified dose-response by Medicaid-share tercile

Two readings sharpen the headline. First, the only positive cell in the entire 3×5 grid is **High-Medicaid $\times$ Most-Binding (T3 $\times$ Q5 = +0.038, t approximately 1.9)**. This is precisely the policy-targeted cell: facilities most exposed to the floor and most fiscally constrained by Medicaid revenue mix. They appear to have stopped the post-COVID staffing slide, while facilities further from the floor or less Medicaid-dependent reverted toward it. Second, the cross-quintile spread (0.32 $\rightarrow$ 0.40 $\rightarrow$ 0.43) provides a within-tercile complement to the triple-interaction estimate: as Medicaid share rises, the distance between the most- and least-binding facilities' responses widens by 0.11 HPRD across terciles, a magnitude consistent with the TWFE triple slope.

5.5 Pooled dose-response

The pooled (cross-Medicaid) dose-response (Table 3 and the pooled dose-response figure) preserves the same Q1→Q5 gradient but at the population-average Medicaid level: Q1 = −0.343 (SE 0.031), Q2 = −0.202, Q3 = −0.196, Q4 = −0.155, Q5 = +0.026 (SE 0.014). The Q1→Q5 spread is 0.37 HPRD, between the T1 and T3 spreads in the stratified dose-response. The dose-response *itself* — without the Medicaid stratification — is the cleanest visual evidence that the staffing floor bound at the binding margin: Q5 facilities held staffing flat, while Q1 facilities reverted toward the floor.

Table 3. Pooled dose-response by bite quintile (Total HPRD)

Quintile	Bite range (HPRD)	ATT	SE	95% CI
Q1 (most above)	bite_w $\ll$ 0	-0.343	0.031	[-0.403, -0.282]
Q2		-0.202	0.018	[-0.237, -0.167]
Q3		-0.196	0.017	[-0.228, -0.164]
Q4		-0.155	0.014	[-0.182, -0.128]
Q5 (deepest below)	bite_w $\gg$ 0	+0.026	0.014	[-0.001, +0.052]




Figure 4: Bite-by-Medicaid ATT heatmap

Q1 to Q5 spread: 0.37 HPRD (monotone gradient).




Figure 5: Pooled dose-response by pre-policy bite quintile

5.6 Rurality $\times$ bite

Table 5 and the rurality dose-response figure report the dose-response stratified by rurality. Rural facilities (648 CCNs) show a substantially steeper bite gradient than urban facilities (3,467 CCNs):

Table 5. Rurality $\times$ bite

Bite quintile	Rural ATT (n=648 CCNs)	Urban ATT (n=3,467 CCNs)
Q1	-0.431 (0.113)	-0.338 (0.033)
Q2	-0.256 (0.033)	-0.192 (0.020)
Q3	-0.091 (0.035)	-0.215 (0.018)
Q4	-0.087 (0.038)	-0.170 (0.014)
Q5	+0.076 (0.029)	+0.004 (0.015)
Q1 to Q5 spread	0.51	0.34




Figure 6: Dose-response by rurality

The rural Bite-Q5 cell (+0.076) is the strongest positive treatment- effect cell anywhere in the

heterogeneity supplement. Most-binding rural facilities staffed *up* in absolute terms — not merely flat — in response to the mandate. The Bite-Q1 cell magnitudes also widen (rural -0.431 vs urban -0.338 in the least-binding facilities), which is consistent with the Medicaid-share heterogeneity already documented (rural facilities are more Medicaid-dependent on average). The interpretation is that whatever labor-market thinness rural facilities face, the mandate forced compliance at the binding margin even more sharply in rural areas than in urban areas.

5.7 Resident-mix robustness

A natural concern is that the binary `treat × post` effect is driven by case-mix composition rather than the staffing-floor mandate. Table 6 addresses this directly. Adding seven LTCFocus pre-treatment mix covariates (percent White / Black / Hispanic, average resident age, acuity index, occupancy percent, percent Alzheimer’s/dementia) interacted with post leaves the binary `treat × post` essentially unchanged: -0.145 (bare TWFE) $\rightarrow -0.140$ (with bite $\times$ post) $\rightarrow -0.139$ (with all seven mix $\times$ post controls). The 95% CIs overlap nearly perfectly. Of the seven mix covariates, only `pct_black × post` (-0.0020 , $p < 0.001$), `pct_hispanic × post` (-0.0050 , $p < 0.01$), and `acuity_index × post` (-0.022 , $p < 0.05$) are individually statistically significant in spec C — facilities with higher Black/Hispanic resident shares and higher acuity saw larger HPRD declines, again consistent with Medicaid-share-correlated fiscal pressure rather than the floor mandate driving the treated-state secular trend. The binary placebo concern documented in the next section is therefore not a case-mix artifact.

Table 6. Resident-mix robustness (binary `treat × post`)

Specification	Controls added	treat × post coef	SE
A bare	facility + quarter FE	-0.145	0.066
B + bite × post	bite × post added	-0.140	0.068
C + 7 mix × post	+ percent White/Black/Hispanic, age, acuity, occupancy, percent Alz/dem × post	-0.139	0.060

5.8 Pre-trends and event-study

The event-study figures plot the binary CSA event-study coefficients for all four nurse-staffing outcomes. The visual takeaway is that pre-period leads are individually small ($|\text{coef}| \leq 0.03$ HPRD across $e=-8$ to $e=-2$ for all outcomes), while post-period coefficients build to magnitudes of -0.14 to -0.22 by $e=+10$ for Total/CNA HPRD.

The joint pre-trend Wald test, however, rejects parallel-trends for three of four outcomes (Total HPRD $\chi^2(7) = 46.7$, $p < 0.001$; CNA HPRD $\chi^2(7) = 65.4$, $p < 0.001$; LPN HPRD $\chi^2(7) = 42.1$, $p < 0.001$; RN HPRD $\chi^2(7) = 12.0$, $p = 0.10$). The tension between visually small leads and a statistically rejected joint test is a power-driven artifact of the large facility-quarter sample (439,274 obs) — the analytic standard errors are tight enough that small pre-period drift becomes detectable. The substantive question is whether the post-period magnitudes survive the implied bias.




Figure 7: Callaway-Sant'Anna event study, Total Nurse HPRD




Figure 8: Callaway-Sant'Anna event study, RN HPRD




Figure 9: Callaway-Sant'Anna event study, CNA HPRD




Figure 10: Callaway-Sant'Anna event study, LPN HPRD

The Rambachan-Roth (2023) HonestDiD bounds figure answers this question. At $M = 1.5$ — i.e., the worst-case post-period bias is allowed to be 1.5 times the worst-case pre-period deviation — the robust 95% CI on Total HPRD still excludes zero at every event time $e = 2$ through $e = 12$. Only $e = 0$ and $e = 1$ lose significance under the worst-case bound. The post-period effects are therefore an order of magnitude larger than the pre-period drift could plausibly explain.




Figure 11: HonestDiD relative-magnitudes sensitivity bounds

5.9 Goodman-Bacon decomposition

The Goodman-Bacon decomposition figure reports the Bacon decomposition of the TWFE binary specification on Total HPRD. The TWFE-implied beta is **−0.0983**, within rounding distance of the CSA simple aggregate (−0.143). The two non-empty buckets are:

- **Treated vs never-treated:** weight 0.813, average beta = −0.123, weighted contribution −0.102 HPRD.
- **Earlier vs later (clean):** weight 0.187, average beta = +0.010, weighted contribution +0.004 HPRD.

The forbidden “later vs already-treated” bucket is empty in my sample (no overlap windows in which an earlier-treated state would serve as a control for a later-treated state with both still treated). This rules out already-treated-as-control contamination as the explanation for the binary effect. The TWFE estimate is dominated by clean treated-vs-never-treated comparisons; the contamination is not the issue.

5.10 Honest caveat — the binary-spec placebo

The remaining identifying caveat is the most serious one. The placebo-date figure plots the placebo distribution: shifting each state’s effective date by −8 quarters (no actual mandate change) and re-estimating the CSA pooled ATT yields a placebo coefficient of **−0.129 (SE 0.010)** on Total HPRD — economically indistinguishable from the actual estimate of −0.143. The binary average effect is therefore not a credible clean policy estimate. It is partly (perhaps largely) a treated-state-specific secular trend: post-COVID staffing erosion was geographically uneven, and many of the adopting states are also the states with the steepest pandemic-era staffing declines.

This is the principal interpretive caveat for the binary specification in this paper. I address it explicitly rather than minimizing it. The binary effect is reported in Tables 2 and 6 as “what the average obscures,” and the paper’s identifying claims rest on three specifications that the placebo cannot replicate:






Figure 13: Placebo-date distribution

1. The **bite-quintile dose-response** (Section 5.5; Table 3 and the pooled dose-response figure): identified by within-treated-state variation in pre-policy bite, differencing out the treated-state secular trend.
2. The **Medicaid-share triple interaction** (Section 5.3; Table 4B and the Medicaid-share dose-response figure): identified by within-state-quarter variation across facilities, which a state-level placebo shift cannot reproduce.
3. The **rurality $\times$ bite stratification** (Section 5.6; Table 5 and the rurality figure): identified by within-state-quarter variation across urban and rural facilities, again immune to a state-level placebo.

All three of those specifications return signed-in-the-policy-direction results. The binary average obscures meaningful heterogeneity that the within-state designs uncover.

5.11 Direct robustness on the heterogeneity headline

The within-state-quarter identifying claim is logically defensible (state-quarter FE absorbs state-level shifts) but should be tested directly. Five additional robustness checks address the most demanding versions of the placebo, small-G, regional-confounding, single-state-driver, and parallel-trends-on-the-triple concerns.

Direct placebo on the triple. Re-estimating the headline triple under the same $-8q$ t0 shift used for the binary placebo returns **+0.040 (SE 0.066, $p = 0.55$)** under facility + state-quarter FE — null and well below the actual +0.153. The state-quarter FE absorbs the state-level shift the binary placebo represents, leaving the cross-facility heterogeneity slope identified by variation

the placebo cannot reproduce.

State-pairs cluster bootstrap. With $G = 13$ treated states the analytic CRV1 SE on the triple is approximately 40% smaller than a state-level pairs cluster bootstrap (resampling the union of treated + never-treated states with replacement, $B = 200$). The bootstrap returns mean $+0.156$, SE 0.047 (vs analytic 0.034), 95% percentile CI $[+0.081, +0.261]$, two-sided $p = \mathbf{0.005}$.

Region $\times$ quarter FE. Re-estimating with Census region $\times$ quarter FE rather than state $\times$ quarter FE (a strictly weaker absorption that probes whether the within-state-quarter result is sensitive to the West/Northeast concentration of treated states) returns a triple of $\mathbf{+0.173}$ (SE $\mathbf{0.038}$, $p < \mathbf{0.001}$), essentially unchanged from the $+0.153$ headline. The same specification recovers a binary `treat  $\times$  post` of -0.177 (consistent with the treated-state pandemic-era drift documented in §5.10).

Rural Bite-Q5 leave-one-state-out. The rural Bite-Q5 cell ($+0.076$, SE 0.029) is identified from 218 rural CCNs spread across 8 treated states with rural Bite-Q5 facilities (CT, FL, IL, MA, ME, NY, OH, PA). Dropping any single treated state moves the rural Bite-Q5 coefficient between $\mathbf{+0.056}$ (drop ME) and $\mathbf{+0.111}$ (drop NY); all 13 LOO replications remain positive and significant at the 5% level. No single state drives the rural finding.

Event-study triple decomposition. The static pooled triple ($+0.153$) decomposed by event time yields a noisy per-quarter trajectory: pre-period $|\text{coef}| \text{ max} = 0.108$ ($e = -8$); pre-period leads e in $[-7, -2]$ are individually insignificant ($|t| < 1.0$). Post-period coefficients build over time from -0.071 ($e = 0$) toward $+0.123$ ($e = 12$), peaking at $+0.136$ ($e = 6$); the static pooled $+0.153$ is approximately the post-period maximum, reflecting the post-period evolution rather than a uniform treatment effect. The pattern is consistent with a policy whose enforcement matures gradually — facilities adjust staffing toward the floor over multiple quarters rather than instantly at t_0 . The event-study triple is noisier per quarter than the static pooled estimate (per-quarter coefficients are mostly individually insignificant because the cross-facility heterogeneity slope is fragmented across 20 event-time bins, losing power), but the static pooled estimate captures the equilibrium response over the 12 post-period quarters where the mandate is operative.

Baseline-HPRD-decile stress test. Adding baseline-HPRD-decile $\times$ post fixed effects mechanically absorbs much of the bite $\times$ post variation (within-state, baseline HPRD and bite are inversely related). Under this stress test the triple drops to $+0.047$ (SE 0.041 , $p = 0.25$). I document this as a stress test rather than a clean robustness check — the baseline-decile is a coarse version of bite, and adding it interacted with post strips out the identifying variation by construction.

6. Discussion

6.1 Summary of findings

The paper’s central empirical claim is that **state staffing-floor mandates work where they bind and where they matter**. The dose-response gradient ($Q1 \rightarrow Q5 = 0.37$ HPRD spread) is monotone and economically meaningful; it steepens with Medicaid share ($0.32 / 0.40 / 0.43$ HPRD across terciles); the only positive cell in the entire 3×5 grid is the policy-targeted high-Medicaid $\times$ most-binding cell; and rural facilities show a steeper gradient still (0.51 HPRD spread) with the strongest positive cell in the supplement (rural Bite-Q5 = $+0.076$). The binary average effect is contaminated by a treated-state secular trend and is reported honestly as such.




Figure 14: Event-study triple decomposition

A way to summarize the findings for a policy audience: the mandate appears to have stopped the post-COVID staffing slide at exactly the facilities where the slide would otherwise have been worst — high- Medicaid, most-binding, and rural facilities. It did not raise staffing across the board, but it bound where the policy’s logic predicted it should bind.

6.2 Relation to prior literature

The findings extend three established strands of work. First, the state-mandate average-effect literature (Werner et al. 2026; Bowblis 2011; Matsudaira 2014; Park and Stearns 2009) documents modest positive average-effect estimates. I find that the average-effect estimate in my 13-state, 2017Q1–2025Q4 sample is contaminated by treated-state pandemic-era drift, but that the underlying mandate mechanism is recoverable from the heterogeneity slopes. The Werner et al. (2026) +0.18 HPRD average and my heterogeneity patterns are consistent: their average reflects a population of adopting states that was mostly bound by the floor, so the average is identified by the same within-state convergence I recover from the dose-response.

Second, the structural-financing literature (Hackmann 2019) establishes that Medicaid reimbursement is the binding margin on staffing investment. My Medicaid-share triple interaction is the direct reduced-form test of that mechanism applied to a mandate setting: when staffing must rise, it rises most where Medicaid-share-correlated fiscal pressure is highest (because the mandate is binding) and *relative to* what the same facility would have done absent the mandate. My +0.153 HPRD per unit of Medicaid share triple slope maps cleanly onto the Hackmann elasticity: if Medicaid revenue is the binding constraint on quality investment, then a binding mandate forces investment exactly where the constraint binds hardest.

Third, the post-pandemic agency-staffing literature (Bowblis et al. 2024; Office of the Assistant Sec-

retary for Planning and Evaluation 2025) shows that nearly half of facilities now substitute toward agency labor. My staffing-level outcome (Total HPRD) does not separate employee hours from contract hours; an obvious extension is to repeat the heterogeneity analysis on agency-substitution outcomes (`Hrs*_emp` vs `Hrs*_ctr` in my cleaned PBJ panel). I expect the high- Medicaid $\times$ most-binding cell to show the largest agency-substitution share, but that is a follow-up.

6.3 Mechanism

Why do high-Medicaid, most-binding facilities respond most? Two non- mutually-exclusive mechanisms:

- **The fiscal-margin channel.** Medicaid-heavy facilities cannot cross-subsidize staffing investment from commercial revenues. When a binding mandate forces staffing up, the only response margins available are (a) eat the loss, (b) cut other inputs (housekeeping, food service, activities — the Bowlblis and Hyer (2013) substitution margin), or (c) substitute toward agency labor at 50–60% premium cost (Bowlblis et al. 2024). All three margins raise the marginal cost of compliance. Facilities with deeper bites have less compliance room to begin with, and so respond more sharply.
- **The mandate-as-binding-constraint channel.** For low-bite, low- Medicaid facilities, the mandate is essentially a paper exercise. For high-bite, high-Medicaid facilities, it is a binding regulatory requirement. The post-period response distribution should therefore widen across both bite and Medicaid dimensions — which is exactly what the 5×5 ATT grid (the 5×5 ATT grid) shows.

Why does rural respond more sharply? The natural prediction is that thinner rural labor markets (Friedrich and Hackmann 2021) would *blunt* the mandate’s bite — facilities cannot hire what is not in the labor pool. My finding goes the other way: rural Bite-Q5 (+0.076) is the strongest positive cell. Two possibilities are consistent with the data. (i) Rural facilities are more Medicaid-dependent on average, so the financing-margin channel dominates the labor-market channel. (ii) Rural facilities face fewer outside options for agency substitution, so a binding mandate is satisfied through directly- employed FTE expansion rather than agency hours — which raises HPRD on the level I measure. I cannot distinguish these two mechanisms with my current outcome set; the agency-substitution analysis flagged in Section 6.2 is the natural test.

6.4 Policy implications

Three implications follow directly from the heterogeneity findings.

For any future federal staffing rule. The 2024 CMS rule’s bite was strongly correlated with Medicaid census (Chidambaram et al. 2024; Skinner et al. 2025). My results suggest that the mandate *does* bind exactly the facilities the rule was designed to reach, but only in the high-Medicaid segment of the bite distribution. A re-promulgated federal rule (or any state-level revival) should target Medicaid-heavy facilities for technical assistance, advance working-capital support, or rate-cell adjustments, because those are the facilities where the mandate will actually bind and where the fiscal margin is tightest.

For state staffing standards. Average-effect estimates from one state (or even from twenty-two states) cannot be transferred to a new state without accounting for that state’s bite distribution and Medicaid mix. A state with a low-bite, low-Medicaid sector will see no measurable response; a state with a high-bite, high-Medicaid sector will see meaningful staffing improvements but at high

marginal cost.

For Medicaid rate-setting. The triple-interaction finding implies that the marginal cost of mandate compliance is concentrated on Medicaid-heavy facilities. If Medicaid rate cells do not adjust to reflect mandate-driven cost increases, the binding margin will fall on facility margins (the Hackmann mechanism in reverse) and over time on closures, ownership turnover, or service-line changes. This is consistent with the Werner et al. (2026) finding of no detectable closure effect *on average* but is also consistent with selection on closure that future work should test directly.

6.5 Limitations

I acknowledge five binding limitations.

1. **Binary placebo similarity.** The $-8q$ placebo coefficient (approximately -0.13) is close to the actual binary CSA estimate (-0.143). The binary average effect should not be interpreted as a clean policy estimate. I pivot to the within-state-quarter heterogeneity specifications, which directly survive the parallel placebo (placebo triple = $+0.040$, SE 0.066 , $p = 0.55$; §5.11).
2. **13-state cohort, regional concentration.** I restrict to the 13 states with post-2017 numerical-floor adoptions, concentrated in West (CA, OR) and Northeast (MA, NJ, NY, CT, RI, ME, DC, PA). External validity to the 12 states with pre-2017 floors (notably California’s 1999/2003 law studied by 1) is limited. A region $\times$ quarter FE robustness preserves the headline triple ($+0.173$ vs $+0.153$; §5.11).
3. **Small-G inference.** With 13 treated states the analytic CRV1 SE on the triple is approximately 40% smaller than a state-pairs cluster bootstrap (boot SE 0.047 vs analytic 0.034 ; ratio 1.40); I report both (analytic $p < 0.001$, bootstrap $p = 0.005$, $B = 200$; §5.11).
4. **HCRIS Medicaid-days field essentially empty.** The HCRIS-SNF `pct_medicaid_days` is populated for only 0.3% of cost-report rows in the available cost-report files. LTCFocus is the sole operative Medicaid-share source, with HCRIS as fallback for a small CCN tail (1.4% of treated set); 60 CCNs (1.4%) lack any Medicaid-share data and are dropped from the heterogeneity analysis.
5. **Quality outcomes deferred — staffing levels only.** Care Compare deficiency citations and time-varying QM stars are available only as a static April-2026 snapshot in the available public files. The current paper makes claims about staffing **levels** in response to the mandate, *not* about staffing **quality** outcomes (deficiency citations, QM star changes). The natural follow-up applies the heterogeneity framework to the deficiency-citation archive and QM star distribution once the time-varying data is loaded.

6.6 What the average-effect headline obscures

A useful framing for the policy reader is to ask what an analyst would have concluded had only the binary CSA pooled ATT been reported. The answer is “the mandate did not work, and may have made staffing worse” — a coefficient of -0.143 HPRD on Total HPRD has the wrong sign for a binding mandate. That conclusion would be incorrect on two counts. First, the -0.143 is not a clean policy estimate; the placebo specification returns a near-identical -0.13 , consistent with a treated-state secular trend rather than a mandate effect. Second, the binary average necessarily averages across facilities for which the mandate is non-binding (Q1, Q2 facilities already above the floor) and facilities for which it is binding (Q4, Q5 facilities below the floor). When the binding-and-fiscally-constrained subset (Bite-Q5 $\times$ Med-Q5 = $+0.033$; rural Bite-Q5 = $+0.076$) is isolated,

the mandate effect is positive — and it is the only cell in the heterogeneity grid where the policy operates as policy-makers intended.

This framing matters for interpreting the evidence. A reader asking “did state staffing-floor mandates raise nurse staffing?” will get one answer from the binary estimate and a different, more accurate answer from the heterogeneity. The paper’s contribution is to show that the heterogeneity is the right answer, and to discipline the binary interpretation accordingly.

6.7 Reconciling with the average-effect literature

A natural question is how to reconcile my paper’s framing with Werner et al. (2026), the most directly comparable recent study. Werner et al. find a +0.18 HPRD (approximately 5%) average effect across twenty-two states 2010–23 with no detectable closure or net-margin response. My binary CSA pooled estimate, in contrast, is −0.143 HPRD on a restricted thirteen-state, post-2017 cohort.

Three reconciliation points are worth making. First, the samples differ: Werner et al. include the 2010–17 wave of state mandates that predate the post-COVID staffing erosion that contaminates my binary estimate. Their sample’s pre-period contains less of the secular trend that drives the placebo concern in my 13-state cohort. Second, their estimator is a TWFE event-study with cohort-by-event-time heterogeneity controls; ours is CSA with the never-treated comparison group plus a TWFE benchmark. The TWFE benchmark in the TWFE benchmark ($\text{treat} \times \text{post} = -0.145$ with quarter FE) is closer in design to theirs but also shows the sign-flip relative to their +0.18 — again pointing to the post-2020 sample composition as the substantive difference. Third, the reconciliation is not “one estimate is right and the other wrong.” The binary average effect *is* what it is in each sample; the contribution of my paper is to show that in the post-2017 cohort, the binary average is contaminated by treated-state pandemic-era drift, and that the policy-relevant signal is in the heterogeneity. A robust reading of the two papers together is: state mandates raise staffing on average (Werner), and the bulk of that effect is concentrated at the high-Medicaid, most-binding, rural facilities where the mandate’s policy logic predicts it should bind (this paper).

6.8 Future research

Four extensions follow directly from the limitations:

1. **Agency-substitution outcomes.** My cleaned PBJ panel separates employee from contract hours (`Hrs*_emp` vs `Hrs*_ctr`). Re-running the heterogeneity on agency-share outcomes would pin down whether high-Medicaid $\times$ most-binding facilities respond on the directly-employed margin or on the agency margin. The Office of the Assistant Secretary for Planning and Evaluation (2025) and Bowblis et al. (2024) work suggests the agency margin should dominate; I have the data to test it.
2. **Time-varying quality outcomes.** Once the Care Compare deficiency-citation archive is loaded, the same heterogeneity framework can test whether staffing improvements translate into quality improvements at the high-bite $\times$ high-Medicaid cells. This is the obvious successor question to Werner et al. (2013) and
 - 1.
3. **Closure and ownership turnover.** The Werner et al. (2026) finding of no closure effect on average masks heterogeneity. Re-running the heterogeneity framework on closure outcomes (PBJ enrollment exit + Care Compare special-focus transitions) would test for selection-on-closure.

4. **Bootstrap-clustered SE.** A state-level cluster bootstrap on the triple-interaction and dose-response estimates would complete the inference package. The triple-interaction t-statistic of 4.6 leaves substantial cushion against any plausible bootstrap inflation, but the formal exercise is overdue.
-

7. Conclusion

When the federal CMS staffing rule was repealed in late 2025, the nursing-home staffing-mandate debate returned to the states. Existing state-mandate evaluations had established that floors raise average direct-care staffing modestly without obvious closure effects. My analysis takes the next step. Using a 510k-row PBJ facility-quarter panel covering 2017Q1–2025Q4 across 13 adopting states, I estimate the staggered-DiD average effect, the dose-response by pre-policy bite, the Medicaid-share $\times$ bite triple interaction, and the rurality stratification.

The binary average effect is contaminated by a treated-state secular trend — a placebo specification yields a near-identical coefficient — and is therefore not a credible clean policy estimate. The identifying within-state-quarter heterogeneity specifications all run in the policy-expected direction. The dose-response gradient is monotone Q1→Q5; it steepens with Medicaid share; the only positive treatment-effect cell in the 3×5 stratified grid is the high- Medicaid $\times$ most-binding cell; and rural facilities show the steepest gradient of all. The triple interaction `bite  $\times$  Medicaid  $\times$  post` is +0.153 (SE 0.034, $t = 4.6$).

The paper’s policy contribution is to argue that average-effect estimates obscure the heterogeneity that matters most for the design of any future staffing rule. The mandate works where it binds and where it matters — among the most-binding facilities that are most fiscally constrained by Medicaid revenue dependence and in rural markets where labor-market thinness is most likely to constrain staffing investment. Targeting Medicaid-heavy facilities for technical-assistance and rate-cell adjustments, and accounting for the bite distribution when evaluating mandate transferability across states, are the immediate policy implications.

The methodological contribution is to apply the Callaway et al. (2024) continuous-treatment dose-response framework to a setting where the binary specification is contaminated by secular trends — and to show that the within- state-quarter heterogeneity slopes are the cleaner identifying variation. In settings where treated-state secular trends threaten the binary identification, the heterogeneity machinery developed here should generalize beyond nursing-home staffing.

References

- Bhaumik, Deepon, and David C. Grabowski. 2025. “Implementation Challenges of the New Federal Nursing Home Staffing Rules Will Vary Across States.” *Health Affairs Scholar* 3 (2): qxaf009. <https://doi.org/10.1093/haschl/qxaf009>.
- Bhaumik, Deepon, Momotazur Rahman, Emily A. Gadbois, et al. 2026. “Staffing Conditions in US Nursing Homes Before, During, and After the COVID-19 Pandemic.” *Health Affairs* 45 (2): 193–99. <https://doi.org/10.1377/hlthaff.2025.00570>.
- Bohler, Forrest, and Eli Y. Adashi. 2025. “Staffing up: The Nursing Home Minimum Staffing

- Rule.” *Journal of General Internal Medicine* 40 (9): 2126–27. <https://doi.org/10.1007/s11606-025-09360-6>.
- Borusyak, Kirill, Xavier Jaravel, and Jann Spiess. 2024. “Revisiting Event-Study Designs: Robust and Efficient Estimation.” *The Review of Economic Studies* 91 (6): 3253–85. <https://doi.org/10.1093/restud/rdae007>.
- Bowblis, John R. 2011. “Staffing Ratios and Quality: An Analysis of Minimum Direct Care Staffing Requirements for Nursing Homes.” *Health Services Research* 46 (5): 1495–516. <https://doi.org/10.1111/j.1475-6773.2011.01274.x>.
- Bowblis, John R., Christopher S. Brunt, Huiwen Xu, Robert Applebaum, and David C. Grabowski. 2024. “Nursing Homes Increasingly Rely on Staffing Agencies for Direct Care Nursing.” *Health Affairs* 43 (3): 327–35. <https://doi.org/10.1377/hlthaff.2023.01101>.
- Bowblis, John R., and Kathryn Hyer. 2013. “Nursing Home Staffing Requirements and Input Substitution: Effects on Housekeeping, Food Service, and Activities Staff.” *Health Services Research* 48 (4): 1539–50. <https://doi.org/10.1111/1475-6773.12046>.
- Brunt, Christopher S. 2023. “Assessing the Impact of Enforcement and Compliance with Minimum Staffing Standards on the Quality of Care in Nursing Homes: Evidence from the Centers for Medicare and Medicaid Services’ Staff Star Rating Downgrade Policy.” *Health Economics* 32 (2): 235–76. <https://doi.org/10.1002/hec.4619>.
- Callaway, Brantly, Andrew Goodman-Bacon, and Pedro H. C. Sant’Anna. 2024. *Difference-in-Differences with a Continuous Treatment*. NBER Working Paper No. 32117. National Bureau of Economic Research. <https://doi.org/10.3386/w32117>.
- Callaway, Brantly, and Pedro H. C. Sant’Anna. 2021. “Difference-in-Differences with Multiple Time Periods.” *Journal of Econometrics* 225 (2): 200–230. <https://doi.org/10.1016/j.jeconom.2020.12.001>.
- Centers for Medicare & Medicaid Services. 2024. *Medicare and Medicaid Programs: Minimum Staffing Standards for Long-Term Care Facilities and Medicaid Institutional Payment Transparency Reporting — Final Rule (CMS-3442-F)*. Federal Register, 89 FR 40876. <https://www.federalregister.gov/documents/2024/05/10/2024-08273/medicare-and-medicaid-programs-minimum-staffing-standards-for-long-term-care-facilities-and-medicaid>.
- Centers for Medicare & Medicaid Services. 2025. *Medicare and Medicaid Programs: Repeal of Minimum Staffing Standards for Long-Term Care Facilities*. Federal Register, Document 2025-21792. <https://www.federalregister.gov/documents/2025/12/03/2025-21792/medicare-and-medicaid-programs-repeal-of-minimum-staffing-standards-for-long-term-care-facilities>.
- Chaisemartin, Clément de, and Xavier D’Haultfoeuille. 2020. “Two-Way Fixed Effects Estimators with Heterogeneous Treatment Effects.” *American Economic Review* 110 (9): 2964–96. <https://doi.org/10.1257/aer.20181169>.
- Chidambaram, Priya, Alice Burns, Tricia Neuman, and Robin Rudowitz. 2024. *A Closer Look at*

the Final Nursing Facility Rule and Which Facilities Might Meet New Staffing Requirements. KFF Issue Brief. <https://www.kff.org/medicaid/a-closer-look-at-the-final-nursing-facility-rule-and-which-facilities-might-meet-new-staffing-requirements/>.

Friedrich, Benjamin U., and Martin B. Hackmann. 2021. “The Returns to Nursing: Evidence from a Parental-Leave Program.” *The Review of Economic Studies* 88 (5): 2308–43. <https://doi.org/10.1093/restud/rdaa082>.

Gandhi, Ashvin, Huizi Yu, and David C. Grabowski. 2021. “High Nursing Staff Turnover in Nursing Homes Offers Important Quality Information.” *Health Affairs* 40 (3): 384–91. <https://doi.org/10.1377/hlthaff.2020.00957>.

Geng, Fangli, David G. Stevenson, and David C. Grabowski. 2019. “Daily Nursing Home Staffing Levels Highly Variable, Often Below CMS Expectations.” *Health Affairs* 38 (7): 1095–100. <https://doi.org/10.1377/hlthaff.2018.05322>.

Goodman-Bacon, Andrew. 2021. “Difference-in-Differences with Variation in Treatment Timing.” *Journal of Econometrics* 225 (2): 254–77. <https://doi.org/10.1016/j.jeconom.2021.03.014>.

Grabowski, David C., and John R. Bowlblis. 2023. “Minimum-Staffing Rules for U.S. Nursing Homes — Opportunities and Challenges.” *New England Journal of Medicine* 389 (18): 1637–40. <https://doi.org/10.1056/NEJMp2310544>.

Gupta, Atul, Sabrina T. Howell, Constantine Yannelis, and Abhinav Gupta. 2024. “Owner Incentives and Performance in Healthcare: Private Equity Investment in Nursing Homes.” *The Review of Financial Studies* 37 (4): 1029–77. <https://doi.org/10.1093/rfs/hhad082>.

Hackmann, Martin B. 2019. “Incentivizing Better Quality of Care: The Role of Medicaid and Competition in the Nursing Home Industry.” *American Economic Review* 109 (5): 1684–716. <https://doi.org/10.1257/aer.20151057>.

Jun, Hankyung, and David C. Grabowski. 2024. “Nursing Home Staffing: Share of Immigrant Certified Nursing Assistants Grew as US-Born Staff Numbers Fell, 2010–21.” *Health Affairs* 43 (1): 108–17. <https://doi.org/10.1377/hlthaff.2023.00881>.

Konetzka, R. Tamara, Sally C. Stearns, and Jeongyoung Park. 2008. “The Staffing-Outcomes Relationship in Nursing Homes.” *Health Services Research* 43 (3): 1025–42. <https://doi.org/10.1111/j.1475-6773.2007.00803.x>.

Matsudaira, Jordan D. 2014. “Government Regulation and the Quality of Healthcare: Evidence from Minimum Staffing Legislation for Nursing Homes.” *Journal of Human Resources* 49 (1): 32–72.

McGarry, Brian E., David C. Grabowski, and Michael L. Barnett. 2020. “Severe Staffing and Personal Protective Equipment Shortages Faced by Nursing Homes During the COVID-19 Pandemic.” *Health Affairs* 39 (10): 1812–21. <https://doi.org/10.1377/hlthaff.2020.01269>.

Medicaid and CHIP Payment and Access Commission. 2024. *Compendium: State Policies Related*

to Nursing Facility Staffing. MACPAC compendium. <https://www.macpac.gov/publication/state-policies-related-to-nursing-facility-staffing/>.

New York State Department of Health. 2023. *Nursing Home Minimum Staffing and Direct Resident Care Spending — Public Health Law §2895-b and 10 NYCRR §415.13*. New York State Department of Health policy guidance. https://www.health.ny.gov/facilities/nursing/minimum_staffing/.

Office of the Assistant Secretary for Planning and Evaluation. 2025. *Use of Contract Staff in Nursing Homes Remains High After the COVID-19 Pandemic*. ASPE Issue Brief, U.S. Department of Health and Human Services. <https://aspe.hhs.gov/reports/use-contract-staff-nhs-remains-high-after-covid-19-pandemic>.

Park, Jeongyoung, and Sally C. Stearns. 2009. “Effects of State Minimum Staffing Standards on Nursing Home Staffing and Quality of Care.” *Health Services Research* 44 (1): 56–78. <https://doi.org/10.1111/j.1475-6773.2008.00906.x>.

Rambachan, Ashesh, and Jonathan Roth. 2023. “A More Credible Approach to Parallel Trends.” *The Review of Economic Studies* 90 (5): 2555–91. <https://doi.org/10.1093/restud/rdad018>.

Shen, Karen, Brian E. McGarry, and Ashvin D. Gandhi. 2023. “Health Care Staff Turnover and Quality of Care at Nursing Homes.” *JAMA Internal Medicine* 183 (11): 1247–54. <https://doi.org/10.1001/jamainternmed.2023.5225>.

Skinner, Robert J., Jiani Yu, Mark A. Unruh, et al. 2025. “State Variation in Staffing and Characteristics of Nursing Homes Most Impacted by New Federal Standards.” *Health Affairs Scholar* 3 (8): qxaf154. <https://doi.org/10.1093/haschl/qxaf154>.

Sun, Liyang, and Sarah Abraham. 2021. “Estimating Dynamic Treatment Effects in Event Studies with Heterogeneous Treatment Effects.” *Journal of Econometrics* 225 (2): 175–99. <https://doi.org/10.1016/j.jeconom.2020.09.006>.

Werner, Rachel M., Xinwei Chen, Norma B. Coe, and Andrew R. Olenski. 2026. “State Nursing Home Minimum Staffing Mandates: Increased Staff Levels, Minimal Impact on Finances and Closures, 2010–23.” *Health Affairs* 45 (3): 269–77. <https://doi.org/10.1377/hlthaff.2025.01223>.

Werner, Rachel M., R. Tamara Konetzka, and Daniel Polsky. 2013. “The Effect of Pay-for-Performance in Nursing Homes: Evidence from State Medicaid Programs.” *Health Services Research* 48 (4): 1393–414. <https://doi.org/10.1111/1475-6773.12035>.