

Suspension vs. Termination: Medicaid Continuity Policies During Incarceration and Effects on Post-Release Coverage, Employment, and Mortality

Jonathan Palisoc, University of Michigan, School of Public Health, Dept of Health Management & Policy

Abstract

Objective: To estimate the effect of states suspending Medicaid enrollment during incarceration versus terminating on post-release Medicaid enrollment, employment, earnings, and mortality. Under suspension, coverage is paused during incarceration and reactivated at release without a new application; under termination, eligibility is ended entirely and the individual must reapply after release from prison, navigating documentation requirements, processing delays, and bureaucratic complexity at a time when they are least equipped to do so.

Data Sources: Census Bureau Criminal Justice Administrative Records System (CJARS) Justice Outcomes Explorer, administratively linked aggregate outcomes for state prison release cohorts across 31 states, 2000-2020; and a novel state-year panel of suspension and termination policies.

Study Design: Staggered difference-in-differences exploiting the rollout of suspension across states between 2005 and 2022 and event study models. Outcomes are release-cohort level and therefore estimations are intent-to-treat (ITT) on the full release cohort.

Principal Findings: After controlling for Medicaid expansion, which was frequently co-timed with adoption of suspension, states saw absolute increases in Medicaid enrollment among justice-involved individuals of 11.1% one year after release; while the increase is consistent across all estimators, the limited between-state variation limits precise detection and most models are not statistically significant. Raw, uncontrolled estimations including Medicaid expansion effects project increases of 22.2-25.8 pp. Subgroup point estimates are broadly similar across sex, race, and age, other than a smaller estimate for women, consistent with their higher baseline enrollment. There were no effects of suspension vs.

termination policies on formal employment, earnings, mortality, or recidivism, though the design limits precise detection of these outcomes.

Keywords: Medicaid, incarceration, administrative burden, reentry, difference-in-differences, CJARS

1. Introduction

Each year, close to 10 million individuals are released from local, county, state and federal prisons in the United States (Carson and Kluckow, 2023; Zeng, 2023). The weeks right after release are the most dangerous period of these individuals' lives: justice-involved individuals are 12.7 times more likely to die in the first two weeks after release than the general population, and drug overdose is the leading cause (Binswanger et al., 2007). Over 80 percent of people in state prisons have at least one chronic condition, and release strips away the limited health infrastructure the correctional system provides (Mallik-Kane and Visser, 2008). While Medicaid is a potential source of post-release coverage, a longstanding federal inmate payment exclusion generally bars federal Medicaid payment for services delivered during incarceration. However, the statute does not require states to *disenroll* individuals during incarceration; as such, states have two options: either suspend coverage and automatically reactivate post-release, or terminate it and require individuals to re-apply once they are able.

This paper asks two questions regarding that difference. The first is whether suspension raised post-release Medicaid enrollment at one, three, and five years after release, and what was the magnitude of this effect after accounting for the co-timed ACA Medicaid expansion. The second is whether any coverage difference is carried through to downstream outcomes of direct policy interest: if post-release individuals who more easily regained coverage saw subsequent improvements in formal employment, earnings, recidivism, and mortality. This second question rests on a common argument among proponents of greater continuity in healthcare treatment for this population: as rates of substance use disorder and serious mental illness among people leaving prison are far above the general adult population, and Medicaid is the payment vehicle for medication for opioid use disorder and outpatient mental health

treatment (Wen et al., 2017), better access to behavioral health treatment could in turn improve health, employment, and recidivism (Aslim et al., 2022; Deza et al., 2024).

To answer both, I construct a state-year panel of suspension and termination policies from 2000 to 2025 and link it to the Census Bureau's Criminal Justice Administrative Records System (CJARS) Justice Outcomes Explorer (JOE), which reports administratively linked outcomes for state prison release cohorts. I estimate staggered difference-in-differences models on state by release-cohort cells, so the estimates are policy-level intent-to-treat associations for the whole release cohort, as the data is pooled and I cannot directly observe individual-level suspension or termination. The paper adds to the literature in two ways: 1) it provides the first multi-state evidence on suspension policy and post-release Medicaid enrollment at one, three, and five years, where the existing quasi-experimental literature is limited and mostly comprised of single-state studies (Section 3.1); and 2) it extends the evidence beyond the proximal outcome of enrollment to distal policy-relevant outcomes of formal employment, earnings, and mortality, using linkages to W-2 records and the Census Numident, which no prior study of this policy has had.

The question is also timely. Effective January 1, 2026, states may not terminate Medicaid or CHIP eligibility for an individual because that individual is an inmate of a public institution, under Section 205 of the Consolidated Appropriations Act, 2024 (P.L. 118-42), which amended Section 1902(a)(84) of the Social Security Act (CMS, 2025). The requirement also removes the cross-state variation that made evaluation possible, so the adoption era studied here is the evidence base that remains, and most of it has been studied one state at a time (Section 3.1). While suspension is already widely acknowledged to cause increases in post-release Medicaid enrollment, I use this paper to study its persistence and whether its effects extend to secondary outcomes the requirement is meant to improve (health, income, recidivism, etc). States that had adopted suspension saw a 22.2 percentage point higher one-year post-release Medicaid enrollment rate than states still terminating, and that contrast is stable across estimators, samples, and timing choices. However, suspension and Medicaid expansion policies happened together in most states, so the raw contrast reflects both policies. Accounting for expansion cuts the estimate in half

to approximately 11 percentage points - still a positive effect but not statistically distinguishable from zero.

2. Background

2.1 The Medicaid Inmate Exclusion Policy

Medicaid is a joint federal-state health insurance program enacted under Title XIX of the Social Security Act in 1965: states administer their own Medicaid programs subject to federal requirements and receive federal matching funds. Eligibility, covered services, and provider payment rates vary substantially across states due to their flexibility in program administration, but federal statute imposes a significant set of structural constraints - one being the Medicaid Inmate Exclusion Policy (MIEP).

The inmate exclusion is codified at Section 1905(a)(31)(A) of the Social Security Act (42 U.S.C. § 1396d(a)(31)(A)) and prohibits federal Medicaid reimbursement for health care services delivered to any “inmate of a public institution,” with a narrow exception for inpatient stays of at least 24 hours (MACPAC, 2018). CMS has interpreted “inmate of a public institution” to include individuals held involuntarily in state and federal prisons, county and local jails, and certain juvenile facilities; it does not include individuals on probation, parole, or home detention. The original rationale was that states and localities, as operators of correctional facilities, bear constitutional responsibility for inmate health care under the Eighth Amendment (*Estelle v. Gamble*, 429 U.S. 97 (1976)), and Congress did not intend for Medicaid to subsidize that obligation.

The exclusion is a payment exclusion, not an eligibility exclusion: federal funds cannot pay for services delivered to an inmate, but incarceration does not itself end eligibility (CMS, 2025). For decades most states chose termination for a number of reasons: (1) inmates were rarely Medicaid-eligible post-release (before ACA expansion), (2) state Medicaid agencies had little financial incentive to maintain eligibility records for individuals who could not generate federal matching funds, and (3) corrections and Medicaid administrative systems were rarely integrated.

As around 600,000 individuals are released from state and federal prisons and an additional ~9 million admissions and releases occur in county and local jails (Carson and Kluckow, 2023; Zeng, 2023), the overwhelming majority of these individuals released from prison were effectively uninsured regardless of prior enrollment in the Medicaid program. Although recent improvements have been made in eligibility redeterminations, even a person who completed a Medicaid re-application on day 1 of reentry could wait as long as 30 to 90 days for a determination.

2.2 Emergence of Suspension Policies

CMS first explicitly authorized suspension as an alternative to termination in a State Medicaid Director Letter issued on May 25, 2004. The letter clarified that states could “suspend rather than terminate” eligibility during incarceration, with coverage reactivated at release without a new application.

Suspension in this context means that the individual’s Medicaid case is held in an inactive but identifiable status during incarceration. In other words, suspension preserves the federal payment bar under the Medicaid Inmate Exclusion Program, but the underlying eligibility record is preserved in the state Medicaid system, the case is automatically reactivated at release, and renewals can, in principle, be processed during the period. On the other hand, termination means the case is closed entirely; at release the individual must submit a new application, requalify on income and other criteria, supply documentation, and wait through a potentially long determination window.

Following the 2004 guidance, states adopted suspension policies in three broad waves. As of January 2025, 46 states suspended Medicaid during incarceration, 3 terminated (Alabama, Georgia, Texas), and 2 have unclear status (KFF, 2025). Figure 1 maps the adoption pattern.

While suspension has been widely adopted, actual implementation has varied across states. Some adopted time-limited suspension (30 days, 90 days, or up to one year, with conversion to termination after that), while others adopted indefinite suspension for the full incarceration period. Scope also varies: some states apply suspension only to state prisons, while others extend it to county and local jails. The operational infrastructure required to make suspension work in practice is substantial - real-time data exchange between corrections agencies and the state Medicaid office to flag admissions, an automated process to

lift the suspension flag at release, and procedures for processing renewals during the period. Bandara et al. (2025) document that several de jure suspension states lack one or more of these components, producing a meaningful gap between the policy on paper and what an individual actually experiences at release. Suspension's effect depends on there being eligibility to preserve, which likely explains why adoption of suspension policies closely tracked with timing of states' ACA expansions. Before the ACA, traditional Medicaid categories excluded most single, non-disabled adults, leaving roughly 80 percent of released individuals uninsured (Albertson et al., 2020); the expansion made an estimated 80 to 90 percent of justice-involved adults eligible in expansion states (Albertson et al., 2020; Camhi, Mistak, and Wachino, 2020) and so markedly enlarged the population entering prison with active coverage.

2.3 2026 Federal Mandate

Section 205 of the Consolidated Appropriations Act, 2024 (P.L. 118-42) amended Section 1902(a)(84) of the Social Security Act to prohibit states from terminating Medicaid or CHIP eligibility solely because an individual is an inmate of a public institution, effective January 1, 2026 (CMS, 2025). CMS guidance emphasizes that states retain a choice among suspension strategies, and that operationalizing the mandate requires data-sharing infrastructure between Medicaid agencies and carceral facilities, mechanisms to identify when a beneficiary becomes an inmate, and coordination on renewals during incarceration. While termination is no longer an option, I argue the results carry significant policy relevance as the termination policy could be reintroduced at the federal level and states vary widely in how fully they operationalize suspension.

2.4 Administrative Burden Framework

The theoretical framework most directly applicable to the suspension-termination distinction is the administrative burden literature, defined by Herd and Moynihan (2018) as “the experience of policy implementation as onerous”. Furthermore, they break burdens into three parts: learning costs (knowing that programs exist and understanding eligibility rules), compliance costs (the time and effort for paperwork and documentation), and psychological costs (stigma, stress, and the emotional toll of navigating government systems).

Termination imposes all three of these burden types on reentering individuals at a moment when their capacity to overcome them is at its lowest. Learning costs are high because individuals who have been incarcerated for months or years may be unaware of changes to Medicaid eligibility rules or application procedures. Compliance costs are extreme because the documentation requirements (proof of identity, residence, and income) involve precisely the things recently released individuals are least likely to have. Psychological costs compound the other two: stigma, frustration, and demoralization weigh on individuals already coping with the social isolation and stigma of incarceration.

When implemented well, suspension may remove many of these frictions: no action, no documentation, no bureaucratic interaction. Renewals and the reactivation step itself can still impose burden, but the overall workload to gain and maintain coverage is significantly reduced. Herd and Moynihan also show that administrative burdens fall hardest on marginalized populations, and recently incarcerated individuals are close to a limiting case: compounding disadvantages (poverty, chronic illness, histories of trauma, stigmatized status) meeting compounding barriers at a moment of acute life transition.

Research on the ACA expansion's effects for justice-involved populations is broadly positive. Aslim et al. (2022) find that expansion reduces recidivism by 11.5 percent, and a companion welfare analysis (Aslim, Mungan, and Yu, 2024) estimates benefit-cost ratios of \$3.45 to \$10.62 per dollar spent. Howell et al. (2023) find an 18 percentage-point coverage increase among criminal-justice-involved adults but no significant change in treatment utilization, and Wen et al. (2017) show that expansion was associated with a 70 percent increase in Medicaid-covered buprenorphine prescriptions.

3. Prior Literature

3.1 Quasi-Experimental Evidence on Suspension and Reentry Coverage

Three prior quasi-experimental studies are relevant to this paper; each examines a different margin of post-release coverage. Gollu and Zapryanova (2022) is the only prior multi-state study of suspension and reentry outcomes: using individual-level administrative prison release data, the authors estimate that suspension reduces the probability of returning to prison within one year by 2.91 percentage points, and

within three years by 4.58 percentage points, with larger effects for Black individuals and repeat offenders. However, they focus on estimating recidivism and do not estimate Medicaid enrollment, employment, or health effects. They also lack linkage to Medicaid enrollment files, tax records, or death records, and rely primarily on cross-sectional variation across states rather than fully exploiting within-state adoption timing.

Packham and Slusky (2024) provide the most directly relevant prior evidence on the mechanism of suspension. Using individual-level criminal justice records linked to Medicaid claims data from a single state (South Carolina), they estimate that the state's combined suspension/simplified re-enrollment policy increased Medicaid enrollment within six months of release by 11.1 percentage points, with corresponding increases in post-release health care utilization. Additionally, Burns et al. (2022) studied Wisconsin's pre-release Medicaid enrollment-assistance program, a related but distinct intervention that simplifies re-enrollment without preserving eligibility through the incarceration period. They document large increases in post-release outpatient care use after program implementation - showing that complementary friction-reduction mechanisms can increase care utilization.

3.2 Coverage Continuity, Mortality, and Crime

Suspension is, in effect, continuous eligibility across the incarceration period, while termination creates automatic involuntary discontinuities in enrollment through disenrollment at the point of incarceration, necessitating justice-involved individuals to complete a full enrollment process after release. That matters because most Medicaid disenrollment occurs due to administrative barriers rather than a change in eligibility (Herd and Moynihan, 2018) - and states adopting continuous eligibility have seen substantial reductions in coverage disruption. Goldman and Sommers (2020) estimate that the ACA expansion reduced churn by roughly half a million adults per year, which amplifies the effect of suspension by enlarging the pool with coverage to preserve.

Regarding mortality, while recent papers have indicated positive effects of Medicaid coverage on reducing mortality (Miller 2023), mortality effects for justice-involved populations have been less clear, perhaps due to the complex, multifaceted nature of post-release life.

Regarding crime, Deza et al. (2024) studied a mass Medicaid disenrollment in Tennessee in 2005 that removed coverage from 190,000 adults and found significant increases in total crime, particularly non-violent crime, suggesting a possible reverse causal effect for those who maintain Medicaid coverage. More broadly, simplifying enrollment reliably raises participation in the Medicaid program: automatic and presumptive eligibility, ex parte renewals, and simplified applications have each been associated with higher enrollment and fewer coverage gaps (Herd and Moynihan, 2025; Sugar et al., 2021), and the Medicaid unwinding after the COVID-19 continuous enrollment provision showed the reverse: millions lost coverage when redetermination resumed, many of them still eligible (KFF, 2025).

4. Data

4.1 The Criminal Justice Administrative Records System

The Criminal Justice Administrative Records System is a joint project of the U.S. Census Bureau and the University of Michigan that harmonizes criminal justice administrative records and links them to Census Bureau survey and administrative data, including tax filings (W-2), federal program participation records, and mortality data from the Census Numident (Finlay, Mueller-Smith, and Papp, 2022). The Justice Outcomes Explorer is a public product derived from CJARS, reporting aggregated outcomes by state, release cohort year, and post-release horizon. Its Social Safety Net module provides the outcome variables:

- **Medicaid enrollment rate:** the share enrolled at the specified post-release horizon (1, 3, or 5 years), through linkage to MSIS/T-MSIS enrollment records, calendar years 2005-2019
- **W-2 employment rate:** the share with any W-2 earnings at the horizon
- **Mean W-2 earnings:** mean W-2 earnings including members with zero earnings
- **SSI receipt rate:** the share receiving Supplemental Security Income, through Social Security Administration linkage, and HUD housing assistance rate (public housing / vouchers, rental assistance)

- **All-cause mortality rate:** the share deceased by the horizon, through the Census Numident file

The JOE corrections module separately provides release cohort counts, recidivism rates, and demographic composition, used for balance tables and composition tests. JOE also publishes the same release-cohort outcomes within one-way subgroup cells, by sex, by race and ethnicity, and by age at release. Two-way subgroups are not supported.

4.2 State Suspension and Termination Policy Panel

I construct a state-year panel classifying each state’s Medicaid eligibility policy for incarcerated individuals from 2000 to 2025, coding each state-year as termination, time-limited suspension, or indefinite suspension. No single existing source covers the full period. I therefore triangulate six independent sources and date each state’s switch to the earliest year they jointly support the policy implementation (KFF, 2025a; MACPAC, 2018, 2023; National Association of Counties, 2015; Gollu and Zapryanova, 2022; McKee et al., 2015; Bandara et al., 2025). The appendix gives the source hierarchy and confidence ratings; one state (Florida) is rated low-confidence, and dropping it lowers the 1-year Medicaid enrollment estimate by 1.3 percentage points (Section 6.6).

Table 1 reports the panel for the 31 CJARS states. Every state began the period under termination, so the design compares states before and after they switched rather than states that always differed. Of the 31 states, 29 adopted suspension between 2005 and 2022 and 2 never did (Georgia and Texas); 27 of the 29 adopted by 2020, the last release cohort JOE observes, while Florida (2022) and Oklahoma (2021) adopted afterward and behave as untreated states throughout. Most adopters chose indefinite suspension, and three CJARS states (Arkansas, Kansas, South Carolina) adopted a time-limited version.

Table 1. State Suspension Policy, Expansion Timing, and Analysis Panel Membership, CJARS States


State	Policy in 2000	Suspension adoption year	Suspension type	ACA expansion year	Medicaid and outcome panel (17)	Recidivism panel (22)
-------	----------------	--------------------------	-----------------	--------------------	---------------------------------	-----------------------

Arizona	Termination	2013	Indefinite	2014	Yes	Yes
Arkansas	Termination	2014	Time-limited, 1 year	2014	No	No
California	Termination	2014	Indefinite	2014	No	Yes
Colorado	Termination	2009	Indefinite	2014	Yes	Yes
Florida ^a	Termination	2022	Indefinite	None	Yes	Yes
Georgia	Termination	Never adopted	Termination throughout	None	No	Yes
Idaho	Termination	2020	Indefinite	2020	No	Yes
Illinois	Termination	2014	Indefinite	2014	Yes	Yes
Indiana	Termination	2015	Indefinite	2015	Yes	No
Iowa	Termination	2008	Indefinite	2014	No	Yes
Kansas	Termination	2018	Time-limited, 1 year	None	Yes	Yes
Maryland	Termination	2014	Indefinite	2014	No	No
Michigan	Termination	2014	Indefinite	2014	No	No
Minnesota	Termination	2007	Indefinite	2014	Yes	Yes
Montana	Termination	2016	Indefinite	2016	Yes	Yes
Nebraska	Termination	2020	Indefinite	2020	Yes	Yes
Nevada	Termination	2014	Indefinite	2014	No	Yes
New Jersey	Termination	2016	Indefinite	2014	No	No
New Mexico	Termination	2015	Indefinite	2014	No	Yes
New York	Termination	2008	Indefinite	2014	Yes	Yes
North Carolina	Termination	2017	Indefinite	2023	Yes	Yes
North Dakota	Termination	2014	Indefinite	2014	No	No
Oklahoma	Termination	2021	Indefinite	2021	Yes	Yes
Oregon	Termination	2005	Indefinite	2014	No	No
Pennsylvania	Termination	2015	Indefinite	2015	Yes	Yes
South Carolina	Termination	2018	Time-limited, 1 year	None	No	No
Texas	Termination	Never adopted	Termination throughout	None	Yes	Yes
Utah	Termination	2020	Indefinite	2020	Yes	Yes
Virginia	Termination	2019	Indefinite	2019	No	No
Washington	Termination	2005	Indefinite	2014	Yes	Yes
Wisconsin	Termination	2014	Indefinite	None	Yes	Yes

Notes: The table covers the 31 states in the CJARS release-cohort data, coded from the state-year policy panel in Section 4.2. “Policy in 2000” is termination for every state, and no CJARS state is coded as suspending before 2005. ACA expansion year is the year the state’s expansion took

effect, coded from KFF, and “None” means the state had not expanded as of the last year of the policy panel. Wisconsin covers adults to 100 percent of the federal poverty level under a waiver and is coded as a non-expansion state throughout. The final two columns indicate whether the state contributes outcome rows to the 17-state outcome panel and to the 22-state reincarceration panel; Indiana appears in the first but not the second. “Florida is the one low-confidence policy coding, and Section 6.6 reports the estimate with it dropped.

Figure 1. State adoption of Medicaid suspension policies, 2000-2025



Notes: Fill color indicates when a state adopted suspension. Bold labels identify the 31 states in the CJARS release-cohort panel, of which Georgia and Texas are never treated; Alabama still terminated as of late 2025 but is outside CJARS. Author’s coding from the state-year policy panel in Section 4.2.


4.3 Release Cohorts, Follow-Up Horizons, and Censoring

Each observation is a state by release-cohort cell, meaning everyone released from a given state’s prisons in a given year, built by CJARS and the Census Bureau from person-level linked administrative records and then aggregated.

I note that two different clocks run at once. I define the **policy clock** as calendar time relative to the state’s suspension adoption, and it defines treatment status and event time; conversely, the **outcome clock** is time since an individual’s own release, and defines the 1-, 3-, and 5-year follow-up points at which JOE measures outcomes. An estimate at horizon h is therefore attached to the policy in place at release, measured h years into each person’s own post-release life rather than h calendar years after a policy

change. Outcomes are point-in-time stocks, so the Medicaid enrollment rate at horizon h is the share of the cohort enrolled at that follow-up point rather than the share ever enrolled during the interval. Figure 2 shows both clocks and the censoring they create.

Figure 2. Two clocks and the observation windows: release cohorts, person-level follow-up horizons, and censoring



Notes: Panel A follows the 2010 release cohort along its own outcome clock, with the 1-, 3-, and 5-year follow-up points measured in 2011, 2013, and 2015. Panel B shows which release cohorts have an observable outcome at each horizon, given that the CJARS MSIS/T-MSIS linkage supports Medicaid measurement only in calendar years 2005–2019; shaded cells are observed and open cells are censored. Author's construction from the JOE data documentation.

The CJARS MSIS/T-MSIS linkage supports Medicaid measurement in calendar years 2005 through 2019, and a horizon- h outcome for release cohort t exists only if $t + h$ falls inside that window. That rule produces three cohort windows, each estimated on its own: release cohorts 2004–2018 at one year, 2002–2016 at three years, and 2000–2014 at five years. Because the most recent cohorts are absent from the

five-year sample, cross-horizon comparisons partly reflect cohort composition, which I check directly in Section 6.2.

4.4 Sample Construction and Panel Membership

I merge the JOE data with the policy panel, the ACA expansion panel, and the state-level controls by state FIPS code and year, and then impose three sample restrictions. State-year cells with fewer than 25 individuals are excluded to satisfy Census Bureau disclosure requirements, which JOE applies at the source, and I restrict to state prison releases. Each Medicaid, employment, and earnings horizon is estimated on the cohort window in Section 4.3, while mortality, housing assistance, and SSI reach back to release cohort 2000 and reincarceration is the only outcome observed through the 2020 cohort. The outcome panels are narrower than the 31-state policy panel, because JOE reports a state's release-cohort outcomes only where the underlying Census Bureau linkage has been built. The same 17 states identify every outcome except reincarceration, which is observed in 22 states (Table 1), and the 20 states outside CJARS, including Alabama, are excluded entirely due to lack of a data-sharing agreement. In short, the sample selects on whether a linkage has been built rather than on policy. Two of the 14 CJARS states without Medicaid outcome rows matter here: Georgia is one of the two never-treated states, and South Carolina is the state Packham and Slusky (2024) study directly. The 17 included and 14 excluded CJARS states look similar on unemployment (5.4 versus 5.8 percent), overdose mortality (14.2 versus 13.5 per 100,000), and poverty (12.7 versus 13.3 percent).

4.5 Outcomes, Treatment, and Estimand

Coverage is the first stage, and the remaining outcomes sit farther along the pathway in Section 1. W-2 employment and mean W-2 earnings would move if untreated substance use disorder or serious mental illness keeps people out of formal work. Reincarceration would move if treatment reduces drug-driven offending and the parole violations that follow relapse, which is the pathway Aslim et al. (2022) and Deza et al. (2024) describe. All-cause mortality is a substantively important downstream outcome, but its

horizons are not aligned with the risk that motivates it. The leading cause of death in the weeks after release is overdose (Binswanger et al., 2007), and a cumulative death rate at one, three, or five years cannot isolate that acute window. HUD housing assistance and SSI receipt are exploratory and are reported in the appendix.

Because JOE reports outcomes for everyone released from state prison in a given state-year, the estimates are intent-to-treat associations for the average outcome of the full release cohort under the de jure policy in force at release. The cohort-level association is diluted by individuals who had no eligibility to preserve at admission, and I cannot observe whether a given person was enrolled before incarceration or reactivated at release. Each state-cohort cell enters unweighted, so the primary estimand is an unweighted state-average association that gives a small release system the same weight as a large one. I designate my primary estimand as unweighted as it answers the state policy question, but also generate a population-weighted secondary estimand by weighting each state-cohort cell by the size of its release cohort (Appendix Table A6). These magnitudes are not comparable to individual-level estimates from single-state claims studies.

4.6 ACA Expansion, Eligibility Thresholds, and State Controls

I code ACA expansion status and expansion effective dates from KFF. Two versions of the expansion indicator appear in the paper, and they answer different questions. As first noted in Section 4.3, the **release-year clock** merges expansion status on the release-cohort year t , so the estimate conditions on the eligibility regime the cohort was released into. The **outcome-year clock** merges it on the outcome-measurement year $t + h$, the regime in force when enrollment is measured. As the horizon lengthens, the differences become more astute, as a cohort released before expansion can have its outcome measured after it. Both versions are reported in Section 6.2; long-horizon interpretation uses the outcome-year clock, and the release-year clock is reported beside it as a separate question. Rebuilding the indicator from the expansion panel's own effective dates reproduces the merged variable in all 561 state-year rows.

For the eligibility-threshold specifications, I build a state-year panel of adult Medicaid income-eligibility limits from KFF State Health Facts covering 2000 through 2018 (KFF, 2026). It carries the limit for working parents in a family of three, which is nearly complete over the estimation window, and the limit for non-disabled adults without dependent children, both as a percent of the federal poverty level. I also include models with state-level controls of the unemployment rate, the age-adjusted drug overdose mortality rate, and the poverty rate (BLS LAUS, CDC WONDER, Census SAIPE), entering only in robustness specifications.

4.7 Summary Statistics

Mean one-year post-release Medicaid enrollment is 32.0 percent in the 1-year sample, with wide variance driven by termination-era observations below 15 percent and post-adoption observations above 50 percent. Adopting states have higher pre-treatment enrollment than the states that never switch inside the panel (21.0 versus 14.4 percent). Table 2 reports the summary statistics.

Table 2. Summary Statistics, Analysis Sample

	N	Mean	SD	Min	Max
Medicaid enrollment, 1-year	212	0.320	0.251	0.078	0.880
Medicaid enrollment, 3-year	207	0.287	0.187	0.091	0.800
Medicaid enrollment, 5-year	201	0.277	0.166	0.093	0.740
W-2 employment, 1-year	212	0.536	0.074	0.350	0.680
W-2 earnings, 1-year (\$)	212	6,909	1,634	3,060	11,600
Mortality, 1-year	264	0.017	0.005	0.006	0.048
HUD housing assistance, 1-year	264	0.011	0.005	0.002	0.031
SSI receipt, 1-year	99	0.043	0.010	0.021	0.060
Post-suspension indicator	561	0.296	0.457	0	1

Notes: The unit of observation is the state by release-cohort cell. Medicaid, employment, HUD, SSI, and mortality outcomes are shares in [0, 1], and earnings are nominal dollars averaged across release-cohort members, including members with zero earnings. Medicaid outcomes are restricted to the 17-state outcome panel and the horizon-specific cohort windows in Section 4.3, mortality and HUD cover release cohorts 2000-2018, and the post-suspension indicator covers all 31 CJARS states and release cohorts 2000-2020. Author's calculations from the CJARS Justice Outcomes Explorer and the state suspension policy panel.

5. Methods

5.1 Identification Strategy

I exploit the staggered adoption of Medicaid suspension policies across states between 2005 and 2022 to estimate the quasi-experimental association between suspension and post-release outcomes. The identifying variation comes from within-state, over-time changes in policy status, separated from time-invariant state characteristics (state fixed effects) and common temporal shocks (year fixed effects). State-years after suspension adoption are classified as “treated,” while state-years under termination are considered as “untreated.” I implement two complementary strategies. The first is a staggered DiD design estimating the average treatment effect of suspension adoption. The second is a heterogeneous DiD testing whether effects are amplified in ACA expansion states. I perform standard parallel trends checks to bolster required assumptions and any causal claims.

5.2 Staggered Difference-in-Differences Estimators

The benchmark two-way fixed effects (TWFE) specification is

$$Y_{st} = \beta \text{Suspension}_{st} + \alpha_s + \lambda_t + X'_{st}\gamma + \varepsilon_{st}$$

where Y_{st} = the outcome for state s in release-cohort year t , Suspension_{st} = a binary indicator equal to one in state-years on or after the state’s suspension adoption year, α_s and λ_t = state and release-cohort-year fixed effects, X_{st} = a vector of time-varying state controls entering only in robustness specifications, ε_{st} = an error term, and β = the parameter of interest, the average change in the cohort-level outcome associated with being under a suspension regime at release. Standard errors are clustered at the state level (Bertrand, Duflo, and Mullainathan, 2004). The expansion-adjusted specifications add a state-specific time-varying post-expansion indicator, $\text{PostExpansion}_{st}$, and the eligibility-threshold specifications add the state’s Medicaid income limit for parents, ParentLimit_{st} , measured in hundreds of percentage points of the federal poverty level. As a single additive indicator cannot fully isolate suspension from an expansion that phased in over time or varied in strength across states, I interpret coefficients as “expansion-adjusted” rather than as net of expansion.

For robustness, I use Goodman-Bacon decomposition (Appendix Figure A9) to check for negative weighting issues, and also utilize the heterogeneity-robust estimators of local-projections difference-in-differences (LP-DiD, Dube et al., 2025) and Callaway and Sant’Anna (2021). I choose LP-DiD for primary robustness as it best accounts for the fact that only Texas contributes Medicaid outcome rows among the never-treated states.

The choice of LP-DiD post window drives the estimator’s standard errors. The default window is the maximum relative year present in the data, which in this panel leaves 4 observations in the one-year pooled regression. I therefore set the post window to five years, matching the event-study window, which puts the same estimate on 73 observations. Table 4 reports both windows, and the appendix gives the window mechanics in full.

5.3 Event Study Analyses

I estimate event study models to assess parallel trends and trace dynamic effects:

$$Y_{st} = \sum_{k=-5, k \neq -1}^5 \beta_k D_{st}^k + \alpha_s + \lambda_t + \varepsilon_{st}$$

where D_{st}^k = an indicator equal to one when release-cohort year t sits k years from state s ’s suspension adoption year g_s , so that $k = t - g_s$, $k = -1$ = the omitted reference period, β_k for $k < 0$ = the pre-treatment leads that test parallel trends individually and jointly, and β_k for $k \geq 0$ = the post-adoption coefficients. The main event study is the expansion-adjusted LP-DiD path, plotted alongside the unadjusted path as the joint-policy-era comparison. Each figure carries a support panel reporting the treated and clean control states at every event time, and I do not show lags beyond five years.

5.4 Inference

While CJARS has data on over 30 states, the relevant panel is limited to 17 clusters, a small comparison group, and common co-timing with Medicaid expansion. To account for these issues, I utilize the wild

cluster bootstrap for TWFE estimates and studentized (bootstrap-t) interval for the two-step expansion-adjusted LP-DiD estimate. I rely on bootstrap rather than analytic standard errors because conventional cluster-robust standard errors understate uncertainty in two distinct ways in my analysis:

First, cluster-robust inference is justified by asymptotics in the number of clusters, and with 17 clusters the variance estimator is biased downward and cluster-t tests reject true nulls well above their nominal size (Cameron, Gelbach, and Miller, 2008). The wild cluster bootstrap with the null imposed addresses this by building the reference distribution for each test statistic from the 17-cluster panel itself rather than from an asymptotic approximation the panel is too small to support.

Second, a conventional standard error for the two-step expansion-adjusted estimate is computed from the second step alone, and therefore treats the estimated first-stage expansion coefficient as known, so none of the first stage's uncertainty reaches the interval, creating an artificially small standard error. I account for this by using the pairs bootstrap, which re-estimates the entire two-step procedure inside every draw, so that first-stage noise propagates into the final interval, and studentizing keeps that interval from depending on any single sample's own noisy standard-error estimate.

Lastly, I also utilize a permutation placebo that reassigns adoption years at random across states; this allows me to observe how often policy timing that is random by construction produces an estimate as large as the one I observe in my analysis.

5.5 Heterogeneity of Treatment Effect Analyses

I interact the suspension indicator with ACA expansion status:

$$Y_{st} = \beta_1 \text{Suspension}_{st} + \beta_2 (\text{Suspension}_{st} \times \text{Expanded}_s) + \alpha_s + \lambda_t + \varepsilon_{st}$$

where Expanded_s = an indicator for states that ever adopted the ACA Medicaid expansion, whose main effect is absorbed by the state fixed effects, β_1 = the suspension estimate in states that never expanded, $\beta_1 + \beta_2$ = the suspension estimate in expansion states, and β_2 = the differential. I supplement this with a


period-split approach that estimates the TWFE model separately for the pre-ACA period (release cohorts 2004-2013 for Medicaid, employment, and earnings) and the post-ACA period (2014-2018), and with split samples among ever-expanded and non-expansion states.

6. Results

6.1 Raw trends

States that adopted suspension saw a 22.2 percentage point higher one-year post-release Medicaid enrollment rate than states still terminating (wild cluster bootstrap $p = 0.001$). Figure 3 plots this increase over time; 3- and 5-year trends are similar and located in the Appendix.

Figure 3. Raw Medicaid enrollment trends by adoption cohort, 1-year horizon



Notes: Mean one-year post-release Medicaid enrollment rate by calendar release-cohort year, with states grouped by the year they adopted suspension. Never-treated states are plotted separately. Group composition is fixed by ultimate adoption cohort, so a change between years is a change in outcomes rather than a change in which states are averaged. The largest upward movements cluster in the 2014-16 window, when most states expanded Medicaid and adopted suspension at the same time. Author's calculations from the CJARS Justice Outcomes Explorer and the state suspension policy panel.

Since suspension policies only preserve Medicaid enrollment people have prior to incarceration, the effect of these policies depend on pre-incarceration Medicaid enrollment rates, which rely on states' eligibility policies. Thus state cohorts who were 'treated' with suspension policies prior to ACA expansion had little coverage to suspend.

6.2 Main results

Table 3 reports the primary estimates: the change in the share of a state's release cohort enrolled in Medicaid one year after release associated with the state having switched from termination to suspension, in levels and net of the co-timed ACA expansion. Table 4 reports the same estimates for 3- and 5-year post-release.

Table 3. Main Results: Medicaid Suspension and 1-Year Post-Release Medicaid Enrollment

Specification	Estimate	SE	Primary p	N entering
<i>Panel A. Joint policy (expansion + suspension)</i>				
(1) TWFE	0.222***	(0.051)	0.001	212
(2) LP-DiD	0.207**	(0.068)	0.031	73
<i>Panel B. Expansion adjusted</i>				
(3) TWFE	0.111	(0.069)	0.301	212
(4) TWFE + parent limit	0.115	(0.068)	0.112	212
(5) LP-DiD	0.137	(0.043)	0.190	73
(6) Stacked DiD	0.127	(0.054)	0.122	577
(7) Callaway-Sant'Anna	0.147	(0.056)	0.293	212
(8) BJS	0.136*	(0.045)	0.085	212

Notes: Outcome is the share of a state's release cohort enrolled in Medicaid one year after release, release cohorts 2004-2018 in 17 states. Every estimate is an unweighted state-average intent-to-treat association with the de jure policy at release (Section 4.5). Coefficients are shares, so the 0.222 in row (1) is 22.2 percentage points. Every specification includes state and release-cohort-year fixed effects, is estimated on a frame of 212 state-cohort observations, and carries state-clustered standard errors in parentheses; the final column reports how many of those observations enter the estimate. Rows (2) and (5) use the LP-DiD estimator of Dube et al. (2025), row (5) first residualizing the outcome on the jointly estimated post-expansion coefficient (Section 5.2). The primary p-value is the wild cluster bootstrap for rows (1) and (3), the permutation placebo for row (2), the analytic cluster-robust p for row (4), and the studentized bootstrap-t for row (5), per Section 5.4 and Table 4. The permutation placebo reassigns adoption years at random across states; 761 of 1,000 reassignments leave a clean comparison and enter the p-value. Analytic p-values for rows (1) through (3) are 0.000, 0.010, and 0.126, and the permutation placebo for row (1) is 0.062 (Table 7). *p<0.10, **p<0.05, ***p<0.01.

Suspension states see a raw Medicaid enrollment increase of 20-22 percentage points one year post release; however, controlling for expansion, which was often co-timed with the policy, reduces the estimate in half to 11.1 points and is not statistically significant. While alternative estimators such as stacked DID and BJS are marginally significant (10.7%, $p = .065$, and 13.6%, $p = .085$) and almost all estimators produce point estimates of approximately 11 percentage points (expansion-controlled), the majority are not statistically distinguishable from zero and have wide confidence intervals, due to the low amount of between-state variation (Appendix Table A12).

A two-step LP-DiD produces an almost identical 11.1 points at the estimator's default window and 13.7 points at the fixed five-year window. Substituting the actual time-varying income-eligibility environment for the binary indicator does not move it either: adding the parent income limit gives 11.5 points, with the limit's own coefficient indistinguishable from zero (Appendix Table A6). The adjusted magnitude is consistent with the closest individual-level evidence, since Packham and Slusky (2024) report an 11.1 point enrollment gain within six months of release under South Carolina's combined suspension and simplified re-enrollment policy, though the estimands differ (Section 4.5).

Table 4 reports the estimates at all three horizons under both expansion clocks. At one year the expansion-adjusted estimate is unchanged at 11.1 points, and at three years it falls from 14.6 to 10.0 points when the indicator moves to the outcome-measurement year. At five years only 4 of 201 rows are coded post-expansion on the release-year clock, so long-horizon interpretation uses the outcome-year clock, which moves the estimate from 19.9 to 10.7 points (Section 4.6).

Table 4. Medicaid Enrollment by Post-Release Horizon, Estimator, and Expansion Adjustment

Estimator	1-year, at t	3-year, at t	5-year, at t	1-year, at t+h	3-year, at t+h	5-year, at t+h
TWFE, no expansion adjustment	0.222*** (0.051)	0.194*** (0.044)	0.207*** (0.058)	-	-	-
TWFE	0.111 (0.069)	0.146*** (0.048)	0.199*** (0.057)	0.111 (0.071)	0.100** (0.039)	0.107** (0.043)
Stacked DiD	0.127** (0.054)	0.163*** (0.049)	0.176** (0.063)	0.128** (0.053)	0.101*** (0.032)	0.091** (0.039)

LP-DiD	0.137 (0.057)	0.191 (0.083)	0.141 (0.089)	0.113 (0.060)	0.132 (0.065)	0.041 (0.063)
Callaway-Sant'Anna	0.147 (0.056)	0.188 (0.063)	0.123 (0.061)	0.125 (0.053)	0.112 (0.054)	0.031 (0.041)
BJS	0.136 (0.045)	0.176 (0.072)	n.e.	0.134 (0.044)	0.116 (0.052)	n.e.

Notes: The outcome is the share of a state's release cohort enrolled in Medicaid at the stated horizon, and every estimate is an unweighted state-average de jure policy-regime intent-to-treat association on the share scale. Each cell is a separate estimate of the horizon-specific effect; the expansion indicator is indexed either to the release-cohort year t or to the outcome-measurement year $t+h$, with the outcome-year clock used for long-horizon interpretation (Section 4.6); the raw TWFE row contains no expansion indicator, so the clock does not apply. The TWFE and stacked rows carry analytic cluster-robust standard errors and their stars follow analytic p-values, per the paper's convention; wild-cluster-bootstrap p-values for every one of those cells are in Appendix Table A13. The three two-step rows (LP-DiD at the fixed five-year post window, Callaway-Sant'Anna, and BJS imputation) residualize the outcome on the jointly estimated post-expansion coefficient; they carry bootstrap standard errors from the stratified state-cluster resampling that re-estimates both steps inside every draw and no significance markers, with their studentized bootstrap-t p-values in Appendix Table A13 and, for the one-year estimate, Table 4. "n.e." means not estimable: the imputation estimator's effective-sample guard suppresses the five-year estimate. TWFE and two-step samples are $N = 212 / 205 / 200$ (17, 17, and 15 states) at the one-, three-, and five-year horizons; the stacked frames have $577 / 571 / 476$ stacked observations on $15 / 15 / 14$ distinct states. Joint-policy-era LP-DiD and Callaway-Sant'Anna estimates, and the common-cohort-window diagnostic, are in Appendix Table A10. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Unadjusted estimates remain positive at three and five years, and a stacked test of the horizon interaction cannot reject equality between the one-year and five-year estimates (coefficient -1.7 points, $p = 0.821$).

Because the five-year sample drops the most recent release cohorts, horizon comparisons could confound persistence with cohort composition. Re-estimating the one-year model on exactly the cells where the five-year outcome is observed gives 16.2 points (Appendix Table A8).

Five-year estimates are significantly weaker. The two-step LP-DiD estimate is 19.0 points at the default window and 14.1 points at the fixed window, its interval covers zero under every interval construction, and it rests on five switching states. In short, the joint-policy-era association shows no clear attenuation at three and five years, while suspension-specific persistence is not identified. JOE also measures cohort coverage rates rather than tracking individuals, so these comparisons are of enrollment aggregates and cannot say whether the same people hold coverage throughout or cycle in and out. All horizon-sensitive estimations also find null results.

6.3 Secondary outcomes: Employment, earnings, and mortality

I do not detect reductions in recidivism due to suspension policies and increased Medicaid coverage.

However, as a number of prior studies find positive effects and data format limits my ability to detect effects under 4-8%, I interpret these findings as an absence of evidence versus definitive proof. I also see

no discernable effects on employment, earnings, and mortality, which I interpret similarly to recidivism as non-determinative.

Table 5 reports the full set of horizons, with all four core outcomes under both estimators. The Medicaid contrast holds at every horizon, while employment, earnings, and mortality move at none of them, and the nominally significant three-year LP-DiD employment and earnings cells rest on small pooled samples (Section 6.6).

Table 5. Horizon Sensitivity, Expansion-Adjusted Estimates, Years 1, 3, and 5 Post-Release

	Medicaid	Employment	Earnings (\$)	Mortality
1-year horizon				
TWFE	0.111 (0.069)	0.015 (0.009)	204.6 (328.1)	0.0011 (0.0008)
Stacked DiD	0.127** (0.054)	0.018** (0.006)	330.5 (293.2)	0.0009 (0.0008)
Callaway-Sant'Anna	0.147 (0.056)	0.024 (0.017)	661.8 (407.0)	0.0008 (0.0009)
BJS	0.136 (0.045)	0.016 (0.006)	331.9 (166.5)	0.0006 (0.0006)
Observations	212	212	212	264
3-year horizon				
TWFE	0.146*** (0.048)	0.005 (0.009)	313.7 (324.7)	0.0015 (0.0013)
Stacked DiD	0.163*** (0.049)	0.006 (0.010)	402.0 (291.6)	0.0015 (0.0011)
Callaway-Sant'Anna	0.188 (0.063)	0.004 (0.028)	-62.2 (576.4)	0.0020 (0.0014)
BJS imputation	0.176 (0.072)	0.002 (0.011)	237.8 (318.9)	0.0012 (0.0009)
Observations	207	207	207	233
5-year horizon				
TWFE	0.199*** (0.057)	-0.005 (0.016)	98.7 (360.0)	0.0011 (0.0013)
Stacked DiD	0.176** (0.063)	-0.003 (0.014)	127.4 (318.5)	0.0016 (0.0014)
Callaway-Sant'Anna	0.123 (0.061)	-0.015 (0.012)	-50.1 (247.5)	0.0013 (0.0014)
Observations	201	201	201	201


Notes: Every estimate is expansion-adjusted. The TWFE and stacked rows estimate the post-expansion indicator jointly (release-year clock, matching Table 4); the Callaway-Sant'Anna and BJS imputation rows residualize the outcome on the jointly estimated post-expansion coefficient and re-estimate both steps inside every draw of the stratified state-cluster bootstrap. Each cell is a separate estimate on the outcome's censoring-consistent cohort window in the 17-state outcome panel (Section 4.3); the five-year window is identified by 15 states. The TWFE and stacked

rows carry analytic cluster-robust standard errors clustered at the state level, with stars from analytic p-values per the paper's convention; the two-step rows carry bootstrap standard errors and no significance markers. Design-based p-values for every cell are in Appendix Table A13: wild cluster bootstrap for the TWFE and stacked rows, studentized bootstrap-t for the two-step rows. "n.e." means the imputation estimator's effective-sample guard suppresses the five-year estimate. The TWFE and two-step rows share the samples in the Observations row; the stacked frames have 577 / 571 / 476 stacked observations on 15 / 15 / 14 distinct states (643 / 593 / 476 for mortality). The stacked expansion-adjusted employment estimate at one year is significant under its wild cluster bootstrap ($p = 0.016$) and the imputation version's bootstrap-t agrees ($p = 0.030$), but the TWFE-adjusted version does not reject (wild cluster bootstrap $p = 0.153$) and the Callaway-Sant'Anna version is far from rejection ($p = 0.498$), so I treat that cell as suggestive rather than as a finding (Section 6.6). The unadjusted TWFE and LP-DiD estimates are in Appendix Table A14. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

6.4 Heterogeneous effects by Medicaid expansion

Figure 4 depicts the expansion-adjusted LP-DiD event study. The post-adoption path rises from -1.6 points at adoption to 24.0 points at five years, and the average post-adoption coefficient is 9.8 points (95 percent interval -1.3 to 20.8), consistent with the pooled two-step estimate of 11.1 points and, like it, not distinguishable from zero. Support falls from 11 treated states at event time -5 to 4 at event time +5, with Texas the only never-treated state throughout (Appendix Table A6).

Figure 4. Expansion-adjusted LP-DiD event study, Medicaid enrollment, 1-year horizon, with event-time support



Notes: The pre-adoption leads are identical with and without the adjustment; no treated state changed expansion status inside its own pre-adoption event window. The upper panel plots event-time coefficients with 95 percent bands from the stratified state-cluster bootstrap in Section 5.4, and the lower panel reports the treated states and clean control states available at each event time. Event time -1 is the omitted reference, and coefficients are long differences against the never-treated and not-yet-treated states. The dotted line in the lower panel marks the three-state

support rule, and exactly one never-treated state contributes at every event time. Author's calculations from the CJARS Justice Outcomes Explorer and the state policy panel.

Pre-period leads are concerning; they are identical with and without the adjustment, at -4.2, -6.1, -8.3, and -8.7 points, since no treated state changed expansion status inside its own pre-adoption window. All four are negative and grow in magnitude as adoption approaches, and the leads at event times -3 and -2 individually exclude zero. A joint Wald test rejects ($\chi^2(4) = 26.6$, $p = 0.00002$, and $\chi^2(4) = 13.8$, $p = 0.008$ under an unstratified covariance). Negative distant leads mean treated states sat below their comparators several years before adoption and were converging upward. If that convergence continued, it would bias the estimate away from zero. I follow Rambachan and Roth (2023) for my sensitivity analysis and find minimal absorption of parallel trends violations.

A state contributes to a within-state treatment coefficient only if its policy switch falls inside the estimation window. In the one-year Medicaid frame, 10 of the 17 states have an observed suspension switch and 7 of those also have an observed expansion switch, so the expansion adjustment is identified by 7 states. In four of the seven (Arizona, Indiana, Montana, Pennsylvania) the two policies are within a year of each other.

Table 6 reports expansion-status heterogeneity. The static interaction cannot separate the two policies either: for Medicaid enrollment the suspension main effect and the interaction are individually indistinguishable from zero once both enter, and the period-split Medicaid estimates stay positive on both sides of 2014.

Table 6. Heterogeneous DiD by ACA Medicaid Expansion Status, 1-Year


	Medicaid	Employment	Earnings	Mortality
Panel A. Static-expansion interaction				
Suspension	0.075 (0.130)	0.024* (0.012)	478.8 (641.6)	0.0018* (0.0010)
Suspension × Expanded	0.195 (0.132)	-0.010 (0.017)	-215.6 (660.6)	-0.0016 (0.0016)
Panel B. Period-split TWFE				

Suspension × Pre-ACA	0.117** (0.050)	0.007 (0.009)	-2.6 (300.1)	-0.0000 (0.0007)
Suspension × Post-ACA	0.147* (0.081)	0.008 (0.008)	-157.3 (300.0)	-0.0013 (0.0035)
<i>Specification details</i>				
State and year fixed effects	Yes	Yes	Yes	Yes
Observations (static)	212	212	212	264
Observations (pre-ACA)	134	134	134	186
Observations (post-ACA)	78	78	78	78

Notes: Cluster-robust standard errors clustered at the state level are in parentheses, and all three panels are estimated on the 17-state outcome panel (Section 4.4). Panel A reports the interaction with time-invariant ACA expansion status as the cross-classifying dimension; the main effect of expansion status is absorbed by the state fixed effects. Panel B reports TWFE coefficients on the suspension indicator estimated separately in the pre- and post-expansion periods. The pre-period is release cohorts 2004-2013 for Medicaid, employment, and earnings and 2000-2013 for mortality; the post-period is 2014-2018 for all four outcomes. Panel C reports the outcome of an attempted time-varying expansion interaction. In the Medicaid sample, every state that adopted both suspension and the expansion did so in an order that makes the interaction collinear with the expansion indicator net of fixed effects, so the coefficient is not separately identified (Section 6.4). *p<0.10, **p<0.05, ***p<0.01.

Figure 5 collects the visual evidence behind this reading. The joint-policy-era LP-DiD paths at the three- and five-year horizons in Panels A and B climb steadily after adoption, while the one-year TWFE path with the time-varying expansion control in Panel C sits near zero. Panel D shows the co-timing directly, with most adopting states' suspension and expansion markers on the same year.

Figure 5. Event studies at the 3- and 5-year horizons, the expansion-controlled event study, and policy co-timing



Notes: Panel D shows each state in the 31-state policy panel as a row, with a marker at its suspension adoption year, a marker at its ACA expansion year, and a connecting segment where the state adopted both policies. Asterisks in Panel D mark the 17 states in the Medicaid outcome panel, a missing marker means the state never adopted or never expanded, and the dashed line marks January 2014. Author's calculations from the CJARS Justice Outcomes Explorer, the state suspension policy panel, and the ACA expansion panel.

6.5 Demographic Subgroup Estimates

Suspension policy related gains are consistent across race, sex and age. Every subgroup with full panel coverage gained roughly 16 to 24 points at one year, with the same pattern at three and five years, and the intervals overlap heavily.

6.6 Robustness

Table 7 reports other robustness checks, which I conduct in three ways: first, I use a number of different estimators to ensure my estimations are consistent across each of them. Second, I check inference by conducting permutation placebos on year of treatment. Third, I conduct leave-one-out tests and drop

certain states to ensure certain outlier states are not driving my findings. Results are generally robust across all estimations and do not significantly vary.

Table 7. Robustness Summary, Medicaid Enrollment, 1-Year Horizon

	Coefficient	SE or interval	N / States	Specification
Panel A. Estimator robustness				
TWFE + time-varying controls	0.230***	(0.041)	198 / 17	Unemployment, overdose, poverty
Cohort-size-weighted TWFE	0.257***	(0.062)	212 / 17	Weighted by release-cohort size
Panel B. Inference robustness				
Permutation placebo p, TWFE	-	0.062	500 permutations	Design-based
Permutation placebo p, LP-DiD	-	0.081	761 admissible of 1,000	Design-based
Panel C. Sample and subsample robustness				
Leave-one-out (range)	[0.206, 0.259]	0.036	-	Across 17 states
Drop Texas (never-treated)	0.209***	(0.052)	198 / 16	Modest attenuation
Drop Kansas (time-limited suspension)	0.225***	(0.056)	197 / 16	Main result unchanged

Notes: Column 1 reports the coefficient on the post-suspension indicator for one-year Medicaid enrollment; column 2 reports the cluster-robust standard error in parentheses, or a p-value or interval where relevant; column 3 reports sample size and cluster count; column 4 describes the specification. All TWFE specifications include state and release-cohort-year fixed effects with standard errors clustered at the state level. Georgia and South Carolina contribute no Medicaid outcome rows and therefore cannot support a leave-one-out check; Texas is the only never-treated state in the Medicaid panel, and dropping it removes its 14 rows.

7. Discussion

7.1 Summary of Findings

Suspension states saw large gains in post-release Medicaid enrollment: roughly 22 percentage points at one year on the raw comparison, and roughly 11 percentage points net of the co-timed ACA expansion, though the adjusted figure is positive but not statistically distinguishable from zero (bootstrap-t interval - 11.0 to +19.2 pp, $p = 0.176$). On the outcome-measurement clock the adjusted estimates are stable at three and five years post-release (10.0 and 10.7 points), and the raw contrast is similar at every horizon. As the

adjusted point estimate is positive at every horizon and under every construction, I argue and read the adjusted results as consistent with a substantial suspension effect that this design cannot establish on its own. The magnitude of coverage gains are similar at three and five years post-release as well. Subgroup estimates are broadly similar across sex, race, and age, with a smaller estimate for women, consistent with their much higher baseline post-release enrollment. I am unable to detect any change in employment, earnings, same-state reincarceration, and mortality.

7.2 Interpretation

The results suggest that administrative policy pathways may matter as much as eligibility policies for this population. An enrollment increase of roughly 11 pp net of expansion, and 22 pp raw, is among the largest effects documented in the administrative burden literature, even at the adjusted end of the range. This is consistent with the prediction that friction reduction produces the largest effects where barriers are greatest. Suspension has the potential to remove all three burden types at once (Section 2.4), and the completeness of that reduction, combined with the vulnerability of the target population, likely explains the large magnitude. However, this interpretation is noted with some degree of caution, since expansion and suspension are co-timed and the design cannot separate them within the co-timed states.

The persistence of the effect at the 3- and 5-year horizons requires extending the standard framework. I frame this persistent effect not as a single barrier removal, but as placement onto a self-sustaining coverage trajectory versus exclusion from one.

The framework also helps explain why the effect is so large relative to other coverage interventions. The ACA expansion changed eligibility; suspension changes access to existing eligibility. For a population where the barriers to completing an application are very substantial (Section 2.4), eliminating the application requirement can produce a coverage gain that is close to half the magnitude of the eligibility expansion itself.

The results also speak to the broader debate about whether administrative burdens are rational disincentives that filter out people who do not value the benefit enough, or arbitrary barriers that exclude people who are fully eligible. An enrollment gain of 11 to 22 pp is more consistent with administrative

barriers deterring eligible participation than with eligibility differences alone, though the aggregate design cannot identify individual eligibility or motives. That reading aligns with Herd and Moynihan's (2025) argument that administrative burdens in safety-net programs function mostly as exclusionary mechanisms rather than efficient screening devices.

7.3 Coverage Is Necessary But Not Sufficient

A central interpretive point from this paper is that strong evidence on coverage take-up does not, on its own, settle whether coverage translates into better health and economic outcomes. The enrollment effect is among the largest documented in the administrative burden literature, yet employment, earnings, reincarceration, and mortality all move imperceptibly.

Three observations fit this pattern. First, the post-release period is dominated by acute mortality risk in which the speed of access matters more than the fact of coverage. Prevention of overdose deaths in the first weeks post-release requires buprenorphine or methadone initiation in those weeks. Insurance reactivation that arrives even a few weeks late, or that arrives without a coordinated handoff to a treating clinician, may not avert the death. Bandara et al. (2025) document that even where suspension is the de jure policy, de facto reactivation is often delayed and rarely accompanied by warm-handoff care coordination, and that fewer than half of states with suspension policies have data-sharing infrastructure adequate to notify Medicaid agencies of an individual's release in real time. In simple terms, keeping Medicaid coverage during and directly after release may not matter if the beneficiary does not even know they are covered or has sufficient trust to utilize the Medicaid system.

Second, the broader Medicaid-mortality literature (Sommers et al., 2012; Borgschulte and Vogler, 2020) typically estimates aggregate relative mortality reductions of 3-6 percent following the ACA expansion among general adult populations over horizons of several years. The justice-involved population at one to five years post-release faces both higher baseline mortality (1.7 percent vs. ~0.5 percent among general adult populations) and a different mix of acute risks. Unfortunately, the study design's minimum detectable effect is too coarse to settle the question (Section 6.4).

Third, the null employment finding has a similar issue: insurance coverage may be necessary to prevent health crises that disrupt work, but it is not sufficient to overcome the labor-market barriers facing people who are returning to society after incarceration, such as employer discrimination, occupational licensing restrictions, parole compliance demands, and skills atrophy. These barriers are not addressed by health insurance and operate through entirely separate channels that this paper does not study.

Suspension substantially increased coverage but did not produce detectable changes in formal employment, earnings, same-state reincarceration, or mortality in this state-cohort design. The pattern suggests that eligibility continuity is a foundation for reentry care, not a substitute for treatment access, provider capacity, or broader reintegration support. Policy and research should treat suspension as a foundational layer on which additional interventions can build - pre-release treatment initiation, care-coordination handoffs, naloxone distribution, transitional employment - rather than as a stand-alone intervention sufficient to close the post-release health and economic gap.

7.4 Policy Implications

The findings support the policy rationale behind the 2026 federal requirement: preserving Medicaid eligibility across the incarceration-to-release transition produces large enrollment gains and is plausibly a foundational input to any reentry-coverage strategy. With the mandate in effect, the central implementation question is no longer whether states formally suspend rather than terminate, but whether they operationalize suspension in ways that produce timely reactivation, successful renewals during incarceration, and connection to care at release. For the states that were still terminating as of late 2025 and for any state with incomplete implementation, these estimates provide a historical benchmark for the enrollment consequences of preserving eligibility: roughly 11 to 22 percentage points at one year, where the lower end nets out the co-timed expansion and is the more defensible forecast for a mandate implemented where expansion status is already fixed. The benchmark should not be extrapolated mechanically, because populations and implementation quality differ. The findings may be most useful for performance monitoring: first-year coverage materially below it would point to operational failure

rather than policy failure. A coverage gain of this size would be expected to raise demand for behavioral health and medication for opioid use disorder services within Medicaid managed care.

The CMS Reentry Section 1115 Demonstration opportunity, introduced in 2023, allows states to cover services for incarcerated individuals in the 30 to 90 days before release; as of mid-2025, 18 states had approval and 9 applications were pending (KFF, 2025; CMS, 2023). Suspension and the waivers are complementary rather than substitutes: suspension preserves eligibility across incarceration, while the waivers expand covered services before release and support the treatment-initiation and care-coordination handoff. States implementing both should in principle see synergistic effects, though the interaction has not yet been evaluated. Three extensions follow naturally from this paper: individual-level CJARS analysis through the federal statistical research data center network, which would replace the cohort-level intent-to-treat estimand with a treatment-on-treated estimand and add power for the mortality and employment outcomes; evaluation of implementation quality as a natural experiment within the mandate era; and identification of the waiver-suspension interaction from the staggered waiver rollout under the now-uniform suspension regime.

7.5 Limitations

7.5.1 ACA co-timing

The co-timing with the ACA expansion is the most important qualification on the magnitude of the estimates. Within the co-timed states, no design can fully separate the two policies. The disaggregation in Section 6.4 identifies the suspension-specific component from the pre-2014 adopters and the never-expanded adopters and implies an expansion-adjusted effect of roughly 11 pp at one year. I report the raw and adjusted figures side by side wherever the estimate is used prescriptively. The preferred event study also carries a systematic negative pre-trend. In the Rambachan and Roth (2023) sensitivity analysis (Section 6.4 and Appendix Figure A10), the joint-policy-era estimate stays distinguishable from zero only while post-adoption violations remain below about a fifth of the largest observed pre-period movement, and the expansion-adjusted robust confidence set contains zero at every bound.

7.5.2 Sample composition

CJARS covers 31 of 51 states, coverage is determined by data-sharing agreements with the Census Bureau, and the outcome regressions are identified by 17 of those. Missing states may differ systematically; if they are disproportionately states with weaker administrative infrastructure, the estimates may not generalize. The sample does include 27 of 29 suspension adopters and 2 of 3 termination states, which gives reasonable coverage of both regimes. Even more constraining is that only two never-treated states appear in the CJARS sample and only Texas contributes Medicaid outcome rows. LP-DiD's not-yet-treated comparators mitigate but do not remove this.

7.5.3 Aggregate estimand and de jure treatment

I use aggregate outcomes rather than individual-level transitions, so my estimates could partly reflect selection rather than the policy effect. Testing pre-determined release-cohort characteristics as placebo outcomes and re-estimating with composition controls attenuates but does not overturn the finding; these checks mitigate the concern without eliminating unobservable compositional change. The design also limits heterogeneity analysis and reduces power, which may explain the inability to detect employment and recidivism effects. Separately, the policy panel captures de jure adoption, while Bandara et al. (2025) document substantial gaps between policy on paper and operational practice. I am unable to account for managed-care penetration, which varies across states and over time and often controls the operational steps after reactivation; state fixed effects absorb time-invariant differences but not differential changes coinciding with adoption.

7.5.4 Observation window and measurement

Medicaid enrollment in CJARS is measurable only in calendar years 2005-2019, tied to the MSIS/T-MSIS linkage window, which prevents observing effects for the most recent adopters and limits sample size; SSI data are available for only a subset of years. I cannot assess impacts beyond five years post-release. W-2 employment captures only formal-sector work reported to the IRS, so informal and off-the-books work is classified as non-employment. Mortality is measured through the Census Numident, which is high-quality but may miss some deaths. JOE recidivism captures same-state contact only, so measured

rates are lower bounds and the nulls are informative about the same-state margin specifically; released individuals could relocate toward suspension states, but release state is determined by facility location rather than by choice, so I expect little bias from that channel.

8. Conclusion

This paper provides the first multi-state evidence on the relationship between Medicaid suspension during incarceration and post-release coverage trajectories. Linking a newly constructed 31-state suspension policy panel to administratively linked CJARS release-cohort outcomes for 17 reporting states, I find a large, persistent intent-to-treat coverage effect: roughly 11 pp 1-year post release net of the co-timed ACA Medicaid expansion, and 22 pp raw. These increases are also seen at three and five years post release, and effects are broadly similar across sex, race, and age. Effects on W-2 employment and earnings are null, same-state reincarceration effects are a precise-enough null that neither confirms nor refutes the modest reductions in prior individual-level work, and mortality estimates are inconclusive and underpowered. These findings arrive at a transitional policy moment. The 2026 federal mandate has since made suspension universal (Section 2.3). They also support a more cautious second message: a strong coverage effect is not, on its own, evidence that the post-release acute-mortality and labor-market gaps will close. More broadly, the results indicate that reducing administrative burdens at critical life transitions, when individuals face the greatest barriers to accessing safety-net programs, can produce substantial improvements in program take-up. The suspension-termination policy distinction has been an important administrative determinant of whether justice-involved individuals returning to society can maintain health coverage through the most vulnerable period of their lives. The lesson extends beyond Medicaid and beyond the justice-involved population: at moments of acute life transition, the friction of a government application can be the difference between coverage and no coverage.

References

Albertson, E. M., Scannell, C., Ashtari, N., & Barnert, E. (2020). Eliminating gaps in Medicaid coverage during reentry after incarceration. *American Journal of Public Health*, 110(3), 317-321.

Anderson, M. L. (2008). Multiple inference and gender differences in the effects of early intervention: A reevaluation of the Abecedarian, Perry Preschool, and Early Training Projects. *Journal of the American Statistical Association*, 103(484), 1481-1495.

Aslim, E. G., Mungan, M. C., Navarro, C. I., & Yu, H. (2022). The effect of public health insurance on criminal recidivism. *Journal of Policy Analysis and Management*, 41(1), 45-91.

Aslim, E. G., Mungan, M. C., & Yu, H. (2024). A welfare analysis of Medicaid and recidivism. *Health Economics*, 33(11), 2463-2507.

Baicker, K., Taubman, S. L., Allen, H. L., Bernstein, M., Gruber, J. H., Newhouse, J. P., Schneider, E. C., Wright, B. J., Zaslavsky, A. M., Finkelstein, A. N., and the Oregon Health Study Group. (2013). The Oregon experiment - effects of Medicaid on clinical outcomes. *New England Journal of Medicine*, 368(18), 1713-1722.

Bandara, S., Saloner, B., Maniates, H., Song, M., & Krawczyk, N. (2025). Implementation of carceral Medicaid suspension and enrollment programs. *Health & Justice*, 13(1), 1.

Bertrand, M., Duflo, E., & Mullainathan, S. (2004). How much should we trust differences-in-differences estimates? *Quarterly Journal of Economics*, 119(1), 249-275.

Binswanger, I. A., Stern, M. F., Deyo, R. A., Heagerty, P. J., Cheadle, A., Elmore, J. G., & Koepsell, T. D. (2007). Release from prison - a high risk of death for former inmates. *New England Journal of Medicine*, 356(2), 157-165.

Borgschulte, M., & Vogler, J. (2020). Did the ACA Medicaid expansion save lives? *Journal of Health Economics*, 72, 102333.

Burns, M. E., Cook, S., Brown, L. M., Dague, L., Tyska, S., Hernandez Romero, K., McNamara, C., & Westergaard, R. P. (2022). Association between assistance with Medicaid enrollment and use of health care after incarceration among adults with a history of substance use. *JAMA Network Open*, 5(1), e2142688. doi:10.1001/jamanetworkopen.2021.42688

Callaway, B., & Sant'Anna, P. H. C. (2021). Difference-in-differences with multiple time periods. *Journal of Econometrics*, 225(2), 200-230.

Camhi, N., Mistak, D., & Wachino, V. (2020). Medicaid's evolving role in advancing the health of people involved in the justice system. Commonwealth Fund Issue Brief.

Carson, E. A., & Kluckow, R. (2023). Prisoners in 2022 - statistical tables. Bureau of Justice Statistics.

Centers for Medicare & Medicaid Services (CMS). (2023). HHS releases new guidance to encourage states to apply for new Medicaid Reentry Section 1115 demonstration opportunity. CMS press release, April 17, 2023.

Centers for Medicare & Medicaid Services (CMS). (2025). *Prohibition on termination of enrollment due to incarceration (Division G, Title I, Section 205, of the Consolidated Appropriations Act, 2024)*. CMCS Informational Bulletin, December 23, 2025.

Deza, M., Lu, T., Maclean, J. C., & Ortega, A. (2024). Losing Medicaid and crime. NBER Working Paper No. 32227.

Dube, A., Girardi, D., Jorda, O., & Taylor, A. M. (2023). A local projections approach to difference-in-differences event studies. NBER Working Paper No. 31184.

Finkelstein, A., Taubman, S., Wright, B., Bernstein, M., Gruber, J., Newhouse, J. P., Allen, H., Baicker, K., and the Oregon Health Study Group. (2012). The Oregon Health Insurance Experiment: Evidence from the first year. *Quarterly Journal of Economics*, 127(3), 1057-1106.

Finlay, K., Mueller-Smith, M., & Papp, J. (2022). The Criminal Justice Administrative Records System. *Scientific Data*, 9, 562.

Goldman, A. L., & Sommers, B. D. (2020). Among low-income adults enrolled in Medicaid, churning decreased after the ACA. *Health Affairs*, 39(1), 85-93.

Gollu, G., & Zapryanova, M. (2022). The effect of Medicaid on recidivism: Evidence from Medicaid suspension and termination policies. *Southern Economic Journal*, 89(2), 326-372.

Goodman-Bacon, A. (2021). Difference-in-differences with variation in treatment timing. *Journal of Econometrics*, 225(2), 254-277.

Herd, P., & Moynihan, D. P. (2018). *Administrative Burden: Policymaking by Other Means*. Russell Sage Foundation.

Herd, P., & Moynihan, D. P. (2025). Administrative burdens in the social safety net. *Journal of Economic Perspectives*, 39(1), 129-150.

Howell, B. A., Hawks, L. C., Balasuriya, L., Chang, V. W., Wang, E. A., & Winkelman, T. N. A. (2023). Health insurance and mental health treatment use among adults with criminal legal involvement after Medicaid expansion. *Psychiatric Services*, 74(10), 1019-1026.

Looney, A., & Turner, N. (2018). Work and opportunity before and after incarceration. Brookings Institution.

MACPAC (2018). *Medicaid and the criminal justice system*. MACPAC Issue Brief, July 2018.

Mallik-Kane, K., & Visser, C. (2008). Health and prisoner reentry. Urban Institute.

Olden, A., & Møen, J. (2022). The triple difference estimator. *Econometrics Journal*, 25(3), 531-553.

Ortiz-Villavicencio, M., & Sant'Anna, P. H. C. (2025). Better understanding triple differences estimators. Working paper, arXiv:2505.09942.

Packham, A., & Slusky, D. (2024). Accessing the safety net: How Medicaid affects health and recidivism. NBER Working Paper No. 31971 (originally December 2023; revised June 2024).

Sommers, B. D., Baicker, K., & Epstein, A. M. (2012). Mortality and access to care among adults after state Medicaid expansions. *New England Journal of Medicine*, 367(11), 1025-1034.

Sugar, S., Peters, C., De Lew, N., & Sommers, B. D. (2021). *Medicaid churning and continuity of care: Evidence and policy considerations before and after the COVID-19 pandemic* (ASPE Issue Brief HP-2021-10). Office of the Assistant Secretary for Planning and Evaluation, U.S. Department of Health and Human Services.

Wen, H., Hockenberry, J. M., Borders, T. F., & Druss, B. G. (2017). Impact of Medicaid expansion on Medicaid-covered utilization of buprenorphine for opioid use disorder treatment. *Medical Care*, 55(4), 336-341.

Zeng, Z. (2023). Jail inmates in 2022 - statistical tables. Bureau of Justice Statistics.

Appendix

Estimator and Inference Detail

LP-DiD post window. The pooled LP-DiD average treatment effect collapses the post period into a single long difference and retains a control state-year only if it is still untreated a fixed number of periods ahead. The estimator's default window is the maximum relative year present in the data, which is 13 in the one-year Medicaid frame because Washington adopted in 2005. At that window the pooled one-year regression runs on 4 observations and returns an analytic standard error of 0.018. Setting the window to five years, which matches the event-study window, puts the same estimate on 73 observations. Both are reported wherever an LP-DiD estimate appears. Because the default window is defined by the data, it changes definition whenever the sample changes, which is why the leave-one-out and state-exclusion diagnostics in the state-influence section below use the fixed-window series as the common estimand.

Two-step expansion adjustment. LP-DiD drops a level covariate as collinear inside its long-difference transformation, so the expansion-adjusted LP-DiD estimate is computed in two steps: the outcome is netted of the jointly estimated post-expansion coefficient, and LP-DiD is then applied to the residualized outcome. The conventional standard error from the second step treats the first-stage coefficient as known and understates uncertainty, which is why no analytic p-value is reported for that column.

Resampling strata. The state-cluster pairs bootstrap resamples states with replacement within strata defined by treatment timing (never-treated, not-yet-treated, early adopter, late adopter) and by expansion status, at fixed stratum size, so every draw contains the same mix of state types as the real panel. The whole two-step procedure, first-stage coefficient included, is re-estimated inside each draw. In the leave-out samples of the state-influence section below, the strata are rebuilt on the states that remain, so no draw can resample an excluded state.

Interval constructions. Three intervals are computed from the same draws: the percentile interval, the basic interval centred on the full-sample estimate, and the studentized (bootstrap-t) interval, which divides

each draw by its own within-draw cluster-robust standard error. The studentized interval is the primary construction, for the reasons given in Section 5.4. Bootstrap bias is reported as the mean of the draws less the full-sample estimate.

Failed draws. A draw is unusable when the resampled panel does not identify the first stage or the long difference, and usability is recorded separately for each post window, since the default window can fail on a sub-sample where the fixed window is well behaved. Usable-draw counts are reported alongside every bootstrap interval. Permutation reassignments that leave no clean comparison at the relevant horizon are likewise reported rather than discarded.

Policy Panel Construction

The published suspension trackers are snapshots taken at different moments, they disagree in a handful of states, and none of them dates the transition year that a staggered design requires. The six sources triangulated for the panel are KFF State Health Facts (2019 and 2025 snapshots), two MACPAC publications (a 2018 issue brief and a chapter in the June 2023 Report to Congress), and the NACo (2015) suspension and termination brief. The remaining three are Gollu and Zapryanova's (2022) appendix coding, a KFF and National Health Law Program issue brief on state eligibility policies for individuals moving into and out of incarceration (McKee et al., 2015), and Bandara et al. (2025) for implementation detail. When sources disagree, the hierarchy is state legislation first, then KFF, then MACPAC, then peer-reviewed research. Each state-year carries a coding confidence rating. As of January 2025 the panel classifies 48 states as suspension states, after resolving the two states KFF reports as unclear.

Treatment Timing and Exposure

Treatment is assigned by the policy regime in force in the state in the release-cohort year. Suspension, however, acts during incarceration and typically begins at admission, so a cohort released shortly after adoption may have entered prison and lost coverage under the prior regime. With a mean incarceration duration of 25.1 months, exposure misclassification in the first post-adoption cohorts is likely. Two

features of the policy are not documented well enough to model: none of the policy-panel sources say whether people already in prison on the effective date were moved to suspended status, which is the variable that would determine how much of the observed ramp is exposure maturation, and the sources support an adoption year but not an adoption month, so excluding the adoption-year release cohort is the available substitute.

Table A10 re-estimates the model under the alternative timing codings this implies: treatment lagged one and two years, adoption-year cohorts excluded, Kansas excluded for its time-limited policy, and both expansion clocks. The point estimate is stable across those choices, and lagging treatment raises rather than lowers it. Across the five codings and both clocks, every unadjusted estimate falls between 17.8 and 27.8 points and every adjusted estimate between 7.4 and 19.3 points. The increase with lag length is what exposure maturation would produce. However, the pre-adoption convergence documented in Section 6.4 or a gradual implementation rollout could produce the same pattern, and it is hard to tell which is at work.

Post-Window Sensitivity

The five-year post window reported in the text matches the event-study window. Re-estimating the expansion-adjusted specification at three-, four-, five-, and six-year fixed post windows leaves the estimate positive at every window, rising from 5.4 points at three years to 12.0 at four, 13.7 at five, and 15.2 at six (Appendix Figure A28). Everything except the window is held fixed: the same 212 rows in 17 states, the same first stage, the same comparison pool, the same clustering, and the same bootstrap replications.

That rise is not a clean dynamic path. Widening the window drops treated states from the pooled long difference. Seven switching states reach a three-year horizon, five reach four years, four reach five years, and only Colorado, Minnesota, and Washington reach six. The long difference therefore ends up resting on the three states that already carry most of the separation between the two policies, and the rise in the point estimate comes from that narrowing rather than from the passage of time.

The primary bootstrap-t interval covers zero at the three-, five-, and six-year windows. At four years it does not, running from 1.7 to 20.3 points with $p = 0.016$. In short, the window changes the precision of the estimate rather than its sign. Nor is the five-year window the one that gives the strongest result: at four years the estimate is smaller and its interval narrower.

State Influence in the Expansion Adjustment

Appendix Figure A27 reports all 17 leave-one-out samples. Every state in the frame is dropped in turn, and not just the ten that switch suspension, because a state can move the estimate through the first-stage expansion coefficient or through the comparison pool without switching treatment itself. Illinois and New York are treated throughout the one-year window and contribute nothing to the treatment coefficient, Kansas contributes a single treated cohort year, and five states are never treated inside it. Leaving out one state at a time moves the adjusted TWFE estimate between 4.9 and 14.6 points around the full-sample 11.1 points, a band of roughly 50 percent in either direction, against 20.6 to 25.9 points around 22.2 points for the unadjusted estimate.

The most influential state is Wisconsin, which is not one of the seven that identify the adjustment. Wisconsin adopted suspension in 2014 and never expanded inside the window, so it moves the estimate through the first stage rather than through the treatment coefficient. Dropping it raises the first-stage post-expansion coefficient from 31.5 to 37.8 points. Since the adjusted estimate is the raw estimate net of that coefficient times the expansion path, a larger first-stage coefficient subtracts more. As such, the adjusted estimate depends on the seven states that switch both policies and also on the few that switch suspension without ever expanding, which here are Wisconsin, Kansas, and North Carolina. All 17 leave-one-out estimates are positive at every window, but none produces a wild cluster bootstrap p below 0.05, and only the sample that drops Washington produces a bootstrap-t interval that excludes zero.

As an additional check, I exclude Washington, Minnesota, and Colorado together, to see how much of the adjustment rests on them. That exclusion removes just over half the treated post-adoption observations

and leaves 167 rows in 14 states, with an adjustment identified by four states whose policies are separated by at most a single year. The published estimate cannot be computed on that sample at all: the LP-DiD default post window is the maximum relative year present in the data, and once those three states are gone the long-difference design collapses, with only 286 of 1,000 bootstrap draws estimable.

At the fixed five-year window the estimate roughly doubles, to 26.0 points, with a bootstrap-t interval running from -3.851 to 0.304 on the share scale, too wide to be informative, and $p = 0.250$. That increase does not mean the association is stronger. The first-stage coefficient falls once the long-separation states are gone, so less is netted out, and Arizona is the only remaining state that reaches horizon 5, so the pooled long difference rests on one state's post-adoption path. The wild cluster bootstrap on the window-free TWFE version gives $p = 0.500$. Running the design on only the separated states and their comparison states gives a TWFE estimate of 14.4 points with a wild cluster bootstrap p of 0.051, close to the published estimate and still not distinguishable from zero. In short, these three states do not appear to generate an unusually large coefficient, and they supply much of the variation that separates suspension from expansion. Without them the adjusted estimate is poorly identified.

Subgroup Estimates

The one-year forest plot is Appendix Figure A18, and Appendix Figures A19 and A20 give the three- and five-year versions. Women's mean one-year post-release enrollment rate is 51.5 percent against 29.2 percent for men, so with roughly half of released women already enrolled at the one-year point there is less headroom for an administrative default to add coverage. Eligibility pathways, incarceration patterns, and cohort composition also differ by sex, and the published subgroup cells do not separate those explanations. Two-way subgroup cells are not published, so no subgroup estimate conditions on more than one characteristic at a time.

Secondary-Outcome Detail

Neither housing assistance nor SSI receipt shows a change under TWFE in raw or adjusted form, and the SSI sample of 97 observations is too small to detect one (Appendix Table A12). The nominally significant LP-DiD housing-assistance estimate rests on 31 pooled observations and is not corroborated by TWFE. LP-DiD also gives negative and nominally significant three-year estimates for employment and earnings. Those rest on a pooled regression of 23 observations, are not corroborated by TWFE at any horizon, and do not survive multiple-testing adjustment (Appendix Table A5). Felony and misdemeanor recharge outcomes are observed in 8 states, below the point at which state-clustered inference is reliable, and I do not interpret them (Appendix Table A8).

Appendix Tables

Table A1. Summary Statistics, Full

	N	Mean	SD	Min	Median	Max
Panel A. Primary outcomes, 1-year horizon						
Medicaid enrollment	212	0.320	0.251	0.078	0.200	0.880
W-2 employment	212	0.536	0.074	0.350	0.540	0.680
W-2 earnings (\$)	212	6,909	1,634	3,060	6,905	11,600
Mortality	264	0.017	0.005	0.006	0.016	0.048
Panel B. Longer-horizon and secondary outcomes						
Medicaid enrollment, 3-year	207	0.287	0.187	0.091	0.210	0.800
Medicaid enrollment, 5-year	201	0.277	0.166	0.093	0.210	0.740
HUD housing assistance, 1-year	264	0.011	0.005	0.002	0.010	0.031
SSI receipt, 1-year	99	0.043	0.010	0.021	0.045	0.060
Incarceration period length (months)	378	25.1	6.6	11.5	24.5	41.7
Panel C. State-year controls						
Unemployment rate (%)	561	5.58	2.13	2.20	5.00	13.70
Overdose mortality rate	554	13.9	6.5	2.6	12.7	44.3
Poverty rate (%)	534	13.0	2.9	6.9	12.7	21.4
Panel D. Policy variable						
Post-suspension indicator	561	0.296	0.457	0	0	1

Notes: The unit of observation is the state-year. Medicaid, employment, HUD, SSI, and mortality outcomes are shares in [0, 1]. Earnings are in nominal dollars averaged across release-cohort members. The overdose mortality rate is CDC WONDER age-adjusted deaths per 100,000. The

sample covers 31 states over 2000-2020; Medicaid outcomes are restricted to 17 states with horizon-specific cohort windows implied by the MSIS/T-MSIS linkage (release cohorts 2004-2018 at 1 year, 2002-2016 at 3 years, 2000-2014 at 5 years).

Table A2. Pre-Adoption Balance, Full

	Adopting states, mean (SD)	Non-adopting states, mean (SD)	Difference	Normalized difference
Panel A. Primary outcomes				
Medicaid enrollment, 1-year	0.210 (0.109)	0.144 (0.041)	0.066	0.80
W-2 employment, 1-year	0.552 (0.079)	0.513 (0.063)	0.039	0.55
W-2 earnings, 1-year (\$)	6,814 (1,822)	6,867 (1,432)	-53	-0.03
Mortality, 1-year	0.016 (0.005)	0.017 (0.004)	-0.001	-0.17
Panel B. Release-cohort characteristics				
Incarceration period (months)	22.6 (5.8)	29.9 (3.5)	-7.3	-1.53
Annual release cohort size	10,096 (7,002)	34,644 (23,043)	-24,549	-1.44
Share male	0.879 (0.027)	0.864 (0.022)	0.015	0.62
Share Black	0.302 (0.156)	0.388 (0.078)	-0.086	-0.70
Share Hispanic	0.149 (0.112)	0.133 (0.107)	0.015	0.14
Share young (15-24)	0.197 (0.050)	0.181 (0.022)	0.016	0.42
Panel C. State controls				
Unemployment rate (%)	5.69 (2.23)	5.60 (1.93)	0.10	0.05
Overdose mortality rate	11.8 (5.0)	13.3 (5.6)	-1.5	-0.28
Poverty rate (%)	12.5 (3.0)	15.3 (1.9)	-2.8	-1.10

Notes: Means with standard deviations in parentheses, computed across pre-adoption state-year observations. “Adopting states” comprises all states that adopted suspension during 2005-2022, observed in their pre-treatment years. “Non-adopting states” comprises the four CJARS states with no suspension adoption inside the 2000-2020 panel: Florida and Oklahoma, whose adoption falls after the panel ends, plus the two never-treated states Georgia and Texas. Georgia contributes no Medicaid, employment, earnings, mortality, housing, or SSI rows, so it enters only the state-control rows of Panel C. Normalized difference is the mean difference divided by the pooled standard deviation. Release-cohort demographic shares are computed from JOE subgroup cells using the release-cohort size variable as the denominator within sex, race, and age subgroups. Level differences in all three panels are absorbed by state fixed effects in the main design.

Table A3. Main Results Across Estimators and Outcomes, 1-Year Horizon

	Medicaid	Medicaid (+ controls)	Employment	Earnings	Mortality
Panel A. Two-way fixed effects					
Suspension	0.222***	0.230***	0.016	315.3	0.0005

	(0.051)	(0.041)	(0.011)	(272.3)	(0.0011)
Panel B. LP-DiD (Dube et al., 2025), estimator default window					
Suspension	0.224*** (0.018)	-	0.003 (0.029)	383.1 (339.4)	-0.0017** (0.0006)
Observations entering the estimate	4	-	4	4	31
Panel C. Callaway-Sant'Anna (2021)					
Suspension (aggregated ATT)	0.258*** (0.046)	-	-	-	0.0001 (0.0015)
Panel D. TWFE, cohort-size weighted					
Suspension	0.257*** (0.062)	-	0.023 (0.020)	351.4 (368.5)	0.0020 (0.0012)
Panel E. Inference and specification					
Wild cluster bootstrap p	0.001	-	0.150	0.280	0.652
State fixed effects	Yes	Yes	Yes	Yes	Yes
Release-cohort-year fixed effects	Yes	Yes	Yes	Yes	Yes
Time-varying controls	No	Yes	No	No	No
Observations	212	198	212	212	264
States	17	17	17	17	17

Notes: Cluster-robust standard errors clustered at the state level are in parentheses. Each column is an outcome-specific specification with state and release-cohort-year fixed effects. Panel A is the benchmark two-way fixed effects specification; Panel B uses the LP-DiD estimator of Dube et al. (2025) at the estimator's default data-dependent post window, and the row beneath reports how many observations enter the pooled long-difference regression, which is why the Panel B standard errors are so small (Section 5.2). Panel C uses the Callaway and Sant'Anna (2021) group-time average treatment effect with not-yet-treated controls. Panel D reweights the TWFE specification by the JOE release-cohort size variable for each outcome. Panel E reports wild cluster bootstrap p-values with the null imposed (Webb weights, 9,999 replications), and no bootstrap was run for the controls specification. Medicaid and employment outcomes are shares, earnings are nominal dollars, and mortality is the share deceased within the horizon. The Medicaid, employment, and earnings samples are restricted to 17 states and release cohorts 2004-2018 by the MSIS/T-MSIS linkage window; the mortality column covers the same 17 states over release cohorts 2000-2018. *p<0.10, **p<0.05, ***p<0.01.

Table A4. Heterogeneous DiD by ACA Medicaid Expansion Status

	Medicaid	Employment	Earnings	Mortality
Panel A. Static-expansion interaction				
Suspension	0.075 (0.130)	0.024* (0.012)	478.8 (641.6)	0.0018* (0.0010)
Suspension × Expanded	0.195	-0.010	-215.6	-0.0016

	(0.132)	(0.017)	(660.6)	(0.0016)
Panel B. Period-split TWFE				
Suspension × Pre-ACA	0.117**	0.007	-2.6	-0.0000
	(0.050)	(0.009)	(300.1)	(0.0007)
Suspension × Post-ACA	0.147*	0.008	-157.3	-0.0013
	(0.081)	(0.008)	(300.0)	(0.0035)
Panel C. Time-varying expansion interaction				
Suspension × Expansion (time-varying)	Collinear	Collinear	Collinear	Collinear
Specification details				
State and year fixed effects	Yes	Yes	Yes	Yes
Observations (static)	212	212	212	264
Observations (pre-ACA)	134	134	134	186
Observations (post-ACA)	78	78	78	78

Notes: Cluster-robust standard errors clustered at the state level are in parentheses, and all three panels are estimated on the 17-state outcome panel (Section 4.4). Panel A reports the interaction with time-invariant ACA expansion status as the cross-classifying dimension; the main effect of expansion status is absorbed by the state fixed effects. Panel B reports TWFE coefficients on the suspension indicator estimated separately in the pre- and post-expansion periods. The pre-period is release cohorts 2004-2013 for Medicaid, employment, and earnings and 2000-2013 for mortality; the post-period is 2014-2018 for all four outcomes. Panel C reports the outcome of an attempted time-varying expansion interaction. In the Medicaid sample, every state that adopted both suspension and the expansion did so in an order that makes the interaction collinear with the expansion indicator net of fixed effects, so the coefficient is not separately identified (Section 6.4). *p<0.10, **p<0.05, ***p<0.01.

Table A5. Horizon Sensitivity, Years 1, 3, and 5 Post-Release, TWFE and LP-DiD

	Medicaid	Employment	Earnings (\$)	Mortality
1-year horizon				
TWFE	0.222***	0.016	315.3	0.0005
	(0.051)	(0.011)	(272.3)	(0.0011)
LP-DiD, estimator default window	0.224***	0.003	383.1	-0.0017**
	(0.018)	(0.029)	(339.4)	(0.0006)
LP-DiD observations entering the estimate	4	4	4	31
Observations (TWFE)	212	212	212	264
3-year horizon				
TWFE	0.194***	0.003	281.1	0.0014
	(0.044)	(0.011)	(268.2)	(0.0012)
LP-DiD, estimator default window	0.187***	-0.031***	-746.3**	0.0027
	(0.007)	(0.006)	(319.9)	(0.0015)
LP-DiD observations entering the estimate	23	23	23	42
Observations (TWFE)	205	205	205	231

5-year horizon				
TWFE	0.207***	-0.006	93.8	0.0007
	(0.058)	(0.016)	(347.6)	(0.0014)
LP-DiD, estimator default window	0.196***	-0.022*	-145.8	0.0026
	(0.018)	(0.011)	(200.4)	(0.0016)
LP-DiD observations entering the estimate	48	48	48	48
Observations (TWFE)	200	200	200	200

Notes: Cluster-robust standard errors clustered at the state level are in parentheses. Each cell is a separate regression of the stated outcome on the post-suspension indicator with state and release-cohort-year fixed effects. Each horizon is estimated on its own censoring-consistent cohort window in the 17-state outcome panel (Section 4.3); the five-year window is identified by 15 states. The LP-DiD rows use the estimator's default data-dependent post window, and the row beneath each pair reports how many observations enter the pooled long-difference regression. The negative and nominally significant three-year LP-DiD employment and earnings estimates rest on 23 pooled observations and are not corroborated by TWFE at any horizon. They also run opposite in sign to the positive one-year LP-DiD earnings estimate and carry false-discovery-rate adjusted q-values above 0.5, so they are reported here and not interpreted as findings (Section 6.6). *p<0.10, **p<0.05, ***p<0.01.

Table A6. Robustness Summary, Medicaid Enrollment, 1-Year Horizon

	Coefficient	SE or interval	N / States	Specification
Panel A. Estimator robustness				
TWFE (baseline)	0.222***	(0.051)	212 / 17	Cluster-robust SE
TWFE + time-varying controls	0.230***	(0.041)	198 / 17	Unemployment, overdose, poverty
LP-DiD (Dube et al., 2025), 5-year window	0.207**	(0.068)	73 pooled / 17	Heterogeneity-robust
Callaway-Sant'Anna (2021)	0.258***	(0.046)	- / 17	Not-yet-treated controls
Cohort-size-weighted TWFE	0.257***	(0.062)	212 / 17	Weighted by release-cohort size
Panel B. Inference robustness				
Wild cluster bootstrap p	-	0.001	-	B = 9,999, Webb, null imposed
Permutation placebo p, TWFE	-	0.062	500 permutations	Design-based
Permutation placebo p, LP-DiD	-	0.081	761 admissible of 1,000	Design-based
Panel C. Sample and subsample robustness				
Leave-one-out (range)	[0.206, 0.259]	max deviation 0.036	-	Across 17 states

Drop Texas (never-treated)	0.209***	(0.052)	198 / 16	Modest attenuation
Drop Kansas (time-limited suspension)	0.225***	(0.056)	197 / 16	Main result unchanged
Pre-ACA only (2004-2013)	0.117**	(0.050)	134 / 17	Aligns with Packham-Slusky 11.1 pp
Post-ACA only (2014-2018)	0.147*	(0.081)	78 / 17	Post-expansion period
Panel D. Policy-coding and cell-size robustness				
Drop low-confidence policy codes	0.209***	-	-	1 state dropped
Cell size N >= 50 (vs. N >= 25)	0.222***	(0.051)	-	Identical, all cells >= 50
Panel E. Composition controls (matched sample)				
Baseline on composition-complete cells	0.164***	(0.041)	143 / 17	Cells with non-missing composition
+ release-cohort composition controls	0.139***	(0.043)	143 / 17	Sex, race, age shares
+ state-year covariates	0.152***	(0.041)	130 / 17	Adds unemployment, overdose, poverty
Panel F. Multiple-testing adjustment				
False-discovery-rate q-value	0.222	q = 0.005	-	Across 14-outcome battery

Notes: Column 1 reports the coefficient on the post-suspension indicator for one-year Medicaid enrollment; column 2 reports the cluster-robust standard error in parentheses, or a p-value or interval where relevant; column 3 reports sample size and cluster count; column 4 describes the specification. All TWFE specifications include state and release-cohort-year fixed effects with standard errors clustered at the state level. Georgia and South Carolina contribute no Medicaid outcome rows and therefore cannot support a leave-one-out check; Texas is the only never-treated state in the Medicaid panel, and dropping it removes its 14 rows. The false-discovery-rate q-value is a Benjamini-Hochberg q computed across the full outcome battery (Medicaid, employment, earnings, and mortality at 1, 3, and 5 years; HUD and SSI at 1 year). *p<0.10, **p<0.05, ***p<0.01.

Table A7. Subgroup DiD Estimates by Sex, Race/Ethnicity, and Age Group

Subgroup	Medicaid 1-yr TWFE	Medicaid 1-yr LP- DiD	Medicaid 3-yr TWFE	Medicaid 3-yr LP- DiD	Medicaid 5-yr TWFE	Medicaid 5-yr LP- DiD	Reincarc . 1-yr TWFE	Reincarc . 1-yr LP-DiD
Overall	0.222*** (0.051)	0.224*** (0.018)	0.194*** (0.044)	0.187*** (0.007)	0.207*** (0.058)	0.196*** (0.018)	-0.022 (0.027)	-0.007 (0.007)
Male	0.231*** (0.054)	0.234*** (0.021)	0.199*** (0.044)	0.201*** (0.009)	0.213*** (0.058)	0.198*** (0.017)	-0.023 (0.027)	-0.008 (0.008)

Female	0.159*** (0.039)	0.177*** (0.014)	0.158*** (0.048)	0.152*** (0.012)	0.165** (0.064)	0.170*** (0.018)	-0.010 (0.023)	-0.007** (0.002)
White	0.214*** (0.050)	0.221*** (0.017)	0.190*** (0.043)	0.196*** (0.010)	0.200*** (0.057)	0.201*** (0.017)	-0.023 (0.030)	-0.015* (0.007)
Black	0.233*** (0.051)	0.237*** (0.025)	0.210*** (0.047)	0.204*** (0.012)	0.219*** (0.063)	0.198*** (0.018)	-0.031 (0.029)	0.001 (0.013)
Hispanic	0.218*** (0.051)	0.214*** (0.014)	0.176*** (0.043)	0.184*** (0.009)	0.185** (0.063)	0.148*** (0.020)	-0.029 (0.023)	0.032*** (0.006)
Asian/Pacific Islander ^a	0.176** (0.060)	n.r.	0.148*** (0.037)	0.168** (0.043)	0.126*** (0.036)	0.152*** (0.036)	-0.050 (0.034)	-0.109 (0.104)
American Indian/Alaska Native ^a	0.176*** (0.052)	n.e.	0.201** (0.065)	0.184*** (0.010)	0.237*** (0.070)	0.200*** (0.023)	-0.021 (0.044)	n.e.
Age 15-24	0.227*** (0.053)	0.173*** (0.020)	0.207*** (0.047)	0.163*** (0.010)	0.216*** (0.064)	0.199*** (0.021)	-0.017 (0.030)	-0.012 (0.019)
Age 25-39	0.234*** (0.054)	0.240*** (0.018)	0.198*** (0.046)	0.207*** (0.010)	0.215*** (0.060)	0.200*** (0.020)	-0.024 (0.027)	-0.018* (0.009)
Age 40+	0.204*** (0.047)	0.211*** (0.023)	0.183*** (0.043)	0.176*** (0.006)	0.192*** (0.055)	0.175*** (0.013)	-0.024 (0.026)	0.003 (0.004)

Notes: Each cell reports the coefficient on the post-suspension indicator from a separate subgroup-specific regression, with the cluster-robust standard error clustered at the state level in parentheses. The subgroup panels are built from JOE's published one-way subgroup cells. LP-DiD columns use the estimator's default post window and inherit the small pooled sample described in Section 5.2, so the LP-DiD standard errors in this table understate uncertainty and the LP-DiD significance markers overstate significance. Full-coverage subgroup panels match the overall panel sample sizes (Medicaid 1-year N = 212 and 17 states; 3-year N = 205-207 and 17 states; 5-year N = 200-201 and 15 states; reincarceration 1-year N = 384 and 22 states). The Black and Hispanic panels miss at most a handful of suppressed cells, except the Black five-year panel, which loses one state (14 rather than 15). "n.r." means LP-DiD returns a degenerate standard error and is not reported; "n.e." means LP-DiD estimation fails in the sparse suppressed panel. ^aUnderpowered: Census Bureau disclosure suppression of cells with fewer than 25 individuals removes 23 to 27 percent of state-cohort cells for the Asian/Pacific Islander and American Indian/Alaska Native panels, a non-random set of small cohorts, so these rows should be interpreted with caution. *p<0.10, **p<0.05, ***p<0.01.

Table A8. Charge-Type Recidivism Detail, Felony and Misdemeanor Recharge, 1-Year Horizon

Outcome	Estimator	Estimate	SE	p	N	States
Felony recharge	TWFE	0.009	0.018	0.636	129	8
Felony recharge	LP-DiD	n.e.	-	-	129	8
Felony recharge	Callaway-Sant'Anna	-0.011	0.014	0.413	129	8
Misdemeanor recharge	TWFE	0.046	0.017	0.030	116	8
Misdemeanor recharge	LP-DiD	n.e.	-	-	116	8
Misdemeanor recharge	Callaway-Sant'Anna	0.002	0.027	0.932	116	8

Notes: Release-cohort window 2000-2020. JOE court-based charge outcomes are observed in only 8 states, since JOE court coverage is far thinner than corrections coverage. With fewer than 10 clusters, state-clustered standard errors are fragile and all estimates in this table should be treated as suggestive. "n.e." means not estimable, because the LP-DiD estimator has insufficient clean-control variation in the 8-state subsample. The positive misdemeanor TWFE estimate does not survive the Callaway and Sant'Anna estimator and is not interpreted as evidence of a misdemeanor-increasing effect (Section 6.6). All recidivism outcomes capture same-state criminal-justice contact only and are lower bounds on true recidivism.

Table A9. Expansion-Adjustment Specification Battery, Medicaid Enrollment

Specification	1-year	3-year
Panel A. Time-varying expansion and eligibility-threshold controls		
TWFE, baseline (no control)	0.222*** (0.051)	0.194*** (0.044)
TWFE + parent income limit only	0.212*** (0.053)	0.192*** (0.041)
TWFE + post-expansion control	0.111 (0.069)	0.146*** (0.048)
TWFE + post-expansion + parent income limit	0.115 (0.068)	0.155*** (0.043)
Post-expansion control coefficient	0.315*** (0.089)	0.132* (0.072)
Parent income limit coefficient (per 100% FPL)	0.039 (0.060)	0.060 (0.044)
LP-DiD two-step, expansion only, default window	0.111 [0.053, 0.270]	0.154 [0.120, 0.326]
LP-DiD two-step, expansion + parent income limit	0.111 [0.038, 0.290]	0.155 [0.122, 0.333]
Panel B. Childless-adult income limit, 2011+ subsample		
TWFE + expansion + parent limit, 2011+ (matched comparator)	0.106 (0.081)	0.014 (0.023)
TWFE + expansion + parent + childless-adult limit, 2011+	0.089 (0.068)	0.013 (0.027)
Childless-adult limit coefficient (per 100% FPL)	0.097 (0.076)	0.005 (0.035)
State-years dropped by the 2011+ restriction	92 of 212 (43.4%)	117 of 204 (57.4%)
Panel C. Split samples		
TWFE, ever-expanded states only (11 states)	0.201** (0.077)	0.151** (0.059)
LP-DiD, ever-expanded states only	0.225** (0.033)	0.180*** (0.019)
TWFE, non-expansion states only (6 states)	0.266* (0.110)	0.185*** (0.009)
Panel D. Drop 2014-16 release cohorts		
TWFE	0.143** (0.066)	0.194*** (0.045)
LP-DiD	n.e.	0.089*** (0.016)

Notes: The outcome is the share of a state's release cohort enrolled in Medicaid at the stated horizon. Cluster-robust standard errors clustered at the state level are in parentheses, and square brackets report state-cluster bootstrap 95 percent percentile intervals for the two-step LP-DiD rows (B = 1,000). The governing bootstrap-t construction in Section 5.4 gives wider intervals than these percentile intervals (Table 4). Panels A and B use all 17 states, except that the three-year parent-limit specifications lose one observation to a single missing eligibility cell (N = 204 rather than 205). Panel C splits the panel into 11 ever-expanded and 6 non-expansion states, and Panel D's three-year cells are identified by 14 states after the 2014-16 cohorts are dropped. Panel A adds a time-varying indicator equal to one in state-years with the ACA expansion in effect and the state Medicaid income-eligibility limit for working parents, scaled per 100 percentage points of the federal poverty level. The two-step LP-DiD rows residualize the outcome on the jointly estimated control coefficients and bootstrap the whole procedure. Panel B adds the income limit for non-disabled adults without dependent children, which the published source reports only from 2011 forward; the matched-comparator row applies the same 2011+ restriction without the childless-adult term so that the covariate can be separated from the sample restriction. The parent limit tracks eligibility for parents while most people leaving prison are childless adults, so the childless-adult limit is the closer measure. The ever-expanded split contains no never-treated state, and the non-expansion split rests on 6 states with 8 treated post-period observations at the 1-year horizon and is underpowered. Panel C reports no LP-DiD rows for the non-expansion split, because the six-state split leaves the estimator effectively a single clean comparison, with a degenerate standard error at one year and three treated post-period observations at three years. "n.e." means not estimable because the LP-DiD long differences span the resulting gap in the panel. *p<0.10, **p<0.05, ***p<0.01.

Table A10. Treatment-Timing Battery and the Expansion Clock, Medicaid Enrollment

Timing coding	Horizon	Raw	Adjusted, expansion at t	Adjusted, expansion at t+h	N
Current coding (regime at release)	1-year	0.222*** (0.051)	0.111 (0.069)	0.111 (0.071)	212
Treatment lagged 1 year	1-year	0.257*** (0.044)	0.141** (0.063)	0.144** (0.056)	212
Treatment lagged 2 years	1-year	0.276*** (0.057)	0.162** (0.064)	0.154** (0.058)	212
Adoption-year cohorts excluded	1-year	0.278*** (0.050)	0.157* (0.078)	0.154** (0.073)	202
Kansas excluded (time-limited)	1-year	0.225*** (0.056)	0.118 (0.070)	0.118 (0.072)	197
Current coding (regime at release)	3-year	0.194*** (0.044)	0.146*** (0.048)	0.100** (0.039)	205
Treatment lagged 1 year	3-year	0.218*** (0.060)	0.173** (0.062)	0.115** (0.052)	205
Treatment lagged 2 years	3-year	0.237*** (0.071)	0.193*** (0.064)	0.120** (0.049)	205
Adoption-year cohorts excluded	3-year	0.232*** (0.058)	0.188** (0.065)	0.127** (0.055)	197
Kansas excluded (time-limited)	3-year	0.186*** (0.046)	0.141** (0.048)	0.098** (0.039)	190
Expansion clock reclassification					
Rows coded post-expansion at t / at t+h	1-year	37 / 44	7 rows reclassified (3.3%)		212
Rows coded post-expansion at t / at t+h	3-year	19 / 39	20 rows reclassified (9.7%)		207
Rows coded post-expansion at t / at t+h	5-year	4 / 35	31 rows reclassified (15.4%)		201

Notes: All estimates are TWFE with state and release-cohort-year fixed effects and cluster-robust standard errors clustered at the state level in parentheses. “Raw” omits the expansion indicator. The two adjusted columns add a state-specific time-varying post-expansion indicator, indexed either to the release-cohort year t or to the outcome-measurement year t+h. Kansas is the only time-limited suspension state in the 17-state Medicaid panel. Across all five codings and both clocks, every raw estimate falls between 0.178 and 0.278 and every adjusted TWFE estimate between 0.098 and 0.193; the LP-DiD two-step versions of the same battery range from 0.074 to 0.172. The lower panel reports how many state-cohort rows change expansion status when the indicator is moved from the release year to the outcome-measurement year. *p<0.10, **p<0.05, ***p<0.01.

Table A11. Event-Time Support and Event-Study Coefficients, Medicaid Enrollment, 1-Year Horizon

Event time	Treated states	Clean control states	Never-treated states available	Adjusted coefficient (SE)	Joint-policy-era coefficient (SE)
-5	11	15	1	-0.042 (0.057)	-0.042 (0.057)
-4	10	14	1	-0.061 (0.054)	-0.061 (0.054)
-3	10	15	1	-0.083 (0.025)	-0.083 (0.025)
-2	11	14	1	-0.087 (0.026)	-0.087 (0.026)
0	10	14	1	-0.016 (0.038)	0.074 (0.027)
1	9	14	1	0.000 (0.041)	0.135 (0.052)
2	9	13	1	0.033 (0.049)	0.177 (0.062)
3	8	12	1	0.097 (0.076)	0.212 (0.070)
4	6	11	1	0.233 (0.083)	0.294 (0.082)
5	4	12	1	0.240 (0.119)	0.390 (0.121)
Average post-adoption effect				0.098 (0.056)	0.214 (0.054)
Joint test of leads				chi-square(4) = 26.6, p = 0.00002	same coefficients

Notes: Event time -1 is the omitted reference, and the estimation frame is the one-year Medicaid panel of 212 rows in 17 states. Standard errors come from the stratified state-cluster bootstrap described in Section 5.4 (500 draws), which re-estimates the entire two-step procedure inside each draw. Coefficients carry no significance markers because the bootstrap produces interval estimates rather than per-coefficient p-values; the Figure 4 bands and the joint lead tests here carry the inference. The pre-period leads are identical across the two columns: no treated state changed ACA expansion status inside its own pre-adoption event window, so the adjustment affects only the post-adoption path, and the joint lead test applies to both. An unstratified covariance on the same coefficients gives chi-square(4) = 13.8, p = 0.008. “Clean control states” are never-treated and not-yet-treated states available at that event time.

Table A12. Expansion-Adjusted Secondary Outcomes With Confidence Intervals, All Horizons

Outcome	Horizon	Raw TWFE	Adjusted TWFE (SE)	95% CI	Wild cluster bootstrap p	N	States
W-2 employment	1-year	0.0162	0.0147 (0.0086)	[-0.0035, 0.0330]	0.153	212	17
W-2 employment	3-year	0.0030	0.0051 (0.0092)	[-0.0147, 0.0248]	0.608	205	17
W-2 employment	5-year	-0.0063	-0.0047 (0.0160)	[-0.0393, 0.0299]	0.779	200	15
W-2 earnings (\$)	3-year	281.1	313.7 (324.7)	[-382.7, 1010.1]	0.381	205	17
W-2 earnings (\$)	5-year	93.8	98.7 (360.0)	[-678.9, 876.4]	0.802	200	15
All-cause mortality	1-year	0.0005	0.0011 (0.0008)	[-0.0006, 0.0028]	0.237	264	17

All-cause mortality	3-year	0.0014	0.0015 (0.0013)	[-0.0013, 0.0044]	0.326	231	17
All-cause mortality	5-year	0.0007	0.0011 (0.0013)	[-0.0016, 0.0039]	0.401	200	15
Reincarceration, same state	1-year	-0.0216	0.0034 (0.0222)	[-0.0427, 0.0495]	0.895	384	22
Reincarceration, same state	3-year	-0.0244	-0.0003 (0.0152)	[-0.0320, 0.0314]	0.983	345	21
Reincarceration, same state	5-year	-0.0101	0.0018 (0.0106)	[-0.0203, 0.0239]	0.869	302	21
Reincarceration, felony charge	1-year	0.0091	0.0176 (0.0197)	[-0.0289, 0.0642]	0.495	129	8
HUD housing assistance	1-year	0.0011	0.0005 (0.0009)	[-0.0014, 0.0023]	0.645	264	17
HUD housing assistance	3-year	0.0024	0.0021 (0.0015)	[-0.0012, 0.0053]	0.256	231	17
SSI receipt	1-year	-0.0003	0.0002 (0.0022)	[-0.0045, 0.0049]	0.937	97	17
SSI receipt	3-year	-0.0002	0.0019 (0.0021)	[-0.0026, 0.0065]	0.370	95	17

Notes: All estimates are TWFE with state and release-cohort-year fixed effects; the adjusted column adds a state-specific time-varying post-expansion indicator, indexed to the release-cohort year. Cluster-robust standard errors clustered at the state level are in parentheses, and intervals use the cluster-t critical value. Wild cluster bootstrap p-values use Webb weights with the null imposed (9,999 replications). The expansion adjustment does not change the sign or the significance of any of these estimates. Two-step expansion-adjusted LP-DiD versions of the reincarceration rows turn positive and nominally significant on pooled long-difference regressions of 36 to 62 observations, contradict the TWFE estimates on the same sample, and are not reported as findings for the reason given in Section 5.2.

Table A13. Joint-Policy-Era LP-DiD, Callaway-Sant’Anna, and the Common-Cohort Window, Medicaid

Enrollment

Specification	Horizon	Estimate	SE	N	States
LP-DiD, raw, 5-year post window	1-year	0.207**	(0.068)	212	17
LP-DiD, raw, 5-year post window	3-year	0.198**	(0.080)	205-207	17
LP-DiD, raw, 5-year post window	5-year	0.144*	(0.081)	200-201	15
Callaway-Sant’Anna, raw	1-year	0.258***	(0.046)	212	17
Callaway-Sant’Anna, raw	3-year	0.169***	(0.058)	205-207	17
Callaway-Sant’Anna, raw	5-year	0.116*	(0.066)	200-201	15
TWFE, raw, cohorts with a complete 5-year outcome (2004-2014)	1-year	0.162***	(0.040)	148-149	15
Callaway-Sant’Anna, raw, same cells	1-year	0.187***	(0.051)	148-149	15

Notes: The outcome is the share of a state’s release cohort enrolled in Medicaid at the stated horizon, and every estimate is an unweighted state-average de jure policy-regime intent-to-treat association with no expansion control. Cluster-robust standard errors clustered at the state level are

in parentheses. LP-DiD is the estimator of Dube et al. (2025) with the post window fixed at five years, and Callaway and Sant’Anna (2021) uses not-yet-treated controls. Each horizon is estimated on its own censoring-consistent cohort window (Section 4.3). The last two rows re-estimate the 1-year model on exactly the cells for which the 5-year outcome is observed, so that horizon comparisons are not confounded with cohort composition (Section 6.2). The expansion-adjusted counterparts are in Table 5. *p<0.10, **p<0.05, ***p<0.01.


Table A14. Inference Reconciliation, Expansion-Adjusted 1-Year Medicaid Estimate, All Constructions

	TWFE + expansion	LP-DiD two-step, default window	LP-DiD two-step, 5- year window
Primary inference	Wild cluster bootstrap	Studentized bootstrap-t	Studentized bootstrap-t
Point estimate	0.111	0.111	0.137
Observations entering the estimate	212	4	73
Analytic cluster-robust SE	(0.069)	(0.018)	(0.043)
Analytic p-value	0.126	0.009	-
Stratified bootstrap SE	-	(0.059)	(0.057)
Bootstrap bias (mean of draws less estimate)	-	+0.028	+0.005
95% percentile interval	-	[0.071, 0.266]	[0.050, 0.266]
95% basic (centred) interval	-	[-0.044, 0.152]	[0.007, 0.223]
95% bootstrap-t interval	-	[-0.110, 0.192]	[-0.155, 0.279]
Percentile p-value	-	0.001	0.001
Bootstrap-t p-value	-	0.176	0.190
Wild cluster bootstrap p-value	0.301	-	-
Permutation placebo p-value	-	0.299	-
Usable bootstrap draws (of 1,000)	-	913	999


Notes: The full construction-by-construction version of Table 4. All columns estimate the same object, the one-year expansion-adjusted suspension estimate on release cohorts 2004-2018 in 17 states, with state-level clustering. Estimates and interval limits are shares of the release cohort. The first row names the primary construction in each column, defined in Section 5.4. Draws are unusable when the resampled panel does not identify the first stage or the long difference, and 239 of the 1,000 permutation reassignments leave no clean comparison and are reported rather than dropped.

Appendix Figures

Appendix Figure A1. TWFE event study, Medicaid enrollment, 1-year horizon




Appendix Figure A2. TWFE event study, W-2 employment, 1-year horizon




Appendix Figure A3. TWFE event study, W-2 earnings, 1-year horizon




Appendix Figure A4. TWFE event study, mortality, 1-year horizon




Appendix Figure A5. LP-DiD event study, Medicaid enrollment, 3-year horizon




Appendix Figure A6. LP-DiD event study, Medicaid enrollment, 5-year horizon



Appendix Figure A7. TWFE event study with time-varying expansion control, Medicaid enrollment, 1-year horizon




Appendix Figure A8. Suspension adoption year versus ACA expansion year




Notes: Each point is a state in the 31-state policy panel that adopted both policies. Points on the 45-degree line adopted suspension in the same year they expanded Medicaid.

Appendix Figure A9. Goodman-Bacon decomposition, Medicaid enrollment




Appendix Figure A10. Pre-trend sensitivity, Rambachan-Roth bounds on the joint-policy-era event study




Notes: The bounds shown apply to the unadjusted joint-policy-era path. On the expansion-adjusted path the robust confidence set contains zero at every relative-magnitude bound, including a bound of zero (Section 6.4).


Appendix Figure A11. Placebo distribution, Medicaid enrollment




Appendix Figure A12. Leave-one-out estimates, Medicaid enrollment




Appendix Figure A13. Permutation placebo distribution, LP-DiD




Appendix Figure A14. Coefficient comparison across estimators




Appendix Figure A15. Policy adoption timeline




Appendix Figure A16. Raw Medicaid enrollment trends by adoption cohort, 3-year horizon



Appendix Figure A17. Raw Medicaid enrollment trends by adoption cohort, 5-year horizon




Appendix Figure A18. Subgroup estimates, Medicaid enrollment, 1-year horizon



Notes: Coefficients on the post-suspension indicator from separate regressions estimated within each published JOE subgroup cell, with 95 percent confidence intervals, for the share enrolled in Medicaid one year after release. Hollow markers denote subgroups where Census Bureau disclosure suppression removes a substantial share of state-cohort cells (Appendix Table A7).


Appendix Figure A19. Subgroup estimates, Medicaid enrollment, 3-year horizon



* Underpowered subgroup: reduced cell coverage due to disclosure suppression (N<25); interpret with caution.

Notes: Hollow markers denote underpowered subgroups affected by disclosure suppression (Appendix Table A7).


Appendix Figure A20. Subgroup estimates, Medicaid enrollment, 5-year horizon




* Underpowered subgroup: reduced cell coverage due to disclosure suppression (N<25); interpret with caution.

Notes: Hollow markers denote underpowered subgroups affected by disclosure suppression (Appendix Table A7).


Appendix Figure A21. LP-DiD event study, reincarceration, 1-year horizon



Appendix Figure A22. TWFE event study, reincarceration, 1-year horizon




Appendix Figure A23. Subgroup estimates, reincarceration, 1-year horizon




Notes: Hollow markers denote underpowered subgroups affected by disclosure suppression (Appendix Table A7).


Appendix Figure A24. Placebo distribution, mortality



Appendix Figure A25. Alternative scale for earnings, levels versus log LP-DiD event study



Appendix Figure A26. Joint-policy-era LP-DiD event study, Medicaid enrollment, 1-year horizon



Coefficients are LP-DiD long differences against never-treated and not-yet-treated states, normalised at event time -1 . Bands are 95 percent intervals from a state-cluster bootstrap stratified on treatment timing and expansion status (500 draws), which re-estimates the full event study within each draw. The lower panel reports the states supporting each event time; the dotted line marks the 3-state support rule. Exactly one never-treated state contributes at every event time.


Notes: The unadjusted LP-DiD path, plotted with the same estimator, bootstrap covariance, and support panel as Figure 4. Its post coefficients run from 0.074 at adoption to 0.390 at five years and average 0.214 (SE 0.054). Because suspension and ACA expansion adoption coincide in most panel states, this path measures the combined policy era rather than suspension alone. Author's calculations from the CJARS Justice Outcomes Explorer and the state suspension policy panel.

Appendix Figure A27. State contributions to the expansion-adjusted 1-year Medicaid estimate, all 17

leave-one-out samples

State contributions to the expansion-adjusted one-year Medicaid estimate


Each row drops one state and re-estimates the adjusted specification on the remaining states; the diamond drops Colorado, Minnesota and Washington together



The one-year Medicaid frame has 212 observations across 17 states, of which 10 switch suspension inside the estimation window and 7 also switch ACA expansion inside it, so only those 7 identify the expansion adjustment. Of the 7, Colorado, Minnesota and Washington separate suspension adoption from expansion by five, seven and nine years; Arizona by one; Indiana, Montana and Pennsylvania by none. Left panel intervals are cluster-robust at the state level. Right panel intervals are bootstrap-t, studentized by the within-draw cluster-robust standard error, from a state-cluster pairs bootstrap stratified on treatment timing and expansion status, with the strata rebuilt on each leave-out sample and the whole two-step procedure re-estimated inside every draw. Bootstrap-t limits are heavy tailed in this application; limits falling outside the plotted range are drawn as arrows. The estimator's default post window is the maximum relative year present in the data, so it changes definition from one leave-out sample to the next; those estimates are reported in the accompanying table but are not plotted here, because a coefficient plot across samples requires a common estimand.

Notes: Each row drops one state and re-estimates the expansion-adjusted specification on the states that remain, and the diamond drops Colorado, Minnesota, and Washington together. The left panel is the two-way fixed effects version with 95 percent cluster-robust intervals. The right panel is the two-step LP-DiD version at the fixed five-year post window, with 95 percent bootstrap-t intervals from the stratified state-cluster bootstrap and strata rebuilt on each leave-out sample. Bootstrap-t limits are heavy tailed here, and limits outside the plotted range are drawn as arrows. Estimates at the estimator's default post window are reported in the appendix text rather than plotted, because that window changes definition from one leave-out sample to the next. The separation years in the figure subtitle are upper bounds (Section 4.6).


Appendix Figure A28. Expansion-adjusted 1-year Medicaid estimate at alternative fixed post windows




Each row re-estimates the expansion-adjusted LP-DiD estimate on the same 212-observation, 17-state frame, changing only the post window at which the pooled long difference is taken. The first stage, the sample, the comparison pool and the state-level clustering are held fixed; the first-stage post-expansion coefficient is 0.315 in every row. Intervals are 95 percent bootstrap-t intervals, studentized by the within-draw cluster-robust standard error, from a state-cluster pairs bootstrap stratified on treatment timing and expansion status with 1,000 replications; the same replications are used at every window, so differences across rows reflect the window and not the resampling. The final column reports the two-sided bootstrap-t p-value. A treated state enters the pooled regression at a given window only if it is observed that many years after its own adoption, so widening the window trades pooled observations and identifying states for horizon. The six-year interval runs from -0.47 to 0.98, wider than the plotted range, and is drawn with arrows.

Notes: Each row re-estimates the two-step expansion-adjusted LP-DiD estimate on the same 212-observation, 17-state frame, changing only the fixed post window at which the pooled long difference is taken. Intervals are 95 percent bootstrap-t intervals from the stratified state-cluster bootstrap in Section 5.4, and the same replications are used at every window, so differences across rows reflect the window rather than the resampling. The columns to the right of the plot report the pooled observations, the number of switching states that reach the window, and the two-sided bootstrap-t p-value. The six-year interval extends beyond the plotted range and is drawn with arrows.


Appendix Figure A29. Raw employment trends by adoption cohort, 1-year horizon




Appendix Figure A30. Raw earnings trends by adoption cohort, 1-year horizon




Appendix Figure A31. Raw mortality trends by adoption cohort, 1-year horizon




Appendix Figure A32. Raw reincarceration trends by adoption cohort, 1-year horizon



Appendix Figure A33. LP-DiD event study, W-2 employment, 1-year horizon



Appendix Figure A34. LP-DiD event study, W-2 earnings, 1-year horizon



Appendix Figure A35. LP-DiD event study, mortality, 1-year horizon

