

Suspension vs. Termination: Medicaid Continuity Policies During Incarceration and Effects on Post-Release Coverage, Employment, and Mortality

Jonathan Palisoc

University of Michigan

Abstract

Objective: To estimate the effect of states suspending Medicaid enrollment during incarceration versus terminating on post-release Medicaid enrollment, employment, earnings, and mortality. Under suspension, coverage is paused during incarceration and reactivated at release without a new application; under termination, eligibility is ended entirely and the individual must reapply after release, navigating documentation requirements, processing delays, and bureaucratic complexity at a time when they are least equipped to do so - reentry. The primary outcome is Medicaid enrollment at one, three, and five years post-release.

Data Sources: Census Bureau Criminal Justice Administrative Records System (CJARS) Justice Outcomes Explorer, administratively linked aggregate outcomes for state prison release cohorts across 31 states, 2000-2020; and a novel state-year panel of suspension and termination policies.

Study Design: Staggered difference-in-differences exploiting the rollout of suspension across states between 2005 and 2022 and event study models. Outcomes are release-cohort level and therefore estimations are intent-to-treat (ITT) on the full release cohort.

Principal Findings: Suspension states saw large gains in post-release Medicaid enrollment: 22.2-25.8 percentage points at one year ($p = 0.001$) dependent on estimator. The effects seem to have long-lasting horizons (18.7 pp at three years, 19.6 pp at five years). As suspension adoption was commonly co-timed with the ACA Medicaid expansion (15 of the 25 policy-panel states that

adopted both policies did both in the same year), I also run a model with a time-varying expansion control attenuates the 1-year estimate to 11.1 pp ($p = 0.009$). I test other subgroups and all indicate significant positive effects: ever-expanded states alone (20.1 pp, $p = 0.026$), within non-expansion states alone (26.6 pp, $p = 0.060$), and dropping 2014-16 release cohorts (14.3 pp, $p = 0.046$). Gains were similar across sex, race, and age subgroups (roughly 17-24 pp), with a smaller estimate for women (17.7 pp vs. 23.4 pp for men) consistent with women's much higher baseline post-release enrollment. Conversely, I do not find significant effects on employment, earnings, and mortality, and little to no effects recidivism, although this may be attributable to power issues and the state-level nature of the data. These estimates extend prior single-state evidence (Packham and Slusky, 2024, Burns et al., 2022) to a multi-state, longer-horizon design, and they provide an implementation benchmark for the recent nationwide shift to suspension that took effect at the start of 2026.

Keywords: Medicaid, incarceration, administrative burden, reentry, difference-in-differences, CJARS

1. Introduction

Each year, more than 600,000 individuals are released from state and federal prisons in the United States. The period after release is one of critical vulnerability: the risk of overdose death is 12.7 times that of the general population (Binswanger et al., 2007), and access to health insurance is often a significant barrier. In Medicaid, a state choosing to suspend coverage versus terminate coverage is the difference between immediate coverage or the burden and wait of a brand-new Medicaid application.

The period immediately following release from prison is among the most difficult times for a high-need, high acuity population. I attempt to estimate the effects of these suspension policies on immediate Medicaid enrollment, as well as longer horizons to see if these gains persist. Over 80 percent of individuals in state prisons have at least one chronic health condition, including substance use disorders, serious mental illness, hypertension, and hepatitis C, yet the reentry process strips away the limited health infrastructure that the correctional system provides, often without any bridge to community-based care (Mallik-Kane and Visser, 2008). Whether Medicaid coverage is available on the day of release, or only after weeks or months of navigating an application process, has direct consequences for health and survival.

The Medicaid Inmate Exclusion Policy (MIEP) bars federal Medicaid payment for services delivered during incarceration but does not require states to terminate the underlying eligibility of incarcerated individuals (Section 2.1). Because federal inmates are not enrolled in state Medicaid programs, I limit my findings to state and local inmates. The statute's silence has allowed states to choose between pausing eligibility - reactivating coverage automatically at release - and terminating it; I exploit the staggered timing of this policy choice to identify its causal effect.

Understanding the effects of this policy distinction matters for a number of reasons. First, federal law now requires suspension: effective January 1, 2026, states may no longer terminate Medicaid eligibility on the basis of incarceration. The mandate substantially reduces the cross-state variation in suspension policy that made pre-mandate evaluation possible, so the historical staggered adoption analyzed here is a uniquely valuable source of evidence (Section 2.3). Alongside this, a flipped version of this policy – requiring termination of all coverage for inmates – could generate significant federal and state savings and reduce Medicaid reliance, both

of which have been noted as high priorities in the second Trump administration. In addition, understanding how suspension interacts with the emerging Section 1115 Reentry Demonstration Waivers (Section 2.6) is essential for designing comprehensive reentry coverage strategies. Lastly, the suspension-termination question tests a hypothesis central to the administrative burden literature: whether reducing friction in program access at critical life transitions can produce large improvements in take-up even when eligibility itself is unchanged.

Despite the policy's importance, the existing quasi-experimental evidence base is thin and concentrated in a single state. Gollu and Zapryanova (2022) examine recidivism under suspension in a multi-state cross-sectional design but do not examine the first stage of enrollment, Packham and Slusky (2024) provide the most direct prior evidence on enrollment from South Carolina, and Burns et al. (2022) study Wisconsin's pre-release Medicaid enrollment-assistance program (Section 3.1); I extend the frontier on all three margins: multi-state scope, longer horizons, and administratively linked secondary outcomes.

To do so, I construct a novel state-year panel of suspension and termination policies from 2000 to 2025 by triangulating six independent policy sources (KFF, MACPAC, NACo, Gollu and Zapryanova, and others) and link it to the Census Bureau's Criminal Justice Administrative Records System (CJARS) Justice Outcomes Explorer (JOE), which provides aggregate, administratively linked release-cohort outcomes for state prison systems across 31 states; I then estimate two-way fixed effects (TWFE) and a number of heterogeneity-robust difference-in-difference estimators.

This paper adds to the literature in two main ways. First, I provide the first multi-state evidence on the relationship between Medicaid suspension policies and post-release Medicaid enrollment trajectories, at one, three, and five years after release - the estimated ITT indicates suspension is

associated with an increase of roughly 22 percentage points in 1-year post-release Medicaid enrollment, with similar magnitudes at the 3- and 5-year horizons, and the estimate is robust across the three estimators and a long battery of robustness checks; the effect magnitude net of expansion is plausibly closer to 11 percentage points. Second, the paper extends the evidence base beyond enrollment to formal employment, earnings, mortality, and housing assistance, using administrative linkages to W-2 records and the Census Numident – and I find no evidence of effects on W-2 employment or earnings, and the mortality estimates are exploratory and not robust across estimators.

2. Institutional and Policy Background

2.1 The Medicaid Inmate Exclusion Policy

Medicaid is a joint federal-state health insurance program enacted under Title XIX of the Social Security Act in 1965: states administer their own Medicaid programs subject to federal requirements and receive federal matching funds. Eligibility, covered services, and provider payment rates vary substantially across states due to their flexibility in program administration, but federal statute imposes a significant set of structural constraints - one being the Medicaid Inmate Exclusion Policy (MIEP).

The MIEP was codified in Section 1905(a)(A) of the Social Security Act in 1965 and prohibits federal Medicaid reimbursement for health care services delivered to any “inmate of a public institution,” with a narrow exception for inpatient hospital stays exceeding 24 hours (MACPAC, 2018). CMS has interpreted “inmate of a public institution” to include individuals held involuntarily in state and federal prisons, county and local jails, and certain juvenile facilities; it does not include individuals on probation, parole, or home detention. The original rationale was

that states and localities, as operators of correctional facilities, bear constitutional responsibility for inmate health care under the Eighth Amendment (*Estelle v. Gamble*, 429 U.S. 97 (1976)), and Congress did not intend for Medicaid to subsidize that obligation.

The MIEP is perhaps best understood as a federal payment exclusion rather than an eligibility exclusion; Incarceration does not itself make an individual ineligible for Medicaid, but rather Federal Medicaid funds simply cannot pay for services provided while the individual is an inmate, except in limited circumstances (CMS, 2025). For decades, most states chose termination by default: state Medicaid agencies had little to no financial incentive to maintain eligibility records for individuals who could not generate federal matching funds, and corrections and Medicaid administrative systems were rarely integrated. Additionally, many inmates were not Medicaid eligible post-release before the ACA's Medicaid expansion, further discouraging coordination.

As around 600,000 individuals are released from state and federal prisons and an additional ~9 million admissions and releases occur in county and local jails (Carson and Kluckow, 2023; Zeng, 2023), the overwhelming majority of these individuals released from prison were effectively uninsured regardless of prior enrollment in the Medicaid program. Although recent improvements have been made in eligibility redeterminations, even a person who completed a Medicaid re-application on day 1 of reentry could wait as long as 30 to 90 days for a determination.

2.2 The Emergence of Suspension Policies

CMS first explicitly authorized suspension as an alternative to termination in a State Medicaid Director Letter issued on May 25, 2004. The letter clarified that states could “suspend rather than terminate” eligibility during incarceration, with coverage reactivated at release without a new

application. Suspension in this context means that the individual's Medicaid case is held in an inactive but identified status during incarceration. In summary, suspension preserves the federal payment bar under MIEP, but the underlying eligibility record is preserved, the case is automatically reactivated at release, and renewals can in principle be processed during the period. On the other hand, termination means the case is closed entirely; at release the individual must submit a new application, requalify on income and other criteria, supply documentation, and wait through a potentially long determination window.

Following the 2004 guidance, states adopted suspension in three broad waves. As of January 2025, 46 states suspend Medicaid during incarceration, 3 terminate (Alabama, Georgia, Texas), and 2 have unclear status (KFF, 2025). Figure 1 maps the adoption pattern.

Implementation has varied across states. Some adopted time-limited suspension (30 days, 90 days, or up to one year, with conversion to termination after that), while others adopted indefinite suspension for the full incarceration period. Scope also varies: some states apply suspension only to state prisons, while others extend it to county and local jails. The operational machinery required to make suspension work in practice is substantial - real-time data exchange between corrections agencies and the state Medicaid office to flag admissions, an automated process to lift the suspension flag at release, and procedures for processing renewals during the period.

Bandara et al. (2025) document that several de jure suspension states lack one or more of these components, producing a meaningful gap between the policy on paper and what an individual actually experiences at release.

2.3 The 2026 Federal Mandate

Section 205 of the Consolidated Appropriations Act, 2024 (P.L. 118-42) expanded the federal prohibition on terminating Medicaid and CHIP eligibility to all inmates of public institutions and

mandated effective January 1, 2026 that states must not terminate Medicaid or CHIP eligibility solely because an individual is an inmate of a public institution (CMS, 2025). The mandate reflects bipartisan consensus informed by advocacy from NACo, NCSL, and MACPAC; CMS guidance from December 2025 emphasizes that states retain choice among different suspension strategies, and that operationalizing the mandate requires data-sharing infrastructure between Medicaid agencies and carceral facilities, mechanisms to identify when a beneficiary becomes an inmate, and coordination on renewals during the incarceration period (CMS, 2025). While termination is no longer an option, I argue that the policy question between it and suspension is still relevant for a number of reasons, including but not limited to: (1) the potential for termination to be reintroduced, and (2) the significant variation in operationalization of suspension, which produces widely different implementation outcomes. For the last termination states and imperfect implementers alike, the estimates in this paper provide a coverage benchmark (Section 7.5).

2.4 The Administrative Burden Framework

The theoretical framework most directly applicable to the suspension-termination distinction is the administrative burden literature, defined by Herd and Moynihan (2018) as “the experience of policy implementation as onerous”; furthermore, they break burdens into three parts: learning costs (knowing that programs exist and understanding eligibility rules), compliance costs (the time and effort for paperwork and documentation), and psychological costs (stigma, stress, and the emotional toll of navigating government systems).

Termination imposes all three burden types on reentering individuals at a moment when their capacity to overcome them is at its lowest. Learning costs are high because individuals who have been incarcerated for months or years may be unaware of changes to Medicaid eligibility rules or

application procedures. Compliance costs are extreme because the documentation requirements (proof of identity, residence, and income) are precisely the things recently released individuals are least likely to have. Psychological costs compound the other two: stigma, frustration, and demoralization weigh on individuals already coping with the social isolation of incarceration. Suspension eliminates all of these burden frictions entirely: There are no learning costs, no compliance costs, and minimal psychological costs - no action, no documentation, no bureaucratic interaction. Herd and Moynihan also provide evidence that administrative burdens disproportionately affect marginalized populations, with recently incarcerated individuals being one such case: a population with compounding disadvantages (poverty, chronic illness, histories of trauma, stigmatized status) facing compounding barriers at a moment of acute life transition.

2.5 ACA Medicaid Expansion and Justice-Involved Populations

The ACA Medicaid expansion to adults earning up to 138 percent of the federal poverty level transformed coverage for justice-involved populations. Before the ACA, traditional Medicaid categories excluded the vast majority of single, non-disabled adults who make up the bulk of the incarcerated population, leaving around 80 percent uninsured after release (Albertson et al., 2020). The expansion made an estimated 80 to 90 percent of justice-involved adults eligible in expansion states (Albertson et al., 2020; Camhi, Mistak, and Wachino, 2020).

As such, it is sensible that policy timing of suspension and expansion was often concurrent, as suspension's value depends on individuals having eligibility to preserve. Before the ACA, many incarcerated individuals were not enrolled at admission, which limited suspension's beneficiary pool; expansion subsequently created a much larger share entering prison with active coverage. Research on expansion's effects on the justice-involved is positive among a number of different outcomes: Aslim et al. (2022) find that expansion reduces recidivism by 11.5 percent, and a

companion welfare analysis (Aslim, Mungan, and Yu, 2024) estimates benefit-cost ratios of \$3.45 to \$10.62 per dollar spent. Additionally, Howell et al. (2023) find an 18 percentage-point coverage increase among criminal-justice-involved adults but no significant change in treatment utilization, and Wen et al. (2017) show that expansion was associated with a 70 percent increase in Medicaid-covered buprenorphine prescriptions.

2.6 Section 1115 Reentry Demonstration Waivers

The CMS Reentry Section 1115 Demonstration opportunity, introduced in 2023, allows states to provide Medicaid-covered services to incarcerated individuals in the 30 to 90 days prior to release. As of mid-2025, 18 states had received approval, with 9 additional applications pending (KFF, 2025; CMS, 2023). These waivers and the titular policy are complementary: suspension addresses the continuity of eligibility across incarceration and provision of services immediately post-release, while the waivers ensure care initiation and provider relationships before release.

2.7 The Reentry Period as a Critical Transition

The reentry period has extensive evidence of being one of the most challenging transitions for any person: Binswanger et al. (2007) established that individuals released from Washington State prisons were 12.7 times more likely to die in the first two weeks post-release than the general population, with drug overdose as the leading cause. The barriers to health care during reentry compound each other: lack of insurance, lack of a medical home, geographic distance from facilities, inability to pay for medications, lack of transportation, stigma, and competing demands on time and attention. Health insurance does not remove all of these barriers, but it is a necessary precondition for accessing the health care system at scale, and a stabilizing force that supports success across multiple reentry domains.

On employment, Looney and Turner (2018) show that only 55 percent of formerly incarcerated individuals have any post-release earnings, and median earnings among those who do work fall below full-time minimum wage. The labor market challenges facing returning citizens suggest that health coverage stability, while important, may not be enough to produce employment gains without complementary interventions.

3. Prior Literature

3.1 Quasi-Experimental Evidence on Suspension and Reentry Coverage

Three prior quasi-experimental studies are relevant to this paper; each examines a different margin of the post-release coverage problem. Gollu and Zapryanova (2022) is the only prior multi-state study of suspension and reentry outcomes: using individual-level administrative prison release data, the authors estimate that suspension reduces the probability of returning to prison within one year by 2.91 percentage points and within three years by 4.58 percentage points, with larger effects for Black individuals and repeat offenders. However, they only estimated recidivism, without measuring Medicaid enrollment, employment, or mortality. It also lacks linkage to Medicaid enrollment files, tax records, or death records, and relies primarily on cross-sectional variation across states rather than fully exploiting within-state adoption timing. It also predates recent advances in staggered-adoption DiD (Goodman-Bacon, 2021; Callaway and Sant'Anna, 2021).

Packham and Slusky (2024) provide the most directly relevant prior evidence on the mechanism of suspension. Using individual-level criminal justice records linked to Medicaid claims data from a single state (South Carolina), they estimate that the state's combined suspension/simplified re-enrollment policy increased Medicaid enrollment within six months of

release by 11.1 percentage points, with corresponding increases in post-release health care utilization.

Burns et al. (2022), in *JAMA Network Open*, studied Wisconsin's pre-release Medicaid enrollment-assistance program, a related but distinct intervention that simplifies re-enrollment without preserving eligibility through the incarceration period. They document large increases in post-release outpatient care use after program implementation – showing that complementary friction-reduction mechanisms can increase care utilization.

3.2 Medicaid Coverage Continuity and Churning

Medicaid enrollees are a high-churning population, with typical beneficiaries covered for fewer than 10 months per year and close to 25 percent experience a coverage change within a calendar year. Churning is associated with disruptions in physician care, reduced medication adherence, increased emergency department use, and worse health. Most Medicaid disenrollment occurs not because individuals become ineligible but because they fail to complete redetermination, which is itself a form of administrative burden (Herd and Moynihan, 2018); states that adopted continuous eligibility policies have seen substantial reductions in coverage disruptions, which suggests that administrative friction rather than eligibility change is the primary driver of coverage gaps.

Suspension is, in effect, continuous eligibility across the incarceration period - analogous to policies states have pursued for other populations - while termination creates involuntary disenrollment and a forced return through the full enrollment process. Goldman and Sommers (2020) provide evidence that the ACA Medicaid expansion decreased coverage disruptions, with an estimated half a million fewer adults churning each year, potentially amplifying the value of suspension by increasing the pool of individuals with coverage to preserve.

3.3 Medicaid and Mortality

While the Oregon Health Insurance Experiment found no statistically significant mortality effects (Finkelstein et al., 2012; Baicker et al., 2013), recent quasi-experimental studies have indicated mild to moderate positive effects. Sommers, Baicker, and Epstein (2012) find that pre-ACA state expansions were associated with a 6.1 percent relative reduction in all-cause mortality, and Borgschulte and Vogler (2020) estimate that the ACA expansion reduced mortality by 3.6 percent. For justice-involved populations, the relationship may be particularly strong given the extreme post-release mortality risk: if suspension expedites Medicaid coverage and, through it, access to medication-assisted treatment and mental health services, it could reduce the acute mortality spike. However, the causal chain from coverage to mortality is long, indirect, and likely attenuated at the progression of each secondary step.

3.4 Insurance Loss, Crime, and Administrative Burden

Deza et al. (2024) study a mass Medicaid disenrollment in Tennessee in 2005 that removed coverage from 190,000 adults, finding significant increases in total crime rates, particularly non-violent crime. Studies of numerous public programs consistently find that simplifying enrollment produces meaningful increases in participation (Herd and Moynihan, 2025). In the Medicaid context, automatic and presumptive eligibility policies, ex parte renewals, and simplified application forms have all been associated with increased enrollment and reduced coverage gaps (Herd and Moynihan, 2025; Sugar et al., 2021). The Medicaid unwinding after the end of the COVID-19 continuous enrollment provision demonstrated the consequences of restoring administrative requirements: millions lost coverage when states resumed redetermination, even though many remained eligible (KFF, 2025).

3.6 Literature Gap and Paper Contribution

Read together, the review points to the gaps this paper addresses (Section 1): multi-state evidence, longer horizons, and administratively linked secondary outcomes; additionally, no prior study has formally examined whether the suspension effect is amplified by the ACA expansion (Section 2.5).

4. Data

4.1 Overview

I combine five free, publicly available data sources: (1) the CJARS Justice Outcomes Explorer (JOE) Social Safety Net module for post-release Medicaid enrollment, employment, and mortality outcomes; (2) the JOE CJARS-based module for release cohort counts and demographic composition; (3) a novel state-year suspension/termination policy panel compiled from multiple policy trackers; (4) state-level ACA Medicaid expansion dates from KFF; and (5) state-level control variables from federal statistical agencies.

4.2 CJARS Justice Outcomes Explorer

The Criminal Justice Administrative Records System (CJARS) is a joint project of the U.S. Census Bureau and the University of Michigan that collects, harmonizes, and links criminal justice administrative records to Census Bureau survey and administrative data, including tax filings (W-2), federal program participation records, and mortality data from the Census Numident (Finlay, Mueller-Smith, and Papp, 2022). CJARS is the most comprehensive effort to date to link criminal justice records to economic and health outcomes at the individual level. The Justice Outcomes Explorer (JOE) is a public data product derived from CJARS that reports aggregated outcomes for justice-involved individuals by state, release cohort year, and post-

release time horizon. The JOE Social Safety Net module provides the primary outcome variables:

- **Medicaid enrollment rate:** the share of individuals enrolled at the specified post-release horizon (1, 3, or 5 years), measured through linkage to MSIS/T-MSIS enrollment records. The linkage supports measurement in calendar years 2005-2019, so the release cohorts observed differ by horizon (see Section 4.3).
- **W-2 employment rate:** the share with any W-2 earnings at the specified horizon. This does not capture self-employment, informal employment, or off-the-books work.
- **Mean W-2 earnings:** mean W-2 earnings including individuals with zero earnings, reflecting both the extensive and intensive margins.
- **SSI receipt rate:** the share receiving Supplemental Security Income, measured through SSA linkage. Available for a more limited set of state-years.
- **HUD housing assistance rate:** the share receiving HUD housing assistance (public housing, Housing Choice Vouchers, or project-based rental assistance).
- **All-cause mortality rate:** the share deceased by the specified horizon, measured through the Census Numident file.

Each outcome is measured at 1-, 3-, and 5-year post-release horizons for state prison release cohorts from 2000 to 2020. A key feature of JOE is that it provides state-level aggregates that can be linked to policy variables without requiring restricted microdata access. The JOE CJARS-based module provides release cohort counts, recidivism rates, and demographic composition (sex, race/ethnicity, age, offense type), which I use for balance tables and composition tests.

4.3 Data Structure: Release Cohorts, Follow-Up Horizons, and Censoring

The structure of the JOE data differs from the individual-level panels used in most of the prior literature. Each observation in the analysis file is a state $\times$ release-cohort cell: for example, all individuals released from a given state's prisons in 2010. These cells are not survey tabulations or ecological correlations of independently collected series. CJARS and the Census Bureau construct them from person-level linked administrative records (corrections rosters linked person-by-person to the administrative sources in Section 4.2) and then aggregate to the state-cohort level, applying disclosure protections before publication. The analysis therefore uses aggregate data, but aggregates built from individual-level administrative linkage.

I define two separate timelines for purposes of this analysis: the first being the policy clock, or calendar time relative to the state's suspension adoption, which defines treatment status and event time in the DiD design; the second is the outcome clock, or person-level time since an individual's own release, which defines the 1-, 3-, and 5-year follow-up horizons at which JOE measures outcomes. A release cohort is followed along its own outcome clock. As an example, the 2010 release cohort has its 1-year outcome measured in 2011, its 3-year outcome in 2013, and its 5-year outcome in 2015 (Figure 4, Panel A). An "effect at the 3-year horizon" is therefore not the effect three calendar years after a policy change; it is the effect of the policy regime in place at release on an outcome measured three years into each individual's own post-release life. One more definitional point: the outcomes are point-in-time stocks at the horizon, not cumulative flows - the Medicaid enrollment rate at horizon h is the share of the release cohort enrolled at the follow-up point, not the share ever enrolled during the interval.

CJARS' MSIS/T-MSIS linkage supports Medicaid outcome measurement in calendar years 2005-2019, and a horizon- h outcome for release cohort t is observed only if $t + h$ falls within that

window. The 1-year Medicaid outcome is therefore available for release cohorts 2004-2018, the 3-year outcome for cohorts 2002-2016, and the 5-year outcome for cohorts 2000-2014 (Figure 4, Panel B). Data censoring means no 3-year outcome exists for the 2017-18 release cohorts, and no 5-year outcome exists for the 2015-18 release cohorts. Each horizon is estimated on its own cohort window, and windows are never mixed across horizons. Because the youngest cohorts are mechanically absent from the 5-year sample, cross-horizon comparisons partly reflect cohort composition; Section 6.2 addresses this with a compositional check.

4.4 Sample Construction

I construct the analysis sample in several steps. First, I download the JOE Social Safety Net and CJARS-based modules and merge the JOE data with the state-year suspension policy panel, the ACA expansion panel, and the state-level controls by state FIPS code and year. Next, I impose sample restrictions:

- I exclude state-year cells with fewer than 25 individuals to satisfy Census Bureau disclosure requirements and ensure statistical precision. JOE suppresses cells below this threshold, so in practice this restriction is applied by the data provider.
- I restrict to state prison releases (excluding jail releases) because the CJARS data primarily track state prison systems and because the policy panel is most reliable for state prison policies. Some states apply suspension only to prisons, not jails, and jail release timing is often unpredictable.
- For the primary Medicaid enrollment analysis, I restrict to the horizon-specific cohort windows implied by the MSIS/T-MSIS linkage described in Section 4.3. For employment and mortality analyses, I use the full available range (2000-2020). The Medicaid enrollment measure is point-in-time.

The resulting policy panel covers 561 state-year observations across 31 CJARS states for release cohorts 2000-2020. The outcome panels are narrower, because JOE only reports a state's release-cohort outcomes when the underlying Census Bureau linkage is built for that state. In practice this means the same set of 17 states identifies every outcome: Medicaid enrollment (212 observations at the 1-year horizon, release cohorts 2004-2018, bounded by the MSIS/T-MSIS linkage window); W-2 employment and earnings (212 observations, same window); mortality and HUD housing assistance (264 observations, available back to 2000 from the Census Numident); SSI receipt (99 observations, the narrowest); and recidivism (the only outcome with a broader 22-state footprint). The 14 CJARS states with no JOE outcome rows for the variables in this paper (AR, CA, GA, IA, ID, MD, MI, ND, NJ, NM, NV, OR, SC, VA) are excluded from the outcome regressions. Among them, Georgia is one of the two never-treated states in the policy panel, and South Carolina is the state studied directly by Packham and Slusky (2024). Because mortality, employment, and earnings draw on the same footprint, the analysis is best understood as a 17-state design, the 31-state policy panel providing representativeness context rather than identifying variation.

Representativeness of the 17-state outcome sample. The 17 states identifying the outcome regressions include 16 of the 27 CJARS suspension adopters and one of the two never-treated states (Texas; Georgia is in the 31-state policy panel but absent from the JOE outcome panels for these outcomes). Average state characteristics in the 17-state sample versus the 14 CJARS states with no JOE outcome rows are similar on observable dimensions: pre-treatment unemployment (5.4 percent vs. 5.8 percent), age-adjusted overdose mortality (14.2 vs. 13.5 per 100,000), and poverty rates (12.7 percent vs. 13.3 percent). The 17-state sample is therefore not visibly selected on baseline economic conditions or overdose intensity. What it does select on, mechanically, is

the existence of a built-out Census Bureau linkage, which correlates with state data infrastructure and CJARS partnership rather than with policy or population fundamentals. This is a representativeness caveat I revisit in Section 7.6. The 20 non-CJARS states, including Alabama, the third termination state, are excluded entirely because they have no data-sharing agreements with the Census Bureau.

Population scope. The policy panel covers 31 CJARS states over release cohorts 2000-2020. The Medicaid outcome panel, which also identifies the employment, earnings, mortality, and housing-assistance results, comprises 17 states: Arizona, Colorado, Florida, Illinois, Indiana, Kansas, Minnesota, Montana, Nebraska, New York, North Carolina, Oklahoma, Pennsylvania, Texas, Utah, Washington, and Wisconsin; the recidivism panel covers 22 states. The sample consists of state prison release cohorts, so the results generalize to prison releases in the covered states, not to jail releases, whose churn is faster, whose release timing is less predictable, and for which suspension policies often do not apply.

Estimand. Because JOE reports aggregate outcomes for all individuals released from state prison in a given state-year, the estimates in this paper are best interpreted as intent-to-treat (ITT) effects of state suspension policy on the average Medicaid enrollment rate of the full release cohort, not as individual-level treatment effects on those whose Medicaid was actually suspended. The ITT estimand is still policy-relevant, but it differs in interpretation from the individual-level effect estimated in single-state administrative-claims studies (Packham and Slusky, 2024). I cannot observe whether a given individual was enrolled in Medicaid before incarceration, whether their eligibility was actually suspended rather than allowed to lapse, or whether coverage was reactivated automatically at release. Section 6.1 bounds the implied effect on the subpopulation whose eligibility was actually preserved. If implementation gaps (Bandara

et al., 2025) weaken the link between de jure policy and de facto practice, the cohort-average ITT effect understates the effect of fully implemented suspension. Section 7.6 returns to these interpretive limits.

4.5 State Suspension/Termination Policy Panel

The most original data contribution of this study is a comprehensive state-year panel classifying each state's Medicaid eligibility policy for incarcerated individuals from 2000 to 2025. Each state-year is classified as termination, time-limited suspension, or indefinite suspension, triangulating six independent sources:

1. KFF State Health Facts (2019 and 2025 snapshots)
2. MACPAC Reports to Congress (2018 and 2023 chapters)
3. the NACo (2015) suspension/termination brief, which reported that 38 states still terminated at the time
4. Gollu and Zapryanova's (2022) appendix coding
5. the KFF (2024) State Medicaid Eligibility Issue Brief
6. Bandara et al. (2025), which provides qualitative implementation details

When sources disagree, I apply a hierarchy: (1) state legislation, (2) KFF, (3) MACPAC, (4) peer-reviewed research. I assign each state a coding confidence rating. Only one state has a low-confidence rating (sensitivity check in Section 6.10).

4.6 ACA Expansion Panel and State Controls

I code ACA expansion status from KFF. State-level controls include the unemployment rate (BLS LAUS), the age-adjusted drug overdose mortality rate (CDC WONDER), and the poverty rate (Census SAIPE). These are included in robustness specifications but not the primary model, following the convention that the main DiD estimate should capture the full policy effect.

4.7 Summary Statistics

Twenty-nine states adopted suspension during the sample period (27 in CJARS), with two never-treated states (Georgia, Texas). Within the 17-state outcome panel, Texas is the binding never-treated comparator (Section 4.4).

Table 1. Summary Statistics

	N	Mean	SD	Min	Median	Max
Panel A. Primary outcomes, 1-year horizon						
Medicaid enrollment	212	0.320	0.251	0.078	0.200	0.880
W-2 employment	212	0.536	0.074	0.350	0.540	0.680
W-2 earnings (\$)	212	6,909	1,634	3,060	6,905	11,600
Yearly Mortality	264	0.017	0.005	0.006	0.016	0.048
Panel B. Medicaid at longer horizons						
Medicaid enrollment (3-yr)	207	0.287	0.187	0.091	0.210	0.800
Medicaid enrollment (5-yr)	201	0.277	0.166	0.093	0.210	0.740
Panel C. Cohort and state characteristics						

	N	Mean	SD	Min	Median	Max
istics						
Incarceration period length (mos.)	378	25.1	6.6	11.5	24.5	41.7
Unemployment rate (%)	561	5.58	2.13	2.20	5.00	13.70
Overdose mortality rate	554	13.9	6.5	2.6	12.7	44.3
Poverty rate (%)	534	13.0	2.9	6.9	12.7	21.4

Notes: Unit of observation is the state-year. Medicaid, employment, and mortality outcomes are shares in [0, 1]. Earnings are in nominal dollars. Overdose mortality rate is CDC WONDER age-adjusted deaths per 100,000. The policy panel covers 31 CJARS states over 2000-2020; the JOE outcome panels (Medicaid, employment, earnings, mortality, HUD) are restricted to the same 17 reporting states, with Medicaid bounded further to release cohorts 2004-2018 at the 1-year horizon by the MSIS/T-MSIS linkage window (Section 4.3).

Table 1 reports descriptive statistics. The high variance of the 1-year Medicaid enrollment rate is driven by the combination of pre-suspension states with low enrollment (many termination-era observations have enrollment below 15%) and post-suspension states with high enrollment (many post-adoption observations exceed 50%). The 3-year and 5-year Medicaid enrollment rates sit slightly below the 1-year rate, consistent with some coverage attrition over longer horizons.

Table 2. Pre-Adoption Balance, Adopting vs. Never-Treated States

	Adopting states	Non-adopting states	Difference	Norm. diff.
	Mean (SD)	Mean (SD)		

Panel A.

	Adopting states	Non-adopting states	Difference	Norm. diff.
Primary outcomes (pre-treatment)				
Medicaid enrollment (1-yr)	0.210 (0.109)	0.144 (0.041)	0.066	0.80
W-2 employment (1-yr)	0.552 (0.079)	0.513 (0.063)	0.039	0.55
W-2 earnings (1-yr, \$)	6,814 (1,822)	6,867 (1,432)	-53	-0.03
Mortality (1-yr)	0.016 (0.005)	0.017 (0.004)	-0.001	-0.17
Panel B. Release-cohort characteristics				
Annual release cohort size	10,096 (7,002)	34,644 (23,043)	-24,549	-1.44
Share male	0.879 (0.027)	0.864 (0.022)	0.015	0.62
Share Black	0.302 (0.156)	0.388 (0.078)	-0.086	-0.70
Share Hispanic	0.149 (0.112)	0.133 (0.107)	0.015	0.14
Share young (15-24)	0.197 (0.050)	0.181 (0.022)	0.016	0.42
Incarceration period (months)	22.6 (5.8)	29.9 (3.5)	-7.3	-1.53
Panel C. State-year controls				
Poverty rate (%)	12.5 (3.0)	15.3 (1.9)	-2.8	-1.10
Unemployment rate (%)	5.69 (2.23)	5.60 (1.93)	0.10	0.05
Overdose mortality rate	11.8 (5.0)	13.3 (5.6)	-1.5	-0.28

Notes: Means (standard deviations) across pre-adoption state-year observations. “Adopting

states” comprises all 27 CJARS states that adopted suspension during 2005-2022, observed in

their pre-treatment years; “non-adopting states” comprises Georgia and Texas, the two never-treated states in the CJARS sample, across all years. Normalized difference is the mean difference divided by the pooled standard deviation. Release-cohort demographic shares are computed using the JOE release-cohort cohort-size variable (cohort size at the 5-year W-2 horizon) as the denominator within sex/race/age subgroup cells. Level differences are absorbed by state fixed effects in the main DiD design; the identifying assumption requires parallel trends rather than level equivalence.

The pre-adoption balance comparison (Table 2) compares adopting states in their pre-treatment years to the non-adopting states (Georgia and Texas across all years). Adopting states have higher pre-treatment Medicaid enrollment (21.0% vs. 14.4%), consistent with the possibility that states with larger Medicaid-enrolled populations were more likely to adopt suspension. The remaining baseline differences are absorbed by state fixed effects; the identifying assumption is parallel trends, assessed in Section 6.

5. Methods

5.1 Identification Strategy

I exploit the staggered adoption of Medicaid suspension policies across states between 2005 and 2022 to estimate the quasi-experimental association between suspension and post-release outcomes. The identifying variation comes from within-state, over-time changes in policy status, separated from time-invariant state characteristics (state fixed effects) and common temporal shocks (year fixed effects). State-years after suspension adoption are classified as “treated,” while state-years under termination are considered as “untreated.” I implement two complementary strategies. The first is a staggered DiD design estimating the average treatment

effect of suspension adoption. The second is a heterogeneous DiD testing whether effects are amplified in ACA expansion states. The key identifying assumption is parallel trends: absent suspension adoption, “treated” states would have followed the same outcome trend as “untreated” states. I assess plausibility through event study analyses and pre-treatment coefficient tests.

5.2 Staggered Difference-in-Differences

The standard two-way fixed effects model is:

$$Y_{it} = \alpha_i + \gamma_t + \beta \cdot \text{PostSuspension}_{it} + X_{it} \cdot \delta + \varepsilon_{it}$$

where Y_{it} = the outcome for state i in release cohort year t , α_i and γ_t = state and year fixed effects, $\text{PostSuspension}_{it}$ = a binary indicator for post-adoption years, and X_{it} = a vector of time-varying controls (in robustness specifications). Standard errors are clustered at the state level (Bertrand, Duflo, and Mullainathan, 2004). State fixed effects absorb all time-invariant differences (Medicaid generosity, incarceration practices, demographics, political orientation), and year fixed effects absorb common temporal shocks (national economic cycles, federal policy changes, secular trends). I report TWFE as a benchmark, noting that with heterogeneous treatment effects and staggered adoption, TWFE may be biased due to negative weighting (Goodman-Bacon, 2021). I diagnose the severity through the Goodman-Bacon decomposition (Appendix Figure A10).

The preferred specification uses LP-DiD (Dube et al., 2023), which is heterogeneity-robust and uses only never-treated and not-yet-treated states as comparators at each horizon. The Callaway and Sant’Anna (2021) group-time ATT estimator serves as the primary cross-check. The two estimators address the negative-weighting bias through different routes.

Texas is the only never-treated state with non-missing Medicaid outcome rows in the JOE panel (Section 4.4). As such, the cross-estimator pattern combined with LP-DiD's use of not-yet-treated comparators, and not any single estimator, is what the design leans on. Cluster-robust standard errors, the wild cluster bootstrap ($B = 999$, Rademacher weights), and design-based randomization inference (500 permutations) provide progressively more conservative inference. Randomization inference is conservative in this design and used mainly as a stress test because the number of admissible treatment-timing permutations is limited; randomly permuting adoption timing also serves as the placebo test.

5.3 Event Study Analyses

I estimate event study models to assess parallel trends and trace dynamic effects:

$$Y_{it} = \alpha_i + \gamma_t + \sum_{k \neq -1} \beta_k \cdot 1[t - E_i = k] + \varepsilon_{it}$$

where E_i = the adoption year and the reference period is $k = -1$. Pre-treatment coefficients ($k < -1$) test the parallel trends assumption (individually and jointly); post-treatment coefficients ($k \geq 0$) trace dynamic effects. I present both TWFE and LP-DiD event studies and note differences when they arise.

5.4 Heterogeneity of Treatment Effect Analyses

I interact the suspension indicator with ACA expansion status:

$$Y_{it} = \alpha_i + \gamma_t + \beta_1 \cdot \text{PostSuspension}_{it} + \beta_2 \cdot (\text{PostSuspension}_{it} \times \text{Expanded}_i) + \varepsilon_{it}$$

where β_1 = the suspension effect in non-expansion states, $\beta_1 + \beta_2$ = the total effect in expansion states, and β_2 = the differential. I supplement this with a period-split approach that estimates the TWFE model separately for the pre-ACA (release cohorts 2004-2013) and post-ACA (2014-2018) periods.

5.5 Inference, Robustness and Sensitivity Analyses

All specifications cluster standard errors at the state level. Section 5.2 describes the wild cluster bootstrap and randomization inference procedures that complement parametric inference. I also conduct the following checks, each reported in Section 6.10 unless noted: event-study pre-trends tests, randomization-inference placebos, leave-one-out estimation, policy-coding sensitivity, alternative cell-size thresholds, horizon sensitivity (Section 6.2), the Goodman-Bacon decomposition, controls sensitivity, the ACA-expansion disaggregation (Section 6.9), and subgroup and secondary-outcome panels (Sections 6.3 and 6.6).

6. Results

6.1 Main Results: Medicaid Enrollment

Table 3. Main Results: Medicaid Suspension and Post-Release Outcomes, 1-Year Horizon

	Medicaid	Medicaid (+ controls)	Employment	Earnings	Mortality
	(1)	(2)	(3)	(4)	(5)
Suspension (post-adoption)	0.222***	0.230***	0.016	315.3	0.0005
LP-DiD estimate (same outcome)	(0.051) 0.224***	(0.041) -	(0.011) 0.003	(272.3) 383.3	(0.0011) -0.0017**
Callaway-Sant'Anna estimate	(0.018) 0.258***	-	(0.029) -	(339.2) -	(0.0006) 0.0001
Cohort-size-weighted TWFE	(0.046) 0.257***	-	0.023	351.4	(0.0015) 0.0020*
Wild	(0.062) <0.001	<0.001	(0.020) 0.135	(368.5) 0.285	(0.0012) 0.667

	Medicaid	Medicaid (+ controls)	Employment	Earnings	Mortality
cluster bootstrap <i>p</i>					
State fixed effects	Yes	Yes	Yes	Yes	Yes
Year fixed effects	Yes	Yes	Yes	Yes	Yes
Time-varying controls	No	Yes	No	No	No
Observations	212	198	212	212	264
States	17	17	17	17	31
R ² (within)	0.214	0.268	0.034	0.021	0.003

Notes: Cluster-robust standard errors (clustered at the state level) in parentheses. Each column is an outcome-specific regression of the stated dependent variable on the post-suspension indicator with state and year fixed effects. Rows 1-2 report the TWFE benchmark; the remaining estimate rows report alternative estimators for the same outcome. Column (2) adds time-varying controls (state unemployment rate, age-adjusted overdose mortality rate, and poverty rate). The LP-DiD estimator (Dube et al., 2023) uses only never-treated and not-yet-treated states as comparators at each horizon. The Callaway-Sant’Anna (2021) estimator reports the aggregated group-time average treatment effect using not-yet-treated controls. The cohort-size-weighted TWFE reweights by the JOE release-cohort cohort-size variable for each outcome. Wild cluster bootstrap *p*-values use Rademacher weights, $B = 999$, with state-level clustering. Medicaid and employment outcomes are shares; earnings are in nominal dollars; mortality is a share. The Medicaid, employment, and earnings outcomes are drawn from the same 17-state JOE outcome panel, restricted to release cohorts 2004-2018 by the MSIS/T-MSIS linkage window (Section 4.3); mortality and HUD outcomes are drawn from the same 17 states with a broader 2000-2020 window. Sample sizes therefore differ across columns because of within-state year coverage rather than because more states identify mortality. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Table 3 presents the main results: 1-year post-release Medicaid enrollment rises by 22.2 percentage points under TWFE (SE = 0.051, $p < 0.001$), 22.4 pp under LP-DiD (SE = 0.018, $p = 0.001$), and 25.8 pp under the Callaway-Sant’Anna group-time average treatment effect estimator (SE = 0.046, $p < 0.001$). Adding time-varying controls to TWFE increases the estimate slightly to 23.0 pp ($p < 0.001$). The close agreement across three estimators, including two that are heterogeneity-robust, indicates that the result is not estimator-specific. Wild cluster bootstrap p-values (Section 5.2) continue to reject the null at conventional thresholds for Medicaid enrollment ($p < 0.001$).

The magnitude is large; relative to the sample mean enrollment rate of 32.0 percent, the 22 pp increase represents a 69 percent gain. Among pre-treatment non-adopting states, mean enrollment was only 14.4 percent, and relative to that baseline the 22 pp estimate implies an increase of roughly 1.5 times the baseline level. This is consistent with the administrative burden interpretation: eliminating the requirement to file a new Medicaid application at a moment of acute instability dramatically increases take-up.

Comparison with prior single-state evidence. Table 4 places the present estimates alongside the prior quasi-experimental evidence. The present estimate of 22.4 percentage points at one year, roughly twice Packham and Slusky’s (2024) South Carolina estimate (Table 4), is plausible given two design differences: (i) my outcome window is one year rather than six months, allowing more time for the reactivation channel to show up; and (ii) my estimand is the state-cohort-average ITT effect across 27 adopting states with varied implementation environments, while Packham and Slusky’s estimate reflects South Carolina’s specific implementation. The expansion-adjusted estimate (Section 6.9) converges almost exactly on the South Carolina figure.

Table 4. Comparison with Prior Quasi-Experimental Evidence

Study	Setting	Intervention	Design	Outcome	Estimate
Packham & Slusky (2024)	South Carolina (one state)	Suspension + simplified re-enrollment	RD/event-study, individual-level Medicaid claims	Enrollment within 6 months post-release	+11.1 pp
Burns et al. (2022)	Wisconsin (one state)	Pre-release Medicaid enrollment-assistance program	Pre/post within enrolled cohort	Post-release outpatient care use	Large positive effect on utilization
Gollu & Zapryanova (2022)	Multi-state	Suspension vs. termination (cross-sectional policy variation)	Individual-level recidivism, OLS with controls	Recidivism within 1 / 3 years	-2.9 pp / -4.6 pp
This paper	17-state JOE outcome panel (within 31-state CJARS policy panel), staggered adoption 2005-2022	Suspension (eligibility preservation)	Staggered DiD (TWFE / LP-DiD / CS), state-year aggregates	Medicaid enrollment at 1 / 3 / 5 years	+22.4 / +18.7 / +19.6 pp (LP-DiD); ~+11.1 pp at 1 yr net of ACA expansion

To put the effect size in the context of the broader Medicaid take-up literature: the ACA expansion increased coverage among eligible adults by roughly 10 to 15 percentage points in the typical expansion state (Sommers et al., 2012). Ex parte renewals and related renewal-simplification policies have also been associated with meaningful reductions in churn-related disenrollment (Herd and Moynihan, 2025; Sugar et al., 2021). The 22 pp raw estimate exceeds the typical coverage gains from the ACA expansion and the churn reductions under renewal simplification. This is plausible given that the target population faces substantially greater enrollment barriers than the typical Medicaid-eligible adult, and given that suspension is a more

complete elimination of administrative burden (automatic reactivation with no action required) than a simplified application or renewal process.

The Goodman-Bacon decomposition (Appendix Figure A10) confirms that the problematic “already-treated vs. later-treated” comparisons receive relatively small weight, and their estimates are positive but smaller than the clean comparisons. As such, TWFE bias is modest and the problematic comparisons attenuate rather than reverse the estimate.

Cohort-composition robustness. Because the JOE outcome is measured at the release-cohort level rather than the individual level, a remaining identification concern is whether suspension adoption coincided with shifts in who was released. The composition placebo tests in Section 6.10 show that share male, share Hispanic, and share young are stable across adoption, while share Black shifts modestly. To assess whether this composition shift inflates the headline result, I re-estimate the main Medicaid TWFE model adding release-cohort composition controls (share male, share Black, share Hispanic, share young, mean incarceration period length, and log cohort size) drawn from the JOE CJARS-based module. Composition data are non-missing for 143 of the 212 state-year observations in the Medicaid sample, so the comparison has to be made on the matched sample. Re-estimating the baseline TWFE on the matched sample (no composition controls) yields 16.4 pp (SE = 0.041, $p = 0.002$, $N = 143$). Adding the composition controls attenuates the estimate to 13.9 pp (SE = 0.043, $p = 0.007$, $N = 143$), and further adding the state-year covariates (unemployment, overdose mortality, and poverty rates) yields 15.2 pp (SE = 0.041, $p = 0.003$, $N = 130$). Two readings follow. First, the matched-sample baseline (16.4 pp) is meaningfully smaller than the full-sample baseline (22.2 pp), so part of the headline magnitude reflects state-years for which composition data are unavailable rather than the explanatory power of composition itself. Second, conditional on the matched sample, composition controls produce

a further attenuation of roughly 2.5 pp, consistent with the share-Black shift contributing modestly to, but not driving, the headline result. The composition-controlled estimate remains positive, statistically significant, and large in policy terms.

Plausibility of the 22 pp magnitude. Because the ITT estimand averages over the full release cohort, including cohort members with no Medicaid eligibility to preserve at incarceration, the implied effect on the population whose eligibility was actually preserved is mechanically larger than 22 pp. A simple accounting bound: if half of the release cohort entered incarceration with active Medicaid (consistent with post-expansion coverage rates among low-income justice-involved adults; Albertson et al., 2020), the implied effect on those whose eligibility was suspended would be on the order of 44 pp; if the share is 30 percent, roughly 73 pp. The pre-ACA-only sensitivity (Section 6.10) further anchors the magnitude: in years when fewer cohort members had Medicaid to preserve, the cohort-average ITT effect is correspondingly smaller, as the bound predicts.

6.2 Horizon Analysis: Persistence at Three and Five Years

The next question is whether the 1-year coverage gain persists as release cohorts move deeper into their post-release lives. Table 5 presents the full horizon analysis: all three estimators at all three horizons, each estimated on its own censoring-consistent cohort window (Section 4.3), together with a compositional check described below. The answer is broadly yes: the effects have not faded at longer horizons. TWFE estimates are 22.2 pp at 1 year, 19.4 pp at 3 years (SE = 0.044, $p < 0.001$), and 20.7 pp at 5 years (SE = 0.058, $p = 0.004$). LP-DiD corroborates at both longer horizons: 18.7 pp at 3 years (SE = 0.007, $p < 0.001$) and 19.6 pp at 5 years (SE = 0.018, $p < 0.001$). The Callaway-Sant'Anna estimates are 16.9 pp at 3 years (SE = 0.058, $p = 0.004$) but 11.6 pp at 5 years (SE = 0.066, $p = 0.078$), the smallest and least precise of the nine horizon

estimates. I read the 5-year Callaway-Sant’Anna estimate as a signal that the 5-year panel is the thinnest cut of the data (15 states, release cohorts 2000-2014, with the 2015-18 cohorts censored), not as evidence against persistence, since the other two estimators remain near 20 pp on the same window.

Table 5. Medicaid Enrollment by Post-Release Horizon and Estimator

Horizon (cohort window)	TWFE	LP-DiD	Callaway- Sant’Anna	N	States
1-year (2004– 2018)	0.222*** (0.051)	0.224*** (0.018)	0.258*** (0.046)	212	17
3-year (2002– 2016)	0.194*** (0.044)	0.187*** (0.007)	0.169*** (0.058)	205–207	17
5-year (2000– 2014)	0.207*** (0.058)	0.196*** (0.018)	0.116* (0.066)	200–201	15
1-year, y5- complete cohorts (2004– 2014)	0.162*** (0.040)	0.089*** (0.009)	0.187*** (0.051)	148–149	15

Notes: Cluster-robust standard errors (state level) in parentheses. Each cell is a separate regression of the horizon-specific Medicaid enrollment rate on the post-suspension indicator (TWFE with state and year fixed effects; LP-DiD per Dube et al., 2023, with never-treated and not-yet-treated comparators; Callaway-Sant’Anna, 2021, group-time ATT with not-yet-treated controls). Each horizon is estimated on its own cohort window implied by the MSIS/T-MSIS linkage (Section 4.3); N varies slightly across estimators within a horizon because of estimator-specific sample requirements. The final row re-estimates the 1-year model restricted to exactly

the state-cohort cells for which the 5-year outcome is observed (release cohorts 2004-2014, y5-linked states) - the compositional check discussed in the text. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

The event studies at the longer horizons reinforce this picture. The LP-DiD event studies for the 3- and 5-year outcomes (Figures 5 and 6) show a clear jump at adoption that persists across post-treatment lags, and the corresponding TWFE event studies (Appendix Figures A14-A15) show pre-period leads that are small and individually indistinguishable from zero at both horizons (3-year leads range from -0.010 to $+0.010$; 5-year leads from -0.037 to -0.015 , all with confidence intervals straddling zero). Raw cohort trends at the 3- and 5-year horizons (Appendix Figures A16-A17) show the same visual divergence after adoption as the 1-year trends. The distant leads of the 3-year LP-DiD event study are significantly negative ($t = -5$: -0.160 , 95% CI $[-0.269, -0.052]$; $t = -4$: -0.152 $[-0.253, -0.050]$), an amplified version of the mild downward drift visible in the distant leads of the 1-year LP-DiD figure (Figure 16), with the nearer leads shrinking toward zero ($t = -3$: -0.065 ; $t = -2$: -0.025). The 5-year distant leads are negative but not significant ($t = -5$: -0.107 $[-0.232, +0.019]$). Negative distant leads of this form mean that treated states' outcomes were below their comparators' five years before adoption. Combined with the small TWFE leads on the same samples, this is more consistent with the ACA-era secular drift documented in Section 6.9 than with anticipatory behavior by soon-to-adopt states.

Compositional check. Because the 5-year sample necessarily drops the youngest release cohorts (2015-18), horizon comparisons could confound persistence with cohort composition. The final row of Table 5 re-estimates the 1-year model on exactly the state-cohort cells where the 5-year outcome is observed (release cohorts 2004-2014, 15 states). The restricted 1-year estimates remain positive and statistically significant but are smaller than their full-sample counterparts:

TWFE 16.2 pp (SE = 0.040, $p = 0.001$), LP-DiD 8.9 pp (SE = 0.009, $p < 0.001$), and Callaway-Sant'Anna 18.7 pp (SE = 0.051, $p < 0.001$). Part of the full-sample 1-year magnitude therefore comes from the post-expansion-era cohorts that the 5-year sample cannot include, a finding that anticipates the ACA disaggregation in Section 6.9. On the common cohort window, the 5-year estimates are as large as or larger than the 1-year estimates, so the persistence conclusion is not a compositional artifact.

The persistence raises an interpretive question. If suspension's primary mechanism is removing the administrative burden of reapplication, the effect should attenuate as terminated individuals eventually overcome those barriers. Standard administrative burden theory would predict shrinking impacts over time, as individuals gradually learn the system, gather documentation, and complete applications. The fact that the effect remains roughly 20 percentage points even at 5 years requires explanation.

Three mechanisms can account for the persistence. First, suspension may operate less through temporary "barrier removal" and more through a default enrollment channel. Once an individual is enrolled, continuous eligibility provisions and automatic (*ex parte*) redetermination processes tend to keep them covered as long as they remain eligible. Under termination, the individual must complete an initial application and then successfully navigate every subsequent redetermination, each of which imposes its own burden.

Second, the persistent effect is consistent with a substantial fraction of terminated individuals never regaining Medicaid coverage, not simply enrolling late. For a population facing compounding disadvantages (no stable address, no identification documents, active substance use disorders, cognitive impairment from traumatic brain injury or untreated mental illness), the administrative barriers to a new Medicaid application may be functionally permanent rather than

merely inconvenient. Five years may simply not be long enough for much of this population to independently navigate the enrollment process, particularly without institutional programs designed to assist with reentry.

Third, the JOE data measure coverage rates for release cohorts at each horizon, not the same individuals tracked longitudinally. The 5-year estimate captures a mix of cohort-level coverage rates reflecting cumulative policy exposure across multiple release years, not necessarily the same individuals maintaining coverage for 5 consecutive years. This cross-sectional aggregation means the persistent effect partly reflects each year's new release cohort benefiting from suspension, rather than a single cohort staying enrolled indefinitely.

I check this interpretation with three quantitative tests. First, a formal attenuation test comparing the 1-year effect (22.2 pp) to the 5-year effect (20.7 pp) cannot reject the null of no attenuation ($z = 0.20$, $p = 0.838$); the 1.6 percentage-point difference is well within sampling error. Second, a stacked-horizon regression that estimates the 1-year effect and tests for differential effects at 3 and 5 years confirms this: the 3-year differential is -5.6 pp ($p = 0.335$) and the 5-year differential is -1.7 pp ($p = 0.821$), neither distinguishable from zero. Third, the raw enrollment trends (Figure 12 and Appendix Figures A16-A17) show that the gap between suspension and termination states is stable across release cohort years from 2004 through 2018 at all three horizons - the divergence does not narrow over calendar time. The data are consistent with a permanent coverage-trajectory difference rather than a temporary barrier that terminated individuals eventually overcome. The slight decline from 1 to 3 years may reflect Medicaid churning through annual redetermination (Goldman and Sommers, 2020), and the stabilization between 3 and 5 years suggests churn-induced attrition is concentrated early while a core coverage gain is maintained.

6.3 Subgroup Analyses: Sex, Race, and Age

The distributional question of who gains coverage under suspension can be addressed with JOE's published one-way subgroup cells, which disaggregate each state $\times$ release-cohort outcome by sex (male, female), race/ethnicity (White, Black, Hispanic, Asian/Pacific Islander, American Indian/Alaska Native), and age group at release (15-24, 25-39, 40+). I re-estimate the main DiD design (TWFE and LP-DiD) separately within each subgroup panel. Figure 7 presents the resulting forest plot for 1-year Medicaid enrollment; Appendix Table A7 reports the full estimates, and Appendix Figures A24-A25 present the corresponding 3- and 5-year forest plots. The headline finding is homogeneity: every subgroup with full panel coverage saw large gains, with 1-year enrollment effects in a band of roughly 17-24 pp. By race/ethnicity, LP-DiD estimates are 22.1 pp for White, 23.7 pp for Black, and 21.4 pp for Hispanic release-cohort members (TWFE: 21.4, 23.3, and 21.8 pp). By age, they are 17.3 pp for ages 15-24, 24.0 pp for ages 25-39, and 21.1 pp for ages 40+ (TWFE: 22.7, 23.4, and 20.4 pp). The subgroup-specific LP-DiD event studies (Appendix Figures A22-A23 for the sex split) show the same dynamic profile as the pooled figure: flat near-term leads and an immediate, persistent post-adoption jump. The pattern at the 3- and 5-year horizons is the same, with every full-coverage subgroup between roughly 15 and 22 pp. Suspension does not appear to concentrate its coverage gains in any demographic segment of the release population. This is consistent with the mechanism operating through the administrative default itself rather than through group-specific take-up behavior.

The one systematic difference is by sex. Women saw smaller gains than men at every horizon (LP-DiD 17.7 pp, SE = 0.014, vs. 23.4 pp, SE = 0.021, at one year; TWFE 15.9 pp vs. 23.1 pp, SE = 0.039 and 0.054). A ceiling interpretation fits the levels: in the subgroup panel, women's

mean 1-year post-release enrollment rate is 0.515, versus 0.292 for men (0.409 versus 0.159 in pre-adoption state-years), reflecting women's greater pre-incarceration attachment to Medicaid through pregnancy-related and parental eligibility pathways. With roughly half of released women already enrolled at the 1-year point, there is simply less headroom for an administrative default to add coverage. The proportional gains are, if anything, larger for men.

Two subgroups cannot support reliable inference: Asian/Pacific Islander and American Indian/Alaska Native panels lose 23-26 percent of state-cohort cells to Census Bureau disclosure suppression (cells with fewer than 25 individuals are withheld), and the suppressed cells are a non-random set of small cohorts. Within the observable cells, the 1-year TWFE estimates are positive and of similar magnitude to the full-coverage groups (Asian/Pacific Islander 17.6 pp, $p = 0.011$; American Indian/Alaska Native 17.6 pp, $p = 0.005$), but the LP-DiD estimator either fails outright or returns degenerate standard errors at the 1-year horizon in these sparse panels. I report these estimates for completeness only. The suppression pattern also highlights the continued need for disaggregated data on these groups.

Lastly, subgroup recidivism estimates (Appendix Table A7 and Appendix Figure A26) are null across the board, mirroring the pooled recidivism result in Section 6.6: every subgroup's TWFE estimate is small (between -3.1 and -1.0 pp) and statistically indistinguishable from zero. Two scattered LP-DiD cells reach nominal significance (female -0.7 pp, $p = 0.019$; Hispanic $+3.2$ pp, $p < 0.001$), but they are small, inconsistent with their TWFE counterparts (the Hispanic TWFE estimate is -2.9 pp, opposite in sign), and not robust across estimators.

6.4 Mortality

Mortality estimates are inconclusive across estimators, and the design is underpowered to resolve them. TWFE returns a near-zero estimate (0.0005, $SE = 0.001$, $p = 0.650$), LP-DiD a small

reduction (-0.17 pp, $SE = 0.0006$, $p = 0.029$), Callaway-Sant'Anna essentially zero (0.0001 , $SE = 0.002$), and the cohort-size-weighted TWFE points in the opposite direction ($+0.0020$, $SE = 0.001$). The wild cluster bootstrap p-value (0.667) and the controls-augmented TWFE ($p = 0.553$) further indicate that no single sign or magnitude is robust.

Against a baseline 1-year mortality rate of 1.65 percent in pre-treatment cells, the minimum detectable effect at 80 percent power is around 0.30 percentage points (~ 18 percent of baseline). The design therefore has reasonable power to detect mortality reductions of roughly 18 percent or larger, but it is underpowered for the 6 percent (Sommers, Baicker, and Epstein, 2012) or 3.6 percent (Borgschulte and Vogler, 2020) reductions documented in the broader Medicaid-mortality literature. A TOST equivalence test against a ± 3.6 percent baseline-relative bound does not reject equivalence-failure ($p_{\text{max}} = 0.46$). I therefore treat mortality as exploratory and unresolved: the design can rule out very large effects but cannot adjudicate whether suspension produces the modest mortality gains that the broader literature would predict. Resolving this requires individual-level data with greater statistical power. Section 7.4 develops the “coverage is necessary but not sufficient” interpretation.

6.5 Employment and Earnings

Employment and earnings also see no significant movement: TWFE estimates for 1-year W-2 employment are 1.6 pp ($SE = 0.011$, $p = 0.148$), and LP-DiD estimates are essentially zero (0.3 pp, $SE = 0.029$, $p = 0.916$). Earnings estimates are positive but noisy (TWFE: $\$315$, $SE = \$272$, $p = 0.264$; LP-DiD: $\$383$, $SE = \$339$, $p = 0.341$). The null result is consistent across 1-, 3-, and 5-year horizons (Appendix Table A5).

The null employment result is still informative. It suggests suspension operates primarily through health coverage, its intended channel, rather than through broader economic stabilization - and

that enrollment is not displacing employment (no work-disincentive effect). A few mechanisms explain why a large coverage gain does not translate into employment effects. First, the employment barriers facing returning citizens are not barriers health insurance addresses; Section 7.4 develops this necessary-but-not-sufficient reading. Second, W-2 employment captures only formal-sector work; if suspension-induced coverage enables pursuit of lower-paying formal employment or channels individuals into SSI, net effects could be attenuated. Third, employment effects may emerge at horizons beyond the 5-year window if coverage access enables treatment for chronic conditions that eventually improve work capacity. Lastly, state-level data may simply lack the power to detect modest employment effects: a 1-2 pp effect would be substantively meaningful but hard to detect with 212 clustered observations.

6.6 Recidivism

Recidivism is the outcome for which the prior literature provides the strongest priors. Gollu and Zapryanova (2022) find recidivism reductions in individual-level data (Section 3.1), and the broader Medicaid-and-crime literature (Aslim et al., 2022; Deza et al., 2024) predicts that coverage gains should reduce reoffending. A measurement caveat has to come before the estimates: JOE recidivism outcomes record a new criminal-justice contact only when it occurs in the same state as the index release. Out-of-state reoffending is unobserved, so measured recidivism rates are lower bounds on true recidivism, and the estimates below are informative about the same-state recidivism margin. Differential cross-state migration by policy regime would threaten even that interpretation, a possibility I return to in Section 7.6.

The reincarceration outcome has the broadest coverage footprint in JOE: 22 states and 384 state-cohort observations at the 1-year horizon (release cohorts 2000-2020), with 21 states at the 3-

year (N = 345, cohorts 2000-2019) and 5-year (N = 302-303, cohorts 2000-2017) horizons. Table 6 reports the estimates.

Table 6. Reincarceration and Recharge Outcomes by Horizon and Estimator

Outcome, horizon (cohort window)	TWFE	LP-DiD	Callaway- Sant'Anna	N	States
Reincarceration, 1-yr (2000–2020)	–0.022 (0.027)	–0.007 (0.007)	–0.034** (0.016)	384	22
Reincarceration, 3-yr (2000–2019)	–0.024 (0.026)	0.006 (0.008)	–0.021 (0.015)	345	21
Reincarceration, 5-yr (2000–2017)	–0.010 (0.020)	0.019* (0.010)	0.019* (0.011)	302–303	21
Felony recharge, 1-yr (2000–2020)	0.009 (0.018)	n.e.	–0.011 (0.014)	129	8
Misdemeanor recharge, 1-yr (2000–2020)	0.046** (0.017)	n.e.	0.002 (0.027)	116	8

Notes: Cluster-robust standard errors (state level) in parentheses. Each cell is a separate regression of the stated recidivism rate on the post-suspension indicator (estimators as in Table 5). JOE recidivism outcomes capture only same-state criminal-justice contact, so all rates are lower bounds on true recidivism. “n.e.” = not estimable: in the 8-state charge-type subsample, the LP-DiD estimator has insufficient clean-control variation. The felony and misdemeanor

recharge outcomes are observed in only 8 states (fewer than 10 clusters), so state-clustered inference for those rows is fragile and the estimates should be treated as suggestive at most.

* $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

The results are best characterized as a null that is precise enough to be informative at the state-cohort level. At the 1-year horizon, TWFE yields -2.2 pp (SE = 0.027, $p = 0.43$) and LP-DiD -0.7 pp (SE = 0.007, $p = 0.32$); the only nominally significant estimate is Callaway-Sant'Anna at -3.4 pp (SE = 0.016, $p = 0.033$). At 3 years the estimates are -2.4 pp (TWFE, $p = 0.36$), $+0.6$ pp (LP-DiD, $p = 0.52$), and -2.1 pp (Callaway-Sant'Anna, $p = 0.15$). By 5 years the Callaway-Sant'Anna estimate attenuates fully and flips sign ($+1.9$ pp, $p = 0.078$), as does LP-DiD ($+1.9$ pp, $p = 0.092$). The event studies show no dynamic pattern hiding behind the static nulls: the 1-year LP-DiD leads are flat (ranging from -0.003 to $+0.021$, all confidence intervals straddling zero) and the post-adoption path shows no sustained decline (Figure 9; Appendix Figures A18-A20 present the TWFE version and the 3- and 5-year horizons). The raw trends show no visible divergence at adoption (Figure 8).

Two readings of this null are worth separating. First, in terms of what the design can exclude: the 1-year TWFE confidence interval spans roughly -7.4 to $+3.1$ pp, so same-state reincarceration reductions larger than about 7 pp are ruled out, while the Gollu-Zapryanova individual-level estimates (-2.9 pp at 1 year, -4.6 pp at 3 years) sit comfortably inside the interval. The state-cohort data therefore neither confirm nor refute the prior individual-level findings; they establish that any true effect is modest in absolute terms. Second, in terms of mechanism: a coverage gain of 11-22 pp (Section 6.9) does not translate into a detectable reduction in same-state reincarceration at this level of aggregation. This is consistent with the “coverage is necessary but not sufficient” interpretation developed in Section 7.4 - the behavioral chain from coverage to

treatment engagement to desistance is long, and cohort-level aggregates are a blunt instrument for detecting it.

The charge-type outcomes (felony and misdemeanor recharge at 1 year) are available in only 8 states, below the threshold at which state-clustered inference is reliable. The felony estimates are null across estimators (TWFE +0.9 pp, $p = 0.64$; Callaway-Sant'Anna -1.1 pp, $p = 0.41$; LP-DiD not estimable). The misdemeanor TWFE estimate is positive and nominally significant (+4.6 pp, $SE = 0.017$, $p = 0.030$), but it is not robust to the Callaway-Sant'Anna estimator (+0.2 pp, $p = 0.93$) and rests on fewer than 10 clusters. I do not interpret it as evidence of a misdemeanor-increasing effect. Appendix Table A8 reports the full charge-type detail. The subgroup recidivism estimates reported in Section 6.3 are null across sex, race, and age groups, so the pooled null is not masking offsetting subgroup effects.

6.7 Other Secondary Outcomes: Housing Assistance and SSI

HUD housing assistance shows a small TWFE estimate (0.1 pp, $SE = 0.001$, $p = 0.310$) but a larger LP-DiD estimate (0.8 pp, $p < 0.001$). The LP-DiD result suggests Medicaid coverage may facilitate housing program access, potentially because enrollment establishes a documented connection to the safety net. However, the TWFE-LP-DiD divergence warrants caution.

SSI receipt shows no effect (TWFE: -0.03 pp, $p = 0.909$; $N = 97$). The small sample severely constrains detection, and SSI is a slow-moving outcome: application and determination take months to years.

6.8 Heterogeneous DiD by ACA Expansion

I estimate two complementary specifications. The first, preferred for interpretation, uses time-invariant expansion status: the coefficient on the suspension indicator captures the suspension effect in non-expansion states, and the coefficient on the interaction with ever-expanded status

captures the differential effect for expansion states. The second uses a time-varying expansion indicator equal to one in state-years with the expansion in effect, which more cleanly separates the within-state activation of expansion from the cross-state classification.

Static-expansion interaction (preferred). The suspension-expansion interaction yields a positive Medicaid enrollment coefficient (0.195, SE = 0.132, $p = 0.158$), implying a larger total effect in expansion states ($0.075 + 0.195 = 0.270$) than in non-expansion states (0.075). This is consistent with the eligibility-to-preserve hypothesis (Section 2.5), but neither coefficient achieves conventional significance, likely reflecting the limited power with only two never-treated states, both non-expansion: the interaction is identified primarily by comparing the within-state effect of suspension across expansion and non-expansion adopting states, without the clean between-group comparison that more never-treated states would provide.

Time-varying expansion interaction (identification limit). Replacing the static expansion status with the time-varying indicator exposes a fundamental identification limit in the Medicaid sample. Among the 17 states with non-missing Medicaid 1-year outcomes (release cohorts 2004-2018), no state-year cell is both pre-suspension and ACA-active. Every state that adopted both suspension and the expansion did so in an order that puts ACA activation either before suspension adoption or inside the post-suspension window. The interaction between post-suspension and time-varying expansion is therefore collinear with the expansion indicator once state and year fixed effects are absorbed, and the interaction coefficient is not separately identified. This is itself an informative finding: in the available sample, the policy combination of “ACA expansion, but not yet suspension” is essentially absent, so the heterogeneous-DiD interaction cannot be cleanly identified from within-state activation timing. The static-expansion

specification remains identifiable because it cross-classifies states by expansion status rather than by within-state expansion timing.

The period-split analysis shows TWFE estimates of 11.7 pp (SE = 0.050, $p = 0.035$) pre-ACA (release cohorts 2004-2013, $N = 134$) and 14.7 pp (SE = 0.081, $p = 0.089$) post-ACA (release cohorts 2014-2018, $N = 78$). The 26 percent larger post-ACA point estimate is consistent with the expansion hypothesis, though the difference is not statistically significant.

For other outcomes, the interaction patterns are less informative. Employment shows a positive base effect in non-expansion states (2.4 pp, $p < 0.10$) with a negative interaction (-1.0 pp, $p > 0.10$), suggesting no amplification in expansion states. For mortality, the base effect in non-expansion states is positive (0.18 pp, $p < 0.10$) with a negative interaction (-0.16 pp, $p > 0.10$). These patterns are hard to interpret and may reflect confounding or small-sample artifacts rather than genuine heterogeneity.

6.9 Disaggregating the Suspension Effect from the ACA Medicaid Expansion

Figure 10 plots each state's suspension adoption year against its expansion year. In the full 31-state policy panel, 25 states adopted both policies, the correlation between adoption year and expansion year is 0.698, and 15 of the 25 dual-policy states adopted suspension in the same year they expanded Medicaid, overwhelmingly the 2014-15 expansion wave. The identifying variation net of expansion therefore comes from the states where the two policies are not co-timed: 7 states that adopted suspension before 2014 and 6 adopters that never expanded within the panel window. Within the 17-state Medicaid outcome panel, the corresponding counts are 7 of 13 dual-policy states co-timed, 5 pre-2014 adopters, and 5 never-expanded adopters, plus Texas (never-treated, non-expansion).

Table 7. Disaggregating Suspension from ACA Expansion: Medicaid Enrollment

Specification	1-year	3-year
Panel A. Time-varying expansion control		
TWFE, baseline (no control)	0.222*** (0.051)	0.194*** (0.044)
TWFE + post-expansion control	0.111 (0.069)	0.146*** (0.048)
<i>post-expansion control coefficient</i>	<i>0.315 (0.089)</i>	<i>0.132 (0.072)</i>
LP-DiD, ACA-residualized outcome (two-step)	0.111*** (0.018)	0.154*** (0.007)
Panel B. Split samples		
TWFE, ever-expanded states only (11 states)	0.201** (0.077)	0.151** (0.059)
LP-DiD, ever-expanded states only	0.225** (0.033)	0.180*** (0.019)
TWFE, non-expansion states only (6 states)	0.266* (0.110)	0.185*** (0.009)
Panel C. Drop 2014–16 release cohorts		
TWFE	0.143** (0.066)	0.194*** (0.045)
LP-DiD	n.e.	0.089*** (0.016)
Panel D. Triple difference (from Section 6.8)		
Suspension main effect	0.075 (0.130)	-
Suspension $\times$ ever-expanded	0.195 (0.132)	-

Notes: Cluster-robust standard errors (state level) in parentheses. Panel A adds a time-varying indicator equal to one in state-years with the ACA expansion in effect (state-specific expansion dates); the LP-DiD row residualizes the outcome on the post-expansion indicator (using the jointly estimated TWFE coefficient) before applying the LP-DiD estimator. Panel B splits the Medicaid panel by ever-expanded status. The ever-expanded split contains no never-treated state, so it identifies from adoption-timing variation among eventual adopters only; the non-expansion split rests on 6 states with 8 treated post-period observations at the 1-year horizon (3 at the 3-year horizon) and is underpowered, and its LP-DiD counterpart returns a degenerate (numerically zero) clustered standard error at the 1-year horizon, effectively a single clean

comparison, and is omitted (the 3-year non-expansion LP-DiD estimate, -0.019 , $SE = 0.007$, is reported in the replication tables and is equally underpowered). Panel C drops the 2014-16 release cohorts; the 1-year LP-DiD is not estimable because the long differences span the resulting gap in the panel. Panel D reprints the static interaction from Section 6.8 (1-year outcome only). $*p < 0.10$, $**p < 0.05$, $***p < 0.01$.

Time-varying expansion control. Adding a state-specific time-varying post-expansion indicator to the 1-year TWFE model cuts the suspension coefficient in half, from 0.222 to 0.111 ($SE = 0.069$, $p = 0.126$), with the expansion control itself absorbing a large 0.315 ($SE = 0.089$). The two-step LP-DiD on the ACA-residualized outcome yields an almost identical 0.111 ($SE = 0.018$, $p = 0.009$) - the same magnitude, but statistically significant because the LP-DiD comparison structure is more efficient. At the 3-year horizon the attenuation is milder: 0.194 to 0.146 (TWFE, $p = 0.008$) and 0.154 (residualized LP-DiD, $p < 0.001$). Adding the expansion control also cleans up the event-study pre-period: the distant leads of the baseline 1-year TWFE event study ($t = -5$: 0.277 , 95% CI [0.115 , 0.439]) flatten to insignificance once the post-expansion control is added ($t = -5$: 0.163 [-0.042 , 0.368]; Figure 11), indicating that much of the visible TWFE pre-trend was ACA-driven drift rather than suspension anticipation. The ACA-residualized LP-DiD event-study leads are likewise all statistically indistinguishable from zero (Appendix Figure A21).

Split samples. Among ever-expanded states alone (where identification comes only from adoption-timing variation among eventual adopters, with no never-treated state in the split), the 1-year effect is 20.1 pp (TWFE, $p = 0.026$) and 22.5 pp (LP-DiD, $p = 0.021$). Among the six non-expansion states alone, the TWFE estimate is actually larger, 26.6 pp ($p = 0.060$), but this split rests on eight treated post-period observations and is underpowered; I report it for

transparency rather than weight. The 3-year splits tell the same story (ever-expanded: TWFE 15.1 pp, LP-DiD 18.0 pp; non-expansion TWFE 18.5 pp on three treated post-period observations). That the effect appears within both subsamples, including states where no expansion ever occurred, is the strongest evidence that suspension has a coverage effect that is not reducible to the ACA.

Dropping co-timed cohorts. Excluding the 2014-16 release cohorts, the window in which most co-timed adoptions occurred, yields a 1-year TWFE estimate of 14.3 pp (SE = 0.066, $p = 0.046$), between the raw and expansion-controlled figures. The 3-year TWFE estimate is essentially unchanged at 19.4 pp ($p < 0.001$), while the 3-year LP-DiD falls to 8.9 pp ($p < 0.001$), with the caveat that LP-DiD long differences spanning the dropped years make this specification conservative.

Across every specification in Table 7, the suspension effect remains positive; in no cut of the data does disaggregating the ACA expansion eliminate it. But the magnitude is a range, not a point. The expansion-controlled estimates center on roughly 11 pp at the 1-year horizon (and 15 pp at three years), while the raw estimates center on 22 pp (19 pp). Within the co-timed states the two policies cannot be fully separated; the separation rests on the non-co-timed adopters identified above, both groups showing clear positive effects. I argue that the suspension effect net of expansion is plausibly closer to 11 pp than to 22 pp - very close to both the pre-ACA-only estimate (11.7 pp, Section 6.8) and Packham and Slusky's (2024) South Carolina estimate (11.1 pp) - and that the remaining gains in the raw magnitude came primarily through the interaction of preserved eligibility with the expansion-era increase in the share of the release cohort that had Medicaid to preserve.

6.10 Robustness

I present a comprehensive set of robustness and sensitivity analyses that, overall, provide strong support for the main Medicaid enrollment finding. Appendix Table A6 summarizes the key results.

Raw cohort trends. Figures 12-15 plot raw outcome trends by adoption cohort (early, mid, and late adopters versus never-treated states) over calendar time for each primary outcome. For Medicaid enrollment (Figure 12), treated cohorts diverge sharply upward after their respective adoption dates, while the never-treated states remain flat - visual confirmation of the treatment effect before any regression adjustment. The raw trends for employment (Figure 13) and earnings (Figure 14) show no comparable divergence, consistent with the null findings for these outcomes. The raw mortality trends (Figure 15) are noisy and show no clear pattern of divergence, consistent with the mixed mortality results.

LP-DiD event studies. Figures 16-19 present LP-DiD event study estimates. For Medicaid enrollment (Figure 16), pre-treatment leads ($t = -5$ through $t = -2$) are small and centered near zero, supporting the parallel trends assumption. A clear positive jump at treatment onset persists through the post-treatment lags at roughly 20-25 pp, consistent with an immediate and lasting effect. The employment (Figure 17), earnings (Figure 18), and mortality (Figure 19) event studies show near-zero pre-treatment leads and post-treatment coefficients consistent with the static regression findings.

LP-DiD pre-trends are substantially cleaner than the TWFE event studies (Appendix Figures A1-A4), consistent with the decomposition evidence that TWFE pre-trend contamination arises from problematic already-treated comparisons (Appendix Figure A10). This supports the choice of LP-DiD as the primary estimator.

Goodman-Bacon decomposition. See Appendix Figure A10, discussed in Section 6.1.

Placebo tests. The TWFE randomization-inference p-value for Medicaid enrollment is 0.062 (500 permutations), indicating that the actual estimate lies near the extreme of the placebo distribution but does not clear conventional significance thresholds in this design-based exercise. The TWFE placebo mean is 0.122, reflecting some mechanical association from the staggered structure, while the actual TWFE estimate is 0.222. The LP-DiD randomization-inference exercise is more conservative: the design-based p-value is 0.153, with a placebo mean near zero (0.007; Appendix Figures A7 and A13).

Leave-one-out. The Medicaid enrollment estimate ranges from 0.206 to 0.259 when dropping each state in turn (maximum deviation: 0.036 from the full-sample estimate of 0.222). No single state drives the result (Appendix Figure A8).

Policy coding sensitivity. Dropping the one state with a low-confidence policy coding yields an estimate of 0.209, virtually unchanged from the full-sample estimate of 0.222 (a change of 0.013, or 6%).

Alternative cell size thresholds. Using a minimum cell size of 50 (rather than 25) yields an estimate of 0.222, identical to baseline. All cells in the Medicaid enrollment sample exceed the 50-person threshold, so this check is confirmatory.

Controls sensitivity. Adding time-varying controls (unemployment rate, overdose mortality rate, poverty rate) to TWFE moves the estimate from 0.222 to 0.230. That adding controls moves the estimate away from zero rather than toward it suggests that observable confounders, if anything, attenuate the uncontrolled estimate.

Coefficient comparison. Appendix Figure A6 shows estimates tightly clustered around 22-23 pp across all specifications.

Cohort-size weighting. Reweighting TWFE by the JOE release-cohort cohort-size variable yields 25.7 pp (SE 6.2 pp), slightly larger than the unweighted 22.2 pp. That the population-weighted estimate is at least as large as the unweighted state-average suggests the effect is not concentrated in small idiosyncratic states.

Wild cluster bootstrap. Bootstrap p-values (Section 5.2) are virtually identical to their cluster-robust counterparts across the outcome battery (Table 3; Appendix Table A6), confirming the inference is not an artifact of finite-cluster asymptotics.

Dropping Texas, the binding never-treated state. Dropping Texas attenuates the estimate from 0.222 to 0.209 (SE 0.052, $p = 0.001$). The modest attenuation underscores why the design leans on LP-DiD's not-yet-treated comparators rather than Texas alone (Section 5.2).

Pre-ACA-only restriction. See the period split in Section 6.8: the pre-ACA estimate (11.7 pp) is roughly half the full-sample 22 pp but remains positive and significant, so the result is not entirely an artifact of the post-ACA period (Section 6.9).

Composition placebo outcomes. I test whether suspension adoption is associated with discontinuous changes in pre-determined release-cohort characteristics (sex, race, age, Hispanic share). The estimates on share male (0.003, $p = 0.52$), share Hispanic (-0.009 , $p = 0.31$), and share young (-0.004 , $p = 0.49$) are small and indistinguishable from zero. The estimate on share Black is 0.024 ($p = 0.03$), pointing to a small post-adoption shift in cohort racial composition that may reflect concurrent criminal-justice reforms in adopting states. This is a modest concern that I revisit in the Section 7.6 limitations.

Sharpened q-values across the outcome battery. Anderson (2008) sharpened false-discovery-rate q-values across the full outcome battery (Medicaid, employment, earnings, and mortality at 1/3/5 years; HUD and SSI at 1 year) leave Medicaid 1-year, 3-year, and 5-year all at $q < 0.02$; no

other outcome survives multiple-testing adjustment at 5 percent. This is consistent with Medicaid enrollment being the only outcome the design has clear power to detect.

7. Discussion

7.1 Summary of Findings

Suspension states saw large gains in post-release Medicaid enrollment: 22.4 percentage points at one year (LP-DiD, $p = 0.001$), with similar estimates from TWFE (22.2 pp) and Callaway-Sant'Anna (25.8 pp). The effects have not faded at longer horizons (18.7 pp at three years, 19.6 pp at five years). However, suspension adoption was substantially co-timed with the ACA Medicaid expansion (15 of the 25 policy-panel states that adopted both policies did so in the same year), so I disaggregate the two policies directly. Adding a time-varying expansion control attenuates the 1-year estimate to 11.1 pp (TWFE, $p = 0.126$; ACA-residualized LP-DiD 11.1 pp, $p = 0.009$), while the effect remains positive within ever-expanded states alone, within non-expansion states alone, and when the 2014-16 release cohorts are dropped. I argue the enrollment effects of suspension are plausibly closer to 11 pp than to 22 pp due to expansion co-timing. Gains were similar across sex, race, and age subgroups (roughly 17-24 pp), with a smaller estimate for women consistent with women's much higher baseline post-release enrollment (Section 6.3). Employment and earnings have not kept up with the coverage gains: I find no significant effects on either. Reincarceration effects are null at the state-cohort level, precise enough to rule out reductions larger than about 7 pp at one year but consistent with the modest effects in prior individual-level work (Section 6.6), and mortality estimates are inconclusive across estimators and underpowered for the 3-6 percent reductions in the broader Medicaid-mortality literature.

7.2 Interpretation

The results suggest that the administrative channel (the ease of maintaining through coverage reduction of administrative burden) is as important as the eligibility channel for this population.

The enrollment increase (roughly 22 pp raw, and roughly 11 pp net of the co-timed ACA expansion; Section 6.9) is among the largest effects documented in the administrative burden literature even at the adjusted end of the range, consistent with the prediction that friction reduction produces the largest effects where barriers are greatest. Suspension eliminates all three burden types (Section 2.4). The completeness of this burden reduction, combined with the extreme vulnerability of the target population, likely explains the magnitude.

The persistence of the effect at the 3- and 5-year horizons requires extending the standard framework (Section 6.2 lays out the candidate mechanisms). The persistent effect is better understood not as a single barrier removal but as placement onto a self-sustaining coverage trajectory versus exclusion from one.

The framework also helps explain why the effect is so large relative to other coverage interventions. The ACA expansion changed eligibility; suspension changes access to existing eligibility. For a population where the barriers to completing an application are close to insurmountable (Section 2.4), eliminating the application requirement produces a larger coverage gain than expanding who is eligible.

The results also speak to the broader debate about whether administrative burdens are rational disincentives (which filter out individuals who do not value the benefit enough) or arbitrary barriers (which exclude individuals who are fully eligible and would benefit). An enrollment gain of 11-22 pp suggests that the large majority of individuals deterred by the termination-reapplication process are, in fact, eligible for and would benefit from Medicaid coverage.

Termination does not filter out ineligible individuals; it excludes eligible individuals who cannot overcome administrative barriers. This is consistent with Herd and Moynihan's (2025) argument that administrative burdens in safety-net programs mostly function as exclusionary mechanisms rather than as efficient screening devices.

7.3 Relation to Prior Literature

The Medicaid enrollment finding complements and extends the prior literature in a few ways. Most directly, it fills the mechanism gap in Gollu and Zapryanova (2022). Their finding that suspension reduces recidivism is consistent with a causal chain running from suspension through Medicaid enrollment through health care access through improved health and behavior. Without observing Medicaid enrollment, they could not confirm the first link in this chain. This paper provides that confirmation: suspension does increase Medicaid enrollment, and by a large magnitude. My recidivism estimates are directionally consistent with their findings at the shorter horizons, but the state-cohort design cannot reject zero, and the fair characterization is a precise-enough null (Section 6.6). Individual-level data provide the precision for binary reoffending outcomes that state-level aggregates cannot.

The null employment result is consistent with other evidence that insurance gains alone do not guarantee utilization changes (Howell et al., 2023); I develop this pattern in Section 7.4.

On mortality, the estimates leave unresolved the hypothesis motivated in Section 3.3 - that coverage in the uniquely high-risk post-release window could have outsized benefits; it remains a question for future investigation rather than an established result (Section 6.4).

7.4 Coverage Is Necessary But Not Sufficient

A central interpretive point from this paper, building on the broader Medicaid-mortality literature, is that strong evidence on coverage take-up does not, on its own, settle whether

coverage translates into better health and economic outcomes. The suspension-enrollment effect estimated here (11-22 pp at one year depending on how the co-timed ACA expansion is handled, persistent at five years) is among the largest documented in the administrative burden literature, yet effects on W-2 employment and earnings are null, same-state reincarceration effects are a precise-enough null, and effects on 1-year all-cause mortality are inconclusive across estimators. This pattern fits three observations.

First, the post-release period is dominated by acute mortality risk in which the speed of access matters more than the fact of coverage. Overdose deaths in the first weeks post-release require buprenorphine or methadone initiation in those weeks. Insurance reactivation that arrives even a few weeks late, or that arrives without a coordinated handoff to a treating clinician, may not avert the death. Bandara et al. (2025) document that even where suspension is the *de jure* policy, *de facto* reactivation is often delayed and rarely accompanied by warm-handoff care coordination, and that fewer than half of states with suspension policies have data-sharing infrastructure adequate to notify Medicaid agencies of an individual's release in real time.

Second, the broader Medicaid-mortality literature (Sommers et al., 2012; Borgschulte and Vogler, 2020) typically estimates aggregate mortality reductions of 3-6 percent following the ACA expansion among general adult populations over horizons of several years. The justice-involved population at one to five years post-release faces both higher baseline mortality (1.7 percent vs. ~0.5 percent among general adult populations) and a different mix of acute risks, and the design's minimum detectable effect is too coarse to settle the question (Section 6.4).

Third, the null employment finding has the same character: insurance coverage may be necessary to prevent health crises that disrupt work, but it is not sufficient to overcome the labor-market barriers facing returning citizens - employer discrimination, occupational licensing restrictions,

parole compliance demands, ban-the-box variation, and skills atrophy. These barriers are not addressed by health insurance and operate through entirely separate channels that this paper does not study.

I argue that while suspension solves the coverage problem at reentry, it does not solve the care continuity problem, the MOUD initiation problem, or the labor-market reintegration problem.

Policy and research should treat suspension as a foundational layer on which additional interventions (Section 1115 reentry waivers, MOUD initiation before release, naloxone distribution, transitional employment programs, ban-the-box) can build, rather than as a stand-alone intervention sufficient to close the post-release health and economic gap.

7.5 Policy Implications

Implementation of the 2026 mandate. The findings support the policy rationale behind Section 205 of the Consolidated Appropriations Act, 2024: preserving Medicaid coverage during the incarceration-to-release transition yields large enrollment gains and is likely a foundational input to any successful reentry-coverage strategy. With the federal mandate now in effect (January 1, 2026), the central implementation question is no longer whether states formally suspend rather than terminate, but whether they operationalize suspension in ways that produce timely reactivation, successful renewals during incarceration, and connection to care at release. For the states that were termination states as of late 2025 (Alabama, Georgia, Texas) and for any state with incomplete de facto implementation, these estimates provide a practical benchmark for the coverage gains pre-mandate adoption produced: roughly 11-22 pp ITT at one year, where the lower end nets out the co-timed ACA expansion (Section 6.9) and is the more defensible forecast for a mandate implemented in an environment where expansion status is already fixed. The figure should in any case be treated as an ecological intent-to-treat estimate rather than a

guaranteed forecast. The scale is large in practical terms even at the adjusted end: these three states account for roughly 100,000 annual prison releases, so an 11-22 percentage-point benchmark implies on the order of 11,000-22,000 additional returning citizens maintaining or regaining Medicaid coverage each year if implementation produces gains similar to historical adoption. Implementation planners can use this benchmark for budgeting (the increased enrollment will generate federal matching funds), workforce planning (more enrollees will require provider capacity), and performance monitoring (coverage rates materially below this benchmark in the first year of implementation may indicate operational failures rather than headline-policy failures).

Managed-care organizations and provider capacity. An enrollment increase of 11-22 pp among returning citizens, concentrated in cohorts with elevated rates of substance use disorder, serious mental illness, and chronic conditions, represents a substantial increase in demand for behavioral health, MOUD, and primary care services within Medicaid managed care networks. In the 41 states that contract with managed-care organizations (MCOs) for at least part of their Medicaid program, MCOs are the operational implementers of post-release reactivation: they assign primary care providers, manage prior authorization for MOUD, and bear capacity risk for behavioral health utilization spikes. Two implications follow. First, MCO contracts that did not previously incorporate reentry-coverage spikes into capitation rate-setting will need to do so under the mandate; rate-setting that fails to anticipate the demand increase documented here may produce either MCO underperformance or provider-network access failures. Second, network-adequacy standards for Medicaid should be evaluated specifically for behavioral health and MOUD adequacy in counties with high reentry volume, since uniform statewide standards may mask county-level capacity gaps that bind precisely where suspended coverage is reactivated.

The Howell et al. (2023) null-utilization finding (Section 2.5) underlines that demand-side coverage gains cannot substitute for supply-side capacity.

Complementarity with Section 1115 waivers. Suspension and the pre-release Section 1115 reentry waivers are complementary rather than substitutes: suspension preserves eligibility across incarceration (the administrative margin), while the waivers expand covered services in the 30 to 90 days before release (the clinical margin of treatment initiation and care-coordination handoff; CMS, 2023). States implementing both policies at once should, in principle, see synergistic effects on care use and post-release outcomes, although this interaction has not yet been evaluated empirically.

Coverage continuity beyond reactivation. The slight attenuation at longer horizons (Section 6.2) suggests that suspension alone does not solve Medicaid churning. Combining suspension with continuous-eligibility guarantees (for example, 12 months of guaranteed continuous eligibility post-release) could further improve coverage continuity and amplify downstream health benefits.

Targeted burden reduction at other transitions. The administrative burden lens also has implications outside the criminal-justice context. Analogous high-friction transitions (aging out of foster care, exiting homelessness shelters, leaving inpatient psychiatric care, transitioning from pediatric to adult Medicaid, leaving long-term care facilities) may similarly benefit from automatic enrollment or coverage-preservation defaults at the moments when individual capacity to navigate paperwork is at its lowest.

7.6 Limitations

The analysis is limited to aggregate outcomes rather than individual-level transitions, and estimates could partially reflect selection rather than the policy effect. I address this by 1) testing

for differential pre-trends and testing pre-determined release-cohort characteristics as placebo outcomes (Section 6.10), and 2) re-estimating the main model with release-cohort composition controls (Section 6.1), which attenuate but do not overturn the finding; these checks mitigate the concern but cannot eliminate unobservable compositional changes. The ecological design also prevents heterogeneity analysis by individual characteristics beyond JOE's published subgroup cells and reduces statistical power relative to individual-level data, which may explain the inability to detect effects on employment and recidivism.

State coverage is also incomplete; CJARS covers 31 of 51 states, with coverage determined by state data-sharing agreements with the Census Bureau, and the 20 missing states, including Alabama, may differ systematically from included states. If missing states are disproportionately those with weaker administrative infrastructure (which might make suspension less effective) or different incarcerated populations, the estimates may not generalize to the full set of states.

However, the sample includes 27 of 29 suspension-adopting states and 2 of 3 termination states, which provides reasonable coverage of both policy regimes. Alongside this, only two never-treated states (Georgia, Texas) appear in the CJARS sample, and only Texas contributes Medicaid outcome rows. This constrains the comparison group for both standard and heterogeneity-robust estimators. LP-DiD's not-yet-treated comparators mitigate but do not remove this limitation, which is particularly constraining for the ACA expansion interaction analysis, where both never-treated states are non-expansion states.

The co-timing with the ACA expansion is the single most important qualification on the headline magnitude (Section 6.9). Within the co-timed states, no design can fully separate the two policies. The disaggregation in Section 6.9 identifies the suspension-specific component from the pre-2014 adopters and the never-expanded adopters, and it implies a suspension effect net of

expansion of roughly 11 pp at one year rather than the raw 22 pp. I report the raw and expansion-adjusted figures side by side wherever the estimate is used prescriptively.

Implementation is a further limitation, in two ways. First, the policy panel captures the de jure adoption of suspension (when the policy was formally enacted or announced), but Bandara et al. (2025) document substantial gaps between de jure policy and de facto implementation in the operational machinery described in Section 2.2. This measurement error in the treatment variable would attenuate the estimates toward zero, suggesting the findings may be conservative: the effect of fully and effectively implemented suspension could be larger than what I estimate from de jure adoption. Second, Medicaid managed-care penetration varies substantially across states and over time, and managed-care plans often control the operational steps that matter after reactivation (Section 7.5). State fixed effects absorb time-invariant differences in managed-care infrastructure, but they do not fully address differential changes in managed-care capacity or contract design that coincide with suspension adoption. If suspension states were also improving managed-care reentry operations, the coverage effects could partly reflect those implementation complements rather than eligibility preservation alone.

The observation window is limited. Medicaid enrollment in CJARS is measurable only in calendar years 2005-2019, tied to the MSIS/T-MSIS linkage window (Section 4.3). This prevents observing effects for the most recent adopters (post-2018) and limits the sample size. Similarly, SSI data are available for only a subset of years, severely limiting power for that outcome. I cannot assess impacts beyond 5 years post-release.

I also considered potential SUTVA violations. The stable unit treatment value assumption requires that one state's policy adoption does not affect outcomes in other states, and it could be violated if individuals strategically relocate to suspension states after release or if federal policy

attention to suspension creates spillovers onto termination states. I do not anticipate meaningful bias from this channel, as the release state is largely determined by the location of the correctional facility, not by the individual.

Lastly, measurement has a few limits. W-2 employment captures only formal-sector employment reported to the IRS; individuals in informal, gig, or off-the-books work are classified as non-employed, potentially masking effects on total work activity. Mortality is measured through the Census Numident, which is high-quality but may miss some deaths, particularly among individuals who are undocumented or transient. As noted in Section 6.6, JOE recidivism captures same-state contact only, so measured rates are lower bounds and the nulls are informative about the same-state margin specifically. If suspension policies differentially affect post-release cross-state migration (for example, if terminated individuals are likelier to relocate toward suspension states), the same-state margin could diverge from total recidivism in ways this design cannot detect.

7.7 A Post-2026 Research Agenda

With the federal mandate in effect, the research agenda on Medicaid suspension shifts from cross-state evaluation of whether states should suspend to within-state evaluation of how well they implement. Five research directions follow from the present paper.

1. **Implementation-quality heterogeneity within the mandate era.** CMS guidance from December 2025 describes a menu of operational suspension strategies, ranging from indefinite suspension with automated reactivation to time-limited suspension with manual reapplication processes (CMS, 2025). The newly mandated states and any state with weak prior implementation (Bandara et al., 2025) provide an implementation-quality natural experiment. Synthetic-control or interrupted-time-series designs comparing post-

mandate coverage trajectories to pre-mandate trends, and to a counterfactual reconstructed from suspension states, can identify whether implementation strategy matters.

2. **Section 1115 reentry waivers and suspension interaction.** The independent and joint effects of eligibility-preserving suspension and pre-release service-coverage waivers (Section 2.6) can be identified by the staggered rollout of waivers within the now-uniform suspension regime. The relevant outcomes (care utilization, MOUD initiation, post-release inpatient and ED use) speak directly to whether the “coverage is necessary but not sufficient” interpretation in Section 7.4 is empirically right.
3. **Individual-level CJARS analysis through the FSRDC network.** The most valuable extension is restricted-access CJARS microdata. Individual-level data would (a) replace the release-cohort ITT estimand with a treatment-on-treated estimand among individuals whose Medicaid was actually suspended; (b) enable heterogeneity by offense type, incarceration period length, and, critically, by health condition (substance use disorder, serious mental illness, chronic disease), constructed from linked MSIS/T-MSIS claims diagnoses, a dimension that JOE’s public aggregates cannot support because they publish no clinical variables; (c) substantially increase statistical power for the mortality and employment outcomes that this design cannot resolve; and (d) measure the timing of post-release Medicaid reactivation (day-of-release vs. month 1 vs. month 3), which is the mechanism that connects suspension to acute mortality avoidance.
4. **MOUD initiation and care-coordination interventions layered on top of suspension.** The “coverage is necessary but not sufficient” reading of these results implies that the next reentry-coverage research question is which complementary interventions translate

eligibility into care: in-prison MOUD initiation, warm-handoff care coordination, naloxone distribution at release, transportation vouchers, and transitional case management. Several of these are now testable as experimental or quasi-experimental policy variation under the post-mandate regime.

5. **Long-run impacts beyond five years.** The 1- to 5-year horizons available in JOE cannot speak to long-run effects on health, employment, recidivism, or mortality. As the CJARS-MSIS/T-MSIS linkages extend forward in time, longer-horizon evaluation of pre-mandate suspension cohorts becomes feasible and may reveal effects that are too small or too distal to detect at five years.
-

8. Conclusion

This paper provides the first multi-state evidence on the relationship between Medicaid suspension during incarceration and post-release coverage trajectories. Linking a newly constructed 31-state suspension policy panel to administratively linked CJARS release-cohort outcomes for 17 reporting states, I find a large, persistent intent-to-treat coverage effect: suspension is associated with an increase of roughly 22 percentage points in 1-year post-release Medicaid enrollment, and roughly 11 pp net of the co-timed ACA Medicaid expansion. The effect persists at the 3- and 5-year horizons and is homogeneous across sex, race, and age subgroups. Effects on W-2 employment and earnings are null, and same-state reincarceration effects are a precise-enough null that neither confirms nor refutes the modest reductions reported in prior individual-level work. Mortality estimates are inconclusive and underpowered for the modest reductions documented in the broader Medicaid-mortality literature (Section 6.4).

These findings arrive at a transitional policy moment. The 2026 federal mandate has since made suspension universal (Section 2.3). They also support a more cautious second message: a strong coverage effect is not, on its own, evidence that the post-release acute-mortality and labor-market gaps will close. Coverage is necessary but not sufficient.

More broadly, the results indicate that reducing administrative burdens at critical life transitions, when individuals face the greatest barriers to accessing safety-net programs, can produce substantial improvements in program take-up. The suspension-termination distinction looks like a minor administrative choice buried in the technical details of Medicaid policy, but it has been an important determinant of whether returning citizens maintain health coverage through the most vulnerable period of their lives. The lesson extends beyond Medicaid and beyond the justice-involved population: at moments of acute life transition, the friction of a government application can be the difference between coverage and no coverage. Translating that coverage gain into measurably better health and economic outcomes is the next policy and research challenge.

References

- Albertson, E. M., Scannell, C., Ashtari, N., & Barnert, E. (2020). Eliminating gaps in Medicaid coverage during reentry after incarceration. *American Journal of Public Health*, 110(3), 317-321.
- Anderson, M. L. (2008). Multiple inference and gender differences in the effects of early intervention: A reevaluation of the Abecedarian, Perry Preschool, and Early Training Projects. *Journal of the American Statistical Association*, 103(484), 1481-1495.
- Aslim, E. G., Mungan, M. C., Navarro, C. I., & Yu, H. (2022). The effect of public health insurance on criminal recidivism. *Journal of Policy Analysis and Management*, 41(1), 45-91.

Aslim, E. G., Mungan, M. C., & Yu, H. (2024). A welfare analysis of Medicaid and recidivism. *Health Economics*, 33(11), 2463-2507.

Baicker, K., Taubman, S. L., Allen, H. L., Bernstein, M., Gruber, J. H., Newhouse, J. P., Schneider, E. C., Wright, B. J., Zaslavsky, A. M., Finkelstein, A. N., and the Oregon Health Study Group. (2013). The Oregon experiment - effects of Medicaid on clinical outcomes. *New England Journal of Medicine*, 368(18), 1713-1722.

Bandara, S., Saloner, B., Maniates, H., Song, M., & Krawczyk, N. (2025). Implementation of carceral Medicaid suspension and enrollment programs. *Health & Justice*, 13(1), 1.

Bertrand, M., Duflo, E., & Mullainathan, S. (2004). How much should we trust differences-in-differences estimates? *Quarterly Journal of Economics*, 119(1), 249-275.

Binswanger, I. A., Stern, M. F., Deyo, R. A., Heagerty, P. J., Cheadle, A., Elmore, J. G., & Koepsell, T. D. (2007). Release from prison - a high risk of death for former inmates. *New England Journal of Medicine*, 356(2), 157-165.

Borgschulte, M., & Vogler, J. (2020). Did the ACA Medicaid expansion save lives? *Journal of Health Economics*, 72, 102333.

Burns, M. E., Cook, S., Brown, L. M., Dague, L., Tyska, S., Hernandez Romero, K., McNamara, C., & Westergaard, R. P. (2022). Association between assistance with Medicaid enrollment and use of health care after incarceration among adults with a history of substance use. *JAMA Network Open*, 5(1), e2142688. doi:10.1001/jamanetworkopen.2021.42688

Callaway, B., & Sant'Anna, P. H. C. (2021). Difference-in-differences with multiple time periods. *Journal of Econometrics*, 225(2), 200-230.

Camhi, N., Mistak, D., & Wachino, V. (2020). Medicaid's evolving role in advancing the health of people involved in the justice system. Commonwealth Fund Issue Brief.

Carson, E. A., & Kluckow, R. (2023). Prisoners in 2022 - statistical tables. Bureau of Justice Statistics.

Centers for Medicare & Medicaid Services (CMS). (2023). HHS releases new guidance to encourage states to apply for new Medicaid Reentry Section 1115 demonstration opportunity. CMS press release, April 17, 2023.

Centers for Medicare & Medicaid Services (CMS). (2025). *Prohibition on termination of enrollment due to incarceration (Division G, Title I, Section 205, of the Consolidated Appropriations Act, 2024)*. CMCS Informational Bulletin, December 23, 2025.

Deza, M., Lu, T., Maclean, J. C., & Ortega, A. (2024). Losing Medicaid and crime. NBER Working Paper No. 32227.

Dube, A., Girardi, D., Jorda, O., & Taylor, A. M. (2023). A local projections approach to difference-in-differences event studies. NBER Working Paper No. 31184.

Finkelstein, A., Taubman, S., Wright, B., Bernstein, M., Gruber, J., Newhouse, J. P., Allen, H., Baicker, K., and the Oregon Health Study Group. (2012). The Oregon Health Insurance Experiment: Evidence from the first year. *Quarterly Journal of Economics*, 127(3), 1057-1106.

Finlay, K., Mueller-Smith, M., & Papp, J. (2022). The Criminal Justice Administrative Records System. *Scientific Data*, 9, 562.

Goldman, A. L., & Sommers, B. D. (2020). Among low-income adults enrolled in Medicaid, churning decreased after the ACA. *Health Affairs*, 39(1), 85-93.

Gollu, G., & Zapryanova, M. (2022). The effect of Medicaid on recidivism: Evidence from Medicaid suspension and termination policies. *Southern Economic Journal*, 89(2), 326-372.

Goodman-Bacon, A. (2021). Difference-in-differences with variation in treatment timing. *Journal of Econometrics*, 225(2), 254-277.

Herd, P., & Moynihan, D. P. (2018). *Administrative Burden: Policymaking by Other Means*. Russell Sage Foundation.

Herd, P., & Moynihan, D. P. (2025). Administrative burdens in the social safety net. *Journal of Economic Perspectives*, 39(1), 129-150.

Howell, B. A., Hawks, L. C., Balasuriya, L., Chang, V. W., Wang, E. A., & Winkelman, T. N. A. (2023). Health insurance and mental health treatment use among adults with criminal legal involvement after Medicaid expansion. *Psychiatric Services*, 74(10), 1019-1026.

Looney, A., & Turner, N. (2018). Work and opportunity before and after incarceration. Brookings Institution.

MACPAC (2018). *Medicaid and the criminal justice system*. MACPAC Issue Brief, July 2018.

Mallik-Kane, K., & Visser, C. (2008). Health and prisoner reentry. Urban Institute.

Olden, A., & Møen, J. (2022). The triple difference estimator. *Econometrics Journal*, 25(3), 531-553.

Ortiz-Villavicencio, M., & Sant'Anna, P. H. C. (2025). Better understanding triple differences estimators. Working paper, arXiv:2505.09942.

Packham, A., & Slusky, D. (2024). Accessing the safety net: How Medicaid affects health and recidivism. NBER Working Paper No. 31971 (originally December 2023; revised June 2024).

Sommers, B. D., Baicker, K., & Epstein, A. M. (2012). Mortality and access to care among adults after state Medicaid expansions. *New England Journal of Medicine*, 367(11), 1025-1034.

Sugar, S., Peters, C., De Lew, N., & Sommers, B. D. (2021). *Medicaid churning and continuity of care: Evidence and policy considerations before and after the COVID-19 pandemic* (ASPE Issue Brief HP-2021-10). Office of the Assistant Secretary for Planning and Evaluation, U.S. Department of Health and Human Services.

Wen, H., Hockenberry, J. M., Borders, T. F., & Druss, B. G. (2017). Impact of Medicaid expansion on Medicaid-covered utilization of buprenorphine for opioid use disorder treatment.

Medical Care, 55(4), 336-341.

Zeng, Z. (2023). Jail inmates in 2022 - statistical tables. Bureau of Justice Statistics.

Appendix Tables

Table A1. Summary Statistics (Full)

	N	Mean	SD	Min	Median	Max
Panel A. Primary outcomes, 1-year horizon						
Medicaid enrollment	212	0.320	0.251	0.078	0.200	0.880
W-2 employment	212	0.536	0.074	0.350	0.540	0.680
W-2 earnings (\$)	212	6,909	1,634	3,060	6,905	11,600
Mortality	264	0.017	0.005	0.006	0.016	0.048
Panel B. Longer-horizon and secondary outcomes						
Medicaid enrollment (3-yr)	207	0.287	0.187	0.091	0.210	0.800
Medicaid	201	0.277	0.166	0.093	0.210	0.740

	N	Mean	SD	Min	Median	Max
enrollment (5-yr)						
HUD housing (1-yr)	264	0.011	0.005	0.002	0.010	0.031
SSI receipt (1-yr)	99	0.043	0.010	0.021	0.045	0.060
Incarceration period length (mos.)	378	25.1	6.6	11.5	24.5	41.7
Panel C.						
State-year controls						
Unemployment rate (%)	561	5.58	2.13	2.20	5.00	13.70
Overdose mortality rate	554	13.9	6.5	2.6	12.7	44.3
Poverty rate (%)	534	13.0	2.9	6.9	12.7	21.4
Panel D.						
Policy variable						
Post-suspension indicator	561	0.296	0.457	0	0	1

Notes: Unit of observation is the state-year. Medicaid, employment, HUD, SSI, and mortality outcomes are shares in [0, 1]. Earnings are in nominal dollars averaged across release-cohort members. Overdose mortality rate is CDC WONDER age-adjusted deaths per 100,000. The sample covers 31 states over 2000-2020; the Medicaid outcomes are restricted to 17 states, with horizon-specific cohort windows implied by the MSIS/T-MSIS linkage (release cohorts 2004-2018 at 1 year, 2002-2016 at 3 years, 2000-2014 at 5 years).

Table A2. Pre-Adoption Balance (Full)

	Adopting states	Non-adopting states	Difference	Normalized diff.
	Mean (SD)	Mean (SD)		
Panel A.				
Primary outcomes				
Medicaid enrollment (1-yr)	0.210 (0.109)	0.144 (0.041)	0.066	0.80
W-2 employment (1-yr)	0.552 (0.079)	0.513 (0.063)	0.039	0.55
W-2 earnings (1-yr, \$)	6,814 (1,822)	6,867 (1,432)	-53	-0.03
Mortality (1-yr)	0.016 (0.005)	0.017 (0.004)	-0.001	-0.17
Panel B.				
Release-cohort characteristics				
Incarceration period (months)	22.6 (5.8)	29.9 (3.5)	-7.3	-1.53
Annual release cohort size	10,096 (7,002)	34,644 (23,043)	-24,549	-1.44
Share male	0.879 (0.027)	0.864 (0.022)	0.015	0.62
Share Black	0.302 (0.156)	0.388 (0.078)	-0.086	-0.70
Share Hispanic	0.149 (0.112)	0.133 (0.107)	0.015	0.14
Share young (15-24)	0.197 (0.050)	0.181 (0.022)	0.016	0.42
Panel C.				
State controls				
Unemployment rate (%)	5.69 (2.23)	5.60 (1.93)	0.10	0.05
Overdose mortality rate	11.8 (5.0)	13.3 (5.6)	-1.5	-0.28
Poverty rate (%)	12.5 (3.0)	15.3 (1.9)	-2.8	-1.10

Notes: Means (standard deviations) across pre-adoption state-year observations. “Adopting states” comprises all states that adopted suspension during 2005-2022, observed in their pre-treatment years; “non-adopting states” comprises Georgia and Texas, the two never-treated states in the CJARS sample, across all years. Normalized difference is the mean difference divided by the pooled standard deviation. Release-cohort demographic shares are computed from JOE state-data subgroup cells using the release-cohort cohort-size variable (cohort size at the 5-year W-2 horizon) as the denominator within sex/race/age subgroups. Level differences in Panels A-C are absorbed by state fixed effects in the main DiD design; the identifying assumption requires parallel trends rather than level equivalence.

Table A3. Main Results Across Estimators and Outcomes, 1-Year Horizon

	Medicaid (1)	Medicaid (+controls) (2)	Employment (3)	Earnings (4)	Mortality (5)
Panel A. Two-way fixed effects (TWFE)					
Suspension	0.222*** (0.051)	0.230*** (0.041)	0.016 (0.011)	315.3 (272.3)	0.0005 (0.0011)
Panel B. LP-DiD (Dube et al., 2023)					
Suspension	0.224*** (0.018)	-	0.003 (0.029)	383.3 (339.2)	-0.0017** (0.0006)
Panel C. Callaway — Sant’Anna (2021)					
Suspension (aggregate)	0.258***	-	-	-	0.0001

	Medicaid	Medicaid (+controls)	Employment	Earnings	Mortality
d ATT)	(0.046)				(0.0015)
Panel D. TWFE, cohort- size weighted					
Suspension	0.257*** (0.062)	-	0.023 (0.020)	351.4 (368.5)	0.0020* (0.0012)
Panel E. Inference					
Wild cluster bootstrap p specification details	<0.001	<0.001	0.135	0.285	0.667
State FE	Yes	Yes	Yes	Yes	Yes
Year FE	Yes	Yes	Yes	Yes	Yes
Time-varying controls	No	Yes	No	No	No
Observations	212	198	212	212	264
States	17	17	17	17	31
R ² (within, TWFE)	0.214	0.268	0.034	0.021	0.003
Panel A)					

Notes: Cluster-robust standard errors (state-level) in parentheses. Each column is an outcome-specific specification with state and year fixed effects. Panel A is the benchmark TWFE; Panel B uses the LP-DiD estimator of Dube et al. (2023), which uses only not-yet-treated and never-treated states as comparators at each horizon; Panel C uses the Callaway-Sant’Anna (2021) group-time ATT estimator with not-yet-treated controls; Panel D is TWFE reweighted by the JOE release-cohort cohort-size variable for each outcome; Panel E reports Rademacher wild cluster bootstrap p -values ($B = 999$). Medicaid and employment outcomes are shares; earnings are nominal dollars; mortality is the share deceased within the horizon. The Medicaid sample is

restricted to 17 states (release cohorts 2004-2018 at the 1-year horizon) by the MSIS/T-MSIS linkage window. *p<0.10, **p<0.05, ***p<0.01.

Table A4. Heterogeneous DiD by ACA Medicaid Expansion Status

	Medicaid	Employment	Earnings	Mortality
	(1)	(2)	(3)	(4)
Panel A. Static- expansion interaction (main)				
Suspension	0.075 (0.130)	0.024* (0.012)	479.0 (642.0)	0.0018* (0.0010)
Suspension × Expanded	0.195 (0.132)	−0.010 (0.017)	−216.0 (661.0)	−0.0016 (0.0016)
Panel B. Period-split TWFE				
Suspension × Pre-ACA (2000–2013)	0.117** (0.050)	0.007 (0.009)	−3.0 (300.0)	−0.0000 (0.0007)
Suspension × Post-ACA (2014–2020)	0.147* (0.081)	0.008 (0.008)	−157.0 (300.0)	−0.0013 (0.0035)
Panel C. Time- varying expansion interaction (identification limit)				
Suspension × Expansion (time- varying)	<i>collinear</i>	<i>collinear</i>	<i>collinear</i>	<i>collinear</i>
Specification details				
State FE,	Yes	Yes	Yes	Yes
Year FE				
Observations (static)	212	212	212	264

	Medicaid	Employment	Earnings	Mortality
Observations (pre-ACA)	134	134	134	186
Observations (post-ACA)	78	78	78	78

Notes: Cluster-robust standard errors (state-level) in parentheses. Panel A reports the heterogeneous-DiD interaction with time-invariant ACA expansion status as the cross-classifying dimension. Panel B reports TWFE coefficients on the suspension indicator estimated separately in the pre- and post-ACA-expansion periods. Panel C reports the outcome of an attempted time-varying expansion interaction: in the Medicaid sample, every state that adopted both suspension and the ACA expansion did so in an order that makes the interaction of post-suspension and time-varying expansion collinear with the expansion indicator net of fixed effects, and the interaction coefficient is not separately identified (see Section 6.8). *p<0.10, **p<0.05, ***p<0.01.

Table A5. Horizon Sensitivity (TWFE), Years 1, 3, and 5 Post-Release

	Medicaid	Employment t	Earnings (\$)	Mortality
1-year horizon				
Suspension	0.222*** (0.051)	0.016 (0.011)	315.3 (272.3)	0.0005 (0.0011)
Observations	212	212	212	264
3-year horizon				
Suspension	0.194*** (0.044)	0.003 (0.011)	281.1 (268.3)	0.0014 (0.0012)
Observations	205	205	205	231
5-year horizon				
Suspension	0.207*** (0.058)	-0.006 (0.016)	93.8 (348.0)	0.0007 (0.0014)
Observations	200	200	200	200
Specification details				
State FE, Year FE	Yes	Yes	Yes	Yes

Notes: Cluster-robust standard errors (state-level) in parentheses. Each cell is a separate TWFE regression of the stated outcome on the post-suspension indicator with state and year fixed effects. *p<0.10, **p<0.05, ***p<0.01.

Table A6. Robustness Summary, Medicaid Enrollment (1-Year)

	Coefficient	SE / CI	N / States	Notes
Panel A.				
Estimator robustness				
TWFE (baseline)	0.222***	(0.051)	212 / 17	Cluster-robust SE
TWFE + time-varying controls	0.230***	(0.041)	198 / 17	Unemployment, overdose, poverty
LP-DiD (Dube et al. 2023)	0.224***	(0.018)	-	Heterogeneity-robust
Callaway–Sant’Anna (2021)	0.258***	(0.046)	-	Not-yet-treated controls
Cohort-size-weighted TWFE	0.257***	(0.062)	212 / 17	Weighted by release-cohort size
Panel B.				
Inference robustness				
Wild cluster bootstrap p	-	< 0.001	-	B = 999, Rademacher
Randomization inference p	-	0.062	500 perms	Design-based
Panel C.				
Sample / subsample robustness				
Leave-one-out (range)	[0.206, 0.259]	max dev. 0.036	-	Across 17 states
Drop Georgia (never-treated)	0.222***	(0.051)	212 / 16	Main result unchanged
Drop Texas (never-treated)	0.209**	(0.052)	212 / 16	Modest attenuation
Drop South Carolina	0.222***	(0.051)	212 / 16	Matches Packham-Slusky scope
Pre-ACA only (2005–2013)	0.117**	(0.050)	134 / 17	Aligns with Packham-Slusky 11.1 pp

	Coefficient	SE / CI	N / States	Notes
Post-ACA only (2014–2018)	0.147*	(0.081)	78 / 17	ACA amplification
Panel D. Policy-coding and cell-size robustness				
Drop low-confidence policy codes	0.209***	-	-	1 state dropped
Cell size N >= 50 (vs. N >= 25)	0.222***	(0.051)	-	Identical (all cells >= 50)
Panel E. Multiple-testing adjustment				
Anderson (2008) sharpened q	0.222	q = 0.005	-	Across 14-outcome battery

Notes: Column 1 reports the coefficient on the post-suspension indicator for Medicaid 1-year

enrollment; column 2 reports the cluster-robust standard error in parentheses (or p-value / CI

where relevant); column 3 reports sample size and cluster count; column 4 describes the

specification. All TWFE specifications include state and year fixed effects with cluster-robust SE

clustered at the state level. The Anderson (2008) sharpened q-value is computed across the full

outcome battery (Medicaid, employment, earnings, mortality at 1/3/5 years; HUD and SSI at 1

year). *p<0.10, **p<0.05, ***p<0.01.

Table A7. Subgroup DiD Estimates by Sex, Race/Ethnicity, and Age Group

Subgroup	Medicaid 1-yr TWFE	Medicaid 1-yr LP-DiD	Medicaid 3-yr TWFE	Medicaid 3-yr LP-DiD	Medicaid 5-yr TWFE	Medicaid 5-yr LP-DiD	Reincarceration 1-yr TWFE	Reincarceration 1-yr LP-DiD
Overall	0.222 (0.051)	0.224 (0.018)	0.194 (0.044)	0.187 (0.007)	0.207 (0.058)	0.196 (0.018)	-0.022 (0.027)	-0.007 (0.007)
Male	0.231 (0.054)	0.234 (0.021)	0.199 (0.044)	0.201 (0.009)	0.213 (0.058)	0.198 (0.017)	-0.023 (0.027)	-0.008 (0.008)
Female	0.159 (0.039)	0.177 (0.014)	0.158 (0.048)	0.152 (0.012)	0.165 (0.064)	0.170 (0.018)	-0.010 (0.023)	-0.007 (0.002)

Subgroup	Medicaid 1-yr TWFE	Medicaid 1-yr LP-DiD	Medicaid 3-yr TWFE	Medicaid 3-yr LP-DiD	Medicaid 5-yr TWFE	Medicaid 5-yr LP-DiD	Reincarceration 1-yr TWFE	Reincarceration 1-yr LP-DiD
White	0.214 (0.050)	0.221 (0.017)	0.190 (0.043)	0.196 (0.010)	0.200 (0.057)	0.201 (0.017)	-0.023 (0.030)	-0.015 (0.007)
Black	0.233 (0.051)	0.237 (0.025)	0.210 (0.047)	0.204 (0.012)	0.219 (0.063)	0.198 (0.018)	-0.031 (0.029)	0.001 (0.013)
Hispanic	0.218 (0.051)	0.214 (0.014)	0.176 (0.043)	0.184 (0.009)	0.185 (0.063)	0.148 (0.020)	-0.029 (0.023)	0.032 (0.006)
Asian/Pacific Islander†	0.176 (0.060)	n.r.	0.148 (0.037)	0.168 (0.043)	0.126 (0.036)	0.152 (0.036)	-0.050 (0.034)	-0.109 (0.104)
American Indian/Alaska Native†	0.176 (0.052)	n.e.	0.201 (0.065)	0.184 (0.010)	0.237 (0.070)	0.200 (0.023)	-0.021 (0.044)	n.e.
Age 15–24	0.227 (0.053)	0.173 (0.020)	0.207 (0.047)	0.163 (0.010)	0.216 (0.064)	0.199 (0.021)	-0.017 (0.030)	-0.012 (0.019)
Age 25–39	0.234 (0.054)	0.240 (0.018)	0.198 (0.046)	0.207 (0.010)	0.215 (0.060)	0.200 (0.020)	-0.024 (0.027)	-0.018 (0.009)
Age 40+	0.204 (0.047)	0.211 (0.023)	0.183 (0.043)	0.176 (0.006)	0.192 (0.055)	0.175 (0.013)	-0.024 (0.026)	0.003 (0.004)

Notes: Each cell reports the coefficient on the post-suspension indicator from a separate

subgroup-specific regression, with the cluster-robust standard error (state level) in parentheses.

The subgroup panels are built from JOE's published one-way subgroup cells. Full-coverage

subgroup panels match the overall panel Ns (Medicaid 1-yr: N = 212, 17 states; 3-yr: N = 205-

207, 17 states; 5-yr: N = 200-201, 15 states; reincarceration 1-yr: N = 384, 22 states), with the

Black and Hispanic panels missing at most a handful of suppressed cells. †Underpowered:

Census Bureau disclosure suppression (cells with N < 25) removes 23-26 percent of state-cohort

cells for the Asian/Pacific Islander and American Indian/Alaska Native panels (e.g., Medicaid 1-

yr N = 163 and 160, respectively; reincarceration 1-yr N = 296 and 293), a non-random set of

small cohorts; interpret with caution. “n.r.” = LP-DiD returns a degenerate (non-finite) standard error and is not reported; “n.e.” = LP-DiD estimation fails in the sparse suppressed panel.

Table A8. Charge-Type Recidivism Detail (Felony and Misdemeanor Recharge, 1-Year)

Outcome	Estimator	Estimate	SE	p	N	States
Felony recharge (1-yr)	TWFE	0.009	0.018	0.636	129	8
Felony recharge (1-yr)	LP-DiD	n.e.	-	-	129	8
Felony recharge (1-yr)	Callaway - Sant’Anna	-0.011	0.014	0.413	129	8
Misdemeanor recharge (1-yr)	TWFE	0.046	0.017	0.030	116	8
Misdemeanor recharge (1-yr)	LP-DiD	n.e.	-	-	116	8
Misdemeanor recharge (1-yr)	Callaway - Sant’Anna	0.002	0.027	0.932	116	8

Notes: Release-cohort window 2000-2020. JOE court-based charge outcomes are observed in only 8 states (JOE court coverage is far thinner than corrections coverage). With fewer than 10 clusters, state-clustered standard errors are fragile, and all estimates in this table should be treated as suggestive. “n.e.” = not estimable: the LP-DiD estimator has insufficient clean-control variation in the 8-state subsample. The positive misdemeanor TWFE estimate is not robust to the Callaway-Sant’Anna estimator and is not interpreted as evidence of a misdemeanor-increasing effect (see Section 6.6). All recidivism outcomes capture same-state criminal-justice contact only and are lower bounds on true recidivism.

Exhibits

Figure 1. State adoption of Medicaid suspension policies.




Figure 1

Figure 2. Suspension and the reentry coverage pathway.




Figure 2

Figure 3. Reentry Medicaid policy architecture: suspension and Section 1115 reentry waivers.





Figure 3

Figure 4. Data structure: state $\times$ release-cohort cells, person-level follow-up horizons, and censoring.

Data structure: state × release-cohort cells, person-level follow-up horizons, and censoring

A. One release cohort, three person-level follow-up horizons



B. Observability of each horizon, release cohorts 2000-2020
(MSIS/T-MSIS linkage: Medicaid outcomes measurable in calendar years 2005-2019)




Figure 4

Figure 5. LP-DiD event study: Medicaid enrollment, 3-year horizon.




Figure 5

Figure 6. LP-DiD event study: Medicaid enrollment, 5-year horizon.





Figure 6

Figure 7. Subgroup estimates: Medicaid enrollment, 1-year horizon.



* Underpowered subgroup: reduced cell coverage due to disclosure suppression ($N < 25$); interpret with caution.

Figure 7

Figure 8. Raw reincarceration trends by adoption cohort, 1-year horizon.




Figure 8

Figure 9. LP-DiD event study: reincarceration, 1-year horizon.




Figure 9

Figure 10. Suspension adoption year versus ACA expansion year.




Figure 10

Figure 11. TWFE event study with time-varying ACA expansion control: Medicaid enrollment, 1-year horizon.




Figure 11

Figure 12. Raw Medicaid enrollment trends by adoption cohort, 1-year horizon.




Figure 12

Figure 13. Raw employment trends by adoption cohort, 1-year horizon.




Figure 13

Figure 14. Raw earnings trends by adoption cohort, 1-year horizon.




Figure 14

Figure 15. Raw mortality trends by adoption cohort, 1-year horizon.




Figure 15

Figure 16. LP-DiD event study: Medicaid enrollment, 1-year horizon.




Figure 16

Figure 17. LP-DiD event study: W-2 employment, 1-year horizon.




Figure 17

Figure 18. LP-DiD event study: W-2 earnings, 1-year horizon.




Figure 18

Figure 19. LP-DiD event study: mortality, 1-year horizon.




Figure 19


Figure 20. Implementation pathway under the 2026 federal mandate.



Figure 20


Appendix Figures

Appendix Figure A1. TWFE event study: Medicaid enrollment (1-yr).




A1

Appendix Figure A2. TWFE event study: W-2 employment (1-yr).




A2

Appendix Figure A3. TWFE event study: W-2 earnings (1-yr).




A3

Appendix Figure A4. TWFE event study: mortality (1-yr).




A4

Appendix Figure A5. Policy adoption timeline.




A5

Appendix Figure A6. Coefficient comparison across estimators.




A6

Appendix Figure A7. Placebo distribution: Medicaid enrollment.




A7

Appendix Figure A8. Leave-one-out: Medicaid enrollment.




A8

Appendix Figure A9. Placebo distribution: mortality.




A9

Appendix Figure A10. Goodman-Bacon decomposition: Medicaid enrollment.




A10

Appendix Figure A11. Alternative scale, earnings: levels vs. log LP-DiD event study.




A11

Appendix Figure A12. TWFE vs. LP-DiD overlay: Medicaid enrollment.




A12

Appendix Figure A13. LP-DiD randomization inference.




A13

Appendix Figure A14. TWFE event study: Medicaid enrollment (3-yr). Estimated on treated states only (never-treated states have undefined event time in the TWFE event-study panel); release cohorts 2002-2016.




A14

Appendix Figure A15. TWFE event study: Medicaid enrollment (5-yr). Release cohorts 2000-2014.




A15

Appendix Figure A16. Raw Medicaid enrollment trends by adoption cohort, 3-year horizon.




A16

Appendix Figure A17. Raw Medicaid enrollment trends by adoption cohort, 5-year horizon.




A17

Appendix Figure A18. TWFE event study: reincarceration (1-yr).




A18

Appendix Figure A19. LP-DiD event study: reincarceration (3-yr).




A19

Appendix Figure A20. LP-DiD event study: reincarceration (5-yr).




A20

Appendix Figure A21. LP-DiD event study: Medicaid enrollment (1-yr), ACA-residualized outcome. The outcome is residualized on the time-varying post-expansion indicator before applying the LP-DiD estimator (Section 6.9); all pre-period leads are statistically indistinguishable from zero.




A21

Appendix Figure A22. LP-DiD event study: Medicaid enrollment (1-yr), male subgroup.




A22

Appendix Figure A23. LP-DiD event study: Medicaid enrollment (1-yr), female subgroup.




A23

Appendix Figure A24. Subgroup forest plot: Medicaid enrollment (3-yr). Hollow markers denote underpowered subgroups affected by disclosure suppression (Table A7).



* Underpowered subgroup: reduced cell coverage due to disclosure suppression (N<25); interpret with caution.


Appendix Figure A25. Subgroup forest plot: Medicaid enrollment (5-yr). Hollow markers denote underpowered subgroups affected by disclosure suppression (Table A7).



* Underpowered subgroup: reduced cell coverage due to disclosure suppression (N<25); interpret with caution.

A25

Appendix Figure A26. Subgroup forest plot: reincarceration (1-yr). Hollow markers denote underpowered subgroups affected by disclosure suppression (Table A7).



* Underpowered subgroup: reduced cell coverage due to disclosure suppression (N<25); interpret with caution.