

Did the Medicaid Unwinding Make Adults Sicker? A Precise Population-Level Null From BRFSS and the ACS, 2019–2024

Target journal: Health Affairs (Research Article) **Main text word count:** approximately 3,500 words **Abstract:** 300 words, unstructured per Health Affairs convention **Exhibits:** 5 main-text exhibits (2 tables, 3 figures); 5 online appendix exhibits **Data availability:** All microdata public (BRFSS, IPUMS ACS, CMS Unwinding Data, BLS LAUS, IPUMS CPS ASEC)

Abstract

The federal continuous Medicaid enrollment requirement ended on March 31, 2023. Over the subsequent eighteen months, states disenrolled more than 25 million Medicaid beneficiaries — the largest mass disenrollment in U.S. Medicaid history — and approximately 69 percent of disenrollments were coded as *procedural* rather than substantive ineligibility determinations. States varied sharply in intensity: the cumulative procedural disenrollment share ranged from 41 percent to 92 percent, reflecting deliberate state choices about pace, outreach, and use of federal waivers. We test whether this state-level intensity variation affected adult health and coverage at the population level, using a continuous-treatment difference-in-differences design on state-year Behavioral Risk Factor Surveillance System (BRFSS) aggregates for 2019–2023 (255 state-year observations) and individual-level American Community Survey (ACS) microdata for 2019–2024 (11.0 million adult observations). Across both data sources, no estimate is statistically distinguishable from zero, and coverage estimates uniformly run toward *more* coverage in high-intensity states. We calculate minimum detectable effects and show that the BRFSS state-year panel and the ACS all-adults sample are adequately powered against TennCare-scale (5-percentage-point) coverage effects, while the ACS below-138%-federal-poverty-level subsample is underpowered (minimum detectable effect 9.7 percentage points). State economic controls using Bureau of Labor Statistics unemployment and Current Population Survey prime-age employment-to-population ratios, a placebo outcome (any exercise in the past month), formal Rambachan-Roth relative-magnitudes sensitivity bounds implemented in the R *HonestDiD* package, and a Romano-Wolf step-down multiple-testing correction collectively confirm the null without tightening it further. Three hypotheses remain consistent with the data and cannot be distinguished by BRFSS or ACS: compositional substitution into alternative coverage, subpopulation concentration of harm in high-risk groups diluted by the larger low-risk majority, and sampling-frame exclusion of housing-unstable adults from survey instruments. BRFSS and the ACS cannot distinguish among these hypotheses because both are cross-sectional surveys that cannot follow individuals across coverage transitions; the decisive test requires individual-level Medicaid claims data, which we commit to as a companion paper. For policymakers, population-level survey evidence alone cannot serve as the empirical base for the administrative-burden reform agenda; reforms must be justified on grounds beyond a detectable population-wide health signal.

Policy Context

On March 31, 2023, the continuous enrollment requirement that had protected Medicaid beneficiaries from disenrollment since March 2020 ended. The requirement, established under the Families First Coronavirus Response Act, had been a condition of the enhanced federal matching funds states received during the COVID-19 public health emergency. Over three years of continuous

enrollment, national Medicaid enrollment grew from approximately 71 million to over 94 million. When the protection expired, states were required to redetermine eligibility for this enlarged population. By August 2024, more than 25 million people had been disenrolled, approximately 69 percent for procedural reasons rather than substantive ineligibility.

State implementation varied sharply. The cumulative procedural disenrollment share — the fraction of all disenrollments coded as procedural — ranged from about 41 percent in states that emphasized outreach and ex parte (automatic data-match) renewal to about 92 percent in states that prioritized rapid completion with tighter deadlines and less enrollee assistance. These differences reflected deliberate state choices about pace, outreach, and use of federal waivers. The U.S. Government Accountability Office subsequently documented compliance issues in almost all states, and at least 35 states reinstated coverage for enrollees erroneously disenrolled for administrative reasons.

Procedural disenrollment is the sharp end of administrative burden: an enrollee who cannot navigate the renewal process loses coverage regardless of whether she remains substantively eligible. The concern motivating procedural disenrollment as a policy target is that it produces *unnecessary* coverage losses that disproportionately fall on low-income, minority, and administratively burdened populations. Rumalla and colleagues documented that Black and Hispanic adults were roughly twice as likely as White adults to lose Medicaid coverage because of inability to complete the renewal process during the unwinding. If procedural disenrollment at this scale produced population-level health harm, it would constitute a strong empirical case for federal administrative-burden reform. This paper tests whether such a signal is detectable in standard public-use population surveys.

Study Data and Methods

Data

Our primary sample is a state-year panel of Behavioral Risk Factor Surveillance System (BRFSS) aggregates for adults aged 18–64 during 2019–2023 (255 state-year observations, with four state-years missing because of BRFSS non-participation and retained as such). Outcomes are the share of adults reporting fair or poor self-rated health, the share unable to see a doctor in the past year because of cost, the share reporting any current health insurance coverage, and the mean number of mentally and physically unhealthy days in the past 30. We also report two placebo outcomes: `share_any_exercise` (structurally unrelated to Medicaid) and `share_flu_shot` (partly insurance-linked).

Our secondary sample is the individual-level IPUMS American Community Survey extract for adults aged 19–64 during 2019–2024 (11.0 million observations). ACS outcomes are indicators for Medicaid coverage, uninsurance, any coverage, and private coverage. Subgroup analyses restrict to adults below 138 percent of the federal poverty level (1.34 million observations) and to low-income adults in the 39 states that had implemented ACA Medicaid expansion as of 2019 (900,000 observations).

Our state-level intensity panel is constructed from CMS Unwinding Data reports through August 2024, supplemented with KFF’s public database. We compute three intensity measures: the cumulative procedural share through August 2024 (our primary measure, range 0.41–0.92, median 0.73), the 2023-only annual procedural share (range 0.24–0.93), and the annualized 2023 procedural disenrollment rate per 1,000 pre-unwinding Medicaid enrollees. The three measures are highly correlated (pairwise $r = 0.93$) and produce near-identical results.

Empirical strategy

The primary specification is a continuous-treatment difference-in-differences with state and year fixed effects:

$$y_{st} = \alpha_s + \alpha_t + (\text{Post}_t \times \text{intensity}_s) + \epsilon_{st}$$

where Post_t equals 1 for $t \geq 2023$ and intensity_s is the state’s time-invariant cumulative procedural disenrollment share. Standard errors are clustered at the state level. Under parallel-trends and monotonic-dose assumptions, ϵ_{st} recovers a population-average dose-response effect consistent with Callaway, Goodman-Bacon, and Sant’Anna’s continuous-treatment framework. We additionally report a binary-intensity specification (indicator for above-median intensity), pre-trend placebos on 2019–2022, and an event-study specification with 2022 as the reference year. The ACS specification is analogous at the individual level, weighted by the IPUMS person weight and clustered at state.

A null is uninformative unless adequately powered. We calculate the minimum detectable effect at 80 percent power (two-sided $\alpha = 0.05$) scaled to the 90th-vs-10th-percentile intensity contrast (approximately 0.31) for each outcome and sample. This scales the per-unit-intensity minimum detectable effect to the policy-relevant difference between a low-intensity state such as North Carolina (procedural share 0.41) and a high-intensity state such as Texas (procedural share 0.88). We benchmark the resulting minimum detectable effects against the 5-percentage-point Medicaid coverage loss DeLeire documented for TennCare, the 3.94-percent-per-quarter prescription-fill decline Rome and colleagues documented for the unwinding, and depression and self-rated health effects from the Oregon Health Insurance Experiment.

Robustness

We report four additional robustness tests. First, we re-estimate the main difference-in-differences with state economic controls. We report both a specification using Bureau of Labor Statistics Local Area Unemployment Statistics only and a specification using BLS unemployment plus a Current Population Survey-derived prime-age (25–54) employment-to-population ratio, the latter replacing a less-accurate labor-force proxy. Second, we estimate the main specification on a placebo outcome (any exercise in the past month) that should be structurally unrelated to Medicaid procedural disenrollment. Third, we apply Rambachan and Roth’s relative-magnitudes honest confidence intervals to the ACS event-study coverage coefficients, using the full joint cluster-robust variance-covariance matrix. We compute these in both a Python manual implementation and a formal R replication using `HonestDiD` 0.2.8 with `method = "C-LF"` and `Mbarvec = c(0.5, 1.0, 2.0)`. Fourth, we apply Romano and Wolf’s step-down procedure to the five-outcome BRFSS primary family using a 199-replication state-clustered bootstrap that imposes the null by centering bootstrap coefficients on the original point estimates.

Study Results

Headline null

Exhibit 1 reports the primary BRFSS continuous-intensity difference-in-differences coefficients. No coefficient reaches conventional significance. The fair/poor health coefficient is -2.0 percentage points per unit of intensity (SE 1.2, $t = -1.70$); mental unhealthy days is -0.70 (SE 0.48); physical unhealthy days is -0.24 (SE 0.36); cost-related care barriers is +0.1 percentage points (SE 1.5); and any coverage is +0.9 percentage points (SE 1.5). The three health coefficients are directionally counterintuitive — slightly *better* self-reported health in higher-intensity states — but the magnitudes

are small and individually insignificant. Coverage goes the “wrong way” as well: higher-intensity states show slightly *more* any-coverage. Pre-trend placebos on 2019–2022 produce coefficients with $|t| < 1.2$ for every outcome, consistent with parallel trends.

The ACS individual-level cross-check confirms the BRFSS null. In the full-adult sample (11.0 million observations), the 2023 Medicaid coefficient is -0.2 percentage points (SE 2.0), the uninsured coefficient is -1.3 (SE 1.3), and the any-coverage coefficient is +1.3 (SE 1.3). In the below-138%-FPL subsample, the Medicaid coefficient is +3.4 percentage points (SE 4.2), the uninsured coefficient is -3.7 (SE 2.9). In the low-income-in-expansion-states subsample, the Medicaid coefficient is +1.5 (SE 4.5), uninsured is -2.9 (SE 3.2). Every coefficient, across every subsample, runs toward more coverage and less uninsurance in higher-intensity states. The BRFSS null is not an instrument artifact; the ACS with 11 million records sees the same null.

The null is adequately powered against TennCare-scale effects

Exhibit 2 reports minimum detectable effects at 80 percent power, scaled to the 90th-vs-10th-percentile intensity contrast. The BRFSS state-year panel’s minimum detectable effect is 1.11 percentage points on any-coverage, 0.83 percentage points on fair/poor health, and 0.26 days on mental unhealthy days — all well below the TennCare-scale benchmarks of 5 percentage points on coverage, 4 percentage points on fair/poor health, and 1.5 days on mental unhealthy days. The ACS all-adults sample has a minimum detectable effect of 4.2 percentage points on Medicaid and 3.5 percentage points on uninsurance, below the TennCare 5-percentage-point benchmark. A TennCare-magnitude population-level shock would be visible in both instruments and is not. The binding constraint on ACS precision is the cross-state intra-cluster correlation, which reduces the effective sample size from 11 million to approximately 2,600–3,200.

The ACS below-138%-FPL subsample is the exception. Its minimum detectable effect on Medicaid coverage is 9.7 percentage points, nearly twice the TennCare benchmark. The low-income null therefore rules out only large effects; it does not rule out a TennCare-sized effect on the most-exposed subpopulation. This is a load-bearing finding: the paper cannot claim that the null holds for the low-income adult population. It holds for the state adult population as a whole and for the all-adults ACS sample.

Robustness

Four robustness tests leave the null intact. First, state economic controls (BLS unemployment plus Current Population Survey prime-age employment-to-population ratio) change the direction-wrong health coefficients by a five-outcome median of approximately 3 percent in magnitude, with no sign flips and no specification reaching significance. Contemporaneous labor-market conditions do not explain the estimates. Second, the placebo exercise outcome is -0.10 percentage points (SE 1.99, 95% confidence interval [-4.0, +3.8]) — tightly centered on zero — and the primary fair/poor health coefficient is roughly 20 times larger in magnitude with a t-statistic roughly 30 times larger in absolute value. A state-level confounder operating through a channel unrelated to Medicaid coverage would have produced comparable movement in the placebo outcome. Third, Rambachan-Roth honest confidence intervals at $M = 1$ (allowing post-treatment pre-trend violations as large as the largest observed pre-treatment deviation) contain zero for every coverage outcome in every ACS sample in both 2023 and 2024, in both the Python manual implementation and the formal R `HonestDiD` 0.2.8 replication on the full joint variance-covariance matrix. Fourth, the Romano-Wolf-adjusted p-values for the five-outcome BRFSS primary family range from 0.518 on fair/poor health

to 0.925 on cost-related barriers. No outcome survives family-wise correction at the 0.05 level or at a relaxed 0.10 level. The family-wise test strengthens the null: the marginal direction-wrong health signals are statistical noise, not hidden effects.

Exhibit 3 plots the BRFSS event-study coefficients by year (with 2022 as the reference) for `share_fair_poor_health` and `share_any_coverage`, alongside the analogous ACS event-study for `has_medicaid` in the below-138%-FPL subsample, with 95 percent confidence bands and Rambachan-Roth relative-magnitudes bounds at $M = 1$. No year-by-intensity coefficient in either panel is distinguishable from zero in the post-treatment period.

Exhibit 4 reports the distribution of state-level unwinding intensity, which is the identifying variation behind the continuous-intensity difference-in-differences. The cumulative procedural disenrollment share ranges from 41 percent (North Carolina) to 92 percent (Texas), with a P10-P90 contrast of roughly 31 percentage points; the distribution is broadly continuous across the 50 states plus the District of Columbia, with no evidence that any single state drives the coefficient. Exhibit 5 summarizes the dose-response structure non-parametrically by assigning states to terciles of intensity and reporting the post-unwinding difference-in-differences for terciles 2 and 3 relative to tercile 1 (the lowest-intensity states). Across all five BRFSS outcomes, neither the tercile-2 nor the tercile-3 contrast is statistically distinguishable from zero, and there is no monotone rising pattern of the kind that would signal a dose-response health effect. The tercile-based specification thereby corroborates the continuous null without relying on the linear-dose functional form.

Discussion

We interpret the null as *precise* rather than *decisive*. Three hypotheses are consistent with the data, and neither BRFSS nor the ACS can discriminate among them.

Compositional substitution. Procedural disenrollees may have transitioned to alternative coverage — employer-sponsored insurance as the labor market recovered, marketplace coverage via the unwinding-related special enrollment period (active March 2023–July 2024), spouse or parent coverage, or re-enrolled Medicaid after GAO-documented state reinstatements. Four pieces of external evidence are consistent with this reading: the McIntyre et al. 2025 finding that the national uninsured rate rose only 0.6 percentage points between 2023 and 2024; the unusually strong 2023 labor market (3.5–4.0 percent unemployment); the GAO’s documentation that at least 35 states reinstated erroneously disenrolled enrollees; and our null on any coverage in both BRFSS and ACS. The hypothesis is *consistent with* the data but is not *identified* by the data. BRFSS and ACS are point-in-time cross-sections that cannot follow individuals across coverage transitions. Notably, our ACS `has_private` coefficients are all insignificant, which is mildly inconsistent with a large employer-coverage-substitution story.

Subpopulation concentration. Health harm may be real but concentrated in a narrow high-risk group — chronically ill adults, adults with serious mental illness or substance use disorders, cognitively impaired adults, or housing-unstable adults. These groups are both a small fraction of the total adult population (so meaningful effects are diluted below the detection threshold in population-weighted averages) and systematically underrepresented in both BRFSS (which requires a working phone and survey participation) and ACS (which requires a stable address). The minimum-detectable-effect analysis is directly load-bearing for this hypothesis: the ACS below-138%-FPL subsample has a 9.7-percentage-point minimum detectable effect on Medicaid coverage — nearly twice the 5-percentage-point TennCare benchmark — so the low-income null cannot rule out TennCare-scale effects on the most-exposed subpopulation. The 9.7-pp bound is precisely the

margin on which a concentrated-harm hypothesis would have to be adjudicated, and our data cannot adjudicate it.

Sampling-frame exclusion. Procedural disenrollment mechanically affects adults who cannot be reached at the address or phone number on file. Housing instability is strongly correlated with being unreachable, and both phone-based BRFSS and address-based ACS systematically undersample housing-unstable adults. The adults most exposed to procedural disenrollment are therefore the adults most likely to be absent from the sampling frame. Under this hypothesis, the null is not an artifact in the narrow sense of survey error; the instruments are structurally blind to the relevant margin.

Several confound stories *are* ruled out. Contemporaneous labor-market conditions do not explain the direction-wrong health estimates (economic-controls specification). State-level confounders operating through non-coverage channels would have moved the placebo exercise outcome, which is tightly centered on zero. Pre-trend violations as large as the largest observed pre-treatment deviation do not overturn the ACS coverage null (R `HonestDiD` 0.2.8 with `method = "C-LF"`). Measurement differences between intensity measures do not matter (the three measures produce near-identical coefficients). The multiple-testing correction absorbs the direction-wrong signals. What remains after these rule-outs is the irreducible fact that cross-sectional survey instruments cannot observe coverage transitions or the housing-unstable, chronically ill, and severely mentally ill adults most exposed to procedural disenrollment.

Limitations

The most fundamental limitation is that BRFSS and the ACS are cross-sectional surveys. They measure coverage at the time of the survey and cannot link an individual’s pre-unwinding Medicaid status to her post-unwinding coverage, health care utilization, or health outcomes. A clean “effect of disenrollment on the disenrolled” requires individual-level claims data with pre/post enrollment linkage. Our cumulative intensity measure averages forward through August 2024 and is applied to 2023 observations, implicitly assuming the unwinding was roughly linear through its first nine months; the annual procedural share robustness check produces the same qualitative pattern. Heterogeneous unwinding modalities (deadline length, outreach intensity, data-match use) are averaged out by the summary intensity measure. Race and ethnicity disparities documented by Rumalla and colleagues are not stratified in our state-year design; follow-up using ACS individual-level data could support race-stratified analysis. Finally, the null is conditional on an unusually strong 2023 labor market and an active federal marketplace special enrollment period; in weaker labor markets or without the special enrollment period, compositional substitution would likely be smaller and observed effects could be larger.

Policy Implications

Three implications follow.

First, at the state adult population level, the 2023 Medicaid unwinding did not produce detectable declines in adult self-reported health or any-coverage rates. The design is adequately powered against TennCare-scale effects on the all-adults margin and against OHIE-scale effects on mental health days. This is an informative null, not an uninformative one, and it is robust across two public-use survey instruments with different sampling methodologies, different outcome batteries, and an order-of-magnitude different sample size.

Second, the informative null does *not* mean procedural disenrollment was harmless. It means that population-level survey evidence cannot identify the relevant effect, because the effect most plausibly concentrates in subpopulations (chronically ill, mentally ill, housing-unstable) that are structurally underrepresented in BRFSS and the ACS, and because disenrollees with alternative coverage pathways — a substantial share, given the strong 2023 labor market — do not change population coverage rates. Administrative-burden reform can be justified on the grounds of Rummalla and colleagues’ racial disparity evidence, the GAO’s documentation of procedural compliance failures, Rome and colleagues’ prescription-fill declines, and the intrinsic normative concern that procedural disenrollment produces unnecessary coverage losses. What population-level survey evidence cannot do is adjudicate a downstream population health signal.

Third, the decisive test is an individual-level analysis that follows actually-disenrolled adults across coverage transitions and measures their utilization and health directly. We commit to such an analysis using T-MSIS / Transformed Medicaid Statistical Information System Analytic Files 2019–2024 as a companion paper, with a Data Use Agreement application to CMS ResDAC under way. That analysis can measure re-enrollment probabilities at 30, 60, 90, and 180 days; preventive, chronic, and emergency utilization trajectories for disenrollees stratified by reason; and subpopulation-specific effects for chronic-condition, mental-illness, and substance-use-disorder cohorts. The present paper and the companion T-MSIS paper together can answer what the policy question requires: which disenrollees were harmed, how much, and through what mechanisms.

For the current administrative-burden reform conversation, the immediate implication is that state procedural-disenrollment intensity is not a summary statistic that tracks measurable population health harm within the first 12 to 18 months of an unwinding episode. Reform advocates should frame the case on unnecessary coverage losses, distributional fairness, and individual-level harm in claims data, rather than on a population-level survey signal that the best available instruments cannot detect.

Exhibits

Exhibit 1. BRFSS and ACS primary difference-in-differences coefficients, state-clustered standard errors. *Source:* `analysis/tables/main_did_summary.csv` (BRFSS); `analysis/tables/acs_individual_results.csv` (ACS).

Exhibit 2. Minimum detectable effects at 80 percent power, scaled to 90th-vs-10th-percentile intensity contrast, with TennCare / Rome / OHIE benchmarks. *Source:* `analysis/tables/mde_and_power.csv`.

Exhibit 3. BRFSS and ACS event-study coefficients with 95 percent pointwise confidence bands and Rambachan-Roth relative-magnitudes bounds at $M = 1$. Top panel: BRFSS year-by-intensity coefficients for `share_fair_poor_health` and `share_any_coverage`, 2019–2023, with 2022 as the reference year. Bottom panel: ACS `has_medicaid`, below-138%-FPL subsample, 2019–2024. *Sources:* `analysis/figures/brfss_event_study_main.png` (top panel); `analysis/figures/acs_event_study_honestdid_below_138fpl.png` (bottom panel). *Note:* The BRFSS structure is continuous-dose with a single post-treatment period (2023); under this structure the Callaway-Sant’Anna, Sun-Abraham, and de Chaisemartin-D’Haultfoeuille staggered-adoption estimators are not applicable, and Goodman-Bacon decomposition is likewise not applicable (there are no 2×2 problematic comparisons). Year-by-intensity interactions with state and year fixed effects and state-clustered (CRV1) standard errors are the correct specification

for this design. Pointwise 95 percent confidence intervals are reported; Rambachan-Roth relative-magnitudes bounds address the concern that pointwise intervals are inappropriate for interpreting the full dynamic path.

Exhibit 4. Distribution of state-level unwinding intensity (cumulative procedural disenrollment share through August 2024). Panel A: ranked horizontal bar chart of all 50 states plus the District of Columbia, with P10, median, and P90 reference lines. Panel B: histogram of the same distribution, shading the P10-P90 policy-relevant dose contrast used to translate difference-in-differences coefficients into comparative-statics magnitudes in Exhibit 2. *Source:* `analysis/figures/intensity_distribution.png`; underlying values in `analysis/tables/intensity_distribution.csv`.

Exhibit 5. Tercile dose-response: post-unwinding difference-in-differences coefficients for tercile 2 and tercile 3 (high-intensity) states relative to tercile 1 (low-intensity, reference), across the five BRFSS primary outcomes, with state-clustered 95 percent confidence intervals. *Source:* `analysis/figures/dose_response_terciles.png`; underlying coefficients in `analysis/tables/dose_response_terciles.csv`.

Online Appendix Exhibits

- **Appendix Exhibit A1.** BRFSS raw outcome trends by above-/below-median unwinding intensity, 2019–2023. *Source:* `analysis/figures/raw_trends_*.png`.
- **Appendix Exhibit A2.** BRFSS difference-in-differences with state economic controls (BLS unemployment only; BLS plus CPS prime-age EPOP). *Sources:* `analysis/tables/economic_controls_summary.csv`, `analysis/tables/economic_controls_cps_summary.csv`.
- **Appendix Exhibit A3.** BRFSS placebo outcome difference-in-differences (exercise, flu shot). *Source:* `analysis/tables/placebo_outcome_summary.csv`.
- **Appendix Exhibit A4.** Romano-Wolf step-down multiple-testing correction for the five-outcome BRFSS primary family. *Source:* `analysis/tables/romano_wolf_paper9.csv`.
- **Appendix Exhibit A5.** R HonestDiD 0.2.8 formal replication panel, 24 sample $\times$ outcome $\times$ post-year rows. *Sources:* `analysis/tables/r_honestdid_acs_results.csv`, `analysis/tables/python_r_honestdid_acs_comparison.csv`.
- **Appendix Exhibit A6.** BRFSS event-study coefficients for the three share outcomes and the two days outcomes (full five-outcome panel; the main text shows the two headline outcomes). *Sources:* `analysis/figures/brfss_event_study_shares_app.png`, `analysis/figures/brfss_event_study_days_app.png`; coefficient file at `analysis/tables/brfss_event_`

Notes

1. Finkelstein A, Taubman S, Wright B, et al. The Oregon Health Insurance Experiment: evidence from the first year. *Q J Econ.* 2012;127(3):1057–1106. doi:10.1093/qje/qjs020.
2. Baicker K, Taubman SL, Allen HL, et al. The Oregon experiment — effects of Medicaid on clinical outcomes. *N Engl J Med.* 2013;368(18):1713–1722. doi:10.1056/NEJMsa1212321.
3. Miller S, Johnson N, Wherry LR. Medicaid and mortality: new evidence from linked survey and administrative data. *Q J Econ.* 2021;136(3):1783–1829. doi:10.1093/qje/qjab004.
4. DeLeire T. The effect of disenrollment from Medicaid on employment, insurance coverage, health and health care utilization. NBER Working Paper 24899. 2018. doi:10.3386/w24899.

5. McIntyre A, Morein M, Kim J, Figueroa JF, Sommers BD. US coverage changes during Medicaid unwinding in 2023. *JAMA Health Forum.* 2025;6(10):e253887. doi:10.1001/jamahealthforum.2025.3887.
6. Rome BN, Han J, McIntyre A, Kesselheim AS, Sommers BD. Changes in medication use during Medicaid continuous enrollment and unwinding. *JAMA Health Forum.* 2026;7(1):e255890. doi:10.1001/jamahealthforum.2025.5890.
7. Rumalla KC, Nelson DB, McConnell KJ, Zhu JM. Racial and ethnic disparities in Medicaid disenrollment after the end of the COVID-19 public health emergency. *JAMA Intern Med.* 2024;184(8):987–989. doi:10.1001/jamainternmed.2024.1503.
8. Dasgupta K, Solomon KT. The effect of ending the pandemic-related mandate of continuous Medicaid coverage on health insurance coverage and economic well-being. *Health Serv Res.* 2026;61(1):1–11. doi:10.1111/1475-6773.70021.
9. Herd P, Moynihan DP. *Administrative Burden: Policymaking by Other Means.* Russell Sage Foundation; 2018.
10. Kenney GM, Haley JM, Wang R, Buettgens M, Anderson N. State variation in Medicaid and CHIP unwinding for children and adults as of November 2023. Urban Institute; 2024. [REPORT — pre-submission verification required.]
11. U.S. Government Accountability Office. Medicaid: Federal oversight of state eligibility redeterminations should reflect lessons learned after COVID-19. GAO-24-106883. 2024. [REPORT — pre-submission verification required.]
12. Callaway B, Goodman-Bacon A, Sant’Anna PHC. Difference-in-differences with a continuous treatment. 2021. [METHODS — not in verified bibliography; pre-submission verification required.]
13. Rambachan A, Roth J. A more credible approach to parallel trends. *Rev Econ Stud.* 2023. [METHODS — not in verified bibliography; pre-submission verification required.]
14. Romano JP, Wolf M. Exact and approximate stepdown methods for multiple hypothesis testing. *J Am Stat Assoc.* 2005;100(469):94–108. [METHODS — not in verified bibliography; pre-submission verification required.]
15. HonestDiD R package, version 0.2.8. Rambachan A, Roth J. 2023. <https://github.com/asheshrambachan/HonestDiD>. [SOFTWARE — pre-submission verification required.]